

Using GNU Fortran

For GCC version 9.2.0

(BLD = 1649)

The gfortran team

Published by the Free Software Foundation
51 Franklin Street, Fifth Floor
Boston, MA 02110-1301, USA

Copyright © 1999-2019 Free Software Foundation, Inc.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with the Invariant Sections being “Funding Free Software”, the Front-Cover Texts being (a) (see below), and with the Back-Cover Texts being (b) (see below). A copy of the license is included in the section entitled “GNU Free Documentation License”.

(a) The FSF’s Front-Cover Text is:

A GNU Manual

(b) The FSF’s Back-Cover Text is:

You have freedom to copy and modify this GNU Manual, like GNU software. Copies published by the Free Software Foundation raise funds for GNU development.

Short Contents

1	Introduction	1
	Invoking GNU Fortran	
2	GNU Fortran Command Options	7
3	Runtime: Influencing runtime behavior with environment variables	33
	Language Reference	
4	Fortran standards status	39
5	Compiler Characteristics	45
6	Extensions	51
7	Mixed-Language Programming	73
8	Coarray Programming	87
9	Intrinsic Procedures	111
10	Intrinsic Modules	287
	Contributing	293
	GNU General Public License	297
	GNU Free Documentation License	309
	Funding Free Software	317
	Option Index	319
	Keyword Index	321

Table of Contents

1	Introduction	1
1.1	About GNU Fortran	1
1.2	GNU Fortran and GCC	2
1.3	Preprocessing and conditional compilation	2
1.4	GNU Fortran and G77	3
1.5	Project Status	3
1.6	Standards	4
1.6.1	Varying Length Character Strings	4
	Part I: Invoking GNU Fortran	5
2	GNU Fortran Command Options	7
2.1	Option summary	7
2.2	Options controlling Fortran dialect	8
2.3	Enable and customize preprocessing	13
2.4	Options to request or suppress errors and warnings	15
2.5	Options for debugging your program or GNU Fortran	20
2.6	Options for directory search	22
2.7	Influencing the linking step	22
2.8	Influencing runtime behavior	22
2.9	Options for code generation conventions	23
2.10	Options for interoperability with other languages	30
2.11	Environment variables affecting <code>gfortran</code>	31
3	Runtime: Influencing runtime behavior with environment variables	33
3.1	TMPDIR—Directory for scratch files	33
3.2	GFORTRAN_STDIN_UNIT—Unit number for standard input	33
3.3	GFORTRAN_STDOUT_UNIT—Unit number for standard output	33
3.4	GFORTRAN_STDERR_UNIT—Unit number for standard error	33
3.5	GFORTRAN_UNBUFFERED_ALL—Do not buffer I/O on all units	33
3.6	GFORTRAN_UNBUFFERED_PRECONNECTED—Do not buffer I/O on preconnected units	33
3.7	GFORTRAN_SHOW_LOCUS—Show location for runtime errors	33
3.8	GFORTRAN_OPTIONAL_PLUS—Print leading + where permitted	34
3.9	GFORTRAN_LIST_SEPARATOR—Separator for list output	34
3.10	GFORTRAN_CONVERT_UNIT—Set endianness for unformatted I/O	34
3.11	GFORTRAN_ERROR_BACKTRACE—Show backtrace on run-time errors	35
3.12	GFORTRAN_FORMATTED_BUFFER_SIZE—Set buffer size for formatted I/O	35

3.13	GFORTTRAN_UNFORMATTED_BUFFER_SIZE—Set buffer size for unformatted I/O.....	35
------	---	----

Part II: Language Reference 37

4 Fortran standards status 39

4.1	Fortran 2003 status	39
4.2	Fortran 2008 status	40
4.3	Status of Fortran 2018 support	42
4.3.1	TS 29113 Status (Further Interoperability with C)	42
4.3.2	TS 18508 Status (Additional Parallel Features)	43

5 Compiler Characteristics 45

5.1	KIND Type Parameters	45
5.2	Internal representation of LOGICAL variables	45
5.3	Evaluation of logical expressions	46
5.4	MAX and MIN intrinsics with REAL NaN arguments	46
5.5	Thread-safety of the runtime library	46
5.6	Data consistency and durability	47
5.7	Files opened without an explicit ACTION= specifier	48
5.8	File operations on symbolic links	48
5.9	File format of unformatted sequential files	48
5.10	Asynchronous I/O	49

6 Extensions 51

6.1	Extensions implemented in GNU Fortran	51
6.1.1	Old-style kind specifications	51
6.1.2	Old-style variable initialization	51
6.1.3	Extensions to namelist	52
6.1.4	X format descriptor without count field	53
6.1.5	Commas in FORMAT specifications	53
6.1.6	Missing period in FORMAT specifications	53
6.1.7	I/O item lists	53
6.1.8	Q exponent-letter	53
6.1.9	BOZ literal constants	53
6.1.10	Real array indices	54
6.1.11	Unary operators	54
6.1.12	Implicitly convert LOGICAL and INTEGER values	54
6.1.13	Hollerith constants support	55
6.1.14	Cray pointers	55
6.1.15	CONVERT specifier	57
6.1.16	OpenMP	58
6.1.17	OpenACC	58
6.1.18	Argument list functions %VAL, %REF and %LOC	59
6.1.19	Read/Write after EOF marker	59
6.1.20	STRUCTURE and RECORD	59
6.1.21	UNION and MAP	62

6.1.22	Type variants for integer intrinsics.....	64
6.1.23	AUTOMATIC and STATIC attributes	65
6.1.24	Extended math intrinsics.....	66
6.1.25	Form feed as whitespace.....	66
6.1.26	TYPE as an alias for PRINT	66
6.1.27	%LOC as an rvalue	66
6.1.28	.XOR. operator	67
6.1.29	Bitwise logical operators	67
6.1.30	Extended I/O specifiers	67
6.1.31	Legacy PARAMETER statements.....	69
6.1.32	Default exponents.....	69
6.2	Extensions not implemented in GNU Fortran	69
6.2.1	ENCODE and DECODE statements.....	69
6.2.2	Variable FORMAT expressions	70
6.2.3	Alternate complex function syntax.....	70
6.2.4	Volatile COMMON blocks.....	70
6.2.5	OPEN (... NAME=).....	70
6.2.6	Q edit descriptor.....	70
7	Mixed-Language Programming.....	73
7.1	Interoperability with C.....	73
7.1.1	Intrinsic Types.....	73
7.1.2	Derived Types and struct.....	73
7.1.3	Interoperable Global Variables.....	74
7.1.4	Interoperable Subroutines and Functions.....	74
7.1.5	Working with Pointers	76
7.1.6	Further Interoperability of Fortran with C.....	78
7.2	GNU Fortran Compiler Directives.....	79
7.2.1	ATTRIBUTES directive.....	79
7.2.2	UNROLL directive	80
7.2.3	BUILTIN directive	80
7.3	Non-Fortran Main Program.....	80
7.3.1	_gfortran_set_args — Save command-line arguments... 81	
7.3.2	_gfortran_set_options — Set library option flags..... 81	
7.3.3	_gfortran_set_convert — Set endian conversion..... 82	
7.3.4	_gfortran_set_record_marker — Set length of record markers	83
7.3.5	_gfortran_set_fpe — Enable floating point exception traps	83
7.3.6	_gfortran_set_max_subrecord_length — Set subrecord length	84
7.4	Naming and argument-passing conventions	84
7.4.1	Naming conventions.....	84
7.4.2	Argument passing conventions.....	85

8	Coarray Programming	87
8.1	Type and enum ABI Documentation	87
8.1.1	caf_token_t	87
8.1.2	caf_register_t	87
8.1.3	caf_deregister_t	87
8.1.4	caf_reference_t	87
8.1.5	caf_team_t	89
8.2	Function ABI Documentation	89
8.2.1	_gfortran_caf_init — Initiation function	90
8.2.2	_gfortran_caf_finish — Finalization function	90
8.2.3	_gfortran_caf_this_image — Querying the image number	90
8.2.4	_gfortran_caf_num_images — Querying the maximal number of images	91
8.2.5	_gfortran_caf_image_status — Query the status of an image	91
8.2.6	_gfortran_caf_failed_images — Get an array of the indexes of the failed images	91
8.2.7	_gfortran_caf_stopped_images — Get an array of the indexes of the stopped images	92
8.2.8	_gfortran_caf_register — Registering coarrays	92
8.2.9	_gfortran_caf_deregister — Deregistering coarrays	93
8.2.10	_gfortran_caf_is_present — Query whether an allocatable or pointer component in a derived type coarray is allocated	94
8.2.11	_gfortran_caf_send — Sending data from a local image to a remote image	94
8.2.12	_gfortran_caf_get — Getting data from a remote image	95
8.2.13	_gfortran_caf_sendget — Sending data between remote images	96
8.2.14	_gfortran_caf_send_by_ref — Sending data from a local image to a remote image with enhanced referencing options	97
8.2.15	_gfortran_caf_get_by_ref — Getting data from a remote image using enhanced references	98
8.2.16	_gfortran_caf_sendget_by_ref — Sending data between remote images using enhanced references on both sides	99
8.2.17	_gfortran_caf_lock — Locking a lock variable	101
8.2.18	_gfortran_caf_unlock — Unlocking a lock variable	101
8.2.19	_gfortran_caf_event_post — Post an event	102
8.2.20	_gfortran_caf_event_wait — Wait that an event occurred	102
8.2.21	_gfortran_caf_event_query — Query event count	103
8.2.22	_gfortran_caf_sync_all — All-image barrier	103
8.2.23	_gfortran_caf_sync_images — Barrier for selected images	104
8.2.24	_gfortran_caf_sync_memory — Wait for completion of segment-memory operations	104

8.2.25	<code>_gfortran_caf_error_stop</code> — Error termination with exit code.....	105
8.2.26	<code>_gfortran_caf_error_stop_str</code> — Error termination with string.....	105
8.2.27	<code>_gfortran_caf_fail_image</code> — Mark the image failed and end its execution.....	105
8.2.28	<code>_gfortran_caf_atomic_define</code> — Atomic variable assignment.....	105
8.2.29	<code>_gfortran_caf_atomic_ref</code> — Atomic variable reference.....	106
8.2.30	<code>_gfortran_caf_atomic_cas</code> — Atomic compare and swap.....	106
8.2.31	<code>_gfortran_caf_atomic_op</code> — Atomic operation.....	107
8.2.32	<code>_gfortran_caf_co_broadcast</code> — Sending data to all images.....	107
8.2.33	<code>_gfortran_caf_co_max</code> — Collective maximum reduction.....	108
8.2.34	<code>_gfortran_caf_co_min</code> — Collective minimum reduction.....	108
8.2.35	<code>_gfortran_caf_co_sum</code> — Collective summing reduction.....	109
8.2.36	<code>_gfortran_caf_co_reduce</code> — Generic collective reduction.....	109
9	Intrinsic Procedures	111
9.1	Introduction to intrinsic procedures.....	111
9.2	<code>ABORT</code> — Abort the program.....	111
9.3	<code>ABS</code> — Absolute value.....	112
9.4	<code>ACCESS</code> — Checks file access modes.....	113
9.5	<code>ACHAR</code> — Character in ASCII collating sequence.....	113
9.6	<code>ACOS</code> — Arccosine function.....	114
9.7	<code>ACOSD</code> — Arccosine function, degrees.....	115
9.8	<code>ACOSH</code> — Inverse hyperbolic cosine function.....	115
9.9	<code>ADJUSTL</code> — Left adjust a string.....	116
9.10	<code>ADJUSTR</code> — Right adjust a string.....	116
9.11	<code>AIMAG</code> — Imaginary part of complex number.....	117
9.12	<code>AINT</code> — Truncate to a whole number.....	117
9.13	<code>ALARM</code> — Execute a routine after a given delay.....	118
9.14	<code>ALL</code> — All values in <i>MASK</i> along <i>DIM</i> are true.....	119
9.15	<code>ALLOCATED</code> — Status of an allocatable entity.....	120
9.16	<code>AND</code> — Bitwise logical AND.....	120
9.17	<code>ANINT</code> — Nearest whole number.....	121
9.18	<code>ANY</code> — Any value in <i>MASK</i> along <i>DIM</i> is true.....	122
9.19	<code>ASIN</code> — Arcsine function.....	123
9.20	<code>ASIND</code> — Arcsine function, degrees.....	123
9.21	<code>ASINH</code> — Inverse hyperbolic sine function.....	124
9.22	<code>ASSOCIATED</code> — Status of a pointer or pointer/target pair....	124
9.23	<code>ATAN</code> — Arctangent function.....	126

9.24	ATAND — Arctangent function, degrees	126
9.25	ATAN2 — Arctangent function	127
9.26	ATAN2D — Arctangent function, degrees	128
9.27	ATANH — Inverse hyperbolic tangent function	129
9.28	ATOMIC_ADD — Atomic ADD operation	129
9.29	ATOMIC_AND — Atomic bitwise AND operation	130
9.30	ATOMIC_CAS — Atomic compare and swap	131
9.31	ATOMIC_DEFINE — Setting a variable atomically	131
9.32	ATOMIC_FETCH_ADD — Atomic ADD operation with prior fetch	132
9.33	ATOMIC_FETCH_AND — Atomic bitwise AND operation with prior fetch	133
9.34	ATOMIC_FETCH_OR — Atomic bitwise OR operation with prior fetch	133
9.35	ATOMIC_FETCH_XOR — Atomic bitwise XOR operation with prior fetch	134
9.36	ATOMIC_OR — Atomic bitwise OR operation	135
9.37	ATOMIC_REF — Obtaining the value of a variable atomically ..	136
9.38	ATOMIC_XOR — Atomic bitwise OR operation	136
9.39	BACKTRACE — Show a backtrace	137
9.40	BESSEL_J0 — Bessel function of the first kind of order 0	137
9.41	BESSEL_J1 — Bessel function of the first kind of order 1	138
9.42	BESSEL_JN — Bessel function of the first kind	138
9.43	BESSEL_Y0 — Bessel function of the second kind of order 0 ..	139
9.44	BESSEL_Y1 — Bessel function of the second kind of order 1 ..	140
9.45	BESSEL_YN — Bessel function of the second kind	140
9.46	BGE — Bitwise greater than or equal to	141
9.47	BGT — Bitwise greater than	141
9.48	BIT_SIZE — Bit size inquiry function	142
9.49	BLE — Bitwise less than or equal to	142
9.50	BLT — Bitwise less than	143
9.51	BTEST — Bit test function	143
9.52	C_ASSOCIATED — Status of a C pointer	144
9.53	C_F_POINTER — Convert C into Fortran pointer	144
9.54	C_F_PROCPTR — Convert C into Fortran procedure pointer	145
9.55	C_FUNLOC — Obtain the C address of a procedure	146
9.56	C_LOC — Obtain the C address of an object	146
9.57	C_SIZEOF — Size in bytes of an expression	147
9.58	CEILING — Integer ceiling function	148
9.59	CHAR — Character conversion function	148
9.60	CHDIR — Change working directory	149
9.61	CHMOD — Change access permissions of files	150
9.62	CMPLX — Complex conversion function	150
9.63	CO_BROADCAST — Copy a value to all images the current set of images	151
9.64	CO_MAX — Maximal value on the current set of images	152
9.65	CO_MIN — Minimal value on the current set of images	153

9.66	CO_REDUCE — Reduction of values on the current set of images	153
9.67	CO_SUM — Sum of values on the current set of images.....	155
9.68	COMMAND_ARGUMENT_COUNT — Get number of command line arguments.....	155
9.69	COMPILER_OPTIONS — Options passed to the compiler.....	156
9.70	COMPILER_VERSION — Compiler version string.....	156
9.71	COMPLEX — Complex conversion function.....	157
9.72	CONJG — Complex conjugate function.....	157
9.73	COS — Cosine function.....	158
9.74	COSD — Cosine function, degrees.....	159
9.75	COSH — Hyperbolic cosine function.....	159
9.76	COTAN — Cotangent function.....	160
9.77	COTAND — Cotangent function, degrees.....	160
9.78	COUNT — Count function.....	161
9.79	CPU_TIME — CPU elapsed time in seconds.....	162
9.80	CSHIFT — Circular shift elements of an array.....	163
9.81	CTIME — Convert a time into a string.....	163
9.82	DATE_AND_TIME — Date and time subroutine.....	164
9.83	DBLE — Double conversion function.....	165
9.84	DCMPLX — Double complex conversion function.....	166
9.85	DIGITS — Significant binary digits function.....	166
9.86	DIM — Positive difference.....	167
9.87	DOT_PRODUCT — Dot product function.....	167
9.88	DPROD — Double product function.....	168
9.89	DREAL — Double real part function.....	169
9.90	DSHIFTL — Combined left shift.....	169
9.91	DSHIFTR — Combined right shift.....	170
9.92	DTIME — Execution time subroutine (or function).....	170
9.93	EOSHIFT — End-off shift elements of an array.....	171
9.94	EPSILON — Epsilon function.....	172
9.95	ERF — Error function.....	173
9.96	ERFC — Error function.....	173
9.97	ERFC_SCALED — Error function.....	174
9.98	ETIME — Execution time subroutine (or function).....	174
9.99	EVENT_QUERY — Query whether a coarray event has occurred	175
9.100	EXECUTE_COMMAND_LINE — Execute a shell command.....	176
9.101	EXIT — Exit the program with status.....	177
9.102	EXP — Exponential function.....	177
9.103	EXPONENT — Exponent function.....	178
9.104	EXTENDS_TYPE_OF — Query dynamic type for extension....	178
9.105	FDATE — Get the current time as a string.....	179
9.106	FGET — Read a single character in stream mode from stdin	179
9.107	FGETC — Read a single character in stream mode.....	180
9.108	FINDLOC — Search an array for a value.....	181
9.109	FLOOR — Integer floor function.....	182

9.110	FLUSH — Flush I/O unit(s)	183
9.111	FNUM — File number function	183
9.112	FPUT — Write a single character in stream mode to stdout ..	184
9.113	FPUTC — Write a single character in stream mode	185
9.114	FRACTION — Fractional part of the model representation ...	186
9.115	FREE — Frees memory	186
9.116	FSEEK — Low level file positioning subroutine	186
9.117	FSTAT — Get file status	188
9.118	FTELL — Current stream position	188
9.119	GAMMA — Gamma function	189
9.120	GERROR — Get last system error message	189
9.121	GETARG — Get command line arguments	190
9.122	GET_COMMAND — Get the entire command line	191
9.123	GET_COMMAND_ARGUMENT — Get command line arguments ...	191
9.124	GETCWD — Get current working directory	192
9.125	GETENV — Get an environmental variable	193
9.126	GET_ENVIRONMENT_VARIABLE — Get an environmental variable	193
9.127	GETGID — Group ID function	194
9.128	GETLOG — Get login name	194
9.129	GETPID — Process ID function	195
9.130	GETUID — User ID function	195
9.131	GMTIME — Convert time to GMT info	196
9.132	HOSTNM — Get system host name	196
9.133	HUGE — Largest number of a kind	197
9.134	HYPOT — Euclidean distance function	197
9.135	IACHAR — Code in ASCII collating sequence	198
9.136	IALL — Bitwise AND of array elements	198
9.137	IAND — Bitwise logical and	199
9.138	IANY — Bitwise OR of array elements	200
9.139	IARGC — Get the number of command line arguments	201
9.140	IBCLR — Clear bit	201
9.141	IBITS — Bit extraction	202
9.142	IBSET — Set bit	202
9.143	ICHAR — Character-to-integer conversion function	203
9.144	IDATE — Get current local time subroutine (day/month/year)	204
9.145	IEOR — Bitwise logical exclusive or	204
9.146	IERRNO — Get the last system error number	205
9.147	IMAGE_INDEX — Function that converts a cosubscript to an image index	205
9.148	INDEX — Position of a substring within a string	206
9.149	INT — Convert to integer type	207
9.150	INT2 — Convert to 16-bit integer type	207
9.151	INT8 — Convert to 64-bit integer type	208
9.152	IOR — Bitwise logical or	208
9.153	IPARITY — Bitwise XOR of array elements	209
9.154	IRAND — Integer pseudo-random number	210

9.155	IS_CONTIGUOUS — Test whether an array is contiguous.....	210
9.156	IS_IOSTAT_END — Test for end-of-file value	211
9.157	IS_IOSTAT_EOR — Test for end-of-record value.....	211
9.158	ISATTY — Whether a unit is a terminal device.....	212
9.159	ISHFT — Shift bits	212
9.160	ISHFTC — Shift bits circularly	213
9.161	ISNAN — Test for a NaN.....	214
9.162	ITIME — Get current local time subroutine (hour/minutes/seconds)	214
9.163	KILL — Send a signal to a process	215
9.164	KIND — Kind of an entity.....	215
9.165	LBOUND — Lower dimension bounds of an array	216
9.166	LCOBOUND — Lower codimension bounds of an array.....	216
9.167	LEADZ — Number of leading zero bits of an integer.....	217
9.168	LEN — Length of a character entity	217
9.169	LEN_TRIM — Length of a character entity without trailing blank characters	218
9.170	LGE — Lexical greater than or equal	218
9.171	LGT — Lexical greater than	219
9.172	LINK — Create a hard link.....	220
9.173	LLE — Lexical less than or equal	220
9.174	LLT — Lexical less than.....	221
9.175	LNBLNK — Index of the last non-blank character in a string..	221
9.176	LOC — Returns the address of a variable	222
9.177	LOG — Natural logarithm function	222
9.178	LOG10 — Base 10 logarithm function.....	223
9.179	LOG_GAMMA — Logarithm of the Gamma function	223
9.180	LOGICAL — Convert to logical type.....	224
9.181	LONG — Convert to integer type.....	224
9.182	LSHIFT — Left shift bits.....	225
9.183	LSTAT — Get file status.....	225
9.184	LTIME — Convert time to local time info.....	226
9.185	MALLOC — Allocate dynamic memory	227
9.186	MASKL — Left justified mask	227
9.187	MASKR — Right justified mask.....	228
9.188	MATMUL — matrix multiplication	228
9.189	MAX — Maximum value of an argument list.....	229
9.190	MAXEXPONENT — Maximum exponent of a real kind.....	229
9.191	MAXLOC — Location of the maximum value within an array..	230
9.192	MAXVAL — Maximum value of an array	231
9.193	MCLOCK — Time function	231
9.194	MCLOCK8 — Time function (64-bit)	232
9.195	MERGE — Merge variables	232
9.196	MERGE_BITS — Merge of bits under mask	233
9.197	MIN — Minimum value of an argument list	233
9.198	MINEXPONENT — Minimum exponent of a real kind	234
9.199	MINLOC — Location of the minimum value within an array..	234
9.200	MINVAL — Minimum value of an array	235

9.201	MOD — Remainder function	236
9.202	MODULO — Modulo function	237
9.203	MOVE_ALLOC — Move allocation from one object to another	237
9.204	MVBITS — Move bits from one integer to another	238
9.205	NEAREST — Nearest representable number	239
9.206	NEW_LINE — New line character	239
9.207	NINT — Nearest whole number	240
9.208	NORM2 — Euclidean vector norms	240
9.209	NOT — Logical negation	241
9.210	NULL — Function that returns an disassociated pointer	241
9.211	NUM_IMAGES — Function that returns the number of images	242
9.212	OR — Bitwise logical OR	243
9.213	PACK — Pack an array into an array of rank one	243
9.214	PARITY — Reduction with exclusive OR	244
9.215	PERROR — Print system error message	245
9.216	POPCNT — Number of bits set	245
9.217	POPPAR — Parity of the number of bits set	246
9.218	PRECISION — Decimal precision of a real kind	246
9.219	PRESENT — Determine whether an optional dummy argument is specified	247
9.220	PRODUCT — Product of array elements	247
9.221	RADIX — Base of a model number	248
9.222	RAN — Real pseudo-random number	248
9.223	RAND — Real pseudo-random number	248
9.224	RANDOM_INIT — Initialize a pseudo-random number generator	249
9.225	RANDOM_NUMBER — Pseudo-random number	250
9.226	RANDOM_SEED — Initialize a pseudo-random number sequence	250
9.227	RANGE — Decimal exponent range	251
9.228	RANK — Rank of a data object	252
9.229	REAL — Convert to real type	252
9.230	RENAME — Rename a file	253
9.231	REPEAT — Repeated string concatenation	254
9.232	RESHAPE — Function to reshape an array	254
9.233	RRSPACING — Reciprocal of the relative spacing	255
9.234	RSHIFT — Right shift bits	255
9.235	SAME_TYPE_AS — Query dynamic types for equality	256
9.236	SCALE — Scale a real value	256
9.237	SCAN — Scan a string for the presence of a set of characters	256
9.238	SECNDS — Time function	257
9.239	SECOND — CPU time function	258
9.240	SELECTED_CHAR_KIND — Choose character kind	258
9.241	SELECTED_INT_KIND — Choose integer kind	259
9.242	SELECTED_REAL_KIND — Choose real kind	259

9.243	SET_EXPONENT — Set the exponent of the model	260
9.244	SHAPE — Determine the shape of an array	261
9.245	SHIFTA — Right shift with fill	262
9.246	SHIFTL — Left shift	262
9.247	SHIFTR — Right shift	262
9.248	SIGN — Sign copying function	263
9.249	SIGNAL — Signal handling subroutine (or function)	264
9.250	SIN — Sine function	264
9.251	SIND — Sine function, degrees	265
9.252	SINH — Hyperbolic sine function	266
9.253	SIZE — Determine the size of an array	266
9.254	SIZEOF — Size in bytes of an expression	267
9.255	SLEEP — Sleep for the specified number of seconds	267
9.256	SPACING — Smallest distance between two numbers of a given type	268
9.257	SPREAD — Add a dimension to an array	268
9.258	SQRT — Square-root function	269
9.259	SRAND — Reinitialize the random number generator	269
9.260	STAT — Get file status	270
9.261	STORAGE_SIZE — Storage size in bits	271
9.262	SUM — Sum of array elements	272
9.263	SYMLNK — Create a symbolic link	272
9.264	SYSTEM — Execute a shell command	273
9.265	SYSTEM_CLOCK — Time function	273
9.266	TAN — Tangent function	274
9.267	TAND — Tangent function, degrees	275
9.268	TANH — Hyperbolic tangent function	276
9.269	THIS_IMAGE — Function that returns the cosubscript index of this image	276
9.270	TIME — Time function	277
9.271	TIME8 — Time function (64-bit)	278
9.272	TINY — Smallest positive number of a real kind	278
9.273	TRAILZ — Number of trailing zero bits of an integer	279
9.274	TRANSFER — Transfer bit patterns	279
9.275	TRANSPOSE — Transpose an array of rank two	280
9.276	TRIM — Remove trailing blank characters of a string	280
9.277	TTYNAM — Get the name of a terminal device	281
9.278	UBOUND — Upper dimension bounds of an array	281
9.279	UCOBOUND — Upper codimension bounds of an array	282
9.280	UMASK — Set the file creation mask	282
9.281	UNLINK — Remove a file from the file system	282
9.282	UNPACK — Unpack an array of rank one into an array	283
9.283	VERIFY — Scan a string for characters not a given set	284
9.284	XOR — Bitwise logical exclusive OR	284

10	Intrinsic Modules	287
10.1	ISO_FORTRAN_ENV.....	287
10.2	ISO_C_BINDING.....	289
10.3	IEEE modules: IEEE_EXCEPTIONS, IEEE_ARITHMETIC, and IEEE_FEATURES.....	290
10.4	OpenMP Modules OMP_LIB and OMP_LIB_KINDS.....	291
10.5	OpenACC Module OPENACC.....	291
	Contributing	293
	Contributors to GNU Fortran	293
	Projects	294
	Proposed Extensions	294
	Compiler extensions:	294
	Environment Options	295
	GNU General Public License	297
	GNU Free Documentation License	309
	ADDENDUM: How to use this License for your documents	316
	Funding Free Software	317
	Option Index	319
	Keyword Index	321

1 Introduction

This manual documents the use of `gfortran`, the GNU Fortran compiler. You can find in this manual how to invoke `gfortran`, as well as its features and incompatibilities.

The GNU Fortran compiler front end was designed initially as a free replacement for, or alternative to, the Unix `f95` command; `gfortran` is the command you will use to invoke the compiler.

1.1 About GNU Fortran

The GNU Fortran compiler supports the Fortran 77, 90 and 95 standards completely, parts of the Fortran 2003, 2008 and 2018 standards, and several vendor extensions. The development goal is to provide the following features:

- Read a user's program, stored in a file and containing instructions written in Fortran 77, Fortran 90, Fortran 95, Fortran 2003, Fortran 2008 or Fortran 2018. This file contains *source code*.
- Translate the user's program into instructions a computer can carry out more quickly than it takes to translate the instructions in the first place. The result after compilation of a program is *machine code*, code designed to be efficiently translated and processed by a machine such as your computer. Humans usually are not as good writing machine code as they are at writing Fortran (or C++, Ada, or Java), because it is easy to make tiny mistakes writing machine code.
- Provide the user with information about the reasons why the compiler is unable to create a binary from the source code. Usually this will be the case if the source code is flawed. The Fortran 90 standard requires that the compiler can point out mistakes to the user. An incorrect usage of the language causes an *error message*.

The compiler will also attempt to diagnose cases where the user's program contains a correct usage of the language, but instructs the computer to do something questionable. This kind of diagnostics message is called a *warning message*.

- Provide optional information about the translation passes from the source code to machine code. This can help a user of the compiler to find the cause of certain bugs which may not be obvious in the source code, but may be more easily found at a lower level compiler output. It also helps developers to find bugs in the compiler itself.
- Provide information in the generated machine code that can make it easier to find bugs in the program (using a debugging tool, called a *debugger*, such as the GNU Debugger `gdb`).
- Locate and gather machine code already generated to perform actions requested by statements in the user's program. This machine code is organized into *modules* and is located and *linked* to the user program.

The GNU Fortran compiler consists of several components:

- A version of the `gcc` command (which also might be installed as the system's `cc` command) that also understands and accepts Fortran source code. The `gcc` command is the *driver* program for all the languages in the GNU Compiler Collection (GCC); With `gcc`, you can compile the source code of any language for which a front end is available in GCC.

- The `gfortran` command itself, which also might be installed as the system's `f95` command. `gfortran` is just another driver program, but specifically for the Fortran compiler only. The difference with `gcc` is that `gfortran` will automatically link the correct libraries to your program.
- A collection of run-time libraries. These libraries contain the machine code needed to support capabilities of the Fortran language that are not directly provided by the machine code generated by the `gfortran` compilation phase, such as intrinsic functions and subroutines, and routines for interaction with files and the operating system.
- The Fortran compiler itself, (`f951`). This is the GNU Fortran parser and code generator, linked to and interfaced with the GCC backend library. `f951` “translates” the source code to assembler code. You would typically not use this program directly; instead, the `gcc` or `gfortran` driver programs will call it for you.

1.2 GNU Fortran and GCC

GNU Fortran is a part of GCC, the *GNU Compiler Collection*. GCC consists of a collection of front ends for various languages, which translate the source code into a language-independent form called *GENERIC*. This is then processed by a common middle end which provides optimization, and then passed to one of a collection of back ends which generate code for different computer architectures and operating systems.

Functionally, this is implemented with a driver program (`gcc`) which provides the command-line interface for the compiler. It calls the relevant compiler front-end program (e.g., `f951` for Fortran) for each file in the source code, and then calls the assembler and linker as appropriate to produce the compiled output. In a copy of GCC which has been compiled with Fortran language support enabled, `gcc` will recognize files with `‘.f’`, `‘.for’`, `‘.ftn’`, `‘.f90’`, `‘.f95’`, `‘.f03’` and `‘.f08’` extensions as Fortran source code, and compile it accordingly. A `gfortran` driver program is also provided, which is identical to `gcc` except that it automatically links the Fortran runtime libraries into the compiled program.

Source files with `‘.f’`, `‘.for’`, `‘.fpp’`, `‘.ftn’`, `‘.F’`, `‘.FOR’`, `‘.FPP’`, and `‘.FTN’` extensions are treated as fixed form. Source files with `‘.f90’`, `‘.f95’`, `‘.f03’`, `‘.f08’`, `‘.F90’`, `‘.F95’`, `‘.F03’` and `‘.F08’` extensions are treated as free form. The capitalized versions of either form are run through preprocessing. Source files with the lower case `‘.fpp’` extension are also run through preprocessing.

This manual specifically documents the Fortran front end, which handles the programming language's syntax and semantics. The aspects of GCC which relate to the optimization passes and the back-end code generation are documented in the GCC manual; see [Section “Introduction” in *Using the GNU Compiler Collection \(GCC\)*](#). The two manuals together provide a complete reference for the GNU Fortran compiler.

1.3 Preprocessing and conditional compilation

Many Fortran compilers including GNU Fortran allow passing the source code through a C preprocessor (CPP; sometimes also called the Fortran preprocessor, FPP) to allow for conditional compilation. In the case of GNU Fortran, this is the GNU C Preprocessor in the traditional mode. On systems with case-preserving file names, the preprocessor is automatically invoked if the filename extension is `‘.F’`, `‘.FOR’`, `‘.FTN’`, `‘.fpp’`, `‘.FPP’`, `‘.F90’`,

‘.F95’, ‘.F03’ or ‘.F08’. To manually invoke the preprocessor on any file, use ‘-cpp’, to disable preprocessing on files where the preprocessor is run automatically, use ‘-nocpp’.

If a preprocessed file includes another file with the Fortran `INCLUDE` statement, the included file is not preprocessed. To preprocess included files, use the equivalent preprocessor statement `#include`.

If GNU Fortran invokes the preprocessor, `__GFORTRAN__` is defined. The macros `__GNUC__`, `__GNUC_MINOR__` and `__GNUC_PATCHLEVEL__` can be used to determine the version of the compiler. See [Section “Overview” in *The C Preprocessor*](#) for details.

GNU Fortran supports a number of `INTEGER` and `REAL` kind types in addition to the kind types required by the Fortran standard. The availability of any given kind type is architecture dependent. The following pre-defined preprocessor macros can be used to conditionally include code for these additional kind types: `__GFC_INT_1__`, `__GFC_INT_2__`, `__GFC_INT_8__`, `__GFC_INT_16__`, `__GFC_REAL_10__`, and `__GFC_REAL_16__`.

While CPP is the de-facto standard for preprocessing Fortran code, Part 3 of the Fortran 95 standard (ISO/IEC 1539-3:1998) defines Conditional Compilation, which is not widely used and not directly supported by the GNU Fortran compiler. You can use the program `coco` to preprocess such files (<http://www.daniellnagle.com/coco.html>).

1.4 GNU Fortran and G77

The GNU Fortran compiler is the successor to `g77`, the Fortran 77 front end included in GCC prior to version 4. It is an entirely new program that has been designed to provide Fortran 95 support and extensibility for future Fortran language standards, as well as providing backwards compatibility for Fortran 77 and nearly all of the GNU language extensions supported by `g77`.

1.5 Project Status

As soon as `gfortran` can parse all of the statements correctly, it will be in the “larva” state. When we generate code, the “puppa” state. When `gfortran` is done, we’ll see if it will be a beautiful butterfly, or just a big bug....

—Andy Vaught, April 2000

The start of the GNU Fortran 95 project was announced on the GCC homepage in March 18, 2000 (even though Andy had already been working on it for a while, of course).

The GNU Fortran compiler is able to compile nearly all standard-compliant Fortran 95, Fortran 90, and Fortran 77 programs, including a number of standard and non-standard extensions, and can be used on real-world programs. In particular, the supported extensions include OpenMP, Cray-style pointers, some old vendor extensions, and several Fortran 2003 and Fortran 2008 features, including TR 15581. However, it is still under development and has a few remaining rough edges. There also is initial support for OpenACC. Note that this is an experimental feature, incomplete, and subject to change in future versions of GCC. See <https://gcc.gnu.org/wiki/OpenACC> for more information.

At present, the GNU Fortran compiler passes the [NIST Fortran 77 Test Suite](#), and produces acceptable results on the [LAPACK Test Suite](#). It also provides respectable performance on the [Polyhedron Fortran compiler benchmarks](#) and the [Livermore Fortran Kernels test](#).

It has been used to compile a number of large real-world programs, including the **HARMONIE** and **HIRLAM** weather forecasting code and the **Tonto** quantum chemistry package; see <https://gcc.gnu.org/wiki/GfortranApps> for an extended list.

Among other things, the GNU Fortran compiler is intended as a replacement for G77. At this point, nearly all programs that could be compiled with G77 can be compiled with GNU Fortran, although there are a few minor known regressions.

The primary work remaining to be done on GNU Fortran falls into three categories: bug fixing (primarily regarding the treatment of invalid code and providing useful error messages), improving the compiler optimizations and the performance of compiled code, and extending the compiler to support future standards—in particular, Fortran 2003, Fortran 2008 and Fortran 2018.

1.6 Standards

The GNU Fortran compiler implements ISO/IEC 1539:1997 (Fortran 95). As such, it can also compile essentially all standard-compliant Fortran 90 and Fortran 77 programs. It also supports the ISO/IEC TR-15581 enhancements to allocatable arrays.

GNU Fortran also have a partial support for ISO/IEC 1539-1:2004 (Fortran 2003), ISO/IEC 1539-1:2010 (Fortran 2008), the Technical Specification **Further Interoperability of Fortran with C** (ISO/IEC TS 29113:2012). Full support of those standards and future Fortran standards is planned. The current status of the support is can be found in the [Section 4.1 \[Fortran 2003 status\]](#), [page 39](#), [Section 4.2 \[Fortran 2008 status\]](#), [page 40](#) and [Section 4.3 \[Fortran 2018 status\]](#), [page 42](#) sections of the documentation.

Additionally, the GNU Fortran compilers supports the OpenMP specification (version 4.0 and most of the features of the 4.5 version, <http://openmp.org/wp/openmp-specifications/>). There also is initial support for the OpenACC specification (targeting version 2.0, <http://www.openacc.org/>). Note that this is an experimental feature, incomplete, and subject to change in future versions of GCC. See <https://gcc.gnu.org/wiki/OpenACC> for more information.

1.6.1 Varying Length Character Strings

The Fortran 95 standard specifies in Part 2 (ISO/IEC 1539-2:2000) varying length character strings. While GNU Fortran currently does not support such strings directly, there exist two Fortran implementations for them, which work with GNU Fortran. They can be found at http://www.fortran.com/iso_varying_string.f95 and at ftp://ftp.nag.co.uk/sc22wg5/ISO_VARYING_STRING/.

Deferred-length character strings of Fortran 2003 supports part of the features of `ISO_VARYING_STRING` and should be considered as replacement. (Namely, allocatable or pointers of the type `character(len=:)`.)

Part I: Invoking GNU Fortran

2 GNU Fortran Command Options

The `gfortran` command supports all the options supported by the `gcc` command. Only options specific to GNU Fortran are documented here.

See [Section “GCC Command Options”](#) in *Using the GNU Compiler Collection (GCC)*, for information on the non-Fortran-specific aspects of the `gcc` command (and, therefore, the `gfortran` command).

All GCC and GNU Fortran options are accepted both by `gfortran` and by `gcc` (as well as any other drivers built at the same time, such as `g++`), since adding GNU Fortran to the GCC distribution enables acceptance of GNU Fortran options by all of the relevant drivers.

In some cases, options have positive and negative forms; the negative form of ‘`-ffoo`’ would be ‘`-fno-foo`’. This manual documents only one of these two forms, whichever one is not the default.

2.1 Option summary

Here is a summary of all the options specific to GNU Fortran, grouped by type. Explanations are in the following sections.

Fortran Language Options

See [Section 2.2 \[Options controlling Fortran dialect\]](#), page 8.

```
-fall-intrinsics -fbackslash -fcray-pointer -fd-lines-as-code
-fd-lines-as-comments
-fdec -fdec-structure -fdec-intrinsic-ints -fdec-static -fdec-math
-fdec-include -fdefault-double-8 -fdefault-integer-8 -fdefault-real-8
-fdefault-real-10 -fdefault-real-16 -fdollar-ok -ffixed-line-length-n
-ffixed-line-length-none -fpad-source -ffree-form -ffree-line-length-n
-ffree-line-length-none -fimplicit-none -finteger-4-integer-8
-fmax-identifier-length -fmodule-private -ffixed-form -fno-range-check
-fopenacc -fopenmp -freal-4-real-10 -freal-4-real-16 -freal-4-real-8
-freal-8-real-10 -freal-8-real-16 -freal-8-real-4 -std=std -ftest-forall-
temp
```

Preprocessing Options

See [Section 2.3 \[Enable and customize preprocessing\]](#), page 13.

```
-A-question[=answer] -Aquestion=answer -C -CC -Dmacro[=defn] -H -P
-Umacro -cpp -dD -dI -dM -dN -dU -fworking-directory -imultilib dir
-iprefix file -iquote -isysroot dir -isystem dir -nocpp -nostdinc
-undef
```

Error and Warning Options

See [Section 2.4 \[Options to request or suppress errors and warnings\]](#), page 15.

```
-Waliasing -Wall -Wampersand -Wargument-mismatch -Warray-bounds -Wc-binding-
type -Wcharacter-truncation -Wconversion
-Wdo-subscript -Wfunction-elimination -Wimplicit-interface
-Wimplicit-procedure -Wintrinsic-shadow -Wuse-without-only -Wintrinsics-std
-Wline-truncation -Wno-align-commons -Wno-tabs -Wreal-q-constant
-Wsurprising -Wunderflow -Wunused-parameter -Wrealloc-lhs
-Wrealloc-lhs-all -Wfrontend-loop-interchange -Wtarget-lifetime
-fmax-errors=n -fsyntax-only -pedantic -pedantic-errors
```

Debugging Options

See [Section 2.5 \[Options for debugging your program or GNU Fortran\]](#), page 20.

```
-fbacktrace -fdump-fortran-optimized -fdump-fortran-original
-fdump-fortran-global -fdump-parse-tree -ffpe-trap=list
-ffpe-summary=list
```

Directory Options

See [Section 2.6 \[Options for directory search\]](#), page 22.

```
-Idir -Jdir -fintrinsic-modules-path dir
```

Link Options

See [Section 2.7 \[Options for influencing the linking step\]](#), page 22.

```
-static-libgfortran
```

Runtime Options

See [Section 2.8 \[Options for influencing runtime behavior\]](#), page 22.

```
-fconvert=conversion -fmax-subrecord-length=length
-frecord-marker=length -fsign-zero
```

Interoperability Options

See [Section 2.10 \[Options for interoperability\]](#), page 30.

```
-fc-prototypes -fc-prototypes-external
```

Code Generation Options

See [Section 2.9 \[Options for code generation conventions\]](#), page 23.

```
-faggressive-function-elimination -fblas-matmul-limit=n
-fbounds-check -ftail-call-workaround -ftail-call-workaround=n
-fcheck-array-temporaries
-fcheck=<all|array-temps|bounds|do|mem|pointer|recursion>
-fcoarray=<none|single|lib> -fexternal-blas -ff2c -ffrontend-loop-interchange
-ffrontend-optimize
-finit-character=n -finit-integer=n -finit-local-zero
-finit-derived
-finit-logical=<true|false> -finit-real=<zero|inf|-inf|nan|snan>
-finline-matmul-limit=n
-fmax-array-constructor=n -fmax-stack-var-size=n -fno-align-commons
-fno-automatic -fno-protect-parens -fno-underscoring
-fsecond-underscore -fpack-derived -frealloc-lhs -frecursive
-frepack-arrays -fshort-enums -fstack-arrays
```

2.2 Options controlling Fortran dialect

The following options control the details of the Fortran dialect accepted by the compiler:

-ffree-form

-ffixed-form

Specify the layout used by the source file. The free form layout was introduced in Fortran 90. Fixed form was traditionally used in older Fortran programs. When neither option is specified, the source form is determined by the file extension.

-fall-intrinsics

This option causes all intrinsic procedures (including the GNU-specific extensions) to be accepted. This can be useful with ‘-std=f95’ to force standard-compliance but get access to the full range of intrinsics available with **gfortran**.

As a consequence, ‘`-Wintrinsics-std`’ will be ignored and no user-defined procedure with the same name as any intrinsic will be called except when it is explicitly declared `EXTERNAL`.

`-fd-lines-as-code`

`-fd-lines-as-comments`

Enable special treatment for lines beginning with `d` or `D` in fixed form sources. If the ‘`-fd-lines-as-code`’ option is given they are treated as if the first column contained a blank. If the ‘`-fd-lines-as-comments`’ option is given, they are treated as comment lines.

`-fdec` DEC compatibility mode. Enables extensions and other features that mimic the default behavior of older compilers (such as DEC). These features are non-standard and should be avoided at all costs. For details on GNU Fortran’s implementation of these extensions see the full documentation.

Other flags enabled by this switch are: ‘`-fdollar-ok`’ ‘`-fcray-pointer`’ ‘`-fdec-structure`’ ‘`-fdec-intrinsic-ints`’ ‘`-fdec-static`’ ‘`-fdec-math`’

If ‘`-fd-lines-as-code`’/‘`-fd-lines-as-comments`’ are unset, then ‘`-fdec`’ also sets ‘`-fd-lines-as-comments`’.

`-fdec-structure`

Enable DEC `STRUCTURE` and `RECORD` as well as `UNION`, `MAP`, and dot (`.'`) as a member separator (in addition to `%`). This is provided for compatibility only; Fortran 90 derived types should be used instead where possible.

`-fdec-intrinsic-ints`

Enable B/I/J/K kind variants of existing integer functions (e.g. `BIAND`, `IIAND`, `JIAND`, etc...). For a complete list of intrinsics see the full documentation.

`-fdec-math`

Enable legacy math intrinsics such as `COTAN` and degree-valued trigonometric functions (e.g. `TAND`, `ATAND`, etc...) for compatibility with older code.

`-fdec-static`

Enable DEC-style `STATIC` and `AUTOMATIC` attributes to explicitly specify the storage of variables and other objects.

`-fdec-include`

Enable parsing of `INCLUDE` as a statement in addition to parsing it as `INCLUDE` line. When parsed as `INCLUDE` statement, `INCLUDE` does not have to be on a single line and can use line continuations.

`-fdollar-ok`

Allow ‘`$`’ as a valid non-first character in a symbol name. Symbols that start with ‘`$`’ are rejected since it is unclear which rules to apply to implicit typing as different vendors implement different rules. Using ‘`$`’ in `IMPLICIT` statements is also rejected.

`-fbackslash`

Change the interpretation of backslashes in string literals from a single backslash character to “C-style” escape characters. The following combinations are

expanded `\a`, `\b`, `\f`, `\n`, `\r`, `\t`, `\v`, `\\`, and `\0` to the ASCII characters alert, backspace, form feed, newline, carriage return, horizontal tab, vertical tab, backslash, and NUL, respectively. Additionally, `\xnn`, `\unnnn` and `\Unnnnnnnnn` (where each *n* is a hexadecimal digit) are translated into the Unicode characters corresponding to the specified code points. All other combinations of a character preceded by `\` are unexpanded.

-fmodule-private

Set the default accessibility of module entities to `PRIVATE`. Use-associated entities will not be accessible unless they are explicitly declared as `PUBLIC`.

-ffixed-line-length-n

Set column after which characters are ignored in typical fixed-form lines in the source file, and, unless `-fno-pad-source`, through which spaces are assumed (as if padded to that length) after the ends of short fixed-form lines.

Popular values for *n* include 72 (the standard and the default), 80 (card image), and 132 (corresponding to “extended-source” options in some popular compilers). *n* may also be `'none'`, meaning that the entire line is meaningful and that continued character constants never have implicit spaces appended to them to fill out the line. `'-ffixed-line-length-0'` means the same thing as `'-ffixed-line-length-none'`.

-fno-pad-source

By default fixed-form lines have spaces assumed (as if padded to that length) after the ends of short fixed-form lines. This is not done either if `'-ffixed-line-length-0'`, `'-ffixed-line-length-none'` or if `'-fno-pad-source'` option is used. With any of those options continued character constants never have implicit spaces appended to them to fill out the line.

-ffree-line-length-n

Set column after which characters are ignored in typical free-form lines in the source file. The default value is 132. *n* may be `'none'`, meaning that the entire line is meaningful. `'-ffree-line-length-0'` means the same thing as `'-ffree-line-length-none'`.

-fmax-identifier-length=n

Specify the maximum allowed identifier length. Typical values are 31 (Fortran 95) and 63 (Fortran 2003 and Fortran 2008).

-fimplicit-none

Specify that no implicit typing is allowed, unless overridden by explicit `IMPLICIT` statements. This is the equivalent of adding `implicit none` to the start of every procedure.

-fcray-pointer

Enable the Cray pointer extension, which provides C-like pointer functionality.

-fopenacc

Enable the OpenACC extensions. This includes OpenACC `!$acc` directives in free form and `c$acc`, `*$acc` and `!$acc` directives in fixed form, `!$` conditional

compilation sentinels in free form and `c$`, `*$` and `!$` sentinels in fixed form, and when linking arranges for the OpenACC runtime library to be linked in.

Note that this is an experimental feature, incomplete, and subject to change in future versions of GCC. See <https://gcc.gnu.org/wiki/OpenACC> for more information.

-fopenmp Enable the OpenMP extensions. This includes OpenMP `!$omp` directives in free form and `c$omp`, `*$omp` and `!$omp` directives in fixed form, `!$` conditional compilation sentinels in free form and `c$`, `*$` and `!$` sentinels in fixed form, and when linking arranges for the OpenMP runtime library to be linked in. The option `'-fopenmp'` implies `'-frecursive'`.

-fno-range-check

Disable range checking on results of simplification of constant expressions during compilation. For example, GNU Fortran will give an error at compile time when simplifying `a = 1. / 0.` With this option, no error will be given and `a` will be assigned the value `+Infinity`. If an expression evaluates to a value outside of the relevant range of `[-HUGE():HUGE()]`, then the expression will be replaced by `-Inf` or `+Inf` as appropriate. Similarly, `DATA i/Z'FFFFFFFF'/` will result in an integer overflow on most systems, but with `'-fno-range-check'` the value will “wrap around” and `i` will be initialized to `-1` instead.

-fdefault-integer-8

Set the default integer and logical types to an 8 byte wide type. This option also affects the kind of integer constants like `42`. Unlike `'-finteger-4-integer-8'`, it does not promote variables with explicit kind declaration.

-fdefault-real-8

Set the default real type to an 8 byte wide type. This option also affects the kind of non-double real constants like `1.0`. This option promotes the default width of `DOUBLE PRECISION` and double real constants like `1.d0` to 16 bytes if possible. If `-fdefault-double-8` is given along with `fdefault-real-8`, `DOUBLE PRECISION` and double real constants are not promoted. Unlike `'-freal-4-real-8'`, `fdefault-real-8` does not promote variables with explicit kind declarations.

-fdefault-real-10

Set the default real type to an 10 byte wide type. This option also affects the kind of non-double real constants like `1.0`. This option promotes the default width of `DOUBLE PRECISION` and double real constants like `1.d0` to 16 bytes if possible. If `-fdefault-double-8` is given along with `fdefault-real-10`, `DOUBLE PRECISION` and double real constants are not promoted. Unlike `'-freal-4-real-10'`, `fdefault-real-10` does not promote variables with explicit kind declarations.

-fdefault-real-16

Set the default real type to an 16 byte wide type. This option also affects the kind of non-double real constants like `1.0`. This option promotes the default width of `DOUBLE PRECISION` and double real constants like `1.d0` to 16 bytes if possible. If `-fdefault-double-8` is given along with `fdefault-real-16`, `DOUBLE PRECISION` and double real constants are not promoted. Unlike

`'-freal-4-real-16'`, `fdefault-real-16` does not promote variables with explicit kind declarations.

`-fdefault-double-8`

Set the `DOUBLE PRECISION` type and double real constants like `1.d0` to an 8 byte wide type. Do nothing if this is already the default. This option prevents `'-fdefault-real-8'`, `'-fdefault-real-10'`, and `'-fdefault-real-16'`, from promoting `DOUBLE PRECISION` and double real constants like `1.d0` to 16 bytes.

`-finteger-4-integer-8`

Promote all `INTEGER(KIND=4)` entities to an `INTEGER(KIND=8)` entities. If `KIND=8` is unavailable, then an error will be issued. This option should be used with care and may not be suitable for your codes. Areas of possible concern include calls to external procedures, alignment in `EQUIVALENCE` and/or `COMMON`, generic interfaces, `BOZ` literal constant conversion, and I/O. Inspection of the intermediate representation of the translated Fortran code, produced by `'-fdump-tree-original'`, is suggested.

`-freal-4-real-8`

`-freal-4-real-10`

`-freal-4-real-16`

`-freal-8-real-4`

`-freal-8-real-10`

`-freal-8-real-16`

Promote all `REAL(KIND=M)` entities to `REAL(KIND=N)` entities. If `REAL(KIND=N)` is unavailable, then an error will be issued. All other real kind types are unaffected by this option. These options should be used with care and may not be suitable for your codes. Areas of possible concern include calls to external procedures, alignment in `EQUIVALENCE` and/or `COMMON`, generic interfaces, `BOZ` literal constant conversion, and I/O. Inspection of the intermediate representation of the translated Fortran code, produced by `'-fdump-tree-original'`, is suggested.

`-std=std` Specify the standard to which the program is expected to conform, which may be one of `'f95'`, `'f2003'`, `'f2008'`, `'f2018'`, `'gnu'`, or `'legacy'`. The default value for `std` is `'gnu'`, which specifies a superset of the latest Fortran standard that includes all of the extensions supported by GNU Fortran, although warnings will be given for obsolete extensions not recommended for use in new code. The `'legacy'` value is equivalent but without the warnings for obsolete extensions, and may be useful for old non-standard programs. The `'f95'`, `'f2003'`, `'f2008'`, and `'f2018'` values specify strict conformance to the Fortran 95, Fortran 2003, Fortran 2008 and Fortran 2018 standards, respectively; errors are given for all extensions beyond the relevant language standard, and warnings are given for the Fortran 77 features that are permitted but obsolescent in later standards. The deprecated option `'-std=f2008ts'` acts as an alias for `'-std=f2018'`. It is only present for backwards compatibility with earlier gfortran versions and should not be used any more.

`-ftest-forall-temp`

Enhance test coverage by forcing most forall assignments to use temporary.

2.3 Enable and customize preprocessing

Preprocessor related options. See section [Section 1.3 \[Preprocessing and conditional compilation\]](#), page 2 for more detailed information on preprocessing in `gfortran`.

- `-cpp`
- `-nocpp` Enable preprocessing. The preprocessor is automatically invoked if the file extension is `.fpp`, `.FPP`, `.F`, `.FOR`, `.FTN`, `.F90`, `.F95`, `.F03` or `.F08`. Use this option to manually enable preprocessing of any kind of Fortran file.
 To disable preprocessing of files with any of the above listed extensions, use the negative form: `-nocpp`.
 The preprocessor is run in traditional mode. Any restrictions of the file-format, especially the limits on line length, apply for preprocessed output as well, so it might be advisable to use the `-ffree-line-length-none` or `-ffixed-line-length-none` options.
- `-dM` Instead of the normal output, generate a list of `#define` directives for all the macros defined during the execution of the preprocessor, including predefined macros. This gives you a way of finding out what is predefined in your version of the preprocessor. Assuming you have no file `foo.f90`, the command
 `touch foo.f90; gfortran -cpp -E -dM foo.f90`
 will show all the predefined macros.
- `-dD` Like `-dM` except in two respects: it does not include the predefined macros, and it outputs both the `#define` directives and the result of preprocessing. Both kinds of output go to the standard output file.
- `-dN` Like `-dD`, but emit only the macro names, not their expansions.
- `-dU` Like `-dD` except that only macros that are expanded, or whose definedness is tested in preprocessor directives, are output; the output is delayed until the use or test of the macro; and `#undef` directives are also output for macros tested but undefined at the time.
- `-dI` Output `#include` directives in addition to the result of preprocessing.
- `-fworking-directory`
 Enable generation of linemarkers in the preprocessor output that will let the compiler know the current working directory at the time of preprocessing. When this option is enabled, the preprocessor will emit, after the initial linemarker, a second linemarker with the current working directory followed by two slashes. GCC will use this directory, when it is present in the preprocessed input, as the directory emitted as the current working directory in some debugging information formats. This option is implicitly enabled if debugging information is enabled, but this can be inhibited with the negated form `-fno-working-directory`. If the `-P` flag is present in the command line, this option has no effect, since no `#line` directives are emitted whatsoever.
- `-idirafter dir`
 Search *dir* for include files, but do it after all directories specified with `-I` and the standard system directories have been exhausted. *dir* is treated as a

system include directory. If *dir* begins with `=`, then the `=` will be replaced by the `sysroot` prefix; see `--sysroot` and `-isysroot`.

`-imultilib dir`

Use *dir* as a subdirectory of the directory containing target-specific C++ headers.

`-iprefix prefix`

Specify *prefix* as the prefix for subsequent `-iwithprefix` options. If the *prefix* represents a directory, you should include the final `'/'`.

`-isysroot dir`

This option is like the `--sysroot` option, but applies only to header files. See the `--sysroot` option for more information.

`-iquote dir`

Search *dir* only for header files requested with `#include "file"`; they are not searched for `#include <file>`, before all directories specified by `-I` and before the standard system directories. If *dir* begins with `=`, then the `=` will be replaced by the `sysroot` prefix; see `--sysroot` and `-isysroot`.

`-isystem dir`

Search *dir* for header files, after all directories specified by `-I` but before the standard system directories. Mark it as a system directory, so that it gets the same special treatment as is applied to the standard system directories. If *dir* begins with `=`, then the `=` will be replaced by the `sysroot` prefix; see `--sysroot` and `-isysroot`.

`-nostdinc`

Do not search the standard system directories for header files. Only the directories you have specified with `-I` options (and the directory of the current file, if appropriate) are searched.

`-undef`

Do not predefine any system-specific or GCC-specific macros. The standard predefined macros remain defined.

`-Apredicate=answer`

Make an assertion with the predicate *predicate* and answer *answer*. This form is preferred to the older form `-A predicate(answer)`, which is still supported, because it does not use shell special characters.

`-A-predicate=answer`

Cancel an assertion with the predicate *predicate* and answer *answer*.

`-C`

Do not discard comments. All comments are passed through to the output file, except for comments in processed directives, which are deleted along with the directive.

You should be prepared for side effects when using `-C`; it causes the preprocessor to treat comments as tokens in their own right. For example, comments appearing at the start of what would be a directive line have the effect of turning that line into an ordinary source line, since the first token on the line is no longer a `'#'`.

Warning: this currently handles C-Style comments only. The preprocessor does not yet recognize Fortran-style comments.

- CC** Do not discard comments, including during macro expansion. This is like ‘-C’, except that comments contained within macros are also passed through to the output file where the macro is expanded.
- In addition to the side-effects of the ‘-C’ option, the ‘-CC’ option causes all C++-style comments inside a macro to be converted to C-style comments. This is to prevent later use of that macro from inadvertently commenting out the remainder of the source line. The ‘-CC’ option is generally used to support lint comments.
- Warning: this currently handles C- and C++-Style comments only. The preprocessor does not yet recognize Fortran-style comments.
- Dname** Predefine name as a macro, with definition 1.
- Dname=definition**
- The contents of *definition* are tokenized and processed as if they appeared during translation phase three in a ‘#define’ directive. In particular, the definition will be truncated by embedded newline characters.
- If you are invoking the preprocessor from a shell or shell-like program you may need to use the shell’s quoting syntax to protect characters such as spaces that have a meaning in the shell syntax.
- If you wish to define a function-like macro on the command line, write its argument list with surrounding parentheses before the equals sign (if any). Parentheses are meaningful to most shells, so you will need to quote the option. With sh and csh, -D‘name(args...)=definition’ works.
- ‘-D’ and ‘-U’ options are processed in the order they are given on the command line. All -imacros file and -include file options are processed after all -D and -U options.
- H** Print the name of each header file used, in addition to other normal activities. Each name is indented to show how deep in the ‘#include’ stack it is.
- P** Inhibit generation of linemarkers in the output from the preprocessor. This might be useful when running the preprocessor on something that is not C code, and will be sent to a program which might be confused by the linemarkers.
- Uname** Cancel any previous definition of *name*, either built in or provided with a ‘-D’ option.

2.4 Options to request or suppress errors and warnings

Errors are diagnostic messages that report that the GNU Fortran compiler cannot compile the relevant piece of source code. The compiler will continue to process the program in an attempt to report further errors to aid in debugging, but will not produce any compiled output.

Warnings are diagnostic messages that report constructions which are not inherently erroneous but which are risky or suggest there is likely to be a bug in the program. Unless ‘-Werror’ is specified, they do not prevent compilation of the program.

You can request many specific warnings with options beginning ‘-W’, for example ‘-Wimplicit’ to request warnings on implicit declarations. Each of these specific warning

options also has a negative form beginning ‘-Wno-’ to turn off warnings; for example, ‘-Wno-implicit’. This manual lists only one of the two forms, whichever is not the default.

These options control the amount and kinds of errors and warnings produced by GNU Fortran:

-fmax-errors=*n*

Limits the maximum number of error messages to *n*, at which point GNU Fortran bails out rather than attempting to continue processing the source code. If *n* is 0, there is no limit on the number of error messages produced.

-fsyntax-only

Check the code for syntax errors, but do not actually compile it. This will generate module files for each module present in the code, but no other output file.

-Wpedantic

-pedantic

Issue warnings for uses of extensions to Fortran. ‘-pedantic’ also applies to C-language constructs where they occur in GNU Fortran source files, such as use of ‘\e’ in a character constant within a directive like `#include`.

Valid Fortran programs should compile properly with or without this option. However, without this option, certain GNU extensions and traditional Fortran features are supported as well. With this option, many of them are rejected.

Some users try to use ‘-pedantic’ to check programs for conformance. They soon find that it does not do quite what they want—it finds some nonstandard practices, but not all. However, improvements to GNU Fortran in this area are welcome.

This should be used in conjunction with ‘-std=f95’, ‘-std=f2003’, ‘-std=f2008’ or ‘-std=f2018’.

-pedantic-errors

Like ‘-pedantic’, except that errors are produced rather than warnings.

-Wall

Enables commonly used warning options pertaining to usage that we recommend avoiding and that we believe are easy to avoid. This currently includes ‘-Waliasing’, ‘-Wampersand’, ‘-Wconversion’, ‘-Wsurprising’, ‘-Wc-binding-type’, ‘-Wintrinsics-std’, ‘-Wtabs’, ‘-Wintrinsic-shadow’, ‘-Wline-truncation’, ‘-Wtarget-lifetime’, ‘-Winteger-division’, ‘-Wreal-q-constant’, ‘-Wunused’ and ‘-Wundefined-do-loop’.

-Waliasing

Warn about possible aliasing of dummy arguments. Specifically, it warns if the same actual argument is associated with a dummy argument with `INTENT(IN)` and a dummy argument with `INTENT(OUT)` in a call with an explicit interface.

The following example will trigger the warning.

```
interface
  subroutine bar(a,b)
    integer, intent(in) :: a
    integer, intent(out) :: b
```



```

        end subroutine
    end interface
    integer :: a

    call bar(a,a)

```

-Wampersand

Warn about missing ampersand in continued character constants. The warning is given with ‘-Wampersand’, ‘-pedantic’, ‘-std=f95’, ‘-std=f2003’, ‘-std=f2008’ and ‘-std=f2018’. Note: With no ampersand given in a continued character constant, GNU Fortran assumes continuation at the first non-comment, non-whitespace character after the ampersand that initiated the continuation.

-Wargument-mismatch

Warn about type, rank, and other mismatches between formal parameters and actual arguments to functions and subroutines. These warnings are recommended and thus enabled by default.

-Warray-temporaries

Warn about array temporaries generated by the compiler. The information generated by this warning is sometimes useful in optimization, in order to avoid such temporaries.

-Wc-binding-type

Warn if the a variable might not be C interoperable. In particular, warn if the variable has been declared using an intrinsic type with default kind instead of using a kind parameter defined for C interoperability in the intrinsic `ISO_C_Binding` module. This option is implied by ‘-Wall’.

-Wcharacter-truncation

Warn when a character assignment will truncate the assigned string.

-Wline-truncation

Warn when a source code line will be truncated. This option is implied by ‘-Wall’. For free-form source code, the default is ‘-Werror=line-truncation’ such that truncations are reported as error.

-Wconversion

Warn about implicit conversions that are likely to change the value of the expression after conversion. Implied by ‘-Wall’.

-Wconversion-extra

Warn about implicit conversions between different types and kinds. This option does *not* imply ‘-Wconversion’.

-Wextra

Enables some warning options for usages of language features which may be problematic. This currently includes ‘-Wcompare-reals’, ‘-Wunused-parameter’ and ‘-Wdo-subscript’.

-Wfrontend-loop-interchange

Enable warning for loop interchanges performed by the ‘-ffrontend-loop-interchange’ option.

-Wimplicit-interface

Warn if a procedure is called without an explicit interface. Note this only checks that an explicit interface is present. It does not check that the declared interfaces are consistent across program units.

-Wimplicit-procedure

Warn if a procedure is called that has neither an explicit interface nor has been declared as **EXTERNAL**.

-Winteger-division

Warn if a constant integer division truncates its result. As an example, $3/5$ evaluates to 0.

-Wintrinsics-std

Warn if **gfortran** finds a procedure named like an intrinsic not available in the currently selected standard (with ‘-std’) and treats it as **EXTERNAL** procedure because of this. ‘-fall-intrinsics’ can be used to never trigger this behavior and always link to the intrinsic regardless of the selected standard.

-Wreal-q-constant

Produce a warning if a real-literal-constant contains a **q** exponent-letter.

-Wsurprising

Produce a warning when “suspicious” code constructs are encountered. While technically legal these usually indicate that an error has been made.

This currently produces a warning under the following circumstances:

- An **INTEGER SELECT** construct has a **CASE** that can never be matched as its lower value is greater than its upper value.
- A **LOGICAL SELECT** construct has three **CASE** statements.
- A **TRANSFER** specifies a source that is shorter than the destination.
- The type of a function result is declared more than once with the same type. If ‘-pedantic’ or standard-conforming mode is enabled, this is an error.
- A **CHARACTER** variable is declared with negative length.

-Wtabs

By default, tabs are accepted as whitespace, but tabs are not members of the Fortran Character Set. For continuation lines, a tab followed by a digit between 1 and 9 is supported. ‘-Wtabs’ will cause a warning to be issued if a tab is encountered. Note, ‘-Wtabs’ is active for ‘-pedantic’, ‘-std=f95’, ‘-std=f2003’, ‘-std=f2008’, ‘-std=f2018’ and ‘-Wall’.

-Wundefined-do-loop

Warn if a **DO** loop with step either 1 or -1 yields an underflow or an overflow during iteration of an induction variable of the loop. This option is implied by ‘-Wall’.

-Wunderflow

Produce a warning when numerical constant expressions are encountered, which yield an **UNDERFLOW** during compilation. Enabled by default.

-Wintrinsic-shadow

Warn if a user-defined procedure or module procedure has the same name as an intrinsic; in this case, an explicit interface or `EXTERNAL` or `INTRINSIC` declaration might be needed to get calls later resolved to the desired intrinsic/procedure. This option is implied by `'-Wall'`.

-Wuse-without-only

Warn if a `USE` statement has no `ONLY` qualifier and thus implicitly imports all public entities of the used module.

-Wunused-dummy-argument

Warn about unused dummy arguments. This option is implied by `'-Wall'`.

-Wunused-parameter

Contrary to `gcc`'s meaning of `'-Wunused-parameter'`, `gfortran`'s implementation of this option does not warn about unused dummy arguments (see `'-Wunused-dummy-argument'`), but about unused `PARAMETER` values. `'-Wunused-parameter'` is implied by `'-Wextra'` if also `'-Wunused'` or `'-Wall'` is used.

-Walign-commons

By default, `gfortran` warns about any occasion of variables being padded for proper alignment inside a `COMMON` block. This warning can be turned off via `'-Wno-align-commons'`. See also `'-falign-commons'`.

-Wfunction-elimination

Warn if any calls to impure functions are eliminated by the optimizations enabled by the `'-ffrontend-optimize'` option. This option is implied by `'-Wextra'`.

-Wrealloc-lhs

Warn when the compiler might insert code to for allocation or reallocation of an allocatable array variable of intrinsic type in intrinsic assignments. In hot loops, the Fortran 2003 reallocation feature may reduce the performance. If the array is already allocated with the correct shape, consider using a whole-array array-spec (e.g. `(:,:,:)`) for the variable on the left-hand side to prevent the reallocation check. Note that in some cases the warning is shown, even if the compiler will optimize reallocation checks away. For instance, when the right-hand side contains the same variable multiplied by a scalar. See also `'-frealloc-lhs'`.

-Wrealloc-lhs-all

Warn when the compiler inserts code to for allocation or reallocation of an allocatable variable; this includes scalars and derived types.

-Wcompare-reals

Warn when comparing real or complex types for equality or inequality. This option is implied by `'-Wextra'`.

-Wtarget-lifetime

Warn if the pointer in a pointer assignment might be longer than the its target. This option is implied by `'-Wall'`.

-Wzerotrip

Warn if a DO loop is known to execute zero times at compile time. This option is implied by ‘-Wall’.

-Wdo-subscript

Warn if an array subscript inside a DO loop could lead to an out-of-bounds access even if the compiler cannot prove that the statement is actually executed, in cases like

```
real a(3)
do i=1,4
  if (condition(i)) then
    a(i) = 1.2
  end if
end do
```

This option is implied by ‘-Wextra’.

-Werror Turns all warnings into errors.

See Section “Options to Request or Suppress Errors and Warnings” in *Using the GNU Compiler Collection (GCC)*, for information on more options offered by the GBE shared by gfortran, gcc and other GNU compilers.

Some of these have no effect when compiling programs written in Fortran.

2.5 Options for debugging your program or GNU Fortran

GNU Fortran has various special options that are used for debugging either your program or the GNU Fortran compiler.

-fdump-fortran-original

Output the internal parse tree after translating the source program into internal representation. This option is mostly useful for debugging the GNU Fortran compiler itself. The output generated by this option might change between releases. This option may also generate internal compiler errors for features which have only recently been added.

-fdump-fortran-optimized

Output the parse tree after front-end optimization. Mostly useful for debugging the GNU Fortran compiler itself. The output generated by this option might change between releases. This option may also generate internal compiler errors for features which have only recently been added.

-fdump-parse-tree

Output the internal parse tree after translating the source program into internal representation. Mostly useful for debugging the GNU Fortran compiler itself. The output generated by this option might change between releases. This option may also generate internal compiler errors for features which have only recently been added. This option is deprecated; use **-fdump-fortran-original** instead.

-fdump-fortran-global

Output a list of the global identifiers after translating into middle-end representation. Mostly useful for debugging the GNU Fortran compiler itself. The output generated by this option might change between releases. This option

may also generate internal compiler errors for features which have only recently been added.

-ffpe-trap=list

Specify a list of floating point exception traps to enable. On most systems, if a floating point exception occurs and the trap for that exception is enabled, a SIGFPE signal will be sent and the program being aborted, producing a core file useful for debugging. *list* is a (possibly empty) comma-separated list of the following exceptions: ‘invalid’ (invalid floating point operation, such as `SQRT(-1.0)`), ‘zero’ (division by zero), ‘overflow’ (overflow in a floating point operation), ‘underflow’ (underflow in a floating point operation), ‘inexact’ (loss of precision during operation), and ‘denormal’ (operation performed on a denormal value). The first five exceptions correspond to the five IEEE 754 exceptions, whereas the last one (‘denormal’) is not part of the IEEE 754 standard but is available on some common architectures such as x86.

The first three exceptions (‘invalid’, ‘zero’, and ‘overflow’) often indicate serious errors, and unless the program has provisions for dealing with these exceptions, enabling traps for these three exceptions is probably a good idea.

If the option is used more than once in the command line, the lists will be joined: ‘**ffpe-trap=list1 ffpe-trap=list2**’ is equivalent to **ffpe-trap=list1,list2**.

Note that once enabled an exception cannot be disabled (no negative form).

Many, if not most, floating point operations incur loss of precision due to rounding, and hence the **ffpe-trap=inexact** is likely to be uninteresting in practice.

By default no exception traps are enabled.

-ffpe-summary=list

Specify a list of floating-point exceptions, whose flag status is printed to `ERROR_UNIT` when invoking `STOP` and `ERROR STOP`. *list* can be either ‘none’, ‘all’ or a comma-separated list of the following exceptions: ‘invalid’, ‘zero’, ‘overflow’, ‘underflow’, ‘inexact’ and ‘denormal’. (See ‘-ffpe-trap’ for a description of the exceptions.)

If the option is used more than once in the command line, only the last one will be used.

By default, a summary for all exceptions but ‘inexact’ is shown.

-fno-backtrace

When a serious runtime error is encountered or a deadly signal is emitted (segmentation fault, illegal instruction, bus error, floating-point exception, and the other POSIX signals that have the action ‘core’), the Fortran runtime library tries to output a backtrace of the error. **-fno-backtrace** disables the backtrace generation. This option only has influence for compilation of the Fortran main program.

See Section “Options for Debugging Your Program or GCC” in *Using the GNU Compiler Collection (GCC)*, for more information on debugging options.

2.6 Options for directory search

These options affect how GNU Fortran searches for files specified by the `INCLUDE` directive and where it searches for previously compiled modules.

It also affects the search paths used by `cpp` when used to preprocess Fortran source.

-I*dir* These affect interpretation of the `INCLUDE` directive (as well as of the `#include` directive of the `cpp` preprocessor).

Also note that the general behavior of `-I` and `INCLUDE` is pretty much the same as of `-I` with `#include` in the `cpp` preprocessor, with regard to looking for `'header.gcc'` files and other such things.

This path is also used to search for `'mod'` files when previously compiled modules are required by a `USE` statement.

See [Section “Options for Directory Search” in *Using the GNU Compiler Collection \(GCC\)*](#), for information on the `-I` option.

-J*dir* This option specifies where to put `'mod'` files for compiled modules. It is also added to the list of directories to searched by an `USE` statement.

The default is the current directory.

-fintrinsic-modules-path *dir*

This option specifies the location of pre-compiled intrinsic modules, if they are not in the default location expected by the compiler.

2.7 Influencing the linking step

These options come into play when the compiler links object files into an executable output file. They are meaningless if the compiler is not doing a link step.

-static-libgfortran

On systems that provide `'libgfortran'` as a shared and a static library, this option forces the use of the static version. If no shared version of `'libgfortran'` was built when the compiler was configured, this option has no effect.

2.8 Influencing runtime behavior

These options affect the runtime behavior of programs compiled with GNU Fortran.

-fconvert=*conversion*

Specify the representation of data for unformatted files. Valid values for conversion are: `'native'`, the default; `'swap'`, swap between big- and little-endian; `'big-endian'`, use big-endian representation for unformatted files; `'little-endian'`, use little-endian representation for unformatted files.

This option has an effect only when used in the main program. The `CONVERT` specifier and the `GFORTTRAN_CONVERT_UNIT` environment variable override the default specified by `-fconvert`.

-frecord-marker=*length*

Specify the length of record markers for unformatted files. Valid values for *length* are 4 and 8. Default is 4. *This is different from previous versions of*

gfortran, which specified a default record marker length of 8 on most systems. If you want to read or write files compatible with earlier versions of **gfortran**, use `'-frecord-marker=8'`.

-fmax-subrecord-length=length

Specify the maximum length for a subrecord. The maximum permitted value for length is 2147483639, which is also the default. Only really useful for use by the **gfortran** testsuite.

-fsign-zero

When enabled, floating point numbers of value zero with the sign bit set are written as negative number in formatted output and treated as negative in the **SIGN** intrinsic. `'-fno-sign-zero'` does not print the negative sign of zero values (or values rounded to zero for I/O) and regards zero as positive number in the **SIGN** intrinsic for compatibility with Fortran 77. The default is `'-fsign-zero'`.

2.9 Options for code generation conventions

These machine-independent options control the interface conventions used in code generation.

Most of them have both positive and negative forms; the negative form of `'-ffoo'` would be `'-fno-foo'`. In the table below, only one of the forms is listed—the one which is not the default. You can figure out the other form by either removing `'no-'` or adding it.

-fno-automatic

Treat each program unit (except those marked as **RECURSIVE**) as if the **SAVE** statement were specified for every local variable and array referenced in it. Does not affect common blocks. (Some Fortran compilers provide this option under the name `'-static'` or `'-save'`.) The default, which is `'-fautomatic'`, uses the stack for local variables smaller than the value given by `'-fmax-stack-var-size'`. Use the option `'-frecursive'` to use no static memory.

Local variables or arrays having an explicit **SAVE** attribute are silently ignored unless the `'-pedantic'` option is added.

-ff2c

Generate code designed to be compatible with code generated by **g77** and **f2c**. The calling conventions used by **g77** (originally implemented in **f2c**) require functions that return type default **REAL** to actually return the C type **double**, and functions that return type **COMPLEX** to return the values via an extra argument in the calling sequence that points to where to store the return value. Under the default GNU calling conventions, such functions simply return their results as they would in GNU C—default **REAL** functions return the C type **float**, and **COMPLEX** functions return the GNU C type **complex**. Additionally, this option implies the `'-fsecond-underscore'` option, unless `'-fno-second-underscore'` is explicitly requested.

This does not affect the generation of code that interfaces with the **libgfortran** library.

Caution: It is not a good idea to mix Fortran code compiled with `'-ff2c'` with code compiled with the default `'-fno-f2c'` calling conventions as, calling

COMPLEX or default REAL functions between program parts which were compiled with different calling conventions will break at execution time.

Caution: This will break code which passes intrinsic functions of type default REAL or COMPLEX as actual arguments, as the library implementations use the ‘-fno-f2c’ calling conventions.

-fno-underscoring

Do not transform names of entities specified in the Fortran source file by appending underscores to them.

With ‘-funderscoring’ in effect, GNU Fortran appends one underscore to external names with no underscores. This is done to ensure compatibility with code produced by many UNIX Fortran compilers.

Caution: The default behavior of GNU Fortran is incompatible with f2c and g77, please use the ‘-ff2c’ option if you want object files compiled with GNU Fortran to be compatible with object code created with these tools.

Use of ‘-fno-underscoring’ is not recommended unless you are experimenting with issues such as integration of GNU Fortran into existing system environments (vis-à-vis existing libraries, tools, and so on).

For example, with ‘-funderscoring’, and assuming that `j()` and `max_count()` are external functions while `my_var` and `lvar` are local variables, a statement like

```
I = J() + MAX_COUNT (MY_VAR, LVAR)
```

is implemented as something akin to:

```
i = j_() + max_count__(&my_var__, &lvar);
```

With ‘-fno-underscoring’, the same statement is implemented as:

```
i = j() + max_count(&my_var, &lvar);
```

Use of ‘-fno-underscoring’ allows direct specification of user-defined names while debugging and when interfacing GNU Fortran code with other languages.

Note that just because the names match does *not* mean that the interface implemented by GNU Fortran for an external name matches the interface implemented by some other language for that same name. That is, getting code produced by GNU Fortran to link to code produced by some other compiler using this or any other method can be only a small part of the overall solution—getting the code generated by both compilers to agree on issues other than naming can require significant effort, and, unlike naming disagreements, linkers normally cannot detect disagreements in these other areas.

Also, note that with ‘-fno-underscoring’, the lack of appended underscores introduces the very real possibility that a user-defined external name will conflict with a name in a system library, which could make finding unresolved-reference bugs quite difficult in some cases—they might occur at program run time, and show up only as buggy behavior at run time.

In future versions of GNU Fortran we hope to improve naming and linking issues so that debugging always involves using the names as they appear in the source, even if the names as seen by the linker are mangled to prevent accidental linking between procedures with incompatible interfaces.

-fsecond-underscore

By default, GNU Fortran appends an underscore to external names. If this option is used GNU Fortran appends two underscores to names with underscores and one underscore to external names with no underscores. GNU Fortran also appends two underscores to internal names with underscores to avoid naming collisions with external names.

This option has no effect if ‘-fno-underscoring’ is in effect. It is implied by the ‘-ff2c’ option.

Otherwise, with this option, an external name such as MAX_COUNT is implemented as a reference to the link-time external symbol `max_count__`, instead of `max_count_`. This is required for compatibility with g77 and f2c, and is implied by use of the ‘-ff2c’ option.

-fcoarray=<keyword>

- ‘none’ Disable coarray support; using coarray declarations and image-control statements will produce a compile-time error. (Default)
- ‘single’ Single-image mode, i.e. `num_images()` is always one.
- ‘lib’ Library-based coarray parallelization; a suitable GNU Fortran co-array library needs to be linked.

-fcheck=<keyword>

Enable the generation of run-time checks; the argument shall be a comma-delimited list of the following keywords. Prefixing a check with ‘no-’ disables it if it was activated by a previous specification.

- ‘all’ Enable all run-time test of ‘-fcheck’.

‘array-temps’

Warns at run time when for passing an actual argument a temporary array had to be generated. The information generated by this warning is sometimes useful in optimization, in order to avoid such temporaries.

Note: The warning is only printed once per location.

- ‘bounds’ Enable generation of run-time checks for array subscripts and against the declared minimum and maximum values. It also checks array indices for assumed and deferred shape arrays against the actual allocated bounds and ensures that all string lengths are equal for character array constructors without an explicit typespec.

Some checks require that ‘-fcheck=bounds’ is set for the compilation of the main program.

Note: In the future this may also include other forms of checking, e.g., checking substring references.

- ‘do’ Enable generation of run-time checks for invalid modification of loop iteration variables.

- `'mem'` Enable generation of run-time checks for memory allocation. Note: This option does not affect explicit allocations using the `ALLOCATE` statement, which will be always checked.
- `'pointer'` Enable generation of run-time checks for pointers and allocatables.
- `'recursion'` Enable generation of run-time checks for recursively called subroutines and functions which are not marked as recursive. See also `'-frecursive'`. Note: This check does not work for OpenMP programs and is disabled if used together with `'-frecursive'` and `'-fopenmp'`.

Example: Assuming you have a file `'foo.f90'`, the command

```
gfortran -fcheck=all,no-array-temps foo.f90
```

will compile the file with all checks enabled as specified above except warnings for generated array temporaries.

`-fbounds-check`

Deprecated alias for `'-fcheck=bounds'`.

`-ftail-call-workaround`

`-ftail-call-workaround=n`

Some C interfaces to Fortran codes violate the gfortran ABI by omitting the hidden character length arguments as described in See [Section 7.4.2 \[Argument passing conventions\], page 85](#). This can lead to crashes because pushing arguments for tail calls can overflow the stack.

To provide a workaround for existing binary packages, this option disables tail call optimization for gfortran procedures with character arguments. With `'-ftail-call-workaround=2'` tail call optimization is disabled in all gfortran procedures with character arguments, with `'-ftail-call-workaround=1'` or equivalent `'-ftail-call-workaround'` only in gfortran procedures with character arguments that call implicitly prototyped procedures.

Using this option can lead to problems including crashes due to insufficient stack space.

It is *very strongly* recommended to fix the code in question. The `'-fc-prototypes-external'` option can be used to generate prototypes which conform to gfortran's ABI, for inclusion in the source code.

Support for this option will likely be withdrawn in a future release of gfortran.

The negative form, `'-fno-tail-call-workaround'` or equivalent `'-ftail-call-workaround=0'`, can be used to disable this option.

Default is currently `'-ftail-call-workaround'`, this will change in future releases.

`-fcheck-array-temporaries`

Deprecated alias for `'-fcheck=array-temps'`.

-fmax-array-constructor=n

This option can be used to increase the upper limit permitted in array constructors. The code below requires this option to expand the array at compile time.

```

program test
implicit none
integer j
integer, parameter :: n = 100000
integer, parameter :: i(n) = (/ (2*j, j = 1, n) /)
print '(10(I0,1X))', i
end program test

```

Caution: This option can lead to long compile times and excessively large object files.

The default value for *n* is 65535.

-fmax-stack-var-size=n

This option specifies the size in bytes of the largest array that will be put on the stack; if the size is exceeded static memory is used (except in procedures marked as RECURSIVE). Use the option **-frecursive** to allow for recursive procedures which do not have a RECURSIVE attribute or for parallel programs. Use **-fno-automatic** to never use the stack.

This option currently only affects local arrays declared with constant bounds, and may not apply to all character variables. Future versions of GNU Fortran may improve this behavior.

The default value for *n* is 32768.

-fstack-arrays

Adding this option will make the Fortran compiler put all arrays of unknown size and array temporaries onto stack memory. If your program uses very large local arrays it is possible that you will have to extend your runtime limits for stack memory on some operating systems. This flag is enabled by default at optimization level **-Ofast** unless **-fmax-stack-var-size** is specified.

-fpack-derived

This option tells GNU Fortran to pack derived type members as closely as possible. Code compiled with this option is likely to be incompatible with code compiled without this option, and may execute slower.

-frepack-arrays

In some circumstances GNU Fortran may pass assumed shape array sections via a descriptor describing a noncontiguous area of memory. This option adds code to the function prologue to repack the data into a contiguous block at runtime.

This should result in faster accesses to the array. However it can introduce significant overhead to the function call, especially when the passed data is noncontiguous.

-fshort-enums

This option is provided for interoperability with C code that was compiled with the **-fshort-enums** option. It will make GNU Fortran choose the smallest

INTEGER kind a given enumerator set will fit in, and give all its enumerators this kind.

`-fexternal-blas`

This option will make `gfortran` generate calls to BLAS functions for some matrix operations like `MATMUL`, instead of using our own algorithms, if the size of the matrices involved is larger than a given limit (see `'-fblas-matmul-limit'`). This may be profitable if an optimized vendor BLAS library is available. The BLAS library will have to be specified at link time.

`-fblas-matmul-limit=n`

Only significant when `'-fexternal-blas'` is in effect. Matrix multiplication of matrices with size larger than (or equal to) n will be performed by calls to BLAS functions, while others will be handled by `gfortran` internal algorithms. If the matrices involved are not square, the size comparison is performed using the geometric mean of the dimensions of the argument and result matrices.

The default value for n is 30.

`-finline-matmul-limit=n`

When front-end optimization is active, some calls to the `MATMUL` intrinsic function will be inlined. This may result in code size increase if the size of the matrix cannot be determined at compile time, as code for both cases is generated. Setting `-finline-matmul-limit=0` will disable inlining in all cases. Setting this option with a value of n will produce inline code for matrices with size up to n . If the matrices involved are not square, the size comparison is performed using the geometric mean of the dimensions of the argument and result matrices. The default value for n is 30. The `-fblas-matmul-limit` can be used to change this value.

`-frecursive`

Allow indirect recursion by forcing all local arrays to be allocated on the stack. This flag cannot be used together with `'-fmax-stack-var-size='` or `'-fno-automatic'`.

`-finit-local-zero`

`-finit-derived`

`-finit-integer=n`

`-finit-real=<zero|inf|-inf|nan|snan>`

`-finit-logical=<true|false>`

`-finit-character=n`

The `'-finit-local-zero'` option instructs the compiler to initialize local INTEGER, REAL, and COMPLEX variables to zero, LOGICAL variables to false, and CHARACTER variables to a string of null bytes. Finer-grained initialization options are provided by the `'-finit-integer=n'`, `'-finit-real=<zero|inf|-inf|nan|snan>'` (which also initializes the real and imaginary parts of local COMPLEX variables), `'-finit-logical=<true|false>'`, and `'-finit-character=n'` (where n is an ASCII character value) options.

With `'-finit-derived'`, components of derived type variables will be initialized according to these flags. Components whose type is not covered by an explicit `'-finit-*` flag will be treated as described above with `'-finit-local-zero'`.

These options do not initialize

- objects with the `POINTER` attribute
- allocatable arrays
- variables that appear in an `EQUIVALENCE` statement.

(These limitations may be removed in future releases).

Note that the `'-finit-real=nan'` option initializes `REAL` and `COMPLEX` variables with a quiet NaN. For a signalling NaN use `'-finit-real=snan'`; note, however, that compile-time optimizations may convert them into quiet NaN and that trapping needs to be enabled (e.g. via `'-ffpe-trap'`).

The `'-finit-integer'` option will parse the value into an integer of type `INTEGER(kind=C_LONG)` on the host. Said value is then assigned to the integer variables in the Fortran code, which might result in wraparound if the value is too large for the kind.

Finally, note that enabling any of the `'-finit-*` options will silence warnings that would have been emitted by `'-Wuninitialized'` for the affected local variables.

`-falign-commons`

By default, `gfortran` enforces proper alignment of all variables in a `COMMON` block by padding them as needed. On certain platforms this is mandatory, on others it increases performance. If a `COMMON` block is not declared with consistent data types everywhere, this padding can cause trouble, and `'-fno-align-commons'` can be used to disable automatic alignment. The same form of this option should be used for all files that share a `COMMON` block. To avoid potential alignment issues in `COMMON` blocks, it is recommended to order objects from largest to smallest.

`-fno-protect-parens`

By default the parentheses in expression are honored for all optimization levels such that the compiler does not do any re-association. Using `'-fno-protect-parens'` allows the compiler to reorder `REAL` and `COMPLEX` expressions to produce faster code. Note that for the re-association optimization `'-fno-signed-zeros'` and `'-fno-trapping-math'` need to be in effect. The parentheses protection is enabled by default, unless `'-Ofast'` is given.

`-frealloc-lhs`

An allocatable left-hand side of an intrinsic assignment is automatically (re)allocated if it is either unallocated or has a different shape. The option is enabled by default except when `'-std=f95'` is given. See also `'-Wrealloc-lhs'`.

`-faggressive-function-elimination`

Functions with identical argument lists are eliminated within statements, regardless of whether these functions are marked `PURE` or not. For example, in

```
a = f(b,c) + f(b,c)
```

there will only be a single call to `f`. This option only works if `'-ffrontend-optimize'` is in effect.

-ffrontend-optimize

This option performs front-end optimization, based on manipulating parts the Fortran parse tree. Enabled by default by any ‘-O’ option except ‘-O0’ and ‘-Og’. Optimizations enabled by this option include:

- inlining calls to `MATMUL`,
- elimination of identical function calls within expressions,
- removing unnecessary calls to `TRIM` in comparisons and assignments,
- replacing `TRIM(a)` with `a(1:LEN_TRIM(a))` and
- short-circuiting of logical operators (`.AND.` and `.OR.`).

It can be deselected by specifying ‘-fno-frontend-optimize’.

-ffrontend-loop-interchange

Attempt to interchange loops in the Fortran front end where profitable. Enabled by default by any ‘-O’ option. At the moment, this option only affects `FORALL` and `DO CONCURRENT` statements with several forall triplets.

See Section “Options for Code Generation Conventions” in *Using the GNU Compiler Collection (GCC)*, for information on more options offered by the GBE shared by `gfortran`, `gcc`, and other GNU compilers.

2.10 Options for interoperability with other languages

-fc-prototypes

This option will generate C prototypes from `BIND(C)` variable declarations, types and procedure interfaces and writes them to standard output. `ENUM` is not yet supported.

The generated prototypes may need inclusion of an appropriate header, such as `<stdint.h>` or `<stdlib.h>`. For types which are not specified using the appropriate kind from the `iso_c_binding` module, a warning is added as a comment to the code.

For function pointers, a pointer to a function returning `int` without an explicit argument list is generated.

Example of use:

```
$ gfortran -fc-prototypes -fsyntax-only foo.f90 > foo.h
```

where the C code intended for interoperating with the Fortran code then uses `#include "foo.h"`.

-fc-prototypes-external

This option will generate C prototypes from external functions and subroutines and write them to standard output. This may be useful for making sure that C bindings to Fortran code are correct. This option does not generate prototypes for `BIND(C)` procedures, use ‘-fc-prototypes’ for that.

The generated prototypes may need inclusion of an appropriate header, such as `<stdint.h>` or `<stdlib.h>`.

This is primarily meant for legacy code to ensure that existing C bindings match what `gfortran` emits. The generated C prototypes should be correct for the

current version of the compiler, but may not match what other compilers or earlier versions of **gfortran** need. For new developments, use of the **BIND(C)** features is recommended.

Example of use:

```
$ gfortran -fc-prototypes-external -fsyntax-only foo.f > foo.h
```

where the C code intended for interoperating with the Fortran code then uses `#include "foo.h"`.

2.11 Environment variables affecting gfortran

The **gfortran** compiler currently does not make use of any environment variables to control its operation above and beyond those that affect the operation of **gcc**.

See Section “Environment Variables Affecting GCC” in *Using the GNU Compiler Collection (GCC)*, for information on environment variables.

See Chapter 3 [Runtime], page 33, for environment variables that affect the run-time behavior of programs compiled with GNU Fortran.

3 Runtime: Influencing runtime behavior with environment variables

The behavior of the `gfortran` can be influenced by environment variables.

Malformed environment variables are silently ignored.

3.1 TMPDIR—Directory for scratch files

When opening a file with `STATUS='SCRATCH'`, GNU Fortran tries to create the file in one of the potential directories by testing each directory in the order below.

1. The environment variable `TMPDIR`, if it exists.
2. On the MinGW target, the directory returned by the `GetTempPath` function. Alternatively, on the Cygwin target, the `TMP` and `TEMP` environment variables, if they exist, in that order.
3. The `P_tmpdir` macro if it is defined, otherwise the directory `'/tmp'`.

3.2 GFORTRAN_STDIN_UNIT—Unit number for standard input

This environment variable can be used to select the unit number preconnected to standard input. This must be a positive integer. The default value is 5.

3.3 GFORTRAN_STDOUT_UNIT—Unit number for standard output

This environment variable can be used to select the unit number preconnected to standard output. This must be a positive integer. The default value is 6.

3.4 GFORTRAN_STDERR_UNIT—Unit number for standard error

This environment variable can be used to select the unit number preconnected to standard error. This must be a positive integer. The default value is 0.

3.5 GFORTRAN_UNBUFFERED_ALL—Do not buffer I/O on all units

This environment variable controls whether all I/O is unbuffered. If the first letter is 'y', 'Y' or '1', all I/O is unbuffered. This will slow down small sequential reads and writes. If the first letter is 'n', 'N' or '0', I/O is buffered. This is the default.

3.6 GFORTRAN_UNBUFFERED_PRECONNECTED—Do not buffer I/O on preconnected units

The environment variable named `GFORTRAN_UNBUFFERED_PRECONNECTED` controls whether I/O on a preconnected unit (i.e. `STDOUT` or `STDERR`) is unbuffered. If the first letter is 'y', 'Y' or '1', I/O is unbuffered. This will slow down small sequential reads and writes. If the first letter is 'n', 'N' or '0', I/O is buffered. This is the default.

3.7 GFORTRAN_SHOW_LOCUS—Show location for runtime errors

If the first letter is 'y', 'Y' or '1', filename and line numbers for runtime errors are printed. If the first letter is 'n', 'N' or '0', do not print filename and line numbers for runtime errors. The default is to print the location.

3.8 GFORTRAN_OPTIONAL_PLUS—Print leading + where permitted

If the first letter is ‘y’, ‘Y’ or ‘1’, a plus sign is printed where permitted by the Fortran standard. If the first letter is ‘n’, ‘N’ or ‘O’, a plus sign is not printed in most cases. Default is not to print plus signs.

3.9 GFORTRAN_LIST_SEPARATOR—Separator for list output

This environment variable specifies the separator when writing list-directed output. It may contain any number of spaces and at most one comma. If you specify this on the command line, be sure to quote spaces, as in

```
$ GFORTRAN_LIST_SEPARATOR=' , ' ./a.out
```

when `a.out` is the compiled Fortran program that you want to run. Default is a single space.

3.10 GFORTRAN_CONVERT_UNIT—Set endianness for unformatted I/O

By setting the `GFORTRAN_CONVERT_UNIT` variable, it is possible to change the representation of data for unformatted files. The syntax for the `GFORTRAN_CONVERT_UNIT` variable is:

```
GFORTRAN_CONVERT_UNIT: mode | mode ';' exception | exception ;
mode: 'native' | 'swap' | 'big_endian' | 'little_endian' ;
exception: mode ':' unit_list | unit_list ;
unit_list: unit_spec | unit_list unit_spec ;
unit_spec: INTEGER | INTEGER '-' INTEGER ;
```

The variable consists of an optional default mode, followed by a list of optional exceptions, which are separated by semicolons from the preceding default and each other. Each exception consists of a format and a comma-separated list of units. Valid values for the modes are the same as for the `CONVERT` specifier:

NATIVE Use the native format. This is the default.

SWAP Swap between little- and big-endian.

LITTLE_ENDIAN Use the little-endian format for unformatted files.

BIG_ENDIAN Use the big-endian format for unformatted files.

A missing mode for an exception is taken to mean **BIG_ENDIAN**. Examples of values for `GFORTRAN_CONVERT_UNIT` are:

`'big_endian'` Do all unformatted I/O in big-endian mode.

`'little_endian;native:10-20,25'` Do all unformatted I/O in little-endian mode, except for units 10 to 20 and 25, which are in native format.

`'10-20'` Units 10 to 20 are big-endian, the rest is native.

Setting the environment variables should be done on the command line or via the `export` command for `sh`-compatible shells and via `setenv` for `csh`-compatible shells.

Example for `sh`:

```
$ gfortran foo.f90
$ GFORTRAN_CONVERT_UNIT='big_endian;native:10-20' ./a.out
```

Example code for `csh`:

```
% gfortran foo.f90
% setenv GFORTRAN_CONVERT_UNIT 'big_endian;native:10-20'
% ./a.out
```

Using anything but the native representation for unformatted data carries a significant speed overhead. If speed in this area matters to you, it is best if you use this only for data that needs to be portable.

See [Section 6.1.15 \[CONVERT specifier\]](#), page 57, for an alternative way to specify the data representation for unformatted files. See [Section 2.8 \[Runtime Options\]](#), page 22, for setting a default data representation for the whole program. The `CONVERT` specifier overrides the `-fconvert` compile options.

Note that the values specified via the `GFORTRAN_CONVERT_UNIT` environment variable will override the `CONVERT` specifier in the open statement. This is to give control over data formats to users who do not have the source code of their program available.

3.11 GFORTRAN_ERROR_BACKTRACE—Show backtrace on run-time errors

If the `GFORTRAN_ERROR_BACKTRACE` variable is set to `'y'`, `'Y'` or `'1'` (only the first letter is relevant) then a backtrace is printed when a serious run-time error occurs. To disable the backtracing, set the variable to `'n'`, `'N'`, `'0'`. Default is to print a backtrace unless the `-fno-backtrace` compile option was used.

3.12 GFORTRAN_FORMATTED_BUFFER_SIZE—Set buffer size for formatted I/O

The `GFORTRAN_FORMATTED_BUFFER_SIZE` environment variable specifies buffer size in bytes to be used for formatted output. The default value is 8192.

3.13 GFORTRAN_UNFORMATTED_BUFFER_SIZE—Set buffer size for unformatted I/O

The `GFORTRAN_UNFORMATTED_BUFFER_SIZE` environment variable specifies buffer size in bytes to be used for unformatted output. The default value is 131072.

Part II: Language Reference

4 Fortran standards status

4.1 Fortran 2003 status

GNU Fortran supports several Fortran 2003 features; an incomplete list can be found below. See also the [wiki page](#) about Fortran 2003.

- Procedure pointers including procedure-pointer components with `PASS` attribute.
- Procedures which are bound to a derived type (type-bound procedures) including `PASS`, `PROCEDURE` and `GENERIC`, and operators bound to a type.
- Abstract interfaces and type extension with the possibility to override type-bound procedures or to have deferred binding.
- Polymorphic entities (“`CLASS`”) for derived types and unlimited polymorphism (“`CLASS(*)`”) – including `SAME_TYPE_AS`, `EXTENDS_TYPE_OF` and `SELECT TYPE` for scalars and arrays and finalization.
- Generic interface names, which have the same name as derived types, are now supported. This allows one to write constructor functions. Note that Fortran does not support static constructor functions. For static variables, only default initialization or structure-constructor initialization are available.
- The `ASSOCIATE` construct.
- Interoperability with C including enumerations,
- In structure constructors the components with default values may be omitted.
- Extensions to the `ALLOCATE` statement, allowing for a type-specification with type parameter and for allocation and initialization from a `SOURCE=` expression; `ALLOCATE` and `DEALLOCATE` optionally return an error message string via `ERRMSG=`.
- Reallocation on assignment: If an intrinsic assignment is used, an allocatable variable on the left-hand side is automatically allocated (if unallocated) or reallocated (if the shape is different). Currently, scalar deferred character length left-hand sides are correctly handled but arrays are not yet fully implemented.
- Deferred-length character variables and scalar deferred-length character components of derived types are supported. (Note that array-valued components are not yet implemented.)
- Transferring of allocations via `MOVE_ALLOC`.
- The `PRIVATE` and `PUBLIC` attributes may be given individually to derived-type components.
- In pointer assignments, the lower bound may be specified and the remapping of elements is supported.
- For pointers an `INTENT` may be specified which affect the association status not the value of the pointer target.
- Intrinsic `command_argument_count`, `get_command`, `get_command_argument`, and `get_environment_variable`.
- Support for Unicode characters (ISO 10646) and UTF-8, including the `SELECTED_CHAR_KIND` and `NEW_LINE` intrinsic functions.

- Support for binary, octal and hexadecimal (BOZ) constants in the intrinsic functions INT, REAL, CMPLX and DBLE.
- Support for namelist variables with allocatable and pointer attribute and nonconstant length type parameter.
- Array constructors using square brackets. That is, [...] rather than (/.../). Type-specification for array constructors like (/ some-type :: ... /).
- Extensions to the specification and initialization expressions, including the support for intrinsics with real and complex arguments.
- Support for the asynchronous input/output.
- FLUSH statement.
- IOMSG= specifier for I/O statements.
- Support for the declaration of enumeration constants via the ENUM and ENUMERATOR statements. Interoperability with gcc is guaranteed also for the case where the -fshort-enums command line option is given.
- TR 15581:
 - ALLOCATABLE dummy arguments.
 - ALLOCATABLE function results
 - ALLOCATABLE components of derived types
- The OPEN statement supports the ACCESS='STREAM' specifier, allowing I/O without any record structure.
- Namelist input/output for internal files.
- Minor I/O features: Rounding during formatted output, using of a decimal comma instead of a decimal point, setting whether a plus sign should appear for positive numbers. On systems where strtod honours the rounding mode, the rounding mode is also supported for input.
- The PROTECTED statement and attribute.
- The VALUE statement and attribute.
- The VOLATILE statement and attribute.
- The IMPORT statement, allowing to import host-associated derived types.
- The intrinsic modules ISO_FORTRAN_ENVIRONMENT is supported, which contains parameters of the I/O units, storage sizes. Additionally, procedures for C interoperability are available in the ISO_C_BINDING module.
- USE statement with INTRINSIC and NON_INTRINSIC attribute; supported intrinsic modules: ISO_FORTRAN_ENV, ISO_C_BINDING, OMP_LIB and OMP_LIB_KINDS, and OPENACC.
- Renaming of operators in the USE statement.

4.2 Fortran 2008 status

The latest version of the Fortran standard is ISO/IEC 1539-1:2010, informally known as Fortran 2008. The official version is available from International Organization for Standardization (ISO) or its national member organizations. The the final draft (FDIS) can be downloaded free of charge from <http://www.nag.co.uk/sc22wg5/links.html>. Fortran is developed by the Working Group 5 of Sub-Committee 22 of the Joint Technical

Committee 1 of the International Organization for Standardization and the International Electrotechnical Commission (IEC). This group is known as **WG5**.

The GNU Fortran compiler supports several of the new features of Fortran 2008; the [wiki](#) has some information about the current Fortran 2008 implementation status. In particular, the following is implemented.

- The `-std=f2008` option and support for the file extensions `.f08` and `.F08`.
- The `OPEN` statement now supports the `NEWUNIT=` option, which returns a unique file unit, thus preventing inadvertent use of the same unit in different parts of the program.
- The `g0` format descriptor and unlimited format items.
- The mathematical intrinsics `ASINH`, `ACOSH`, `ATANH`, `ERF`, `ERFC`, `GAMMA`, `LOG_GAMMA`, `BESSEL_J0`, `BESSEL_J1`, `BESSEL_JN`, `BESSEL_Y0`, `BESSEL_Y1`, `BESSEL_YN`, `HYPOT`, `NORM2`, and `ERFC_SCALED`.
- Using complex arguments with `TAN`, `SINH`, `COSH`, `TANH`, `ASIN`, `ACOS`, and `ATAN` is now possible; `ATAN(Y,X)` is now an alias for `ATAN2(Y,X)`.
- Support of the `PARITY` intrinsic functions.
- The following bit intrinsics: `LEADZ` and `TRAILZ` for counting the number of leading and trailing zero bits, `POPCNT` and `POPPAR` for counting the number of one bits and returning the parity; `BGE`, `BGT`, `BLE`, and `BLT` for bitwise comparisons; `DSHIFTL` and `DSHIFTR` for combined left and right shifts, `MASKL` and `MASKR` for simple left and right justified masks, `MERGE_BITS` for a bitwise merge using a mask, `SHIFTA`, `SHIFTL` and `SHIFTR` for shift operations, and the transformational bit intrinsics `IALL`, `IANY` and `IPARITY`.
- Support of the `EXECUTE_COMMAND_LINE` intrinsic subroutine.
- Support for the `STORAGE_SIZE` intrinsic inquiry function.
- The `INT{8,16,32}` and `REAL{32,64,128}` kind type parameters and the array-valued named constants `INTEGER_KINDS`, `LOGICAL_KINDS`, `REAL_KINDS` and `CHARACTER_KINDS` of the intrinsic module `ISO_FORTRAN_ENV`.
- The module procedures `C_SIZEOF` of the intrinsic module `ISO_C_BINDINGS` and `COMPILER_VERSION` and `COMPILER_OPTIONS` of `ISO_FORTRAN_ENV`.
- Coarray support for serial programs with `-fcoarray=single` flag and experimental support for multiple images with the `-fcoarray=lib` flag.
- Submodules are supported. It should be noted that `MODULEs` do not produce the `smod` file needed by the descendent `SUBMODULEs` unless they contain at least one `MODULE PROCEDURE` interface. The reason for this is that `SUBMODULEs` are useless without `MODULE PROCEDUREs`. See <http://j3-fortran.org/doc/meeting/207/15-209.txt> for a discussion and a draft interpretation. Adopting this interpretation has the advantage that code that does not use submodules does not generate `smod` files.
- The `DO CONCURRENT` construct is supported.
- The `BLOCK` construct is supported.
- The `STOP` and the new `ERROR STOP` statements now support all constant expressions. Both show the signals which were signaling at termination.
- Support for the `CONTIGUOUS` attribute.
- Support for `ALLOCATE` with `MOLD`.

- Support for the **IMPURE** attribute for procedures, which allows for **ELEMENTAL** procedures without the restrictions of **PURE**.
- Null pointers (including `NULL()`) and not-allocated variables can be used as actual argument to optional non-pointer, non-allocatable dummy arguments, denoting an absent argument.
- Non-pointer variables with **TARGET** attribute can be used as actual argument to **POINTER** dummies with **INTENT(IN)**.
- Pointers including procedure pointers and those in a derived type (pointer components) can now be initialized by a target instead of only by `NULL`.
- The **EXIT** statement (with construct-name) can be now be used to leave not only the **DO** but also the **ASSOCIATE**, **BLOCK**, **IF**, **SELECT CASE** and **SELECT TYPE** constructs.
- Internal procedures can now be used as actual argument.
- Minor features: obsolesce diagnostics for **ENTRY** with `'-std=f2008'`; a line may start with a semicolon; for internal and module procedures **END** can be used instead of **END SUBROUTINE** and **END FUNCTION**; **SELECTED_REAL_KIND** now also takes a **RADIX** argument; intrinsic types are supported for **TYPE**(*intrinsic-type-spec*); multiple type-bound procedures can be declared in a single **PROCEDURE** statement; implied-shape arrays are supported for named constants (**PARAMETER**).

4.3 Status of Fortran 2018 support

- **ERROR STOP** in a **PURE** procedure An **ERROR STOP** statement is permitted in a **PURE** procedure.
- **IMPLICIT NONE** with a spec-list Support the **IMPLICIT NONE** statement with an `implicit-none-spec-list`.
- Behavior of **INQUIRE** with the **RECL=** specifier

The behavior of the **INQUIRE** statement with the **RECL=** specifier now conforms to Fortran 2018.

4.3.1 TS 29113 Status (Further Interoperability with C)

GNU Fortran supports some of the new features of the Technical Specification (TS) 29113 on Further Interoperability of Fortran with C. The [wiki](#) has some information about the current TS 29113 implementation status. In particular, the following is implemented.

See also [Section 7.1.6 \[Further Interoperability of Fortran with C\]](#), page 78.

- The **OPTIONAL** attribute is allowed for dummy arguments of **BIND(C)** procedures.
- The **RANK** intrinsic is supported.
- GNU Fortran's implementation for variables with **ASYNCHRONOUS** attribute is compatible with TS 29113.
- Assumed types (**TYPE(*)**).
- Assumed-rank (**DIMENSION(..)**).
- **ISO_Fortran_binding** (now in Fortran 2018 18.4) is implemented such that conversion of the array descriptor for assumed type or assumed rank arrays is done in the library. The include file `ISO_Fortran_binding.h` is can be found in `~prefix/lib/gcc/$target/$version`.

4.3.2 TS 18508 Status (Additional Parallel Features)

GNU Fortran supports the following new features of the Technical Specification 18508 on Additional Parallel Features in Fortran:

- The new atomic `ADD`, `CAS`, `FETCH` and `ADD/OR/XOR`, `OR` and `XOR` intrinsics.
- The `CO_MIN` and `CO_MAX` and `SUM` reduction intrinsics. And the `CO_BROADCAST` and `CO_REDUCE` intrinsic, except that those do not support polymorphic types or types with allocatable, pointer or polymorphic components.
- Events (`EVENT POST`, `EVENT WAIT`, `EVENT_QUERY`)
- Failed images (`FAIL IMAGE`, `IMAGE_STATUS`, `FAILED_IMAGES`, `STOPPED_IMAGES`)

5 Compiler Characteristics

This chapter describes certain characteristics of the GNU Fortran compiler, that are not specified by the Fortran standard, but which might in some way or another become visible to the programmer.

5.1 KIND Type Parameters

The KIND type parameters supported by GNU Fortran for the primitive data types are:

INTEGER 1, 2, 4, 8*, 16*, default: 4**

LOGICAL 1, 2, 4, 8*, 16*, default: 4**

REAL 4, 8, 10*, 16*, default: 4***

COMPLEX 4, 8, 10*, 16*, default: 4***

DOUBLE PRECISION
 4, 8, 10*, 16*, default: 8***

CHARACTER
 1, 4, default: 1

* not available on all systems

** unless `-fdefault-integer-8` is used

*** unless `-fdefault-real-8` is used (see [Section 2.2 \[Fortran Dialect Options\]](#), page 8)

The KIND value matches the storage size in bytes, except for COMPLEX where the storage size is twice as much (or both real and imaginary part are a real value of the given size). It is recommended to use the [Section 9.240 \[SELECTED_CHAR_KIND\]](#), page 258, [Section 9.241 \[SELECTED_INT_KIND\]](#), page 259 and [Section 9.242 \[SELECTED_REAL_KIND\]](#), page 259 intrinsics or the INT8, INT16, INT32, INT64, REAL32, REAL64, and REAL128 parameters of the ISO_FORTRAN_ENV module instead of the concrete values. The available kind parameters can be found in the constant arrays CHARACTER_KINDS, INTEGER_KINDS, LOGICAL_KINDS and REAL_KINDS in the [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287 module. For C interoperability, the kind parameters of the [Section 10.2 \[ISO_C_BINDING\]](#), page 289 module should be used.

5.2 Internal representation of LOGICAL variables

The Fortran standard does not specify how variables of LOGICAL type are represented, beyond requiring that LOGICAL variables of default kind have the same storage size as default INTEGER and REAL variables. The GNU Fortran internal representation is as follows.

A LOGICAL(KIND=N) variable is represented as an INTEGER(KIND=N) variable, however, with only two permissible values: 1 for .TRUE. and 0 for .FALSE.. Any other integer value results in undefined behavior.

See also [Section 7.4.2 \[Argument passing conventions\]](#), page 85 and [Section 7.1 \[Interoperability with C\]](#), page 73.

5.3 Evaluation of logical expressions

The Fortran standard does not require the compiler to evaluate all parts of an expression, if they do not contribute to the final result. For logical expressions with `.AND.` or `.OR.` operators, in particular, GNU Fortran will optimize out function calls (even to impure functions) if the result of the expression can be established without them. However, since not all compilers do that, and such an optimization can potentially modify the program flow and subsequent results, GNU Fortran throws warnings for such situations with the `'-Wfunction-elimination'` flag.

5.4 MAX and MIN intrinsics with REAL NaN arguments

The Fortran standard does not specify what the result of the `MAX` and `MIN` intrinsics are if one of the arguments is a NaN. Accordingly, the GNU Fortran compiler does not specify that either, as this allows for faster and more compact code to be generated. If the programmer wishes to take some specific action in case one of the arguments is a NaN, it is necessary to explicitly test the arguments before calling `MAX` or `MIN`, e.g. with the `IEEE_IS_NAN` function from the intrinsic module `IEEE_ARITHMETIC`.

5.5 Thread-safety of the runtime library

GNU Fortran can be used in programs with multiple threads, e.g. by using OpenMP, by calling OS thread handling functions via the `ISO_C_BINDING` facility, or by GNU Fortran compiled library code being called from a multi-threaded program.

The GNU Fortran runtime library, (`libgfortran`), supports being called concurrently from multiple threads with the following exceptions.

During library initialization, the C `getenv` function is used, which need not be thread-safe. Similarly, the `getenv` function is used to implement the `GET_ENVIRONMENT_VARIABLE` and `GETENV` intrinsics. It is the responsibility of the user to ensure that the environment is not being updated concurrently when any of these actions are taking place.

The `EXECUTE_COMMAND_LINE` and `SYSTEM` intrinsics are implemented with the `system` function, which need not be thread-safe. It is the responsibility of the user to ensure that `system` is not called concurrently.

For platforms not supporting thread-safe POSIX functions, further functionality might not be thread-safe. For details, please consult the documentation for your operating system.

The GNU Fortran runtime library uses various C library functions that depend on the locale, such as `strtod` and `snprintf`. In order to work correctly in locale-aware programs that set the locale using `setlocale`, the locale is reset to the default “C” locale while executing a formatted `READ` or `WRITE` statement. On targets supporting the POSIX 2008 per-thread locale functions (e.g. `newlocale`, `uselocale`, `freelocale`), these are used and thus the global locale set using `setlocale` or the per-thread locales in other threads are not affected. However, on targets lacking this functionality, the global `LC_NUMERIC` locale is set to “C” during the formatted I/O. Thus, on such targets it's not safe to call `setlocale` concurrently from another thread while a Fortran formatted I/O operation is in progress. Also, other threads doing something dependent on the `LC_NUMERIC` locale might not work correctly if a formatted I/O operation is in progress in another thread.

5.6 Data consistency and durability

This section contains a brief overview of data and metadata consistency and durability issues when doing I/O.

With respect to durability, GNU Fortran makes no effort to ensure that data is committed to stable storage. If this is required, the GNU Fortran programmer can use the intrinsic `FNUM` to retrieve the low level file descriptor corresponding to an open Fortran unit. Then, using e.g. the `ISO_C_BINDING` feature, one can call the underlying system call to flush dirty data to stable storage, such as `fsync` on POSIX, `_commit` on MingW, or `fcntl(fd, F_FULLSYNC, 0)` on Mac OS X. The following example shows how to call `fsync`:

```
! Declare the interface for POSIX fsync function
interface
  function fsync (fd) bind(c,name="fsync")
    use iso_c_binding, only: c_int
    integer(c_int), value :: fd
    integer(c_int) :: fsync
  end function fsync
end interface

! Variable declaration
integer :: ret

! Opening unit 10
open (10,file="foo")

! ...
! Perform I/O on unit 10
! ...

! Flush and sync
flush(10)
ret = fsync(fnum(10))

! Handle possible error
if (ret /= 0) stop "Error calling FSYNC"
```

With respect to consistency, for regular files GNU Fortran uses buffered I/O in order to improve performance. This buffer is flushed automatically when full and in some other situations, e.g. when closing a unit. It can also be explicitly flushed with the `FLUSH` statement. Also, the buffering can be turned off with the `GFORTRAN_UNBUFFERED_ALL` and `GFORTRAN_UNBUFFERED_PRECONNECTED` environment variables. Special files, such as terminals and pipes, are always unbuffered. Sometimes, however, further things may need to be done in order to allow other processes to see data that GNU Fortran has written, as follows.

The Windows platform supports a relaxed metadata consistency model, where file metadata is written to the directory lazily. This means that, for instance, the `dir` command can show a stale size for a file. One can force a directory metadata update by closing the unit, or by calling `_commit` on the file descriptor. Note, though, that `_commit` will force all dirty data to stable storage, which is often a very slow operation.

The Network File System (NFS) implements a relaxed consistency model called open-to-close consistency. Closing a file forces dirty data and metadata to be flushed to the server, and opening a file forces the client to contact the server in order to revalidate cached data. `fsync` will also force a flush of dirty data and metadata to the server. Similar to `open` and

`close`, acquiring and releasing `fcntl` file locks, if the server supports them, will also force cache validation and flushing dirty data and metadata.

5.7 Files opened without an explicit `ACTION=` specifier

The Fortran standard says that if an `OPEN` statement is executed without an explicit `ACTION=` specifier, the default value is processor dependent. GNU Fortran behaves as follows:

1. Attempt to open the file with `ACTION='READWRITE'`
2. If that fails, try to open with `ACTION='READ'`
3. If that fails, try to open with `ACTION='WRITE'`
4. If that fails, generate an error

5.8 File operations on symbolic links

This section documents the behavior of GNU Fortran for file operations on symbolic links, on systems that support them.

- Results of `INQUIRE` statements of the “inquire by file” form will relate to the target of the symbolic link. For example, `INQUIRE(FILE="foo",EXIST=ex)` will set `ex` to `.true.` if `foo` is a symbolic link pointing to an existing file, and `.false.` if `foo` points to a non-existing file (“dangling” symbolic link).
- Using the `OPEN` statement with a `STATUS="NEW"` specifier on a symbolic link will result in an error condition, whether the symbolic link points to an existing target or is dangling.
- If a symbolic link was connected, using the `CLOSE` statement with a `STATUS="DELETE"` specifier will cause the symbolic link itself to be deleted, not its target.

5.9 File format of unformatted sequential files

Unformatted sequential files are stored as logical records using record markers. Each logical record consists of one or more subrecords.

Each subrecord consists of a leading record marker, the data written by the user program, and a trailing record marker. The record markers are four-byte integers by default, and eight-byte integers if the `'-fmax-subrecord-length=8'` option (which exists for backwards compatibility only) is in effect.

The representation of the record markers is that of unformatted files given with the `'-fconvert'` option, the [Section 6.1.15 \[CONVERT specifier\]](#), [page 57](#) in an open statement or the [Section 3.10 \[GFORTTRAN_CONVERT_UNIT\]](#), [page 34](#) environment variable.

The maximum number of bytes of user data in a subrecord is 2147483639 (2 GiB - 9) for a four-byte record marker. This limit can be lowered with the `'-fmax-subrecord-length'` option, although this is rarely useful. If the length of a logical record exceeds this limit, the data is distributed among several subrecords.

The absolute of the number stored in the record markers is the number of bytes of user data in the corresponding subrecord. If the leading record marker of a subrecord contains a negative number, another subrecord follows the current one. If the trailing record marker contains a negative number, then there is a preceding subrecord.

In the most simple case, with only one subrecord per logical record, both record markers contain the number of bytes of user data in the record.

The format for unformatted sequential data can be duplicated using unformatted stream, as shown in the example program for an unformatted record containing a single subrecord:

```
program main
  use iso_fortran_env, only: int32
  implicit none
  integer(int32) :: i
  real, dimension(10) :: a, b
  call random_number(a)
  open (10,file='test.dat',form='unformatted',access='stream')
  inquire (iolength=i) a
  write (10) i, a, i
  close (10)
  open (10,file='test.dat',form='unformatted')
  read (10) b
  if (all (a == b)) print *, 'success!'
end program main
```

5.10 Asynchronous I/O

Asynchronous I/O is supported if the program is linked against the POSIX thread library. If that is not the case, all I/O is performed as synchronous. On systems which do not support pthread condition variables, such as AIX, I/O is also performed as synchronous.

On some systems, such as Darwin or Solaris, the POSIX thread library is always linked in, so asynchronous I/O is always performed. On other systems, such as Linux, it is necessary to specify `-pthread`, `-lpthread` or `-fopenmp` during the linking step.

6 Extensions

The two sections below detail the extensions to standard Fortran that are implemented in GNU Fortran, as well as some of the popular or historically important extensions that are not (or not yet) implemented. For the latter case, we explain the alternatives available to GNU Fortran users, including replacement by standard-conforming code or GNU extensions.

6.1 Extensions implemented in GNU Fortran

GNU Fortran implements a number of extensions over standard Fortran. This chapter contains information on their syntax and meaning. There are currently two categories of GNU Fortran extensions, those that provide functionality beyond that provided by any standard, and those that are supported by GNU Fortran purely for backward compatibility with legacy compilers. By default, `-std=gnu` allows the compiler to accept both types of extensions, but to warn about the use of the latter. Specifying either `-std=f95`, `-std=f2003`, `-std=f2008`, or `-std=f2018` disables both types of extensions, and `-std=legacy` allows both without warning. The special compile flag `-fdec` enables additional compatibility extensions along with those enabled by `-std=legacy`.

6.1.1 Old-style kind specifications

GNU Fortran allows old-style kind specifications in declarations. These look like:

```
TYPESPEC*size x,y,z
```

where `TYPESPEC` is a basic type (`INTEGER`, `REAL`, etc.), and where `size` is a byte count corresponding to the storage size of a valid kind for that type. (For `COMPLEX` variables, `size` is the total size of the real and imaginary parts.) The statement then declares `x`, `y` and `z` to be of type `TYPESPEC` with the appropriate kind. This is equivalent to the standard-conforming declaration

```
TYPESPEC(k) x,y,z
```

where `k` is the kind parameter suitable for the intended precision. As kind parameters are implementation-dependent, use the `KIND`, `SELECTED_INT_KIND` and `SELECTED_REAL_KIND` intrinsics to retrieve the correct value, for instance `REAL*8 x` can be replaced by:

```
INTEGER, PARAMETER :: dbl = KIND(1.0d0)
REAL(KIND=dbl) :: x
```

6.1.2 Old-style variable initialization

GNU Fortran allows old-style initialization of variables of the form:

```
INTEGER i/1/,j/2/
REAL x(2,2) /3*0.,1./
```

The syntax for the initializers is as for the `DATA` statement, but unlike in a `DATA` statement, an initializer only applies to the variable immediately preceding the initialization. In other words, something like `INTEGER I,J/2,3/` is not valid. This style of initialization is only allowed in declarations without double colons (`::`); the double colons were introduced in Fortran 90, which also introduced a standard syntax for initializing variables in type declarations.

Examples of standard-conforming code equivalent to the above example are:

```

! Fortran 90
  INTEGER :: i = 1, j = 2
  REAL :: x(2,2) = RESHAPE((/0.,0.,0.,1./),SHAPE(x))
! Fortran 77
  INTEGER i, j
  REAL x(2,2)
  DATA i/1/, j/2/, x/3*0.,1./

```

Note that variables which are explicitly initialized in declarations or in DATA statements automatically acquire the SAVE attribute.

6.1.3 Extensions to namelist

GNU Fortran fully supports the Fortran 95 standard for namelist I/O including array qualifiers, substrings and fully qualified derived types. The output from a namelist write is compatible with namelist read. The output has all names in upper case and indentation to column 1 after the namelist name. Two extensions are permitted:

Old-style use of '\$' instead of '&'

```

$MYNML
  X(:)%Y(2) = 1.0 2.0 3.0
  CH(1:4) = "abcd"
$END

```

It should be noted that the default terminator is '/' rather than '&END'.

Querying of the namelist when inputting from stdin. After at least one space, entering '?' sends to stdout the namelist name and the names of the variables in the namelist:

```

?

&mynml
  x
  x%y
  ch
&end

```

Entering '=?' outputs the namelist to stdout, as if WRITE(*,NML = mynml) had been called:

```

=?

&MYNML
  X(1)%Y= 0.000000 , 1.000000 , 0.000000 ,
  X(2)%Y= 0.000000 , 2.000000 , 0.000000 ,
  X(3)%Y= 0.000000 , 3.000000 , 0.000000 ,
  CH=abcd, /

```

To aid this dialog, when input is from stdin, errors send their messages to stderr and execution continues, even if IOSTAT is set.

PRINT namelist is permitted. This causes an error if '-std=f95' is used.

```

PROGRAM test_print
  REAL, dimension (4) :: x = (/1.0, 2.0, 3.0, 4.0/)
  NAMELIST /mynml/ x
  PRINT mynml
END PROGRAM test_print

```

Expanded namelist reads are permitted. This causes an error if '-std=f95' is used. In the following example, the first element of the array will be given the value 0.00 and the two succeeding elements will be given the values 1.00 and 2.00.

```
&MYNML
  X(1,1) = 0.00 , 1.00 , 2.00
/
```

When writing a namelist, if no `DELIM=` is specified, by default a double quote is used to delimit character strings. If `-std=F95`, `F2003`, or `F2008`, etc, the `delim` status is set to 'none'. Defaulting to quotes ensures that namelists with character strings can be subsequently read back in accurately.

6.1.4 X format descriptor without count field

To support legacy codes, GNU Fortran permits the count field of the `X` edit descriptor in `FORMAT` statements to be omitted. When omitted, the count is implicitly assumed to be one.

```
      PRINT 10, 2, 3
10     FORMAT (I1, X, I1)
```

6.1.5 Commas in FORMAT specifications

To support legacy codes, GNU Fortran allows the comma separator to be omitted immediately before and after character string edit descriptors in `FORMAT` statements.

```
      PRINT 10, 2, 3
10     FORMAT ('F00=' I1' BAR=' I2)
```

6.1.6 Missing period in FORMAT specifications

To support legacy codes, GNU Fortran allows missing periods in format specifications if and only if `'-std=legacy'` is given on the command line. This is considered non-conforming code and is discouraged.

```
      REAL :: value
      READ(*,10) value
10     FORMAT ('F4')
```

6.1.7 I/O item lists

To support legacy codes, GNU Fortran allows the input item list of the `READ` statement, and the output item lists of the `WRITE` and `PRINT` statements, to start with a comma.

6.1.8 Q exponent-letter

GNU Fortran accepts real literal constants with an exponent-letter of `Q`, for example, `1.23Q45`. The constant is interpreted as a `REAL(16)` entity on targets that support this type. If the target does not support `REAL(16)` but has a `REAL(10)` type, then the real-literal-constant will be interpreted as a `REAL(10)` entity. In the absence of `REAL(16)` and `REAL(10)`, an error will occur.

6.1.9 BOZ literal constants

Besides decimal constants, Fortran also supports binary (`b`), octal (`o`) and hexadecimal (`z`) integer constants. The syntax is: `'prefix quote digits quote'`, where the prefix is either `b`, `o` or `z`, quote is either `'` or `"` and the digits are for binary 0 or 1, for octal between 0 and 7, and for hexadecimal between 0 and `F`. (Example: `b'01011101'`.)

Up to Fortran 95, `BOZ` literals were only allowed to initialize integer variables in `DATA` statements. Since Fortran 2003 `BOZ` literals are also allowed as argument of `REAL`, `DBLE`,

INT and CMPLX; the result is the same as if the integer BOZ literal had been converted by TRANSFER to, respectively, real, double precision, integer or complex. As GNU Fortran extension the intrinsic procedures FLOAT, DFLOAT, COMPLEX and DCMPLX are treated alike.

As an extension, GNU Fortran allows hexadecimal BOZ literal constants to be specified using the X prefix, in addition to the standard Z prefix. The BOZ literal can also be specified by adding a suffix to the string, for example, Z'ABC' and 'ABC'Z are equivalent.

Furthermore, GNU Fortran allows using BOZ literal constants outside DATA statements and the four intrinsic functions allowed by Fortran 2003. In DATA statements, in direct assignments, where the right-hand side only contains a BOZ literal constant, and for old-style initializers of the form integer i /o'0173'/, the constant is transferred as if TRANSFER had been used; for COMPLEX numbers, only the real part is initialized unless CMPLX is used. In all other cases, the BOZ literal constant is converted to an INTEGER value with the largest decimal representation. This value is then converted numerically to the type and kind of the variable in question. (For instance, real :: r = b'0000001' + 1 initializes r with 2.0.) As different compilers implement the extension differently, one should be careful when doing bitwise initialization of non-integer variables.

Note that initializing an INTEGER variable with a statement such as DATA i/Z'FFFFFFFF'/ will give an integer overflow error rather than the desired result of -1 when i is a 32-bit integer on a system that supports 64-bit integers. The '-fno-range-check' option can be used as a workaround for legacy code that initializes integers in this manner.

6.1.10 Real array indices

As an extension, GNU Fortran allows the use of REAL expressions or variables as array indices.

6.1.11 Unary operators

As an extension, GNU Fortran allows unary plus and unary minus operators to appear as the second operand of binary arithmetic operators without the need for parenthesis.

```
X = Y * -Z
```

6.1.12 Implicitly convert LOGICAL and INTEGER values

As an extension for backwards compatibility with other compilers, GNU Fortran allows the implicit conversion of LOGICAL values to INTEGER values and vice versa. When converting from a LOGICAL to an INTEGER, .FALSE. is interpreted as zero, and .TRUE. is interpreted as one. When converting from INTEGER to LOGICAL, the value zero is interpreted as .FALSE. and any nonzero value is interpreted as .TRUE..

```
LOGICAL :: l
l = 1

INTEGER :: i
i = .TRUE.
```

However, there is no implicit conversion of INTEGER values in if-statements, nor of LOGICAL or INTEGER values in I/O operations.

6.1.13 Hollerith constants support

GNU Fortran supports Hollerith constants in assignments, function arguments, and `DATA` and `ASSIGN` statements. A Hollerith constant is written as a string of characters preceded by an integer constant indicating the character count, and the letter `H` or `h`, and stored in bitwise fashion in a numeric (`INTEGER`, `REAL`, or `complex`) or `LOGICAL` variable. The constant will be padded or truncated to fit the size of the variable in which it is stored.

Examples of valid uses of Hollerith constants:

```
complex*16 x(2)
data x /16Habcdefghijklmnop, 16Hqrstuvwxyz012345/
x(1) = 16HABCDEFGHJKLMNOP
call foo (4h abc)
```

Invalid Hollerith constants examples:

```
integer*4 a
a = 8H12345678 ! Valid, but the Hollerith constant will be truncated.
a = 0H          ! At least one character is needed.
```

In general, Hollerith constants were used to provide a rudimentary facility for handling character strings in early Fortran compilers, prior to the introduction of `CHARACTER` variables in Fortran 77; in those cases, the standard-compliant equivalent is to convert the program to use proper character strings. On occasion, there may be a case where the intent is specifically to initialize a numeric variable with a given byte sequence. In these cases, the same result can be obtained by using the `TRANSFER` statement, as in this example.

```
INTEGER(KIND=4) :: a
a = TRANSFER ("abcd", a)      ! equivalent to: a = 4Habcd
```

6.1.14 Cray pointers

Cray pointers are part of a non-standard extension that provides a C-like pointer in Fortran. This is accomplished through a pair of variables: an integer "pointer" that holds a memory address, and a "pointee" that is used to dereference the pointer.

Pointer/pointee pairs are declared in statements of the form:

```
pointer ( <pointer> , <pointee> )
```

or,

```
pointer ( <pointer1> , <pointee1> ), ( <pointer2> , <pointee2> ), ...
```

The pointer is an integer that is intended to hold a memory address. The pointee may be an array or scalar. If an assumed-size array is permitted within the scoping unit, a pointee can be an assumed-size array. That is, the last dimension may be left unspecified by using a `*` in place of a value. A pointee cannot be an assumed shape array. No space is allocated for the pointee.

The pointee may have its type declared before or after the pointer statement, and its array specification (if any) may be declared before, during, or after the pointer statement. The pointer may be declared as an integer prior to the pointer statement. However, some machines have default integer sizes that are different than the size of a pointer, and so the following code is not portable:

```
integer ipt
pointer (ipt, iarr)
```

If a pointer is declared with a kind that is too small, the compiler will issue a warning; the resulting binary will probably not work correctly, because the memory addresses stored

in the pointers may be truncated. It is safer to omit the first line of the above example; if explicit declaration of `ipt`'s type is omitted, then the compiler will ensure that `ipt` is an integer variable large enough to hold a pointer.

Pointer arithmetic is valid with Cray pointers, but it is not the same as C pointer arithmetic. Cray pointers are just ordinary integers, so the user is responsible for determining how many bytes to add to a pointer in order to increment it. Consider the following example:

```
real target(10)
real pointee(10)
pointer (ipt, pointee)
ipt = loc (target)
ipt = ipt + 1
```

The last statement does not set `ipt` to the address of `target(1)`, as it would in C pointer arithmetic. Adding 1 to `ipt` just adds one byte to the address stored in `ipt`.

Any expression involving the pointee will be translated to use the value stored in the pointer as the base address.

To get the address of elements, this extension provides an intrinsic function `LOC()`. The `LOC()` function is equivalent to the `&` operator in C, except the address is cast to an integer type:

```
real ar(10)
pointer(ipt, arpte(10))
real arpte
ipt = loc(ar) ! Makes arpte is an alias for ar
arpte(1) = 1.0 ! Sets ar(1) to 1.0
```

The pointer can also be set by a call to the `MALLOC` intrinsic (see [Section 9.185 \[MALLOC\]](#), page 227).

Cray pointees often are used to alias an existing variable. For example:

```
integer target(10)
integer iarr(10)
pointer (ipt, iarr)
ipt = loc(target)
```

As long as `ipt` remains unchanged, `iarr` is now an alias for `target`. The optimizer, however, will not detect this aliasing, so it is unsafe to use `iarr` and `target` simultaneously. Using a pointee in any way that violates the Fortran aliasing rules or assumptions is illegal. It is the user's responsibility to avoid doing this; the compiler works under the assumption that no such aliasing occurs.

Cray pointers will work correctly when there is no aliasing (i.e., when they are used to access a dynamically allocated block of memory), and also in any routine where a pointee is used, but any variable with which it shares storage is not used. Code that violates these rules may not run as the user intends. This is not a bug in the optimizer; any code that violates the aliasing rules is illegal. (Note that this is not unique to GNU Fortran; any Fortran compiler that supports Cray pointers will “incorrectly” optimize code with illegal aliasing.)

There are a number of restrictions on the attributes that can be applied to Cray pointers and pointees. Pointees may not have the `ALLOCATABLE`, `INTENT`, `OPTIONAL`, `DUMMY`, `TARGET`, `INTRINSIC`, or `POINTER` attributes. Pointers may not have the `DIMENSION`, `POINTER`, `TARGET`, `ALLOCATABLE`, `EXTERNAL`, or `INTRINSIC` attributes, nor may they be function results. Pointees may not occur in more than one pointer statement. A pointee cannot be a pointer. Pointees cannot occur in equivalence, common, or data statements.

A Cray pointer may also point to a function or a subroutine. For example, the following excerpt is valid:

```
implicit none
external sub
pointer (subptr,subpte)
external subpte
subptr = loc(sub)
call subpte()
[...]
subroutine sub
[...]
end subroutine sub
```

A pointer may be modified during the course of a program, and this will change the location to which the pointee refers. However, when pointees are passed as arguments, they are treated as ordinary variables in the invoked function. Subsequent changes to the pointer will not change the base address of the array that was passed.

6.1.15 CONVERT specifier

GNU Fortran allows the conversion of unformatted data between little- and big-endian representation to facilitate moving of data between different systems. The conversion can be indicated with the `CONVERT` specifier on the `OPEN` statement. See [Section 3.10 \[GFORTRAN_CONVERT_UNIT\]](#), page 34, for an alternative way of specifying the data format via an environment variable.

Valid values for `CONVERT` are:

`CONVERT='NATIVE'` Use the native format. This is the default.

`CONVERT='SWAP'` Swap between little- and big-endian.

`CONVERT='LITTLE_ENDIAN'` Use the little-endian representation for unformatted files.

`CONVERT='BIG_ENDIAN'` Use the big-endian representation for unformatted files.

Using the option could look like this:

```
open(file='big.dat',form='unformatted',access='sequential', &
      convert='big_endian')
```

The value of the conversion can be queried by using `INQUIRE(CONVERT=ch)`. The values returned are `'BIG_ENDIAN'` and `'LITTLE_ENDIAN'`.

`CONVERT` works between big- and little-endian for `INTEGER` values of all supported kinds and for `REAL` on IEEE systems of kinds 4 and 8. Conversion between different “extended double” types on different architectures such as m68k and x86_64, which GNU Fortran supports as `REAL(KIND=10)` and `REAL(KIND=16)`, will probably not work.

Note that the values specified via the `GFORTRAN_CONVERT_UNIT` environment variable will override the `CONVERT` specifier in the open statement. This is to give control over data formats to users who do not have the source code of their program available.

Using anything but the native representation for unformatted data carries a significant speed overhead. If speed in this area matters to you, it is best if you use this only for data that needs to be portable.

6.1.16 OpenMP

OpenMP (Open Multi-Processing) is an application programming interface (API) that supports multi-platform shared memory multiprocessing programming in C/C++ and Fortran on many architectures, including Unix and Microsoft Windows platforms. It consists of a set of compiler directives, library routines, and environment variables that influence run-time behavior.

GNU Fortran strives to be compatible to the [OpenMP Application Program Interface v4.5](#).

To enable the processing of the OpenMP directive `!$omp` in free-form source code; the `c$omp`, `*$omp` and `!$omp` directives in fixed form; the `!$` conditional compilation sentinels in free form; and the `c$`, `*$` and `!$` sentinels in fixed form, `gfortran` needs to be invoked with the `'-fopenmp'`. This also arranges for automatic linking of the GNU Offloading and Multi Processing Runtime Library [Section “libgomp” in GNU Offloading and Multi Processing Runtime Library](#).

The OpenMP Fortran runtime library routines are provided both in a form of a Fortran 90 module named `omp_lib` and in a form of a Fortran `include` file named `'omp_lib.h'`.

An example of a parallelized loop taken from Appendix A.1 of the OpenMP Application Program Interface v2.5:

```
SUBROUTINE A1(N, A, B)
  INTEGER I, N
  REAL B(N), A(N)
  !$OMP PARALLEL DO !I is private by default
  DO I=2,N
    B(I) = (A(I) + A(I-1)) / 2.0
  ENDDO
  !$OMP END PARALLEL DO
END SUBROUTINE A1
```

Please note:

- `'-fopenmp'` implies `'-frecursive'`, i.e., all local arrays will be allocated on the stack. When porting existing code to OpenMP, this may lead to surprising results, especially to segmentation faults if the stacksize is limited.
- On glibc-based systems, OpenMP enabled applications cannot be statically linked due to limitations of the underlying pthreads-implementation. It might be possible to get a working solution if `-Wl,--whole-archive -lpthread -Wl,--no-whole-archive` is added to the command line. However, this is not supported by `gcc` and thus not recommended.

6.1.17 OpenACC

OpenACC is an application programming interface (API) that supports offloading of code to accelerator devices. It consists of a set of compiler directives, library routines, and environment variables that influence run-time behavior.

GNU Fortran strives to be compatible to the [OpenACC Application Programming Interface v2.0](#).

To enable the processing of the OpenACC directive `!$acc` in free-form source code; the `c$acc`, `*$acc` and `!$acc` directives in fixed form; the `!$` conditional compilation sentinels in free form; and the `c$`, `*$` and `!$` sentinels in fixed form, `gfortran` needs to be invoked with

the ‘`-fopenacc`’. This also arranges for automatic linking of the GNU Offloading and Multi Processing Runtime Library Section “libgomp” in *GNU Offloading and Multi Processing Runtime Library*.

The OpenACC Fortran runtime library routines are provided both in a form of a Fortran 90 module named `openacc` and in a form of a Fortran `include` file named ‘`openacc_lib.h`’.

Note that this is an experimental feature, incomplete, and subject to change in future versions of GCC. See <https://gcc.gnu.org/wiki/OpenACC> for more information.

6.1.18 Argument list functions %VAL, %REF and %LOC

GNU Fortran supports argument list functions `%VAL`, `%REF` and `%LOC` statements, for backward compatibility with g77. It is recommended that these should be used only for code that is accessing facilities outside of GNU Fortran, such as operating system or windowing facilities. It is best to constrain such uses to isolated portions of a program—portions that deal specifically and exclusively with low-level, system-dependent facilities. Such portions might well provide a portable interface for use by the program as a whole, but are themselves not portable, and should be thoroughly tested each time they are rebuilt using a new compiler or version of a compiler.

`%VAL` passes a scalar argument by value, `%REF` passes it by reference and `%LOC` passes its memory location. Since gfortran already passes scalar arguments by reference, `%REF` is in effect a do-nothing. `%LOC` has the same effect as a Fortran pointer.

An example of passing an argument by value to a C subroutine `foo`:

```
C
C prototype      void foo_ (float x);
C
      external foo
      real*4 x
      x = 3.14159
      call foo (%VAL (x))
      end
```

For details refer to the g77 manual <https://gcc.gnu.org/onlinedocs/gcc-3.4.6/g77/index.html#Top>.

Also, `c_by_val.f` and its partner `c_by_val.c` of the GNU Fortran testsuite are worth a look.

6.1.19 Read/Write after EOF marker

Some legacy codes rely on allowing `READ` or `WRITE` after the EOF file marker in order to find the end of a file. GNU Fortran normally rejects these codes with a run-time error message and suggests the user consider `BACKSPACE` or `REWIND` to properly position the file before the EOF marker. As an extension, the run-time error may be disabled using `-std=legacy`.

6.1.20 STRUCTURE and RECORD

Record structures are a pre-Fortran-90 vendor extension to create user-defined aggregate data types. Support for record structures in GNU Fortran can be enabled with the ‘`-fdec-structure`’ compile flag. If you have a choice, you should instead use Fortran 90’s “derived types”, which have a different syntax.

In many cases, record structures can easily be converted to derived types. To convert, replace `STRUCTURE /structure-name/` by `TYPE type-name`. Additionally, replace `RECORD`

/structure-name/ by `TYPE(type-name)`. Finally, in the component access, replace the period (.) by the percent sign (%).

Here is an example of code using the non portable record structure syntax:

```
! Declaring a structure named 'item' and containing three fields:
! an integer ID, an description string and a floating-point price.
STRUCTURE /item/
  INTEGER id
  CHARACTER(LEN=200) description
  REAL price
END STRUCTURE

! Define two variables, an single record of type 'item'
! named 'pear', and an array of items named 'store_catalog'
RECORD /item/ pear, store_catalog(100)

! We can directly access the fields of both variables
pear.id = 92316
pear.description = "juicy D'Anjou pear"
pear.price = 0.15
store_catalog(7).id = 7831
store_catalog(7).description = "milk bottle"
store_catalog(7).price = 1.2

! We can also manipulate the whole structure
store_catalog(12) = pear
print *, store_catalog(12)
```

This code can easily be rewritten in the Fortran 90 syntax as following:

```
! 'STRUCTURE /name/ ... END STRUCTURE' becomes
! 'TYPE name ... END TYPE'
TYPE item
  INTEGER id
  CHARACTER(LEN=200) description
  REAL price
END TYPE

! 'RECORD /name/ variable' becomes 'TYPE(name) variable'
TYPE(item) pear, store_catalog(100)

! Instead of using a dot (.) to access fields of a record, the
! standard syntax uses a percent sign (%)
pear%id = 92316
pear%description = "juicy D'Anjou pear"
pear%price = 0.15
store_catalog(7)%id = 7831
store_catalog(7)%description = "milk bottle"
store_catalog(7)%price = 1.2
```

```
! Assignments of a whole variable do not change
store_catalog(12) = pear
print *, store_catalog(12)
```

GNU Fortran implements STRUCTURES like derived types with the following rules and exceptions:

- Structures act like derived types with the `SEQUENCE` attribute. Otherwise they may contain no specifiers.
- Structures may contain a special field with the name `%FILL`. This will create an anonymous component which cannot be accessed but occupies space just as if a component of the same type was declared in its place, useful for alignment purposes. As an example, the following structure will consist of at least sixteen bytes:

```
structure /padded/
  character(4) start
  character(8) %FILL
  character(4) end
end structure
```

- Structures may share names with other symbols. For example, the following is invalid for derived types, but valid for structures:

```
structure /header/
  ! ...
end structure
record /header/ header
```

- Structure types may be declared nested within another parent structure. The syntax is:

```
structure /type-name/
  ...
  structure [ /<type-name>/ ] <field-list>
  ...
```

The type name may be omitted, in which case the structure type itself is anonymous, and other structures of the same type cannot be instantiated. The following shows some examples:

```
structure /appointment/
  ! nested structure definition: app_time is an array of two 'time'
  structure /time/ app_time (2)
    integer(1) hour, minute
  end structure
  character(10) memo
end structure
```

```
! The 'time' structure is still usable
record /time/ now
now = time(5, 30)
```

```
...
```

```
structure /appointment/
```

```

! anonymous nested structure definition
structure start, end
  integer(1) hour, minute
end structure
character(10) memo
end structure

```

- Structures may contain UNION blocks. For more detail see the section on [Section 6.1.21 \[UNION and MAP\]](#), page 62.
- Structures support old-style initialization of components, like those described in [Section 6.1.2 \[Old-style variable initialization\]](#), page 51. For array initializers, an initializer may contain a repeat specification of the form `<literal-integer> * <constant-initializer>`. The value of the integer indicates the number of times to repeat the constant initializer when expanding the initializer list.

6.1.21 UNION and MAP

Unions are an old vendor extension which were commonly used with the non-standard [Section 6.1.20 \[STRUCTURE and RECORD\]](#), page 59 extensions. Use of UNION and MAP is automatically enabled with ‘`-fdec-structure`’.

A UNION declaration occurs within a structure; within the definition of each union is a number of MAP blocks. Each MAP shares storage with its sibling maps (in the same union), and the size of the union is the size of the largest map within it, just as with unions in C. The major difference is that component references do not indicate which union or map the component is in (the compiler gets to figure that out).

Here is a small example:

```

structure /myunion/
union
  map
    character(2) w0, w1, w2
  end map
  map
    character(6) long
  end map
end union
end structure


record /myunion/ rec
! After this assignment...
rec.long = 'hello!'

! The following is true:
! rec.w0 === 'he'
! rec.w1 === 'll'
! rec.w2 === 'o!'

```

The two maps share memory, and the size of the union is ultimately six bytes:

0	1	2	3	4	5	6	Byte offset



Following is an example mirroring the layout of an Intel x86_64 register:

```

structure /reg/
  union ! U0
    ! rax
    map
      character(16) rx
    end map
    map
      character(8) rh      ! rah
      union ! U1
        map
          character(8) rl   ! ral
        end map
        map
          character(8) ex   ! eax
        end map
        map
          character(4) eh   ! eah
          union ! U2
            map
              character(4) el ! eal
            end map
            map
              character(4) x  ! ax
            end map
            map
              character(2) h  ! ah
              character(2) l  ! al
            end map
          end union
        end map
      end union
    end map
  end union
end structure
record /reg/ a

! After this assignment...
a.rx      =      'AAAAAAAA.BBB.C.D'

! The following is true:
a.rx === 'AAAAAAAA.BBB.C.D'

```

```

a.rh === 'AAAAAAAA'
a.rl === '.BBB.C.D'
a.ex === '.BBB.C.D'
a.eh === '.BBB'
a.el === '.C.D'
a.x  === '.C.D'
a.h  === '.C'
a.l  === '.D'

```

6.1.22 Type variants for integer intrinsics

Similar to the D/C prefixes to real functions to specify the input/output types, GNU Fortran offers B/I/J/K prefixes to integer functions for compatibility with DEC programs. The types implied by each are:

```

B - INTEGER(kind=1)
I - INTEGER(kind=2)
J - INTEGER(kind=4)
K - INTEGER(kind=8)

```

GNU Fortran supports these with the flag ‘-fdec-intrinsic-ints’. Intrinsics for which prefixed versions are available and in what form are noted in [Chapter 9 \[Intrinsic Procedures\]](#), page 111. The complete list of supported intrinsics is here:

Intrinsic	B	I	J	K
Section 9.3 [ABS] , page 112	BABS	IIABS	JIABS	KIABS
Section 9.51 [BTEST] , page 143	BBTEST	BITEST	BJTEST	BKTEST
Section 9.137 [IAND] , page 199	BIAND	IIAND	JIAND	KIAND
Section 9.140 [IBCLR] , page 201	BBCLR	IIBCLR	JIBCLR	KIBCLR
Section 9.141 [IBITS] , page 202	BBITS	IIBITS	JIBITS	KIBITS
Section 9.142 [IBSET] , page 202	BBSET	IIBSET	JIBSET	KIBSET
Section 9.145 [IEOR] , page 204	BIEOR	IIEOR	JIEOR	KIEOR
Section 9.152 [IOR] , page 208	BIOR	IIOR	JIOR	KIOR

Section 9.159 [ISHFT], page 212	BSHFT	IISHFT	JISHFT	KISHFT
Section 9.160 [ISHFTC], page 213	BSHFTC	IISHFTC	JISHFTC	KISHFTC
Section 9.201 [MOD], page 236	BMOD	IMOD	JMOD	KMOD
Section 9.209 [NOT], page 241	BNOT	INOT	JNOT	KNOT
Section 9.229 [REAL], page 252	--	FLOATI	FLOATJ	FLOATK

6.1.23 AUTOMATIC and STATIC attributes

With ‘-fdec-static’ GNU Fortran supports the DEC extended attributes **STATIC** and **AUTOMATIC** to provide explicit specification of entity storage. These follow the syntax of the Fortran standard **SAVE** attribute.

STATIC is exactly equivalent to **SAVE**, and specifies that an entity should be allocated in static memory. As an example, **STATIC** local variables will retain their values across multiple calls to a function.

Entities marked **AUTOMATIC** will be stack automatic whenever possible. **AUTOMATIC** is the default for local variables smaller than ‘-fmax-stack-var-size’, unless ‘-fno-automatic’ is given. This attribute overrides ‘-fno-automatic’, ‘-fmax-stack-var-size’, and blanket **SAVE** statements.

Examples:

```

subroutine f
  integer, automatic :: i  ! automatic variable
  integer x, y             ! static variables
  save
  ...
endsubroutine

subroutine f
  integer a, b, c, x, y, z
  static :: x
  save y
  automatic z, c
  ! a, b, c, and z are automatic
  ! x and y are static
endsubroutine

! Compiled with -fno-automatic
subroutine f
  integer a, b, c, d
  automatic :: a
  ! a is automatic; b, c, and d are static

```

endsubroutine

6.1.24 Extended math intrinsics

GNU Fortran supports an extended list of mathematical intrinsics with the compile flag ‘-fdec-math’ for compatibility with legacy code. These intrinsics are described fully in [Chapter 9 \[Intrinsic Procedures\]](#), [page 111](#) where it is noted that they are extensions and should be avoided whenever possible.

Specifically, ‘-fdec-math’ enables the [Section 9.76 \[COTAN\]](#), [page 160](#) intrinsic, and trigonometric intrinsics which accept or produce values in degrees instead of radians. Here is a summary of the new intrinsics:

Radians

[Section 9.6 \[ACOS\]](#), [page 114](#)
[Section 9.19 \[ASIN\]](#), [page 123](#)
[Section 9.23 \[ATAN\]](#), [page 126](#)
[Section 9.25 \[ATAN2\]](#), [page 127](#)
[Section 9.73 \[COS\]](#), [page 158](#)
[Section 9.76 \[COTAN\]](#), [page 160*](#)
[Section 9.250 \[SIN\]](#), [page 264](#)
[Section 9.266 \[TAN\]](#), [page 274](#)

* Enabled with ‘-fdec-math’.

Degrees

[Section 9.7 \[ACOSD\]](#), [page 115*](#)
[Section 9.20 \[ASIND\]](#), [page 123*](#)
[Section 9.24 \[ATAND\]](#), [page 126*](#)
[Section 9.26 \[ATAN2D\]](#), [page 128*](#)
[Section 9.74 \[COSD\]](#), [page 159*](#)
[Section 9.77 \[COTAND\]](#), [page 160*](#)
[Section 9.251 \[SIND\]](#), [page 265*](#)
[Section 9.267 \[TAND\]](#), [page 275*](#)

For advanced users, it may be important to know the implementation of these functions. They are simply wrappers around the standard radian functions, which have more accurate builtin versions. These functions convert their arguments (or results) to degrees (or radians) by taking the value modulus 360 (or 2π) and then multiplying it by a constant radian-to-degree (or degree-to-radian) factor, as appropriate. The factor is computed at compile-time as $180/\pi$ (or $\pi/180$).

6.1.25 Form feed as whitespace

Historically, legacy compilers allowed insertion of form feed characters (‘\f’, ASCII 0xC) at the beginning of lines for formatted output to line printers, though the Fortran standard does not mention this. GNU Fortran supports the interpretation of form feed characters in source as whitespace for compatibility.

6.1.26 TYPE as an alias for PRINT

For compatibility, GNU Fortran will interpret TYPE statements as PRINT statements with the flag ‘-fdec’. With this flag asserted, the following two examples are equivalent:

```
TYPE *, 'hello world'
PRINT *, 'hello world'
```

6.1.27 %LOC as an rvalue

Normally %LOC is allowed only in parameter lists. However the intrinsic function LOC does the same thing, and is usable as the right-hand-side of assignments. For compatibility, GNU Fortran supports the use of %LOC as an alias for the builtin LOC with ‘-std=legacy’. With this feature enabled the following two examples are equivalent:

```

integer :: i, l
l = %loc(i)
call sub(l)
integer :: i
call sub(%loc(i))

```

6.1.28 .XOR. operator

GNU Fortran supports `.XOR.` as a logical operator with `-std=legacy` for compatibility with legacy code. `.XOR.` is equivalent to `.NEQV..` That is, the output is true if and only if the inputs differ.

6.1.29 Bitwise logical operators

With `-fdec`, GNU Fortran relaxes the type constraints on logical operators to allow integer operands, and performs the corresponding bitwise operation instead. This flag is for compatibility only, and should be avoided in new code. Consider:

```

INTEGER :: i, j
i = z'33'
j = z'cc'
print *, i .AND. j

```

In this example, compiled with `-fdec`, GNU Fortran will replace the `.AND.` operation with a call to the intrinsic [Section 9.137 \[IAND\], page 199](#) function, yielding the bitwise-and of `i` and `j`.

Note that this conversion will occur if at least one operand is of integral type. As a result, a logical operand will be converted to an integer when the other operand is an integer in a logical operation. In this case, `.TRUE.` is converted to 1 and `.FALSE.` to 0.

Here is the mapping of logical operator to bitwise intrinsic used with `-fdec`:

Operator	Intrinsic	Bitwise operation
<code>.NOT.</code>	Section 9.209 [NOT], page 241	complement
<code>.AND.</code>	Section 9.137 [IAND], page 199	intersection
<code>.OR.</code>	Section 9.152 [IOR], page 208	union
<code>.NEQV.</code>	Section 9.145 [IEOR], page 204	exclusive or
<code>.EQV.</code>	Section 9.209 [NOT], page 241 (Section 9.145 [IEOR], page 204)	complement of exclusive or

6.1.30 Extended I/O specifiers

GNU Fortran supports the additional legacy I/O specifiers `CARRIAGECONTROL`, `READONLY`, and `SHARE` with the compile flag `-fdec`, for compatibility.

CARRIAGECONTROL

The `CARRIAGECONTROL` specifier allows a user to control line termination settings between output records for an I/O unit. The specifier has no meaning for

readonly files. When `CARRIAGECONTROL` is specified upon opening a unit for formatted writing, the exact `CARRIAGECONTROL` setting determines what characters to write between output records. The syntax is:

```
OPEN(..., CARRIAGECONTROL=cc)
```

Where `cc` is a character expression that evaluates to one of the following values:

'LIST'	One line feed between records (default)
'FORTRAN'	Legacy interpretation of the first character (see below)
'NONE'	No separator between records

With `CARRIAGECONTROL='FORTRAN'`, when a record is written, the first character of the input record is not written, and instead determines the output record separator as follows:

Leading character	Meaning	Output separating character(s)
'+'	Overprinting	Carriage return only
'-'	New line	Line feed and carriage return
'0'	Skip line	Two line feeds and carriage return
'1'	New page	Form feed and carriage return
'\$'	Prompting	Line feed (no carriage return)
CHAR(0)	Overprinting advance)	(no None

READONLY The `READONLY` specifier may be given upon opening a unit, and is equivalent to specifying `ACTION='READ'`, except that the file may not be deleted on close (i.e. `CLOSE` with `STATUS="DELETE"`). The syntax is:

```
OPEN(..., READONLY)
```

SHARE The `SHARE` specifier allows system-level locking on a unit upon opening it for controlled access from multiple processes/threads. The `SHARE` specifier has several forms:

```
OPEN(..., SHARE=sh)
OPEN(..., SHARED)
OPEN(..., NOSHARED)
```

Where `sh` in the first form is a character expression that evaluates to a value as seen in the table below. The latter two forms are aliases for particular values of `sh`:

Explicit form	Short form	Meaning
SHARE='DENYRW'	NOSHARED	Exclusive (write) lock
SHARE='DENYNONE'	SHARED	Shared (read) lock

In general only one process may hold an exclusive (write) lock for a given file at a time, whereas many processes may hold shared (read) locks for the same file.

The behavior of locking may vary with your operating system. On POSIX systems, locking is implemented with `fcntl`. Consult your corresponding operating system's manual pages for further details. Locking via `SHARE=` is not supported on other systems.

6.1.31 Legacy PARAMETER statements

For compatibility, GNU Fortran supports legacy PARAMETER statements without parentheses with ‘-std=legacy’. A warning is emitted if used with ‘-std=gnu’, and an error is acknowledged with a real Fortran standard flag (‘-std=f95’, etc...). These statements take the following form:

```
implicit real (E)
parameter e = 2.718282
real c
parameter c = 3.0e8
```

6.1.32 Default exponents

For compatibility, GNU Fortran supports a default exponent of zero in real constants with ‘-fdec’. For example, 9e would be interpreted as 9e0, rather than an error.

6.2 Extensions not implemented in GNU Fortran

The long history of the Fortran language, its wide use and broad userbase, the large number of different compiler vendors and the lack of some features crucial to users in the first standards have lead to the existence of a number of important extensions to the language. While some of the most useful or popular extensions are supported by the GNU Fortran compiler, not all existing extensions are supported. This section aims at listing these extensions and offering advice on how best make code that uses them running with the GNU Fortran compiler.

6.2.1 ENCODE and DECODE statements

GNU Fortran does not support the ENCODE and DECODE statements. These statements are best replaced by READ and WRITE statements involving internal files (CHARACTER variables and arrays), which have been part of the Fortran standard since Fortran 77. For example, replace a code fragment like

```
INTEGER*1 LINE(80)
REAL A, B, C
c ... Code that sets LINE
DECODE (80, 9000, LINE) A, B, C
9000 FORMAT (1X, 3(F10.5))
```

with the following:

```
CHARACTER(LEN=80) LINE
REAL A, B, C
c ... Code that sets LINE
READ (UNIT=LINE, FMT=9000) A, B, C
9000 FORMAT (1X, 3(F10.5))
```

Similarly, replace a code fragment like

```
INTEGER*1 LINE(80)
REAL A, B, C
c ... Code that sets A, B and C
ENCODE (80, 9000, LINE) A, B, C
9000 FORMAT (1X, 'OUTPUT IS ', 3(F10.5))
```

with the following:

```
CHARACTER(LEN=80) LINE
REAL A, B, C
```

```

c      ... Code that sets A, B and C
      WRITE (UNIT=LINE, FMT=9000) A, B, C
9000 FORMAT (1X, 'OUTPUT IS ', 3(F10.5))

```

6.2.2 Variable FORMAT expressions

A variable FORMAT expression is format statement which includes angle brackets enclosing a Fortran expression: `FORMAT(I<N>)`. GNU Fortran does not support this legacy extension. The effect of variable format expressions can be reproduced by using the more powerful (and standard) combination of internal output and string formats. For example, replace a code fragment like this:

```

      WRITE(6,20) INT1
20    FORMAT(I<N+1>)

```

with the following:

```

c      Variable declaration
      CHARACTER(LEN=20) FMT
c
c      Other code here...
c
      WRITE(FMT,'("I", I0, ")")') N+1
      WRITE(6,FMT) INT1

```

or with:

```

c      Variable declaration
      CHARACTER(LEN=20) FMT
c
c      Other code here...
c
      WRITE(FMT,*) N+1
      WRITE(6,"(I" // ADJUSTL(FMT) // ")") INT1

```

6.2.3 Alternate complex function syntax

Some Fortran compilers, including `g77`, let the user declare complex functions with the syntax `COMPLEX FUNCTION name*16()`, as well as `COMPLEX*16 FUNCTION name()`. Both are non-standard, legacy extensions. `gfortran` accepts the latter form, which is more common, but not the former.

6.2.4 Volatile COMMON blocks

Some Fortran compilers, including `g77`, let the user declare `COMMON` with the `VOLATILE` attribute. This is invalid standard Fortran syntax and is not supported by `gfortran`. Note that `gfortran` accepts `VOLATILE` variables in `COMMON` blocks since revision 4.3.

6.2.5 OPEN(... NAME=)

Some Fortran compilers, including `g77`, let the user declare `OPEN( ... NAME=)`. This is invalid standard Fortran syntax and is not supported by `gfortran`. `OPEN( ... NAME=)` should be replaced with `OPEN( ... FILE=)`.

6.2.6 Q edit descriptor

Some Fortran compilers provide the `Q` edit descriptor, which transfers the number of characters left within an input record into an integer variable.

A direct replacement of the `Q` edit descriptor is not available in `gfortran`. How to replicate its functionality using standard-conforming code depends on what the intent of the original code is.

Options to replace `Q` may be to read the whole line into a character variable and then counting the number of non-blank characters left using `LEN_TRIM`. Another method may be to use formatted stream, read the data up to the position where the `Q` descriptor occurred, use `INQUIRE` to get the file position, count the characters up to the next `NEW_LINE` and then start reading from the position marked previously.

7 Mixed-Language Programming

This chapter is about mixed-language interoperability, but also applies if one links Fortran code compiled by different compilers. In most cases, use of the C Binding features of the Fortran 2003 standard is sufficient, and their use is highly recommended.

7.1 Interoperability with C

Since Fortran 2003 (ISO/IEC 1539-1:2004(E)) there is a standardized way to generate procedure and derived-type declarations and global variables which are interoperable with C (ISO/IEC 9899:1999). The `bind(C)` attribute has been added to inform the compiler that a symbol shall be interoperable with C; also, some constraints are added. Note, however, that not all C features have a Fortran equivalent or vice versa. For instance, neither C's unsigned integers nor C's functions with variable number of arguments have an equivalent in Fortran.

Note that array dimensions are reversely ordered in C and that arrays in C always start with index 0 while in Fortran they start by default with 1. Thus, an array declaration `A(n,m)` in Fortran matches `A[m][n]` in C and accessing the element `A(i,j)` matches `A[j-1][i-1]`. The element following `A(i,j)` (C: `A[j-1][i-1]`; assuming $i < n$) in memory is `A(i+1,j)` (C: `A[j-1][i]`).

7.1.1 Intrinsic Types

In order to ensure that exactly the same variable type and kind is used in C and Fortran, the named constants shall be used which are defined in the `ISO_C_BINDING` intrinsic module. That module contains named constants for kind parameters and character named constants for the escape sequences in C. For a list of the constants, see [Section 10.2 \[ISO_C_BINDING\]](#), [page 289](#).

For logical types, please note that the Fortran standard only guarantees interoperability between C99's `_Bool` and Fortran's `C_Bool`-kind logicals and C99 defines that `true` has the value 1 and `false` the value 0. Using any other integer value with GNU Fortran's `LOGICAL` (with any kind parameter) gives an undefined result. (Passing other integer values than 0 and 1 to GCC's `_Bool` is also undefined, unless the integer is explicitly or implicitly casted to `_Bool`.)

7.1.2 Derived Types and `struct`

For compatibility of derived types with `struct`, one needs to use the `BIND(C)` attribute in the type declaration. For instance, the following type declaration

```
USE ISO_C_BINDING
TYPE, BIND(C) :: myType
  INTEGER(C_INT) :: i1, i2
  INTEGER(C_SIGNED_CHAR) :: i3
  REAL(C_DOUBLE) :: d1
  COMPLEX(C_FLOAT_COMPLEX) :: c1
  CHARACTER(KIND=C_CHAR) :: str(5)
END TYPE
```

matches the following `struct` declaration in C

```
struct {
  int i1, i2;
```

```

/* Note: "char" might be signed or unsigned. */
signed char i3;
double d1;
float _Complex c1;
char str[5];
} myType;

```

Derived types with the C binding attribute shall not have the `sequence` attribute, type parameters, the `extends` attribute, nor type-bound procedures. Every component must be of interoperable type and kind and may not have the `pointer` or `allocatable` attribute. The names of the components are irrelevant for interoperability.

As there exist no direct Fortran equivalents, neither unions nor structs with bit field or variable-length array members are interoperable.

7.1.3 Interoperable Global Variables

Variables can be made accessible from C using the C binding attribute, optionally together with specifying a binding name. Those variables have to be declared in the declaration part of a `MODULE`, be of interoperable type, and have neither the `pointer` nor the `allocatable` attribute.

```

MODULE m
  USE myType_module
  USE ISO_C_BINDING
  integer(C_INT), bind(C, name="_MyProject_flags") :: global_flag
  type(myType), bind(C) :: tp
END MODULE

```

Here, `_MyProject_flags` is the case-sensitive name of the variable as seen from C programs while `global_flag` is the case-insensitive name as seen from Fortran. If no binding name is specified, as for `tp`, the C binding name is the (lowercase) Fortran binding name. If a binding name is specified, only a single variable may be after the double colon. Note of warning: You cannot use a global variable to access `errno` of the C library as the C standard allows it to be a macro. Use the `IERRNO` intrinsic (GNU extension) instead.

7.1.4 Interoperable Subroutines and Functions

Subroutines and functions have to have the `BIND(C)` attribute to be compatible with C. The dummy argument declaration is relatively straightforward. However, one needs to be careful because C uses call-by-value by default while Fortran behaves usually similar to call-by-reference. Furthermore, strings and pointers are handled differently. Note that in Fortran 2003 and 2008 only explicit size and assumed-size arrays are supported but not assumed-shape or deferred-shape (i.e. `allocatable` or `pointer`) arrays. However, those are allowed since the Technical Specification 29113, see [Section 7.1.6 \[Further Interoperability of Fortran with C\]](#), page 78

To pass a variable by value, use the `VALUE` attribute. Thus, the following C prototype

```
int func(int i, int *j)
```

matches the Fortran declaration

```

integer(c_int) function func(i,j)
  use iso_c_binding, only: c_int
  integer(c_int), VALUE :: i
  integer(c_int) :: j

```

Note that pointer arguments also frequently need the `VALUE` attribute, see [Section 7.1.5 \[Working with Pointers\]](#), page 76.

Strings are handled quite differently in C and Fortran. In C a string is a NUL-terminated array of characters while in Fortran each string has a length associated with it and is thus not terminated (by e.g. NUL). For example, if one wants to use the following C function,

```
#include <stdio.h>
void print_C(char *string) /* equivalent: char string[] */
{
    printf("%s\n", string);
}
```

to print “Hello World” from Fortran, one can call it using

```
use iso_c_binding, only: C_CHAR, C_NULL_CHAR
interface
    subroutine print_c(string) bind(C, name="print_C")
        use iso_c_binding, only: c_char
        character(kind=c_char) :: string(*)
    end subroutine print_c
end interface
call print_c(C_CHAR_"Hello World"//C_NULL_CHAR)
```

As the example shows, one needs to ensure that the string is NUL terminated. Additionally, the dummy argument *string* of `print_C` is a length-one assumed-size array; using `character(len=*)` is not allowed. The example above uses `c_char_"Hello World"` to ensure the string literal has the right type; typically the default character kind and `c_char` are the same and thus `"Hello World"` is equivalent. However, the standard does not guarantee this.

The use of strings is now further illustrated using the C library function `strncpy`, whose prototype is

```
char *strncpy(char *restrict s1, const char *restrict s2, size_t n);
```

The function `strncpy` copies at most *n* characters from string *s2* to *s1* and returns *s1*. In the following example, we ignore the return value:

```
use iso_c_binding
implicit none
character(len=30) :: str, str2
interface
    ! Ignore the return value of strncpy -> subroutine
    ! "restrict" is always assumed if we do not pass a pointer
    subroutine strncpy(dest, src, n) bind(C)
        import
        character(kind=c_char), intent(out) :: dest(*)
        character(kind=c_char), intent(in)  :: src(*)
        integer(c_size_t), value, intent(in) :: n
    end subroutine strncpy
end interface
str = repeat('X',30) ! Initialize whole string with 'X'
call strncpy(str, c_char_"Hello World"//C_NULL_CHAR, &
             len(c_char_"Hello World",kind=c_size_t))
print '(a)', str ! prints: "Hello WorldXXXXXXXXXXXXXXXXXXXX"
end
```

The intrinsic procedures are described in [Chapter 9 \[Intrinsic Procedures\]](#), page 111.

7.1.5 Working with Pointers

C pointers are represented in Fortran via the special opaque derived type `type(c_ptr)` (with private components). Thus one needs to use intrinsic conversion procedures to convert from or to C pointers.

For some applications, using an assumed type (`TYPE(*)`) can be an alternative to a C pointer; see [Section 7.1.6 \[Further Interoperability of Fortran with C\]](#), page 78.

For example,

```
use iso_c_binding
type(c_ptr) :: cptr1, cptr2
integer, target :: array(7), scalar
integer, pointer :: pa(:), ps
cptr1 = c_loc(array(1)) ! The programmer needs to ensure that the
                        ! array is contiguous if required by the C
                        ! procedure
cptr2 = c_loc(scalar)
call c_f_pointer(cptr2, ps)
call c_f_pointer(cptr2, pa, shape=[7])
```

When converting C to Fortran arrays, the one-dimensional `SHAPE` argument has to be passed.

If a pointer is a dummy-argument of an interoperable procedure, it usually has to be declared using the `VALUE` attribute. `void*` matches `TYPE(C_PTR)`, `VALUE`, while `TYPE(C_PTR)` alone matches `void**`.

Procedure pointers are handled analogously to pointers; the C type is `TYPE(C_FUNPTR)` and the intrinsic conversion procedures are `C_F_PROCPTR` and `C_FUNLOC`.

Let us consider two examples of actually passing a procedure pointer from C to Fortran and vice versa. Note that these examples are also very similar to passing ordinary pointers between both languages. First, consider this code in C:

```
/* Procedure implemented in Fortran. */
void get_values (void (*)(double));

/* Call-back routine we want called from Fortran. */
void
print_it (double x)
{
    printf ("Number is %f.\n", x);
}

/* Call Fortran routine and pass call-back to it. */
void
foobar ()
{
    get_values (&print_it);
}
```

A matching implementation for `get_values` in Fortran, that correctly receives the procedure pointer from C and is able to call it, is given in the following `MODULE`:

```
MODULE m
  IMPLICIT NONE

  ! Define interface of call-back routine.
  ABSTRACT INTERFACE
    SUBROUTINE callback (x)
```

```

        USE, INTRINSIC :: ISO_C_BINDING
        REAL(KIND=C_DOUBLE), INTENT(IN), VALUE :: x
    END SUBROUTINE callback
END INTERFACE

CONTAINS

    ! Define C-bound procedure.
    SUBROUTINE get_values (cproc) BIND(C)
        USE, INTRINSIC :: ISO_C_BINDING
        TYPE(C_FUNPTR), INTENT(IN), VALUE :: cproc

        PROCEDURE(callback), POINTER :: proc

        ! Convert C to Fortran procedure pointer.
        CALL C_F_PROCPTR (cproc, proc)

        ! Call it.
        CALL proc (1.0_C_DOUBLE)
        CALL proc (-42.0_C_DOUBLE)
        CALL proc (18.12_C_DOUBLE)
    END SUBROUTINE get_values

END MODULE m

```

Next, we want to call a C routine that expects a procedure pointer argument and pass it a Fortran procedure (which clearly must be interoperable!). Again, the C function may be:

```

int
call_it (int (*func)(int), int arg)
{
    return func (arg);
}

```

It can be used as in the following Fortran code:

```

MODULE m
    USE, INTRINSIC :: ISO_C_BINDING
    IMPLICIT NONE

    ! Define interface of C function.
    INTERFACE
        INTEGER(KIND=C_INT) FUNCTION call_it (func, arg) BIND(C)
            USE, INTRINSIC :: ISO_C_BINDING
            TYPE(C_FUNPTR), INTENT(IN), VALUE :: func
            INTEGER(KIND=C_INT), INTENT(IN), VALUE :: arg
        END FUNCTION call_it
    END INTERFACE

CONTAINS

    ! Define procedure passed to C function.
    ! It must be interoperable!
    INTEGER(KIND=C_INT) FUNCTION double_it (arg) BIND(C)
        INTEGER(KIND=C_INT), INTENT(IN), VALUE :: arg
        double_it = arg + arg
    END FUNCTION double_it

    ! Call C function.

```

```

SUBROUTINE foobar ()
  TYPE(C_FUNPTR) :: cproc
  INTEGER(KIND=C_INT) :: i

  ! Get C procedure pointer.
  cproc = C_FUNLOC (double_it)

  ! Use it.
  DO i = 1_C_INT, 10_C_INT
    PRINT *, call_it (cproc, i)
  END DO
END SUBROUTINE foobar

END MODULE m

```

7.1.6 Further Interoperability of Fortran with C

The Technical Specification ISO/IEC TS 29113:2012 on further interoperability of Fortran with C extends the interoperability support of Fortran 2003 and Fortran 2008. Besides removing some restrictions and constraints, it adds assumed-type (`TYPE(*)`) and assumed-rank (`dimension`) variables and allows for interoperability of assumed-shape, assumed-rank and deferred-shape arrays, including allocatables and pointers.

Note: Currently, GNU Fortran does not use internally the array descriptor (dope vector) as specified in the Technical Specification, but uses an array descriptor with different fields. Assumed type and assumed rank formal arguments are converted in the library to the specified form. The ISO_Fortran_binding API functions (also Fortran 2018 18.4) are implemented in `libgfortran`. Alternatively, the Chasm Language Interoperability Tools, <http://chasm-interop.sourceforge.net/>, provide an interface to GNU Fortran's array descriptor.

The Technical Specification adds the following new features, which are supported by GNU Fortran:

- The `ASYNCHRONOUS` attribute has been clarified and extended to allow its use with asynchronous communication in user-provided libraries such as in implementations of the Message Passing Interface specification.
- Many constraints have been relaxed, in particular for the `C_LOC` and `C_F_POINTER` intrinsics.
- The `OPTIONAL` attribute is now allowed for dummy arguments; an absent argument matches a `NULL` pointer.
- Assumed types (`TYPE(*)`) have been added, which may only be used for dummy arguments. They are unlimited polymorphic but contrary to `CLASS(*)` they do not contain any type information, similar to C's `void *` pointers. Expressions of any type and kind can be passed; thus, it can be used as replacement for `TYPE(C_PTR)`, avoiding the use of `C_LOC` in the caller.

Note, however, that `TYPE(*)` only accepts scalar arguments, unless the `DIMENSION` is explicitly specified. As `DIMENSION(*)` only supports array (including array elements) but no scalars, it is not a full replacement for `C_LOC`. On the other hand, assumed-type assumed-rank dummy arguments (`TYPE(*)`, `DIMENSION(..)`) allow for both scalars and arrays, but require special code on the callee side to handle the array descriptor.

- Assumed-rank arrays (`DIMENSION(...)`) as dummy argument allow that scalars and arrays of any rank can be passed as actual argument. As the Technical Specification does not provide for direct means to operate with them, they have to be used either from the C side or be converted using `C_LOC` and `C_F_POINTER` to scalars or arrays of a specific rank. The rank can be determined using the `RANK` intrinsic.

Currently unimplemented:

- GNU Fortran always uses an array descriptor, which does not match the one of the Technical Specification. The `ISO_Fortran_binding.h` header file and the C functions it specifies are not available.
- Using assumed-shape, assumed-rank and deferred-shape arrays in `BIND(C)` procedures is not fully supported. In particular, C interoperable strings of other length than one are not supported as this requires the new array descriptor.

7.2 GNU Fortran Compiler Directives

7.2.1 ATTRIBUTES directive

The Fortran standard describes how a conforming program shall behave; however, the exact implementation is not standardized. In order to allow the user to choose specific implementation details, compiler directives can be used to set attributes of variables and procedures which are not part of the standard. Whether a given attribute is supported and its exact effects depend on both the operating system and on the processor; see [Section “C Extensions” in *Using the GNU Compiler Collection \(GCC\)*](#) for details.

For procedures and procedure pointers, the following attributes can be used to change the calling convention:

- `CDECL` – standard C calling convention
- `STDCALL` – convention where the called procedure pops the stack
- `FASTCALL` – part of the arguments are passed via registers instead using the stack

Besides changing the calling convention, the attributes also influence the decoration of the symbol name, e.g., by a leading underscore or by a trailing at-sign followed by the number of bytes on the stack. When assigning a procedure to a procedure pointer, both should use the same calling convention.

On some systems, procedures and global variables (module variables and `COMMON` blocks) need special handling to be accessible when they are in a shared library. The following attributes are available:

- `DLEXPOR` – provide a global pointer to a pointer in the DLL
- `DLLIMPORT` – reference the function or variable using a global pointer

For dummy arguments, the `NO_ARG_CHECK` attribute can be used; in other compilers, it is also known as `IGNORE_TKR`. For dummy arguments with this attribute actual arguments of any type and kind (similar to `TYPE(*)`), scalars and arrays of any rank (no equivalent in Fortran standard) are accepted. As with `TYPE(*)`, the argument is unlimited polymorphic and no type information is available. Additionally, the argument may only be passed to dummy arguments with the `NO_ARG_CHECK` attribute and as argument to the `PRESENT` intrinsic function and to `C_LOC` of the `ISO_C_BINDING` module.

Variables with `NO_ARG_CHECK` attribute shall be of assumed-type (`TYPE(*)`; recommended) or of type `INTEGER`, `LOGICAL`, `REAL` or `COMPLEX`. They shall not have the `ALLOCATE`, `CODIMENSION`, `INTENT(OUT)`, `POINTER` or `VALUE` attribute; furthermore, they shall be either scalar or of assumed-size (`dimension(*)`). As `TYPE(*)`, the `NO_ARG_CHECK` attribute requires an explicit interface.

- `NO_ARG_CHECK` – disable the type, kind and rank checking

The attributes are specified using the syntax

```
!GCC$ ATTRIBUTES attribute-list :: variable-list
```

where in free-form source code only whitespace is allowed before `!GCC$` and in fixed-form source code `!GCC$`, `cGCC$` or `*GCC$` shall start in the first column.

For procedures, the compiler directives shall be placed into the body of the procedure; for variables and procedure pointers, they shall be in the same declaration part as the variable or procedure pointer.

7.2.2 UNROLL directive

The syntax of the directive is

```
!GCC$ unroll N
```

You can use this directive to control how many times a loop should be unrolled. It must be placed immediately before a `DO` loop and applies only to the loop that follows. `N` is an integer constant specifying the unrolling factor. The values of 0 and 1 block any unrolling of the loop.

7.2.3 BUILTIN directive

The syntax of the directive is

```
!GCC$ BUILTIN (B) attributes simd FLAGS IF('target')
```

You can use this directive to define which middle-end built-ins provide vector implementations. `B` is name of the middle-end built-in. `FLAGS` are optional and must be either `"(inbranch)"` or `"(notinbranch)"`. `IF` statement is optional and is used to filter multilib ABIs for the built-in that should be vectorized. Example usage:

```
!GCC$ builtin (sinf) attributes simd (notinbranch) if('x86_64')
```

The purpose of the directive is to provide an API among the GCC compiler and the GNU C Library which would define vector implementations of math routines.

7.3 Non-Fortran Main Program

Even if you are doing mixed-language programming, it is very likely that you do not need to know or use the information in this section. Since it is about the internal structure of GNU Fortran, it may also change in GCC minor releases.

When you compile a `PROGRAM` with GNU Fortran, a function with the name `main` (in the symbol table of the object file) is generated, which initializes the `libgfortran` library and then calls the actual program which uses the name `MAIN__`, for historic reasons. If you link GNU Fortran compiled procedures to, e.g., a C or C++ program or to a Fortran program compiled by a different compiler, the `libgfortran` library is not initialized and thus a few intrinsic procedures do not work properly, e.g. those for obtaining the command-line arguments.

Therefore, if your `PROGRAM` is not compiled with GNU Fortran and the GNU Fortran compiled procedures require intrinsics relying on the library initialization, you need to initialize the library yourself. Using the default options, `gfortran` calls `_gfortran_set_args` and `_gfortran_set_options`. The initialization of the former is needed if the called procedures access the command line (and for backtracing); the latter sets some flags based on the standard chosen or to enable backtracing. In typical programs, it is not necessary to call any initialization function.

If your `PROGRAM` is compiled with GNU Fortran, you shall not call any of the following functions. The `libgfortran` initialization functions are shown in C syntax but using C bindings they are also accessible from Fortran.

7.3.1 `_gfortran_set_args` — Save command-line arguments

Description:

`_gfortran_set_args` saves the command-line arguments; this initialization is required if any of the command-line intrinsics is called. Additionally, it shall be called if backtracing is enabled (see `_gfortran_set_options`).

Syntax: `void _gfortran_set_args (int argc, char *argv[])`

Arguments:

<code>argc</code>	number of command line argument strings
<code>argv</code>	the command-line argument strings; <code>argv[0]</code> is the pathname of the executable itself.

Example:

```
int main (int argc, char *argv[])
{
    /* Initialize libgfortran. */
    _gfortran_set_args (argc, argv);
    return 0;
}
```

7.3.2 `_gfortran_set_options` — Set library option flags

Description:

`_gfortran_set_options` sets several flags related to the Fortran standard to be used, whether backtracing should be enabled and whether range checks should be performed. The syntax allows for upward compatibility since the number of passed flags is specified; for non-passed flags, the default value is used. See also see [Section 2.9 \[Code Gen Options\], page 23](#). Please note that not all flags are actually used.

Syntax: `void _gfortran_set_options (int num, int options[])`

Arguments:

<code>num</code>	number of options passed
<code>argv</code>	The list of flag values

option flag list:

<i>option</i> [0]	Allowed standard; can give run-time errors if e.g. an input-output edit descriptor is invalid in a given standard. Possible values are (bitwise or-ed) GFC_STD_F77 (1), GFC_STD_F95_OBS (2), GFC_STD_F95_DEL (4), GFC_STD_F95 (8), GFC_STD_F2003 (16), GFC_STD_GNU (32), GFC_STD_LEGACY (64), GFC_STD_F2008 (128), GFC_STD_F2008_OBS (256), GFC_STD_F2008_TS (512), GFC_STD_F2018 (1024), GFC_STD_F2018_OBS (2048), and GFC_STD_F2018_DEL (4096). Default: GFC_STD_F95_OBS GFC_STD_F95_DEL GFC_STD_F95 GFC_STD_F2003 GFC_STD_F2008 GFC_STD_F2008_TS GFC_STD_F2008_OBS GFC_STD_F77 GFC_STD_F2018 GFC_STD_F2018_OBS GFC_STD_F2018_DEL GFC_STD_GNU GFC_STD_LEGACY.
<i>option</i> [1]	Standard-warning flag; prints a warning to standard error. Default: GFC_STD_F95_DEL GFC_STD_LEGACY.
<i>option</i> [2]	If non zero, enable pedantic checking. Default: off.
<i>option</i> [3]	Unused.
<i>option</i> [4]	If non zero, enable backtracing on run-time errors. Default: off. (Default in the compiler: on.) Note: Installs a signal handler and requires command-line initialization using <code>_gfortran_set_args</code> .
<i>option</i> [5]	If non zero, supports signed zeros. Default: enabled.
<i>option</i> [6]	Enables run-time checking. Possible values are (bitwise or-ed): GFC_RTCHECK_BOUNDS (1), GFC_RTCHECK_ARRAY_TEMPS (2), GFC_RTCHECK_RECURSION (4), GFC_RTCHECK_DO (16), GFC_RTCHECK_POINTER (32). Default: disabled.
<i>option</i> [7]	Unused.
<i>option</i> [8]	Show a warning when invoking STOP and ERROR STOP if a floating-point exception occurred. Possible values are (bitwise or-ed) GFC_FPE_INVALID (1), GFC_FPE_DENORMAL (2), GFC_FPE_ZERO (4), GFC_FPE_OVERFLOW (8), GFC_FPE_UNDERFLOW (16), GFC_FPE_INEXACT (32). Default: None (0). (Default in the compiler: GFC_FPE_INVALID GFC_FPE_DENORMAL GFC_FPE_ZERO GFC_FPE_OVERFLOW GFC_FPE_UNDERFLOW.)

Example:

```
/* Use gfortran 4.9 default options. */
static int options[] = {68, 511, 0, 0, 1, 1, 0, 0, 31};
_gfortran_set_options (9, &options);
```

7.3.3 `_gfortran_set_convert` — Set endian conversion

Description:

`_gfortran_set_convert` set the representation of data for unformatted files.

Syntax: `void _gfortran_set_convert (int conv)`

Arguments:

<i>conv</i>	Endian conversion, GFC_CONVERT_NATIVE GFC_CONVERT_SWAP (1), (2), GFC_CONVERT_LITTLE (3).	possible values: (0, default), GFC_CONVERT_BIG
-------------	---	--

Example:

```
int main (int argc, char *argv[])
{
    /* Initialize libgfortran. */
    _gfortran_set_args (argc, argv);
    _gfortran_set_convert (1);
    return 0;
}
```

7.3.4 `_gfortran_set_record_marker` — Set length of record markers

Description:

`_gfortran_set_record_marker` sets the length of record markers for unformatted files.

Syntax: `void _gfortran_set_record_marker (int val)`

Arguments:

<i>val</i>	Length of the record marker; valid values are 4 and 8. Default is 4.
------------	--

Example:

```
int main (int argc, char *argv[])
{
    /* Initialize libgfortran. */
    _gfortran_set_args (argc, argv);
    _gfortran_set_record_marker (8);
    return 0;
}
```

7.3.5 `_gfortran_set_fpe` — Enable floating point exception traps

Description:

`_gfortran_set_fpe` enables floating point exception traps for the specified exceptions. On most systems, this will result in a SIGFPE signal being sent and the program being aborted.

Syntax: `void _gfortran_set_fpe (int val)`

Arguments:

<i>option[0]</i>	IEEE exceptions. Possible values are (bitwise or-ed) zero (0, default) no trapping, GFC_FPE_INVALID (1), GFC_FPE_DENORMAL (2), GFC_FPE_ZERO (4), GFC_FPE_OVERFLOW (8), GFC_FPE_UNDERFLOW (16), and GFC_FPE_INEXACT (32).
------------------	--

Example:

```
int main (int argc, char *argv[])
{
    /* Initialize libgfortran. */
    ...
}
```

```

_gfortran_set_args (argc, argv);
/* FPE for invalid operations such as SQRT(-1.0). */
_gfortran_set_fpe (1);
return 0;
}

```

7.3.6 `_gfortran_set_max_subrecord_length` — Set subrecord length

Description:

`_gfortran_set_max_subrecord_length` set the maximum length for a subrecord. This option only makes sense for testing and debugging of unformatted I/O.

Syntax: `void _gfortran_set_max_subrecord_length (int val)`

Arguments:

`val` the maximum length for a subrecord; the maximum permitted value is 2147483639, which is also the default.

Example:

```

int main (int argc, char *argv[])
{
    /* Initialize libgfortran. */
    _gfortran_set_args (argc, argv);
    _gfortran_set_max_subrecord_length (8);
    return 0;
}

```

7.4 Naming and argument-passing conventions

This section gives an overview about the naming convention of procedures and global variables and about the argument passing conventions used by GNU Fortran. If a C binding has been specified, the naming convention and some of the argument-passing conventions change. If possible, mixed-language and mixed-compiler projects should use the better defined C binding for interoperability. See [Section 7.1 \[Interoperability with C\]](#), page 73.

7.4.1 Naming conventions

According the Fortran standard, valid Fortran names consist of a letter between A to Z, a to z, digits 0, 1 to 9 and underscores (`_`) with the restriction that names may only start with a letter. As vendor extension, the dollar sign (`$`) is additionally permitted with the option `'-fdollar-ok'`, but not as first character and only if the target system supports it.

By default, the procedure name is the lower-cased Fortran name with an appended underscore (`_`); using `'-fno-underscoring'` no underscore is appended while `-fsecond-underscore` appends two underscores. Depending on the target system and the calling convention, the procedure might be additionally dressed; for instance, on 32bit Windows with `stdcall`, an at-sign `@` followed by an integer number is appended. For the changing the calling convention, see [Section 7.2 \[GNU Fortran Compiler Directives\]](#), page 79.

For common blocks, the same convention is used, i.e. by default an underscore is appended to the lower-cased Fortran name. Blank commons have the name `__BLNK__`.

For procedures and variables declared in the specification space of a module, the name is formed by `__`, followed by the lower-cased module name, `_MOD_`, and the lower-cased Fortran name. Note that no underscore is appended.

7.4.2 Argument passing conventions

Subroutines do not return a value (matching C99's `void`) while functions either return a value as specified in the platform ABI or the result variable is passed as hidden argument to the function and no result is returned. A hidden result variable is used when the result variable is an array or of type `CHARACTER`.

Arguments are passed according to the platform ABI. In particular, complex arguments might not be compatible to a struct with two real components for the real and imaginary part. The argument passing matches the one of C99's `_Complex`. Functions with scalar complex result variables return their value and do not use a by-reference argument. Note that with the `'-ff2c'` option, the argument passing is modified and no longer completely matches the platform ABI. Some other Fortran compilers use `f2c` semantic by default; this might cause problems with interoperability.

GNU Fortran passes most arguments by reference, i.e. by passing a pointer to the data. Note that the compiler might use a temporary variable into which the actual argument has been copied, if required semantically (copy-in/copy-out).

For arguments with `ALLOCATABLE` and `POINTER` attribute (including procedure pointers), a pointer to the pointer is passed such that the pointer address can be modified in the procedure.

For dummy arguments with the `VALUE` attribute: Scalar arguments of the type `INTEGER`, `LOGICAL`, `REAL` and `COMPLEX` are passed by value according to the platform ABI. (As vendor extension and not recommended, using `%VAL()` in the call to a procedure has the same effect.) For `TYPE(C_PTR)` and procedure pointers, the pointer itself is passed such that it can be modified without affecting the caller.

For Boolean (`LOGICAL`) arguments, please note that GCC expects only the integer value 0 and 1. If a GNU Fortran `LOGICAL` variable contains another integer value, the result is undefined. As some other Fortran compilers use `-1` for `.TRUE.`, extra care has to be taken – such as passing the value as `INTEGER`. (The same value restriction also applies to other front ends of GCC, e.g. to GCC's C99 compiler for `_Bool` or GCC's Ada compiler for `Boolean`.)

For arguments of `CHARACTER` type, the character length is passed as a hidden argument at the end of the argument list. For deferred-length strings, the value is passed by reference, otherwise by value. The character length has the C type `size_t` (or `INTEGER(kind=C_SIZE_T)` in Fortran). Note that this is different to older versions of the GNU Fortran compiler, where the type of the hidden character length argument was a C `int`. In order to retain compatibility with older versions, one can e.g. for the following Fortran procedure

```
subroutine fstrlen (s, a)
  character(len=*) :: s
  integer :: a
  print*, len(s)
end subroutine fstrlen
```

define the corresponding C prototype as follows:

```
#if __GNUC__ > 7
typedef size_t fortran_charlen_t;
#else
typedef int fortran_charlen_t;
#endif

void fstrlen_ (char*, int*, fortran_charlen_t);
```

In order to avoid such compiler-specific details, for new code it is instead recommended to use the `ISO_C_BINDING` feature.

Note with C binding, `CHARACTER(len=1)` result variables are returned according to the platform ABI and no hidden length argument is used for dummy arguments; with `VALUE`, those variables are passed by value.

For `OPTIONAL` dummy arguments, an absent argument is denoted by a `NULL` pointer, except for scalar dummy arguments of type `INTEGER`, `LOGICAL`, `REAL` and `COMPLEX` which have the `VALUE` attribute. For those, a hidden Boolean argument (`logical(kind=C_bool),value`) is used to indicate whether the argument is present.

Arguments which are assumed-shape, assumed-rank or deferred-rank arrays or, with `'-fcoarray=lib'`, allocatable scalar coarrays use an array descriptor. All other arrays pass the address of the first element of the array. With `'-fcoarray=lib'`, the token and the offset belonging to nonallocatable coarrays dummy arguments are passed as hidden argument along the character length hidden arguments. The token is an opaque pointer identifying the coarray and the offset is a passed-by-value integer of kind `C_PTRDIFF_T`, denoting the byte offset between the base address of the coarray and the passed scalar or first element of the passed array.

The arguments are passed in the following order

- Result variable, when the function result is passed by reference
- Character length of the function result, if it is of type `CHARACTER` and no C binding is used
- The arguments in the order in which they appear in the Fortran declaration
- The present status for optional arguments with value attribute, which are internally passed by value
- The character length and/or coarray token and offset for the first argument which is a `CHARACTER` or a nonallocatable coarray dummy argument, followed by the hidden arguments of the next dummy argument of such a type

8 Coarray Programming

8.1 Type and enum ABI Documentation

8.1.1 `caf_token_t`

Typedef of type `void *` on the compiler side. Can be any data type on the library side.

8.1.2 `caf_register_t`

Indicates which kind of coarray variable should be registered.

```
typedef enum caf_register_t {
    CAF_REGTYPE_COARRAY_STATIC,
    CAF_REGTYPE_COARRAY_ALLOC,
    CAF_REGTYPE_LOCK_STATIC,
    CAF_REGTYPE_LOCK_ALLOC,
    CAF_REGTYPE_CRITICAL,
    CAF_REGTYPE_EVENT_STATIC,
    CAF_REGTYPE_EVENT_ALLOC,
    CAF_REGTYPE_COARRAY_ALLOC_REGISTER_ONLY,
    CAF_REGTYPE_COARRAY_ALLOC_ALLOCATE_ONLY
}
caf_register_t;
```

The values `CAF_REGTYPE_COARRAY_ALLOC_REGISTER_ONLY` and `CAF_REGTYPE_COARRAY_ALLOC_ALLOCATE_ONLY` are for allocatable components in derived type coarrays only. The first one sets up the token without allocating memory for allocatable component. The latter one only allocates the memory for an allocatable component in a derived type coarray. The token needs to be setup previously by the `REGISTER_ONLY`. This allows to have allocatable components un-allocated on some images. The status whether an allocatable component is allocated on a remote image can be queried by `_caf_is_present` which used internally by the `ALLOCATED` intrinsic.

8.1.3 `caf_deregister_t`

```
typedef enum caf_deregister_t {
    CAF_DEREGTYPE_COARRAY_DEREGISTER,
    CAF_DEREGTYPE_COARRAY_DEALLOCATE_ONLY
}
caf_deregister_t;
```

Allows to specify the type of deregistration of a coarray object. The `CAF_DEREGTYPE_COARRAY_DEALLOCATE_ONLY` flag is only allowed for allocatable components in derived type coarrays.

8.1.4 `caf_reference_t`

The structure used for implementing arbitrary reference chains. A `CAF_REFERENCE_T` allows to specify a component reference or any kind of array reference of any rank supported by gfortran. For array references all kinds as known by the compiler/Fortran standard are supported indicated by a `MODE`.

```

typedef enum caf_ref_type_t {
    /* Reference a component of a derived type, either regular one or an
       allocatable or pointer type. For regular ones idx in caf_reference_t is
       set to -1. */
    CAF_REF_COMPONENT,
    /* Reference an allocatable array. */
    CAF_REF_ARRAY,
    /* Reference a non-allocatable/non-pointer array. I.e., the coarray object
       has no array descriptor associated and the addressing is done
       completely using the ref. */
    CAF_REF_STATIC_ARRAY
} caf_ref_type_t;

typedef enum caf_array_ref_t {
    /* No array ref. This terminates the array ref. */
    CAF_ARR_REF_NONE = 0,
    /* Reference array elements given by a vector. Only for this mode
       caf_reference_t.u.a.dim[i].v is valid. */
    CAF_ARR_REF_VECTOR,
    /* A full array ref (:). */
    CAF_ARR_REF_FULL,
    /* Reference a range on elements given by start, end and stride. */
    CAF_ARR_REF_RANGE,
    /* Only a single item is referenced given in the start member. */
    CAF_ARR_REF_SINGLE,
    /* An array ref of the kind (i:), where i is an arbitrary valid index in the
       array. The index i is given in the start member. */
    CAF_ARR_REF_OPEN_END,
    /* An array ref of the kind (:i), where the lower bound of the array ref
       is given by the remote side. The index i is given in the end member. */
    CAF_ARR_REF_OPEN_START
} caf_array_ref_t;

/* References to remote components of a derived type. */
typedef struct caf_reference_t {
    /* A pointer to the next ref or NULL. */
    struct caf_reference_t *next;
    /* The type of the reference. */
    /* caf_ref_type_t, replaced by int to allow specification in fortran FE. */
    int type;
    /* The size of an item referenced in bytes. I.e. in an array ref this is
       the factor to advance the array pointer with to get to the next item.
       For component refs this gives just the size of the element referenced. */
    size_t item_size;
    union {
        struct {
            /* The offset (in bytes) of the component in the derived type.
               Unused for allocatable or pointer components. */

```



```

    ptrdiff_t offset;
    /* The offset (in bytes) to the caf_token associated with this
       component.  NULL, when not allocatable/pointer ref.  */
    ptrdiff_t caf_token_offset;
} c;
struct {
    /* The mode of the array ref.  See CAF_ARR_REF_*.  */
    /* caf_array_ref_t, replaced by unsigend char to allow specification in
       fortran FE.  */
    unsigned char mode[GFC_MAX_DIMENSIONS];
    /* The type of a static array.  Unset for array's with descriptors.  */
    int static_array_type;
    /* Subscript refs (s) or vector refs (v).  */
    union {
        struct {
            /* The start and end boundary of the ref and the stride.  */
            index_type start, end, stride;
        } s;
        struct {
            /* nvec entries of kind giving the elements to reference.  */
            void *vector;
            /* The number of entries in vector.  */
            size_t nvec;
            /* The integer kind used for the elements in vector.  */
            int kind;
        } v;
    } dim[GFC_MAX_DIMENSIONS];
} a;
} u;
} caf_reference_t;

```

The references make up a single linked list of reference operations. The `NEXT` member links to the next reference or `NULL` to indicate the end of the chain. Component and array refs can be arbitrarily mixed as long as they comply to the Fortran standard.

NOTES The member `STATIC_ARRAY_TYPE` is used only when the `TYPE` is `CAF_REF_STATIC_ARRAY`. The member gives the type of the data referenced. Because no array descriptor is available for a descriptor-less array and type conversion still needs to take place the type is transported here.

At the moment `CAF_ARR_REF_VECTOR` is not implemented in the front end for descriptor-less arrays. The library `caf_single` has untested support for it.

8.1.5 `caf_team_t`

Opaque pointer to represent a team-handle. This type is a stand-in for the future implementation of teams. It is about to change without further notice.

8.2 Function ABI Documentation

8.2.1 `_gfortran_caf_init` — Initialization function

Description:

This function is called at startup of the program before the Fortran main program, if the latter has been compiled with ‘`-fcoarray=lib`’. It takes as arguments the command-line arguments of the program. It is permitted to pass two NULL pointers as argument; if non-NULL, the library is permitted to modify the arguments.

Syntax: `void _gfortran_caf_init (int *argc, char ***argv)`

Arguments:

<code>argc</code>	intent(inout) An integer pointer with the number of arguments passed to the program or NULL.
<code>argv</code>	intent(inout) A pointer to an array of strings with the command-line arguments or NULL.

NOTES The function is modelled after the initialization function of the Message Passing Interface (MPI) specification. Due to the way coarray registration works, it might not be the first call to the library. If the main program is not written in Fortran and only a library uses coarrays, it can happen that this function is never called. Therefore, it is recommended that the library does not rely on the passed arguments and whether the call has been done.

8.2.2 `_gfortran_caf_finish` — Finalization function

Description:

This function is called at the end of the Fortran main program, if it has been compiled with the ‘`-fcoarray=lib`’ option.

Syntax: `void _gfortran_caf_finish (void)`

NOTES For non-Fortran programs, it is recommended to call the function at the end of the main program. To ensure that the shutdown is also performed for programs where this function is not explicitly invoked, for instance non-Fortran programs or calls to the system’s `exit()` function, the library can use a destructor function. Note that programs can also be terminated using the `STOP` and `ERROR STOP` statements; those use different library calls.

8.2.3 `_gfortran_caf_this_image` — Querying the image number

Description:

This function returns the current image number, which is a positive number.

Syntax: `int _gfortran_caf_this_image (int distance)`

Arguments:

<code>distance</code>	As specified for the <code>this_image</code> intrinsic in TS18508. Shall be a non-negative number.
-----------------------	--

NOTES If the Fortran intrinsic `this_image` is invoked without an argument, which is the only permitted form in Fortran 2008, GCC passes 0 as first argument.

8.2.4 `_gfortran_caf_num_images` — Querying the maximal number of images

Description:

This function returns the number of images in the current team, if *distance* is 0 or the number of images in the parent team at the specified distance. If *failed* is -1, the function returns the number of all images at the specified distance; if it is 0, the function returns the number of nonfailed images, and if it is 1, it returns the number of failed images.

Syntax: `int _gfortran_caf_num_images(int distance, int failed)`

Arguments:

distance the distance from this image to the ancestor. Shall be positive.
failed shall be -1, 0, or 1

NOTES This function follows TS18508. If the `num_image` intrinsic has no arguments, then the compiler passes `distance=0` and `failed=-1` to the function.

8.2.5 `_gfortran_caf_image_status` — Query the status of an image

Description:

Get the status of the image given by the id *image* of the team given by *team*. Valid results are zero, for image is ok, `STAT_STOPPED_IMAGE` from the `ISO_FORTRAN_ENV` module to indicate that the image has been stopped and `STAT_FAILED_IMAGE` also from `ISO_FORTRAN_ENV` to indicate that the image has executed a `FAIL IMAGE` statement.

Syntax: `int _gfortran_caf_image_status (int image, caf_team_t * team)`

Arguments:

image the positive scalar id of the image in the current TEAM.
team optional; team on the which the inquiry is to be performed.

NOTES This function follows TS18508. Because team-functionality is not yet implemented a null-pointer is passed for the *team* argument at the moment.

8.2.6 `_gfortran_caf_failed_images` — Get an array of the indexes of the failed images

Description:

Get an array of image indexes in the current *team* that have failed. The array is sorted ascendingly. When *team* is not provided the current team is to be used. When *kind* is provided then the resulting array is of that integer kind else it is of default integer kind. The returns an unallocated size zero array when no images have failed.

Syntax: `int _gfortran_caf_failed_images (caf_team_t * team, int * kind)`

Arguments:

team optional; team on the which the inquiry is to be performed.
image optional; the kind of the resulting integer array.

NOTES This function follows TS18508. Because team-functionality is not yet implemented a null-pointer is passed for the *team* argument at the moment.

8.2.7 `_gfortran_caf_stopped_images` — Get an array of the indexes of the stopped images

Description:

Get an array of image indexes in the current *team* that have stopped. The array is sorted ascendingly. When *team* is not provided the current team is to be used. When *kind* is provided then the resulting array is of that integer kind else it is of default integer kind. The returns an unallocated size zero array when no images have failed.

Syntax: `int _gfortran_caf_stopped_images (caf_team_t * team, int * kind)`

Arguments:

team optional; team on the which the inquiry is to be performed.
image optional; the kind of the resulting integer array.

NOTES This function follows TS18508. Because team-functionality is not yet implemented a null-pointer is passed for the *team* argument at the moment.

8.2.8 `_gfortran_caf_register` — Registering coarrays

Description:

Registers memory for a coarray and creates a token to identify the coarray. The routine is called for both coarrays with **SAVE** attribute and using an explicit **ALLOCATE** statement. If an error occurs and *STAT* is a **NULL** pointer, the function shall abort with printing an error message and starting the error termination. If no error occurs and *STAT* is present, it shall be set to zero. Otherwise, it shall be set to a positive value and, if not-**NULL**, *ERRMSG* shall be set to a string describing the failure. The routine shall register the memory provided in the **DATA**-component of the array descriptor *DESC*, when that component is non-**NULL**, else it shall allocate sufficient memory and provide a pointer to it in the **DATA**-component of *DESC*. The array descriptor has rank zero, when a scalar object is to be registered and the array descriptor may be invalid after the call to `_gfortran_caf_register`. When an array is to be allocated the descriptor persists.

For `CAF_REGTYPE_COARRAY_STATIC` and `CAF_REGTYPE_COARRAY_ALLOC`, the passed size is the byte size requested. For `CAF_REGTYPE_LOCK_STATIC`, `CAF_REGTYPE_LOCK_ALLOC` and `CAF_REGTYPE_CRITICAL` it is the array size or one for a scalar.

When `CAF_REGTYPE_COARRAY_ALLOC_REGISTER_ONLY` is used, then only a token for an allocatable or pointer component is created. The *SIZE* parameter is not used then. On the contrary when `CAF_REGTYPE_COARRAY_ALLOC_ALLOCATE_ONLY` is specified, then the *token* needs to be registered by a previous call with regtype `CAF_REGTYPE_COARRAY_ALLOC_REGISTER_ONLY` and either the memory specified in the *DESC*'s data-ptr is registered or allocate when the data-ptr is **NULL**.

Syntax: `void caf_register (size_t size, caf_register_t type, caf_token_t *token, gfc_descriptor_t *desc, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>size</i>	For normal coarrays, the byte size of the coarray to be allocated; for lock types and event types, the number of elements.
<i>type</i>	one of the <code>caf_register_t</code> types.
<i>token</i>	intent(out) An opaque pointer identifying the coarray.
<i>desc</i>	intent(inout) The (pseudo) array descriptor.
<i>stat</i>	intent(out) For allocatable coarrays, stores the <code>STAT=</code> ; may be NULL
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL
<i>errmsg_len</i>	the buffer size of <code>errmsg</code> .

NOTES

Nonallocatable coarrays have to be registered prior use from remote images. In order to guarantee this, they have to be registered before the main program. This can be achieved by creating constructor functions. That is what GCC does such that also for nonallocatable coarrays the memory is allocated and no static memory is used. The token permits to identify the coarray; to the processor, the token is a nonaliasing pointer. The library can, for instance, store the base address of the coarray in the token, some handle or a more complicated struct. The library may also store the array descriptor *DESC* when its rank is non-zero. For lock types, the value shall only be used for checking the allocation status. Note that for critical blocks, the locking is only required on one image; in the locking statement, the processor shall always pass an image index of one for critical-block lock variables (`CAF_REGTYPE_CRITICAL`). For lock types and critical-block variables, the initial value shall be unlocked (or, respectively, not in critical section) such as the value false; for event types, the initial state should be no event, e.g. zero.

8.2.9 `_gfortran_caf_deregister` — Deregistering coarrays*Description:*

Called to free or deregister the memory of a coarray; the processor calls this function for automatic and explicit deallocation. In case of an error, this function shall fail with an error message, unless the *STAT* variable is not null. The library is only expected to free memory it allocated itself during a call to `_gfortran_caf_register`.

Syntax: `void caf_deregister (caf_token_t *token, caf_deregister_t type, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>token</i>	the token to free.
<i>type</i>	the type of action to take for the coarray. A <code>CAF_DEREGTYPE_COARRAY_DEALLOCATE_ONLY</code> is allowed only for allocatable or pointer components of derived type coarrays. The action only deallocates the local memory without deleting the token.
<i>stat</i>	intent(out) Stores the <code>STAT=</code> ; may be NULL
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL

`errmsg_len` the buffer size of `errmsg`.

NOTES For nonallocatable coarrays this function is never called. If a cleanup is required, it has to be handled via the `finish`, `stop` and `error stop` functions, and via destructors.

8.2.10 `_gfortran_caf_is_present` — Query whether an allocatable or pointer component in a derived type coarray is allocated

Description:

Used to query the coarray library whether an allocatable component in a derived type coarray is allocated on a remote image.

Syntax: `void _gfortran_caf_is_present (caf_token_t token, int image_index, gfc_reference_t *ref)`

Arguments:

<code>token</code>	An opaque pointer identifying the coarray.
<code>image_index</code>	The ID of the remote image; must be a positive number.
<code>ref</code>	A chain of references to address the allocatable or pointer component in the derived type coarray. The object reference needs to be a scalar or a full array reference, respectively.

8.2.11 `_gfortran_caf_send` — Sending data from a local image to a remote image

Description:

Called to send a scalar, an array section or a whole array from a local to a remote image identified by the `image_index`.

Syntax: `void _gfortran_caf_send (caf_token_t token, size_t offset, int image_index, gfc_descriptor_t *dest, caf_vector_t *dst_vector, gfc_descriptor_t *src, int dst_kind, int src_kind, bool may_require_tmp, int *stat)`

Arguments:

<code>token</code>	intent(in) An opaque pointer identifying the coarray.
<code>offset</code>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the coarray.
<code>image_index</code>	intent(in) The ID of the remote image; must be a positive number.
<code>dest</code>	intent(in) Array descriptor for the remote image for the bounds and the size. The <code>base_addr</code> shall not be accessed.
<code>dst_vector</code>	intent(in) If not NULL, it contains the vector subscript of the destination array; the values are relative to the dimension triplet of the <code>dest</code> argument.
<code>src</code>	intent(in) Array descriptor of the local array to be transferred to the remote image
<code>dst_kind</code>	intent(in) Kind of the destination argument
<code>src_kind</code>	intent(in) Kind of the source argument

may_require_tmp intent(in) The variable is **false** when it is known at compile time that the *dest* and *src* either cannot overlap or overlap (fully or partially) such that walking *src* and *dest* in element wise element order (honoring the stride value) will not lead to wrong results. Otherwise, the value is **true**.

stat intent(out) when non-NULL give the result of the operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.

NOTES It is permitted to have *image_index* equal the current image; the memory of the send-to and the send-from might (partially) overlap in that case. The implementation has to take care that it handles this case, e.g. using `memmove` which handles (partially) overlapping memory. If *may_require_tmp* is true, the library might additionally create a temporary variable, unless additional checks show that this is not required (e.g. because walking backward is possible or because both arrays are contiguous and `memmove` takes care of overlap issues).

Note that the assignment of a scalar to an array is permitted. In addition, the library has to handle numeric-type conversion and for strings, padding and different character kinds.

8.2.12 `_gfortran_caf_get` — Getting data from a remote image

Description:

Called to get an array section or a whole array from a remote, image identified by the *image_index*.

Syntax:

```
void _gfortran_caf_get (caf_token_t token, size_t offset, int
image_index, gfc_descriptor_t *src, caf_vector_t *src_vector,
gfc_descriptor_t *dest, int src_kind, int dst_kind, bool may_
require_tmp, int *stat)
```

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>offset</i>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the coarray.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number.
<i>dest</i>	intent(out) Array descriptor of the local array to store the data retrieved from the remote image
<i>src</i>	intent(in) Array descriptor for the remote image for the bounds and the size. The <code>base_addr</code> shall not be accessed.
<i>src_vector</i>	intent(in) If not NULL, it contains the vector subscript of the source array; the values are relative to the dimension triplet of the <i>src</i> argument.
<i>dst_kind</i>	intent(in) Kind of the destination argument
<i>src_kind</i>	intent(in) Kind of the source argument

<i>may_require_tmp</i>	intent(in) The variable is false when it is known at compile time that the <i>dest</i> and <i>src</i> either cannot overlap or overlap (fully or partially) such that walking <i>src</i> and <i>dest</i> in element wise element order (honoring the stride value) will not lead to wrong results. Otherwise, the value is true .
<i>stat</i>	intent(out) When non-NULL give the result of the operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.

NOTES It is permitted to have *image_index* equal the current image; the memory of the send-to and the send-from might (partially) overlap in that case. The implementation has to take care that it handles this case, e.g. using **memmove** which handles (partially) overlapping memory. If *may_require_tmp* is true, the library might additionally create a temporary variable, unless additional checks show that this is not required (e.g. because walking backward is possible or because both arrays are contiguous and **memmove** takes care of overlap issues). Note that the library has to handle numeric-type conversion and for strings, padding and different character kinds.

8.2.13 `_gfortran_caf_sendget` — Sending data between remote images

Description:

Called to send a scalar, an array section or a whole array from a remote image identified by the *src_image_index* to a remote image identified by the *dst_image_index*.

Syntax:

```
void _gfortran_caf_sendget (caf_token_t dst_token, size_t dst_offset, int dst_image_index, gfc_descriptor_t *dest, caf_vector_t *dst_vector, caf_token_t src_token, size_t src_offset, int src_image_index, gfc_descriptor_t *src, caf_vector_t *src_vector, int dst_kind, int src_kind, bool may_require_tmp, int *stat)
```

Arguments:

<i>dst_token</i>	intent(in) An opaque pointer identifying the destination coarray.
<i>dst_offset</i>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the destination coarray.
<i>dst_image_index</i>	intent(in) The ID of the destination remote image; must be a positive number.
<i>dest</i>	intent(in) Array descriptor for the destination remote image for the bounds and the size. The base_addr shall not be accessed.
<i>dst_vector</i>	intent(int) If not NULL, it contains the vector subscript of the destination array; the values are relative to the dimension triplet of the <i>dest</i> argument.
<i>src_token</i>	intent(in) An opaque pointer identifying the source coarray.

<i>src_offset</i>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the source coarray.
<i>src_image_index</i>	intent(in) The ID of the source remote image; must be a positive number.
<i>src</i>	intent(in) Array descriptor of the local array to be transferred to the remote image.
<i>src_vector</i>	intent(in) Array descriptor of the local array to be transferred to the remote image
<i>dst_kind</i>	intent(in) Kind of the destination argument
<i>src_kind</i>	intent(in) Kind of the source argument
<i>may_require_tmp</i>	intent(in) The variable is false when it is known at compile time that the <i>dest</i> and <i>src</i> either cannot overlap or overlap (fully or partially) such that walking <i>src</i> and <i>dest</i> in element wise element order (honoring the stride value) will not lead to wrong results. Otherwise, the value is true .
<i>stat</i>	intent(out) when non-NULL give the result of the operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.

NOTES It is permitted to have the same image index for both *src_image_index* and *dst_image_index*; the memory of the send-to and the send-from might (partially) overlap in that case. The implementation has to take care that it handles this case, e.g. using `memmove` which handles (partially) overlapping memory. If *may_require_tmp* is true, the library might additionally create a temporary variable, unless additional checks show that this is not required (e.g. because walking backward is possible or because both arrays are contiguous and `memmove` takes care of overlap issues).

Note that the assignment of a scalar to an array is permitted. In addition, the library has to handle numeric-type conversion and for strings, padding and different character kinds.

8.2.14 `_gfortran_caf_send_by_ref` — Sending data from a local image to a remote image with enhanced referencing options

Description:

Called to send a scalar, an array section or a whole array from a local to a remote image identified by the *image_index*.

Syntax:

```
void _gfortran_caf_send_by_ref (caf_token_t token, int image_index,
gfc_descriptor_t *src, caf_reference_t *refs, int dst_kind, int
src_kind, bool may_require_tmp, bool dst_reallocatable, int *stat,
int dst_type)
```

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number.

<i>src</i>	intent(in) Array descriptor of the local array to be transferred to the remote image
<i>refs</i>	intent(in) The references on the remote array to store the data given by <i>src</i> . Guaranteed to have at least one entry.
<i>dst_kind</i>	intent(in) Kind of the destination argument
<i>src_kind</i>	intent(in) Kind of the source argument
<i>may_require_tmp</i>	intent(in) The variable is false when it is known at compile time that the <i>dest</i> and <i>src</i> either cannot overlap or overlap (fully or partially) such that walking <i>src</i> and <i>dest</i> in element wise element order (honoring the stride value) will not lead to wrong results. Otherwise, the value is true .
<i>dst_reallocatable</i>	intent(in) Set when the destination is of allocatable or pointer type and the <i>refs</i> will allow reallocation, i.e., the ref is a full array or component ref.
<i>stat</i>	intent(out) When non-NULL give the result of the operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.
<i>dst_type</i>	intent(in) Give the type of the destination. When the destination is not an array, than the precise type, e.g. of a component in a derived type, is not known, but provided here.

NOTES It is permitted to have *image_index* equal the current image; the memory of the send-to and the send-from might (partially) overlap in that case. The implementation has to take care that it handles this case, e.g. using **memmove** which handles (partially) overlapping memory. If *may_require_tmp* is true, the library might additionally create a temporary variable, unless additional checks show that this is not required (e.g. because walking backward is possible or because both arrays are contiguous and **memmove** takes care of overlap issues). Note that the assignment of a scalar to an array is permitted. In addition, the library has to handle numeric-type conversion and for strings, padding and different character kinds.

Because of the more complicated references possible some operations may be unsupported by certain libraries. The library is expected to issue a precise error message why the operation is not permitted.

8.2.15 `_gfortran_caf_get_by_ref` — Getting data from a remote image using enhanced references

Description:

Called to get a scalar, an array section or a whole array from a remote image identified by the *image_index*.

Syntax:

```
void _gfortran_caf_get_by_ref (caf_token_t token, int image_index,
caf_reference_t *refs, gfc_descriptor_t *dst, int dst_kind, int
src_kind, bool may_require_tmp, bool dst_reallocatable, int *stat,
int src_type)
```

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number.
<i>refs</i>	intent(in) The references to apply to the remote structure to get the data.
<i>dst</i>	intent(in) Array descriptor of the local array to store the data transferred from the remote image. May be reallocated where needed and when <i>DST_REALLOCATABLE</i> allows it.
<i>dst_kind</i>	intent(in) Kind of the destination argument
<i>src_kind</i>	intent(in) Kind of the source argument
<i>may_require_tmp</i>	intent(in) The variable is false when it is known at compile time that the <i>dest</i> and <i>src</i> either cannot overlap or overlap (fully or partially) such that walking <i>src</i> and <i>dest</i> in element wise element order (honoring the stride value) will not lead to wrong results. Otherwise, the value is true .
<i>dst_reallocatable</i>	intent(in) Set when <i>DST</i> is of allocatable or pointer type and its refs allow reallocation, i.e., the full array or a component is referenced.
<i>stat</i>	intent(out) When non-NULL give the result of the operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.
<i>src_type</i>	intent(in) Give the type of the source. When the source is not an array, then the precise type, e.g. of a component in a derived type, is not known, but provided here.

NOTES It is permitted to have *image_index* equal the current image; the memory of the send-to and the send-from might (partially) overlap in that case. The implementation has to take care that it handles this case, e.g. using **memmove** which handles (partially) overlapping memory. If *may_require_tmp* is true, the library might additionally create a temporary variable, unless additional checks show that this is not required (e.g. because walking backward is possible or because both arrays are contiguous and **memmove** takes care of overlap issues). Note that the library has to handle numeric-type conversion and for strings, padding and different character kinds.

Because of the more complicated references possible some operations may be unsupported by certain libraries. The library is expected to issue a precise error message why the operation is not permitted.

8.2.16 *_gfortran_caf_sendget_by_ref* — Sending data between remote images using enhanced references on both sides

Description:

Called to send a scalar, an array section or a whole array from a remote image identified by the *src_image_index* to a remote image identified by the *dst_image_index*.

Syntax: `void _gfortran_caf_sendget_by_ref (caf_token_t dst_token, int
dst_image_index, caf_reference_t *dst_refs, caf_token_t src_token,
int src_image_index, caf_reference_t *src_refs, int dst_kind, int
src_kind, bool may_require_tmp, int *dst_stat, int *src_stat, int
dst_type, int src_type)`

Arguments:

<i>dst_token</i>	intent(in) An opaque pointer identifying the destination coarray.
<i>dst_image_index</i>	intent(in) The ID of the destination remote image; must be a positive number.
<i>dst_refs</i>	intent(in) The references on the remote array to store the data given by the source. Guaranteed to have at least one entry.
<i>src_token</i>	intent(in) An opaque pointer identifying the source coarray.
<i>src_image_index</i>	intent(in) The ID of the source remote image; must be a positive number.
<i>src_refs</i>	intent(in) The references to apply to the remote structure to get the data.
<i>dst_kind</i>	intent(in) Kind of the destination argument
<i>src_kind</i>	intent(in) Kind of the source argument
<i>may_require_tmp</i>	intent(in) The variable is false when it is known at compile time that the <i>dest</i> and <i>src</i> either cannot overlap or overlap (fully or partially) such that walking <i>src</i> and <i>dest</i> in element wise element order (honoring the stride value) will not lead to wrong results. Otherwise, the value is true .
<i>dst_stat</i>	intent(out) when non-NULL give the result of the send-operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.
<i>src_stat</i>	intent(out) When non-NULL give the result of the get-operation, i.e., zero on success and non-zero on error. When NULL and an error occurs, then an error message is printed and the program is terminated.
<i>dst_type</i>	intent(in) Give the type of the destination. When the destination is not an array, than the precise type, e.g. of a component in a derived type, is not known, but provided here.
<i>src_type</i>	intent(in) Give the type of the source. When the source is not an array, than the precise type, e.g. of a component in a derived type, is not known, but provided here.

NOTES It is permitted to have the same image index for both *src_image_index* and *dst_image_index*; the memory of the send-to and the send-from might (partially) overlap in that case. The implementation has to take care that it handles this case, e.g. using **memmove** which handles (partially) overlapping memory. If *may_require_tmp* is true, the library might additionally create a temporary variable, unless additional checks show that this is not required (e.g. because

walking backward is possible or because both arrays are contiguous and `memmove` takes care of overlap issues).

Note that the assignment of a scalar to an array is permitted. In addition, the library has to handle numeric-type conversion and for strings, padding and different character kinds.

Because of the more complicated references possible some operations may be unsupported by certain libraries. The library is expected to issue a precise error message why the operation is not permitted.

8.2.17 `_gfortran_caf_lock` — Locking a lock variable

Description:

Acquire a lock on the given image on a scalar locking variable or for the given array element for an array-valued variable. If the *acquired_lock* is `NULL`, the function returns after having obtained the lock. If it is non-`NULL`, then *acquired_lock* is assigned the value `true` (one) when the lock could be obtained and `false` (zero) otherwise. Locking a lock variable which has already been locked by the same image is an error.

Syntax: `void _gfortran_caf_lock (caf_token_t token, size_t index, int image_index, int *acquired_lock, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>index</i>	intent(in) Array index; first array index is 0. For scalars, it is always 0.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number.
<i>acquired_lock</i>	intent(out) If not <code>NULL</code> , it returns whether lock could be obtained.
<i>stat</i>	intent(out) Stores the <code>STAT=</code> ; may be <code>NULL</code> .
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be <code>NULL</code> .
<i>errmsg_len</i>	intent(in) the buffer size of <i>errmsg</i>

NOTES This function is also called for critical blocks; for those, the array index is always zero and the image index is one. Libraries are permitted to use other images for critical-block locking variables.

8.2.18 `_gfortran_caf_unlock` — Unlocking a lock variable

Description:

Release a lock on the given image on a scalar locking variable or for the given array element for an array-valued variable. Unlocking a lock variable which is unlocked or has been locked by a different image is an error.

Syntax: `void _gfortran_caf_unlock (caf_token_t token, size_t index, int image_index, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>index</i>	intent(in) Array index; first array index is 0. For scalars, it is always 0.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number.
<i>stat</i>	intent(out) For allocatable coarrays, stores the STAT=; may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>errmsg_len</i>	intent(in) the buffer size of errmsg

NOTES This function is also called for critical block; for those, the array index is always zero and the image index is one. Libraries are permitted to use other images for critical-block locking variables.

8.2.19 `_gfortran_caf_event_post` — Post an event

Description:

Increment the event count of the specified event variable.

Syntax: `void _gfortran_caf_event_post (caf_token_t token, size_t index, int image_index, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>index</i>	intent(in) Array index; first array index is 0. For scalars, it is always 0.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number; zero indicates the current image, when accessed noncoindexed.
<i>stat</i>	intent(out) Stores the STAT=; may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>errmsg_len</i>	intent(in) the buffer size of errmsg

NOTES This acts like an atomic add of one to the remote image's event variable. The statement is an image-control statement but does not imply sync memory. Still, all preceeding push communications of this image to the specified remote image have to be completed before `event_wait` on the remote image returns.

8.2.20 `_gfortran_caf_event_wait` — Wait that an event occurred

Description:

Wait until the event count has reached at least the specified *until_count*; if so, atomically decrement the event variable by this amount and return.

Syntax: `void _gfortran_caf_event_wait (caf_token_t token, size_t index, int until_count, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
--------------	---

<i>index</i>	intent(in) Array index; first array index is 0. For scalars, it is always 0.
<i>until_count</i>	intent(in) The number of events which have to be available before the function returns.
<i>stat</i>	intent(out) Stores the STAT=; may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>errmsg_len</i>	intent(in) the buffer size of errmsg

NOTES This function only operates on a local coarray. It acts like a loop checking atomically the value of the event variable, breaking if the value is greater or equal the requested number of counts. Before the function returns, the event variable has to be decremented by the requested *until_count* value. A possible implementation would be a busy loop for a certain number of spins (possibly depending on the number of threads relative to the number of available cores) followed by another waiting strategy such as a sleeping wait (possibly with an increasing number of sleep time) or, if possible, a futex wait.

The statement is an image-control statement but does not imply sync memory. Still, all preceeding push communications of this image to the specified remote image have to be completed before `event_wait` on the remote image returns.

8.2.21 `_gfortran_caf_event_query` — Query event count

Description:

Return the event count of the specified event variable.

Syntax: `void _gfortran_caf_event_query (caf_token_t token, size_t index, int image_index, int *count, int *stat)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>index</i>	intent(in) Array index; first array index is 0. For scalars, it is always 0.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number; zero indicates the current image when accessed noncoindexed.
<i>count</i>	intent(out) The number of events currently posted to the event variable.
<i>stat</i>	intent(out) Stores the STAT=; may be NULL.

NOTES The typical use is to check the local event variable to only call `event_wait` when the data is available. However, a coindexed variable is permitted; there is no ordering or synchronization implied. It acts like an atomic fetch of the value of the event variable.

8.2.22 `_gfortran_caf_sync_all` — All-image barrier

Description:

Synchronization of all images in the current team; the program only continues on a given image after this function has been called on all images of the cur-

NOTE A simple implementation could be

`__asm__ __volatile__ (""::"memory")` to prevent code movements.

8.2.25 `_gfortran_caf_error_stop` — Error termination with exit code

Description:

Invoked for an `ERROR STOP` statement which has an integer argument. The function should terminate the program with the specified exit code.

Syntax: `void _gfortran_caf_error_stop (int error)`

Arguments:

`error` intent(in) The exit status to be used.

8.2.26 `_gfortran_caf_error_stop_str` — Error termination with string

Description:

Invoked for an `ERROR STOP` statement which has a string as argument. The function should terminate the program with a nonzero-exit code.

Syntax: `void _gfortran_caf_error_stop (const char *string, size_t len)`

Arguments:

`string` intent(in) the error message (not zero terminated)
`len` intent(in) the length of the string

8.2.27 `_gfortran_caf_fail_image` — Mark the image failed and end its execution

Description:

Invoked for an `FAIL IMAGE` statement. The function should terminate the current image.

Syntax: `void _gfortran_caf_fail_image ()`

NOTES This function follows TS18508.

8.2.28 `_gfortran_caf_atomic_define` — Atomic variable assignment

Description:

Assign atomically a value to an integer or logical variable.

Syntax: `void _gfortran_caf_atomic_define (caf_token_t token, size_t offset, int image_index, void *value, int *stat, int type, int kind)`

Arguments:

`token` intent(in) An opaque pointer identifying the coarray.
`offset` intent(in) By which amount of bytes the actual data is shifted compared to the base address of the coarray.
`image_index` intent(in) The ID of the remote image; must be a positive number; zero indicates the current image when used noncoindexed.

<i>value</i>	intent(in) the value to be assigned, passed by reference
<i>stat</i>	intent(out) Stores the status STAT= and may be NULL.
<i>type</i>	intent(in) The data type, i.e. BT_INTEGER (1) or BT_LOGICAL (2).
<i>kind</i>	intent(in) The kind value (only 4; always int)

8.2.29 `_gfortran_caf_atomic_ref` — Atomic variable reference

Description:

Reference atomically a value of a kind-4 integer or logical variable.

Syntax: `void _gfortran_caf_atomic_ref (caf_token_t token, size_t offset, int image_index, void *value, int *stat, int type, int kind)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>offset</i>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the coarray.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number; zero indicates the current image when used noncoindexed.
<i>value</i>	intent(out) The variable assigned the atomically referenced variable.
<i>stat</i>	intent(out) Stores the status STAT= and may be NULL.
<i>type</i>	the data type, i.e. BT_INTEGER (1) or BT_LOGICAL (2).
<i>kind</i>	The kind value (only 4; always int)

8.2.30 `_gfortran_caf_atomic_cas` — Atomic compare and swap

Description:

Atomic compare and swap of a kind-4 integer or logical variable. Assigns atomically the specified value to the atomic variable, if the latter has the value specified by the passed condition value.

Syntax: `void _gfortran_caf_atomic_cas (caf_token_t token, size_t offset, int image_index, void *old, void *compare, void *new_val, int *stat, int type, int kind)`

Arguments:

<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>offset</i>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the coarray.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number; zero indicates the current image when used noncoindexed.
<i>old</i>	intent(out) The value which the atomic variable had just before the cas operation.
<i>compare</i>	intent(in) The value used for comparison.

<i>new_val</i>	intent(in) The new value for the atomic variable, assigned to the atomic variable, if <code>compare</code> equals the value of the atomic variable.
<i>stat</i>	intent(out) Stores the status <code>STAT=</code> and may be <code>NULL</code> .
<i>type</i>	intent(in) the data type, i.e. <code>BT_INTEGER</code> (1) or <code>BT_LOGICAL</code> (2).
<i>kind</i>	intent(in) The kind value (only 4; always <code>int</code>)

8.2.31 `_gfortran_caf_atomic_op` — Atomic operation

Description:

Apply an operation atomically to an atomic integer or logical variable. After the operation, *old* contains the value just before the operation, which, respectively, adds (`GFC_CAF_ATOMIC_ADD`) atomically the *value* to the atomic integer variable or does a bitwise AND, OR or exclusive OR between the atomic variable and *value*; the result is then stored in the atomic variable.

Syntax: `void _gfortran_caf_atomic_op (int op, caf_token_t token, size_t offset, int image_index, void *value, void *old, int *stat, int type, int kind)`

Arguments:

<i>op</i>	intent(in) the operation to be performed; possible values <code>GFC_CAF_ATOMIC_ADD</code> (1), <code>GFC_CAF_ATOMIC_AND</code> (2), <code>GFC_CAF_ATOMIC_OR</code> (3), <code>GFC_CAF_ATOMIC_XOR</code> (4).
<i>token</i>	intent(in) An opaque pointer identifying the coarray.
<i>offset</i>	intent(in) By which amount of bytes the actual data is shifted compared to the base address of the coarray.
<i>image_index</i>	intent(in) The ID of the remote image; must be a positive number; zero indicates the current image when used noncoindexed.
<i>old</i>	intent(out) The value which the atomic variable had just before the atomic operation.
<i>val</i>	intent(in) The new value for the atomic variable, assigned to the atomic variable, if <code>compare</code> equals the value of the atomic variable.
<i>stat</i>	intent(out) Stores the status <code>STAT=</code> and may be <code>NULL</code> .
<i>type</i>	intent(in) the data type, i.e. <code>BT_INTEGER</code> (1) or <code>BT_LOGICAL</code> (2)
<i>kind</i>	intent(in) the kind value (only 4; always <code>int</code>)

8.2.32 `_gfortran_caf_co_broadcast` — Sending data to all images

Description:

Distribute a value from a given image to all other images in the team. Has to be called collectively.

Syntax: `void _gfortran_caf_co_broadcast (gfc_descriptor_t *a, int source_image, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>a</i>	intent(inout) An array descriptor with the data to be broadcasted (on <i>source_image</i>) or to be received (other images).
<i>source_image</i>	intent(in) The ID of the image from which the data should be broadcasted.
<i>stat</i>	intent(out) Stores the status STAT= and may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>errmsg_len</i>	intent(in) the buffer size of <i>errmsg</i> .

8.2.33 `_gfortran_caf_co_max` — Collective maximum reduction*Description:*

Calculates for each array element of the variable *a* the maximum value for that element in the current team; if *result_image* has the value 0, the result shall be stored on all images, otherwise, only on the specified image. This function operates on numeric values and character strings.

Syntax: `void _gfortran_caf_co_max (gfc_descriptor_t *a, int result_image, int *stat, char *errmsg, int a_len, size_t errmsg_len)`

Arguments:

<i>a</i>	intent(inout) An array descriptor for the data to be processed. On the destination image(s) the result overwrites the old content.
<i>result_image</i>	intent(in) The ID of the image to which the reduced value should be copied to; if zero, it has to be copied to all images.
<i>stat</i>	intent(out) Stores the status STAT= and may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>a_len</i>	intent(in) the string length of argument <i>a</i>
<i>errmsg_len</i>	intent(in) the buffer size of <i>errmsg</i>

NOTES If *result_image* is nonzero, the data in the array descriptor *a* on all images except of the specified one become undefined; hence, the library may make use of this.

8.2.34 `_gfortran_caf_co_min` — Collective minimum reduction*Description:*

Calculates for each array element of the variable *a* the minimum value for that element in the current team; if *result_image* has the value 0, the result shall be stored on all images, otherwise, only on the specified image. This function operates on numeric values and character strings.

Syntax: `void _gfortran_caf_co_min (gfc_descriptor_t *a, int result_image, int *stat, char *errmsg, int a_len, size_t errmsg_len)`

Arguments:

<i>a</i>	intent(inout) An array descriptor for the data to be processed. On the destination image(s) the result overwrites the old content.
<i>result_image</i>	intent(in) The ID of the image to which the reduced value should be copied to; if zero, it has to be copied to all images.
<i>stat</i>	intent(out) Stores the status STAT= and may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>a_len</i>	intent(in) the string length of argument <i>a</i>
<i>errmsg_len</i>	intent(in) the buffer size of <i>errmsg</i>

NOTES If *result_image* is nonzero, the data in the array descriptor *a* on all images except of the specified one become undefined; hence, the library may make use of this.

8.2.35 `_gfortran_caf_co_sum` — Collective summing reduction

Description:

Calculates for each array element of the variable *a* the sum of all values for that element in the current team; if *result_image* has the value 0, the result shall be stored on all images, otherwise, only on the specified image. This function operates on numeric values only.

Syntax: `void _gfortran_caf_co_sum (gfc_descriptor_t *a, int result_image, int *stat, char *errmsg, size_t errmsg_len)`

Arguments:

<i>a</i>	intent(inout) An array descriptor with the data to be processed. On the destination image(s) the result overwrites the old content.
<i>result_image</i>	intent(in) The ID of the image to which the reduced value should be copied to; if zero, it has to be copied to all images.
<i>stat</i>	intent(out) Stores the status STAT= and may be NULL.
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be NULL.
<i>errmsg_len</i>	intent(in) the buffer size of <i>errmsg</i>

NOTES If *result_image* is nonzero, the data in the array descriptor *a* on all images except of the specified one become undefined; hence, the library may make use of this.

8.2.36 `_gfortran_caf_co_reduce` — Generic collective reduction

Description:

Calculates for each array element of the variable *a* the reduction value for that element in the current team; if *result_image* has the value 0, the result shall be stored on all images, otherwise, only on the specified image. The *opr* is a pure function doing a mathematically commutative and associative operation.

The *opr_flags* denote the following; the values are bitwise ored. `GFC_CAF_BYREF` (1) if the result should be returned by reference; `GFC_CAF_HIDDENLEN` (2) whether the result and argument string lengths shall be specified as hidden arguments; `GFC_CAF_ARG_VALUE` (4) whether the arguments shall be passed by value, `GFC_CAF_ARG_DESC` (8) whether the arguments shall be passed by descriptor.

Syntax: `void _gfortran_caf_co_reduce (gfc_descriptor_t *a, void * (*opr) (void *, void *), int opr_flags, int result_image, int *stat, char *errmsg, int a_len, size_t errmsg_len)`

Arguments:

<i>a</i>	intent(inout) An array descriptor with the data to be processed. On the destination image(s) the result overwrites the old content.
<i>opr</i>	intent(in) Function pointer to the reduction function
<i>opr_flags</i>	intent(in) Flags regarding the reduction function
<i>result_image</i>	intent(in) The ID of the image to which the reduced value should be copied to; if zero, it has to be copied to all images.
<i>stat</i>	intent(out) Stores the status <code>STAT=</code> and may be <code>NULL</code> .
<i>errmsg</i>	intent(out) When an error occurs, this will be set to an error message; may be <code>NULL</code> .
<i>a_len</i>	intent(in) the string length of argument <i>a</i>
<i>errmsg_len</i>	intent(in) the buffer size of <i>errmsg</i>

NOTES If *result_image* is nonzero, the data in the array descriptor *a* on all images except of the specified one become undefined; hence, the library may make use of this.

For character arguments, the result is passed as first argument, followed by the result string length, next come the two string arguments, followed by the two hidden string length arguments. With C binding, there are no hidden arguments and by-reference passing and either only a single character is passed or an array descriptor.

9 Intrinsic Procedures

9.1 Introduction to intrinsic procedures

The intrinsic procedures provided by GNU Fortran include all of the intrinsic procedures required by the Fortran 95 standard, a set of intrinsic procedures for backwards compatibility with G77, and a selection of intrinsic procedures from the Fortran 2003 and Fortran 2008 standards. Any conflict between a description here and a description in either the Fortran 95 standard, the Fortran 2003 standard or the Fortran 2008 standard is unintentional, and the standard(s) should be considered authoritative.

The enumeration of the `KIND` type parameter is processor defined in the Fortran 95 standard. GNU Fortran defines the default integer type and default real type by `INTEGER(KIND=4)` and `REAL(KIND=4)`, respectively. The standard mandates that both data types shall have another kind, which have more precision. On typical target architectures supported by `gfortran`, this kind type parameter is `KIND=8`. Hence, `REAL(KIND=8)` and `DOUBLE PRECISION` are equivalent. In the description of generic intrinsic procedures, the kind type parameter will be specified by `KIND=*`, and in the description of specific names for an intrinsic procedure the kind type parameter will be explicitly given (e.g., `REAL(KIND=4)` or `REAL(KIND=8)`). Finally, for brevity the optional `KIND=` syntax will be omitted.

Many of the intrinsic procedures take one or more optional arguments. This document follows the convention used in the Fortran 95 standard, and denotes such arguments by square brackets.

GNU Fortran offers the `'-std=f95'` and `'-std=gnu'` options, which can be used to restrict the set of intrinsic procedures to a given standard. By default, `gfortran` sets the `'-std=gnu'` option, and so all intrinsic procedures described here are accepted. There is one caveat. For a select group of intrinsic procedures, `g77` implemented both a function and a subroutine. Both classes have been implemented in `gfortran` for backwards compatibility with `g77`. It is noted here that these functions and subroutines cannot be intermixed in a given subprogram. In the descriptions that follow, the applicable standard for each intrinsic procedure is noted.

9.2 ABORT — Abort the program

Description:

`ABORT` causes immediate termination of the program. On operating systems that support a core dump, `ABORT` will produce a core dump. It will also print a backtrace, unless `-fno-backtrace` is given.

Standard: GNU extension

Class: Subroutine

Syntax: `CALL ABORT`

Return value:

Does not return.

Example:

```

program test_abort
  integer :: i = 1, j = 2
  if (i /= j) call abort
end program test_abort

```

See also: [Section 9.101 \[EXIT\]](#), page 177, [Section 9.163 \[KILL\]](#), page 215, [Section 9.39 \[BACKTRACE\]](#), page 137

9.3 ABS — Absolute value

Description:

ABS(A) computes the absolute value of A.

Standard: Fortran 77 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = ABS(A)

Arguments:

A The type of the argument shall be an INTEGER, REAL, or COMPLEX.

Return value:

The return value is of the same type and kind as the argument except the return value is REAL for a COMPLEX argument.

Example:

```

program test_abs
  integer :: i = -1
  real :: x = -1.e0
  complex :: z = (-1.e0,0.e0)
  i = abs(i)
  x = abs(x)
  z = abs(z)
end program test_abs

```

Specific names:

Name	Argument	Return type	Standard
ABS(A)	REAL(4) A	REAL(4)	Fortran 77 and later
CABS(A)	COMPLEX(4) A	REAL(4)	Fortran 77 and later
DABS(A)	REAL(8) A	REAL(8)	Fortran 77 and later
IABS(A)	INTEGER(4) A	INTEGER(4)	Fortran 77 and later
BABS(A)	INTEGER(1) A	INTEGER(1)	GNU extension
IIABS(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JIABS(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KIABS(A)	INTEGER(8) A	INTEGER(8)	GNU extension
ZABS(A)	COMPLEX(8) A	COMPLEX(8)	GNU extension
CDABS(A)	COMPLEX(8) A	COMPLEX(8)	GNU extension

9.4 ACCESS — Checks file access modes

Description:

ACCESS(NAME, MODE) checks whether the file *NAME* exists, is readable, writable or executable. Except for the executable check, ACCESS can be replaced by Fortran 95's INQUIRE.

Standard: GNU extension

Class: Inquiry function

Syntax: RESULT = ACCESS(NAME, MODE)

Arguments:

<i>NAME</i>	Scalar CHARACTER of default kind with the file name. Tailing blank are ignored unless the character achar(0) is present, then all characters up to and excluding achar(0) are used as file name.
<i>MODE</i>	Scalar CHARACTER of default kind with the file access mode, may be any concatenation of "r" (readable), "w" (writable) and "x" (executable), or " " to check for existence.

Return value:

Returns a scalar INTEGER, which is 0 if the file is accessible in the given mode; otherwise or if an invalid argument has been given for MODE the value 1 is returned.

Example:

```

program access_test
  implicit none
  character(len=*), parameter :: file = 'test.dat'
  character(len=*), parameter :: file2 = 'test.dat '//achar(0)
  if(access(file,' ') == 0) print *, trim(file),' is exists'
  if(access(file,'r') == 0) print *, trim(file),' is readable'
  if(access(file,'w') == 0) print *, trim(file),' is writable'
  if(access(file,'x') == 0) print *, trim(file),' is executable'
  if(access(file2,'rwx') == 0) &
    print *, trim(file2),' is readable, writable and executable'
end program access_test

```

Specific names:

See also:

9.5 ACHAR — Character in ASCII collating sequence

Description:

ACHAR(I) returns the character located at position I in the ASCII collating sequence.

Standard: Fortran 77 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: RESULT = ACHAR(I [, KIND])

Arguments:

I The type shall be **INTEGER**.
KIND (Optional) An **INTEGER** initialization expression indicating
the kind parameter of the result.

Return value:

The return value is of type **CHARACTER** with a length of one. If the *KIND* argument is present, the return value is of the specified kind and of the default kind otherwise.

Example:

```
program test_achar
  character c
  c = achar(32)
end program test_achar
```

Note: See [Section 9.143 \[ICHAR\]](#), page 203 for a discussion of converting between numerical values and formatted string representations.

See also: [Section 9.59 \[CHAR\]](#), page 148, [Section 9.135 \[IACHAR\]](#), page 198,
[Section 9.143 \[ICHAR\]](#), page 203

9.6 ACOS — Arccosine function

Description:

ACOS(X) computes the arccosine of *X* (inverse of **COS(X)**).

Standard: Fortran 77 and later, for a complex argument Fortran 2008 or later

Class: Elemental function

Syntax: **RESULT = ACOS(X)**

Arguments:

X The type shall either be **REAL** with a magnitude that is less
than or equal to one - or the type shall be **COMPLEX**.

Return value:

The return value is of the same type and kind as *X*. The real part of the result is in radians and lies in the range $0 \leq \Re \text{acos}(x) \leq \pi$.

Example:

```
program test_acos
  real(8) :: x = 0.866_8
  x = acos(x)
end program test_acos
```

Specific names:

Name	Argument	Return type	Standard
ACOS(X)	REAL(4) X	REAL(4)	Fortran 77 and later
DACOS(X)	REAL(8) X	REAL(8)	Fortran 77 and later

See also: Inverse function: [Section 9.73 \[COS\]](#), page 158 Degrees function: [Section 9.7 \[ACOSD\]](#), page 115

9.7 ACOSD — Arccosine function, degrees

Description:

ACOSD(*X*) computes the arccosine of *X* in degrees (inverse of COSD(*X*)).

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘-fdec-math’

Class: Elemental function

Syntax: RESULT = ACOSD(*X*)

Arguments:

X The type shall either be REAL with a magnitude that is less than or equal to one - or the type shall be COMPLEX.

Return value:

The return value is of the same type and kind as *X*. The real part of the result is in degrees and lies in the range $0 \leq \Re \operatorname{acos}(x) \leq 180$.

Example:

```
program test_acosd
  real(8) :: x = 0.866_8
  x = acosd(x)
end program test_acosd
```

Specific names:

Name	Argument	Return type	Standard
ACOSD(<i>X</i>)	REAL(4) <i>X</i>	REAL(4)	GNU Extension
DACOSD(<i>X</i>)	REAL(8) <i>X</i>	REAL(8)	GNU Extension

See also: Inverse function: [Section 9.74 \[COSD\], page 159](#) Radians function: [Section 9.6 \[ACOS\], page 114](#)

9.8 ACOSH — Inverse hyperbolic cosine function

Description:

ACOSH(*X*) computes the inverse hyperbolic cosine of *X*.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = ACOSH(*X*)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has the same type and kind as *X*. If *X* is complex, the imaginary part of the result is in radians and lies between $0 \leq \Im \operatorname{acosh}(x) \leq \pi$.

Example:

```

PROGRAM test_acosh
  REAL(8), DIMENSION(3) :: x = (/ 1.0, 2.0, 3.0 /)
  WRITE (*,*) ACOSH(x)
END PROGRAM

```

Specific names:

Name	Argument	Return type	Standard
DACOSH(X)	REAL(8) X	REAL(8)	GNU extension

See also: Inverse function: [Section 9.75 \[COSH\], page 159](#)

9.9 ADJUSTL — Left adjust a string

Description:

ADJUSTL(*STRING*) will left adjust a string by removing leading spaces. Spaces are inserted at the end of the string as needed.

Standard: Fortran 90 and later

Class: Elemental function

Syntax: RESULT = ADJUSTL(*STRING*)

Arguments:

STRING The type shall be CHARACTER.

Return value:

The return value is of type CHARACTER and of the same kind as *STRING* where leading spaces are removed and the same number of spaces are inserted on the end of *STRING*.

Example:

```

program test_adjustl
  character(len=20) :: str = '   gfortran'
  str = adjustl(str)
  print *, str
end program test_adjustl

```

See also: [Section 9.10 \[ADJUSTR\], page 116](#), [Section 9.276 \[TRIM\], page 280](#)

9.10 ADJUSTR — Right adjust a string

Description:

ADJUSTR(*STRING*) will right adjust a string by removing trailing spaces. Spaces are inserted at the start of the string as needed.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = ADJUSTR(*STRING*)

Arguments:

STR The type shall be CHARACTER.

Return value:

The return value is of type **CHARACTER** and of the same kind as *STRING* where trailing spaces are removed and the same number of spaces are inserted at the start of *STRING*.

Example:

```
program test_adjustr
  character(len=20) :: str = 'gfortran'
  str = adjustr(str)
  print *, str
end program test_adjustr
```

See also: [Section 9.9 \[ADJUSTL\], page 116](#), [Section 9.276 \[TRIM\], page 280](#)

9.11 AIMAG — Imaginary part of complex number

Description:

AIMAG(Z) yields the imaginary part of complex argument Z. The IMAG(Z) and IMAGPART(Z) intrinsic functions are provided for compatibility with g77, and their use in new code is strongly discouraged.

Standard: Fortran 77 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = AIMAG(Z)

Arguments:

Z The type of the argument shall be **COMPLEX**.

Return value:

The return value is of type **REAL** with the kind type parameter of the argument.

Example:

```
program test_aimag
  complex(4) z4
  complex(8) z8
  z4 = cmplx(1.e0_4, 0.e0_4)
  z8 = cmplx(0.e0_8, 1.e0_8)
  print *, aimag(z4), dimag(z8)
end program test_aimag
```

Specific names:

Name	Argument	Return type	Standard
AIMAG(Z)	COMPLEX Z	REAL	GNU extension
DIMAG(Z)	COMPLEX(8) Z	REAL(8)	GNU extension
IMAG(Z)	COMPLEX Z	REAL	GNU extension
IMAGPART(Z)	COMPLEX Z	REAL	GNU extension

9.12 AINT — Truncate to a whole number

Description:

AINT(A [, KIND]) truncates its argument to a whole number.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = AINT(A [, KIND])`

Arguments:

A The type of the argument shall be `REAL`.
KIND (Optional) An `INTEGER` initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `REAL` with the kind type parameter of the argument if the optional *KIND* is absent; otherwise, the kind type parameter will be given by *KIND*. If the magnitude of *X* is less than one, `AINT(X)` returns zero. If the magnitude is equal to or greater than one then it returns the largest whole number that does not exceed its magnitude. The sign is the same as the sign of *X*.

Example:

```
program test_aint
  real(4) x4
  real(8) x8
  x4 = 1.234E0_4
  x8 = 4.321_8
  print *, aint(x4), dint(x8)
  x8 = aint(x4,8)
end program test_aint
```

Specific names:

Name	Argument	Return type	Standard
<code>AINT(A)</code>	<code>REAL(4) A</code>	<code>REAL(4)</code>	Fortran 77 and later
<code>DINT(A)</code>	<code>REAL(8) A</code>	<code>REAL(8)</code>	Fortran 77 and later

9.13 ALARM — Execute a routine after a given delay

Description:

`ALARM(SECONDS, HANDLER [, STATUS])` causes external subroutine *HANDLER* to be executed after a delay of *SECONDS* by using `alarm(2)` to set up a signal and `signal(2)` to catch it. If *STATUS* is supplied, it will be returned with the number of seconds remaining until any previously scheduled alarm was due to be delivered, or zero if there was no previously scheduled alarm.

Standard: GNU extension

Class: Subroutine

Syntax: `CALL ALARM(SECONDS, HANDLER [, STATUS])`

Arguments:

SECONDS The type of the argument shall be a scalar `INTEGER`. It is `INTENT(IN)`.
HANDLER Signal handler (`INTEGER FUNCTION` or `SUBROUTINE`) or dummy/global `INTEGER` scalar. The scalar values may be either `SIG_IGN=1` to ignore the alarm generated or `SIG_DFL=0` to set the default action. It is `INTENT(IN)`.

STATUS (Optional) *STATUS* shall be a scalar variable of the default INTEGER kind. It is INTENT(OUT).

Example:

```
program test_alarm
  external handler_print
  integer i
  call alarm (3, handler_print, i)
  print *, i
  call sleep(10)
end program test_alarm
```

This will cause the external routine *handler_print* to be called after 3 seconds.

9.14 ALL — All values in *MASK* along *DIM* are true

Description:

ALL(MASK [, DIM]) determines if all the values are true in *MASK* in the array along dimension *DIM*.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = ALL(MASK [, DIM])

Arguments:

MASK The type of the argument shall be LOGICAL and it shall not be scalar.

DIM (Optional) *DIM* shall be a scalar integer with a value that lies between one and the rank of *MASK*.

Return value:

ALL(MASK) returns a scalar value of type LOGICAL where the kind type parameter is the same as the kind type parameter of *MASK*. If *DIM* is present, then ALL(MASK, DIM) returns an array with the rank of *MASK* minus 1. The shape is determined from the shape of *MASK* where the *DIM* dimension is elided.

- (A) ALL(MASK) is true if all elements of *MASK* are true. It also is true if *MASK* has zero size; otherwise, it is false.
- (B) If the rank of *MASK* is one, then ALL(MASK, DIM) is equivalent to ALL(MASK). If the rank is greater than one, then ALL(MASK, DIM) is determined by applying ALL to the array sections.

Example:

```
program test_all
  logical l
  l = all(/.true., .true., .true./)
  print *, l
  call section
contains
  subroutine section
    integer a(2,3), b(2,3)
    a = 1
    b = 1
```

```

        b(2,2) = 2
        print *, all(a .eq. b, 1)
        print *, all(a .eq. b, 2)
    end subroutine section
end program test_all

```

9.15 ALLOCATED — Status of an allocatable entity

Description:

ALLOCATED(ARRAY) and ALLOCATED(SCALAR) check the allocation status of *ARRAY* and *SCALAR*, respectively.

Standard: Fortran 95 and later. Note, the SCALAR= keyword and allocatable scalar entities are available in Fortran 2003 and later.

Class: Inquiry function

Syntax:

```

RESULT = ALLOCATED(ARRAY)
RESULT = ALLOCATED(SCALAR)

```

Arguments:

<i>ARRAY</i>	The argument shall be an ALLOCATABLE array.
<i>SCALAR</i>	The argument shall be an ALLOCATABLE scalar.

Return value:

The return value is a scalar LOGICAL with the default logical kind type parameter. If the argument is allocated, then the result is .TRUE.; otherwise, it returns .FALSE.

Example:

```

program test_allocated
    integer :: i = 4
    real(4), allocatable :: x(:)
    if (.not. allocated(x)) allocate(x(i))
end program test_allocated

```

9.16 AND — Bitwise logical AND

Description:

Bitwise logical AND.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. For integer arguments, programmers should consider the use of the [Section 9.137 \[IAND\]](#), [page 199](#) intrinsic defined by the Fortran standard.

Standard: GNU extension

Class: Function

Syntax: RESULT = AND(I, J)

Arguments:

<i>I</i>	The type shall be either a scalar INTEGER type or a scalar LOGICAL type or a boz-literal-constant.
----------	--

J The type shall be the same as the type of *I* or a boz-literal-constant. *I* and *J* shall not both be boz-literal-constants. If either *I* or *J* is a boz-literal-constant, then the other argument must be a scalar `INTEGER`.

Return value:

The return type is either a scalar `INTEGER` or a scalar `LOGICAL`. If the kind type parameters differ, then the smaller kind type is implicitly converted to larger kind, and the return has the larger kind. A boz-literal-constant is converted to an `INTEGER` with the kind type parameter of the other argument as-if a call to [Section 9.149 \[INT\]](#), page 207 occurred.

Example:

```
PROGRAM test_and
  LOGICAL :: T = .TRUE., F = .FALSE.
  INTEGER :: a, b
  DATA a / Z'F' /, b / Z'3' /

  WRITE (*,*) AND(T, T), AND(T, F), AND(F, T), AND(F, F)
  WRITE (*,*) AND(a, b)
END PROGRAM
```

See also: Fortran 95 elemental function: [Section 9.137 \[IAND\]](#), page 199

9.17 ANINT — Nearest whole number

Description:

`ANINT(A [, KIND])` rounds its argument to the nearest whole number.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = ANINT(A [, KIND])`

Arguments:

A The type of the argument shall be `REAL`.
KIND (Optional) An `INTEGER` initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type real with the kind type parameter of the argument if the optional *KIND* is absent; otherwise, the kind type parameter will be given by *KIND*. If *A* is greater than zero, `ANINT(A)` returns `AIN(X+0.5)`. If *A* is less than or equal to zero then it returns `AIN(X-0.5)`.

Example:

```
program test_anint
  real(4) x4
  real(8) x8
  x4 = 1.234E0_4
  x8 = 4.321_8
  print *, anint(x4), dnint(x8)
  x8 = anint(x4,8)
end program test_anint
```

Specific names:

Name	Argument	Return type	Standard
AINT(A)	REAL(4) A	REAL(4)	Fortran 77 and later
DNINT(A)	REAL(8) A	REAL(8)	Fortran 77 and later

9.18 ANY — Any value in *MASK* along *DIM* is true

Description:

ANY(MASK [, DIM]) determines if any of the values in the logical array *MASK* along dimension *DIM* are .TRUE..

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = ANY(MASK [, DIM])

Arguments:

MASK The type of the argument shall be LOGICAL and it shall not be scalar.

DIM (Optional) *DIM* shall be a scalar integer with a value that lies between one and the rank of *MASK*.

Return value:

ANY(MASK) returns a scalar value of type LOGICAL where the kind type parameter is the same as the kind type parameter of *MASK*. If *DIM* is present, then ANY(MASK, DIM) returns an array with the rank of *MASK* minus 1. The shape is determined from the shape of *MASK* where the *DIM* dimension is elided.

- (A) ANY(MASK) is true if any element of *MASK* is true; otherwise, it is false. It also is false if *MASK* has zero size.
- (B) If the rank of *MASK* is one, then ANY(MASK,DIM) is equivalent to ANY(MASK). If the rank is greater than one, then ANY(MASK,DIM) is determined by applying ANY to the array sections.

Example:

```

program test_any
  logical l
  l = any(/.true., .true., .true./)
  print *, l
  call section
contains
  subroutine section
    integer a(2,3), b(2,3)
    a = 1
    b = 1
    b(2,2) = 2
    print *, any(a .eq. b, 1)
    print *, any(a .eq. b, 2)
  end subroutine section
end program test_any

```

9.19 ASIN — Arcsine function

Description:

ASIN(*X*) computes the arcsine of its *X* (inverse of SIN(*X*)).

Standard: Fortran 77 and later, for a complex argument Fortran 2008 or later

Class: Elemental function

Syntax: RESULT = ASIN(*X*)

Arguments:

X The type shall be either **REAL** and a magnitude that is less than or equal to one - or be **COMPLEX**.

Return value:

The return value is of the same type and kind as *X*. The real part of the result is in radians and lies in the range $-\pi/2 \leq \Re \operatorname{asin}(x) \leq \pi/2$.

Example:

```
program test_asin
  real(8) :: x = 0.866_8
  x = asin(x)
end program test_asin
```

Specific names:

Name	Argument	Return type	Standard
ASIN(<i>X</i>)	REAL(4) <i>X</i>	REAL(4)	Fortran 77 and later
DASIN(<i>X</i>)	REAL(8) <i>X</i>	REAL(8)	Fortran 77 and later

See also: Inverse function: [Section 9.250 \[SIN\]](#), page 264 Degrees function: [Section 9.20 \[ASIND\]](#), page 123

9.20 ASIND — Arcsine function, degrees

Description:

ASIND(*X*) computes the arcsine of its *X* in degrees (inverse of SIND(*X*)).

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘-fdec-math’.

Class: Elemental function

Syntax: RESULT = ASIND(*X*)

Arguments:

X The type shall be either **REAL** and a magnitude that is less than or equal to one - or be **COMPLEX**.

Return value:

The return value is of the same type and kind as *X*. The real part of the result is in degrees and lies in the range $-90 \leq \Re \operatorname{asin}(x) \leq 90$.

Example:

```

program test_asind
  real(8) :: x = 0.866_8
  x = asind(x)
end program test_asind

```

Specific names:

Name	Argument	Return type	Standard
ASIND(X)	REAL(4) X	REAL(4)	GNU Extension
DASIND(X)	REAL(8) X	REAL(8)	GNU Extension

See also: Inverse function: [Section 9.251 \[SIND\]](#), page 265 Radians function: [Section 9.19 \[ASIN\]](#), page 123

9.21 ASINH — Inverse hyperbolic sine function

Description:

ASINH(X) computes the inverse hyperbolic sine of X.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = ASINH(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value is of the same type and kind as X. If X is complex, the imaginary part of the result is in radians and lies between $-\pi/2 \leq \Im \operatorname{asinh}(x) \leq \pi/2$.

Example:

```

PROGRAM test_asinh
  REAL(8), DIMENSION(3) :: x = (/ -1.0, 0.0, 1.0 /)
  WRITE (*,*) ASINH(x)
END PROGRAM

```

Specific names:

Name	Argument	Return type	Standard
DASINH(X)	REAL(8) X	REAL(8)	GNU extension.

See also: Inverse function: [Section 9.252 \[SINH\]](#), page 266

9.22 ASSOCIATED — Status of a pointer or pointer/target pair

Description:

ASSOCIATED(POINTER [, TARGET]) determines the status of the pointer *POINTER* or if *POINTER* is associated with the target *TARGET*.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: RESULT = ASSOCIATED(POINTER [, TARGET])

Arguments:

POINTER *POINTER* shall have the *POINTER* attribute and it can be of any type.

TARGET (Optional) *TARGET* shall be a pointer or a target. It must have the same type, kind type parameter, and array rank as *POINTER*.

The association status of neither *POINTER* nor *TARGET* shall be undefined.

Return value:

ASSOCIATED(POINTER) returns a scalar value of type *LOGICAL(4)*. There are several cases:

- (A) When the optional *TARGET* is not present then
ASSOCIATED(POINTER) is true if *POINTER* is associated with a target; otherwise, it returns false.
- (B) If *TARGET* is present and a scalar target, the result is true if
TARGET is not a zero-sized storage sequence and the target associated with *POINTER* occupies the same storage units. If *POINTER* is disassociated, the result is false.
- (C) If *TARGET* is present and an array target, the result is true if
TARGET and *POINTER* have the same shape, are not zero-sized arrays, are arrays whose elements are not zero-sized storage sequences, and *TARGET* and *POINTER* occupy the same storage units in array element order. As in case(B), the result is false, if *POINTER* is disassociated.
- (D) If *TARGET* is present and an scalar pointer, the result is true
if *TARGET* is associated with *POINTER*, the target associated with *TARGET* are not zero-sized storage sequences and occupy the same storage units. The result is false, if either *TARGET* or *POINTER* is disassociated.
- (E) If *TARGET* is present and an array pointer, the result is true if
target associated with *POINTER* and the target associated with *TARGET* have the same shape, are not zero-sized arrays, are arrays whose elements are not zero-sized storage sequences, and *TARGET* and *POINTER* occupy the same storage units in array element order. The result is false, if either *TARGET* or *POINTER* is disassociated.

Example:

```

program test_associated
  implicit none
  real, target :: tgt(2) = (/1., 2./)
  real, pointer :: ptr(:)
  ptr => tgt
  if (associated(ptr) .eqv. .false.) call abort
  if (associated(ptr,tgt) .eqv. .false.) call abort
end program test_associated

```

See also: [Section 9.210 \[NULL\], page 241](#)

9.23 ATAN — Arctangent function

Description:

ATAN(X) computes the arctangent of X.

Standard: Fortran 77 and later, for a complex argument and for two arguments Fortran 2008 or later

Class: Elemental function

Syntax:

```
RESULT = ATAN(X)
RESULT = ATAN(Y, X)
```

Arguments:

X The type shall be REAL or COMPLEX; if Y is present, X shall be REAL.

Y shall be of
the same ty-
pe and kind
as X.

Return value:

The return value is of the same type and kind as X. If Y is present, the result is identical to ATAN2(Y,X). Otherwise, it the arcus tangent of X, where the real part of the result is in radians and lies in the range $-\pi/2 \leq \Re \operatorname{atan}(x) \leq \pi/2$.

Example:

```
program test_atan
  real(8) :: x = 2.866_8
  x = atan(x)
end program test_atan
```

Specific names:

Name	Argument	Return type	Standard
ATAN(X)	REAL(4) X	REAL(4)	Fortran 77 and later
DATAN(X)	REAL(8) X	REAL(8)	Fortran 77 and later

See also: Inverse function: [Section 9.266 \[TAN\], page 274](#) Degrees function: [Section 9.24 \[ATAND\], page 126](#)

9.24 ATAND — Arctangent function, degrees

Description:

ATAND(X) computes the arctangent of X in degrees (inverse of [Section 9.267 \[TAND\], page 275](#)).

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘-fdec-math’.

Class: Elemental function

Syntax:

```
RESULT = ATAND(X)
RESULT = ATAND(Y, X)
```

Arguments:

X The type shall be **REAL** or **COMPLEX**; if *Y* is present, *X* shall be **REAL**.

Y shall be of the same type and kind as *X*.

Return value:

The return value is of the same type and kind as *X*. If *Y* is present, the result is identical to **ATAND2**(*Y*,*X*). Otherwise, it is the arcus tangent of *X*, where the real part of the result is in degrees and lies in the range $-90 \leq \Re \operatorname{atand}(x) \leq 90$.

Example:

```
program test_atand
  real(8) :: x = 2.866_8
  x = atand(x)
end program test_atand
```

Specific names:

Name	Argument	Return type	Standard
ATAND (<i>X</i>)	REAL (4) <i>X</i>	REAL (4)	GNU Extension
DATAND (<i>X</i>)	REAL (8) <i>X</i>	REAL (8)	GNU Extension

See also: Inverse function: [Section 9.267 \[TAND\]](#), page 275 Radians function: [Section 9.23 \[ATAN\]](#), page 126

9.25 ATAN2 — Arctangent function

Description:

ATAN2(*Y*, *X*) computes the principal value of the argument function of the complex number $X + iY$. This function can be used to transform from Cartesian into polar coordinates and allows to determine the angle in the correct quadrant.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: **RESULT** = **ATAN2**(*Y*, *X*)

Arguments:

Y The type shall be **REAL**.

X The type and kind type parameter shall be the same as *Y*. If *Y* is zero, then *X* must be nonzero.

Return value:

The return value has the same type and kind type parameter as *Y*. It is the principal value of the complex number $X + iY$. If *X* is nonzero, then it lies in the range $-\pi \leq \operatorname{atan}(x) \leq \pi$. The sign is positive if *Y* is positive. If *Y* is zero,

then the return value is zero if X is strictly positive, π if X is negative and Y is positive zero (or the processor does not handle signed zeros), and $-\pi$ if X is negative and Y is negative zero. Finally, if X is zero, then the magnitude of the result is $\pi/2$.

Example:

```
program test_atan2
  real(4) :: x = 1.e0_4, y = 0.5e0_4
  x = atan2(y,x)
end program test_atan2
```

Specific names:

Name	Argument	Return type	Standard
ATAN2(X , Y)	REAL(4) X , Y	REAL(4)	Fortran 77 and later
DATAN2(X , Y)	REAL(8) X , Y	REAL(8)	Fortran 77 and later

See also: Alias: [Section 9.23 \[ATAN\]](#), page 126 Degrees function: [Section 9.26 \[ATAN2D\]](#), page 128

9.26 ATAN2D — Arctangent function, degrees

Description:

ATAN2D(Y , X) computes the principal value of the argument function of the complex number $X + iY$ in degrees. This function can be used to transform from Cartesian into polar coordinates and allows to determine the angle in the correct quadrant.

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘-fdec-math’.

Class: Elemental function

Syntax: RESULT = ATAN2D(Y , X)

Arguments:

Y	The type shall be REAL.
X	The type and kind type parameter shall be the same as Y . If Y is zero, then X must be nonzero.

Return value:

The return value has the same type and kind type parameter as Y . It is the principal value of the complex number $X + iY$. If X is nonzero, then it lies in the range $-180 \leq \text{atan}(x) \leq 180$. The sign is positive if Y is positive. If Y is zero, then the return value is zero if X is strictly positive, 180 if X is negative and Y is positive zero (or the processor does not handle signed zeros), and -180 if X is negative and Y is negative zero. Finally, if X is zero, then the magnitude of the result is 90.

Example:

```
program test_atan2d
  real(4) :: x = 1.e0_4, y = 0.5e0_4
  x = atan2d(y,x)
end program test_atan2d
```


Specific names:

Name	Argument	Return type	Standard
ATAN2D(X, Y)	REAL(4) X, Y	REAL(4)	GNU Extension
DATAN2D(X, Y)	REAL(8) X, Y	REAL(8)	GNU Extension

See also: Alias: [Section 9.24 \[ATAND\]](#), page 126 Radians function: [Section 9.25 \[ATAN2\]](#), page 127

9.27 ATANH — Inverse hyperbolic tangent function

Description:

ATANH(X) computes the inverse hyperbolic tangent of X.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = ATANH(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has same type and kind as X. If X is complex, the imaginary part of the result is in radians and lies between $-\pi/2 \leq \Im \operatorname{atanh}(x) \leq \pi/2$.

Example:

```
PROGRAM test_atanh
  REAL, DIMENSION(3) :: x = (/ -1.0, 0.0, 1.0 /)
  WRITE (*,*) ATANH(x)
END PROGRAM
```

Specific names:

Name	Argument	Return type	Standard
DATANH(X)	REAL(8) X	REAL(8)	GNU extension

See also: Inverse function: [Section 9.268 \[TANH\]](#), page 276

9.28 ATOMIC_ADD — Atomic ADD operation

Description:

ATOMIC_ADD(ATOM, VALUE) atomically adds the value of VAR to the variable ATOM. When STAT is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed ATOM, if the remote image has stopped, it is assigned the value of ISO_FORTRAN_ENV's STAT_STOPPED_IMAGE and if the remote image has failed, the value STAT_FAILED_IMAGE.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_ADD (ATOM, VALUE [, STAT])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with <code>ATOMIC_INT_KIND</code> kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*]
  call atomic_add (atom[1], this_image())
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.32 \[ATOMIC_FETCH_ADD\]](#), page 132, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.29 \[ATOMIC_AND\]](#), page 130, [Section 9.36 \[ATOMIC_OR\]](#), page 135, [Section 9.38 \[ATOMIC_XOR\]](#), page 136

9.29 ATOMIC_AND — Atomic bitwise AND operation

Description:

`ATOMIC_AND(ATOM, VALUE)` atomically defines *ATOM* with the bitwise AND between the values of *ATOM* and *VALUE*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a co-indexed *ATOM*, if the remote image has stopped, it is assigned the value of `ISO_FORTRAN_ENV`'s `STAT_STOPPED_IMAGE` and if the remote image has failed, the value `STAT_FAILED_IMAGE`.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: `CALL ATOMIC_AND (ATOM, VALUE [, STAT])`

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with <code>ATOMIC_INT_KIND</code> kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*]
  call atomic_and (atom[1], int(b'10100011101'))
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.33 \[ATOMIC_FETCH_AND\]](#), page 133, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.28 \[ATOMIC_ADD\]](#), page 129, [Section 9.36 \[ATOMIC_OR\]](#), page 135, [Section 9.38 \[ATOMIC_XOR\]](#), page 136

9.30 ATOMIC_CAS — Atomic compare and swap

Description:

ATOMIC_CAS compares the variable *ATOM* with the value of *COMPARE*; if the value is the same, *ATOM* is set to the value of *NEW*. Additionally, *OLD* is set to the value of *ATOM* that was used for the comparison. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of ISO_FORTRAN_ENV's STAT_STOPPED_IMAGE and if the remote image has failed, the value STAT_FAILED_IMAGE.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_CAS (ATOM, OLD, COMPARE, NEW [, STAT])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of either integer type with ATOMIC_INT_KIND kind or logical type with ATOMIC_LOGICAL_KIND kind.
<i>OLD</i>	Scalar of the same type and kind as <i>ATOM</i> .
<i>COMPARE</i>	Scalar variable of the same type and kind as <i>ATOM</i> .
<i>NEW</i>	Scalar variable of the same type as <i>ATOM</i> . If kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  logical(atomic_logical_kind) :: atom[*], prev
  call atomic_cas (atom[1], prev, .false., .true.)
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.37 \[ATOMIC_REF\]](#), page 136, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287

9.31 ATOMIC_DEFINE — Setting a variable atomically

Description:

ATOMIC_DEFINE(ATOM, VALUE) defines the variable *ATOM* with the value *VALUE* atomically. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of ISO_FORTRAN_ENV's STAT_STOPPED_IMAGE and if the remote image has failed, the value STAT_FAILED_IMAGE.

Standard: Fortran 2008 and later; with *STAT*, TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_DEFINE (ATOM, VALUE [, STAT])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of either integer type with <code>ATOMIC_INT_KIND</code> kind or logical type with <code>ATOMIC_LOGICAL_KIND</code> kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*]
  call atomic_define (atom[1], this_image())
end program atomic

```

See also: [Section 9.37 \[ATOMIC_REF\]](#), page 136, [Section 9.30 \[ATOMIC_CAS\]](#), page 131, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.28 \[ATOMIC_ADD\]](#), page 129, [Section 9.29 \[ATOMIC_AND\]](#), page 130, [Section 9.36 \[ATOMIC_OR\]](#), page 135, [Section 9.38 \[ATOMIC_XOR\]](#), page 136

9.32 ATOMIC_FETCH_ADD — Atomic ADD operation with prior fetch

Description:

`ATOMIC_FETCH_ADD(ATOM, VALUE, OLD)` atomically stores the value of *ATOM* in *OLD* and adds the value of *VAR* to the variable *ATOM*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of `ISO_FORTRAN_ENV`'s `STAT_STOPPED_IMAGE` and if the remote image has failed, the value `STAT_FAILED_IMAGE`.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: `CALL ATOMIC_FETCH_ADD (ATOM, VALUE, old [, STAT])`

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with <code>ATOMIC_INT_KIND</code> kind. <code>ATOMIC_LOGICAL_KIND</code> kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>OLD</i>	Scalar of the same type and kind as <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*], old

```

```

      call atomic_add (atom[1], this_image(), old)
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.28 \[ATOMIC_ADD\]](#), page 129, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.33 \[ATOMIC_FETCH_AND\]](#), page 133, [Section 9.34 \[ATOMIC_FETCH_OR\]](#), page 133, [Section 9.35 \[ATOMIC_FETCH_XOR\]](#), page 134

9.33 ATOMIC_FETCH_AND — Atomic bitwise AND operation with prior fetch

Description:

ATOMIC_AND(ATOM, VALUE) atomically stores the value of *ATOM* in *OLD* and defines *ATOM* with the bitwise AND between the values of *ATOM* and *VALUE*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of ISO_FORTRAN_ENV's STAT_STOPPED_IMAGE and if the remote image has failed, the value STAT_FAILED_IMAGE.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_FETCH_AND (ATOM, VALUE, OLD [, STAT])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with ATOMIC_INT_KIND kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>OLD</i>	Scalar of the same type and kind as <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*], old
  call atomic_fetch_and (atom[1], int(b'10100011101'), old)
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.29 \[ATOMIC_AND\]](#), page 130, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.32 \[ATOMIC_FETCH_ADD\]](#), page 132, [Section 9.34 \[ATOMIC_FETCH_OR\]](#), page 133, [Section 9.35 \[ATOMIC_FETCH_XOR\]](#), page 134

9.34 ATOMIC_FETCH_OR — Atomic bitwise OR operation with prior fetch

Description:

ATOMIC_OR(ATOM, VALUE) atomically stores the value of *ATOM* in *OLD* and defines *ATOM* with the bitwise OR between the values of *ATOM* and *VALUE*.

When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of *ISO_FORTRAN_ENV*'s *STAT_STOPPED_IMAGE* and if the remote image has failed, the value *STAT_FAILED_IMAGE*.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_FETCH_OR (*ATOM*, *VALUE*, *OLD* [, *STAT*])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with <i>ATOMIC_INT_KIND</i> kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>OLD</i>	Scalar of the same type and kind as <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*], old
  call atomic_fetch_or (atom[1], int(b'10100011101'), old)
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\], page 131](#), [Section 9.36 \[ATOMIC_OR\], page 135](#), [Section 10.1 \[ISO_FORTRAN_ENV\], page 287](#), [Section 9.32 \[ATOMIC_FETCH_ADD\], page 132](#), [Section 9.33 \[ATOMIC_FETCH_AND\], page 133](#), [Section 9.35 \[ATOMIC_FETCH_XOR\], page 134](#)

9.35 ATOMIC_FETCH_XOR — Atomic bitwise XOR operation with prior fetch

Description:

ATOMIC_XOR(*ATOM*, *VALUE*) atomically stores the value of *ATOM* in *OLD* and defines *ATOM* with the bitwise XOR between the values of *ATOM* and *VALUE*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of *ISO_FORTRAN_ENV*'s *STAT_STOPPED_IMAGE* and if the remote image has failed, the value *STAT_FAILED_IMAGE*.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_FETCH_XOR (*ATOM*, *VALUE*, *OLD* [, *STAT*])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with <i>ATOMIC_INT_KIND</i> kind.
-------------	--

<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>OLD</i>	Scalar of the same type and kind as <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*], old
  call atomic_fetch_xor (atom[1], int(b'10100011101'), old)
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.38 \[ATOMIC_XOR\]](#), page 136, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.32 \[ATOMIC_FETCH_ADD\]](#), page 132, [Section 9.33 \[ATOMIC_FETCH_AND\]](#), page 133, [Section 9.34 \[ATOMIC_FETCH_OR\]](#), page 133

9.36 ATOMIC_OR — Atomic bitwise OR operation

Description:

`ATOMIC_OR(ATOM, VALUE)` atomically defines *ATOM* with the bitwise AND between the values of *ATOM* and *VALUE*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a co-indexed *ATOM*, if the remote image has stopped, it is assigned the value of `ISO_FORTRAN_ENV`'s `STAT_STOPPED_IMAGE` and if the remote image has failed, the value `STAT_FAILED_IMAGE`.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: `CALL ATOMIC_OR (ATOM, VALUE [, STAT])`

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with <code>ATOMIC_INT_KIND</code> kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*]
  call atomic_or (atom[1], int(b'10100011101'))
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.34 \[ATOMIC_FETCH_OR\]](#), page 133, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.28 \[ATOMIC_ADD\]](#), page 129, [Section 9.36 \[ATOMIC_OR\]](#), page 135, [Section 9.38 \[ATOMIC_XOR\]](#), page 136

9.37 ATOMIC_REF — Obtaining the value of a variable atomically

Description:

ATOMIC_DEFINE(ATOM, VALUE) atomically assigns the value of the variable *ATOM* to *VALUE*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of ISO_FORTRAN_ENV's STAT_STOPPED_IMAGE and if the remote image has failed, the value STAT_FAILED_IMAGE.

Standard: Fortran 2008 and later; with *STAT*, TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_REF(VALUE, ATOM [, STAT])

Arguments:

<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>ATOM</i>	Scalar coarray or coindexed variable of either integer type with ATOMIC_INT_KIND kind or logical type with ATOMIC_LOGICAL_KIND kind.
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  logical(atomic_logical_kind) :: atom[*]
  logical :: val
  call atomic_ref (atom, .false.)
  ! ...
  call atomic_ref (atom, val)
  if (val) then
    print *, "Obtained"
  end if
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.30 \[ATOMIC-CAS\]](#), page 131, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.32 \[ATOMIC_FETCH_ADD\]](#), page 132, [Section 9.33 \[ATOMIC_FETCH_AND\]](#), page 133, [Section 9.34 \[ATOMIC_FETCH_OR\]](#), page 133, [Section 9.35 \[ATOMIC_FETCH_XOR\]](#), page 134

9.38 ATOMIC_XOR — Atomic bitwise OR operation

Description:

ATOMIC_AND(ATOM, VALUE) atomically defines *ATOM* with the bitwise XOR between the values of *ATOM* and *VALUE*. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value; in particular, for a coindexed *ATOM*, if the remote image has stopped, it is assigned the value of

ISO_FORTRAN_ENV's STAT_STOPPED_IMAGE and if the remote image has failed, the value STAT_FAILED_IMAGE.

Standard: TS 18508 or later

Class: Atomic subroutine

Syntax: CALL ATOMIC_XOR (ATOM, VALUE [, STAT])

Arguments:

<i>ATOM</i>	Scalar coarray or coindexed variable of integer type with ATOMIC_INT_KIND kind.
<i>VALUE</i>	Scalar of the same type as <i>ATOM</i> . If the kind is different, the value is converted to the kind of <i>ATOM</i> .
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  integer(atomic_int_kind) :: atom[*]
  call atomic_xor (atom[1], int(b'10100011101'))
end program atomic

```

See also: [Section 9.31 \[ATOMIC_DEFINE\]](#), page 131, [Section 9.35 \[ATOMIC_FETCH_XOR\]](#), page 134, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287, [Section 9.28 \[ATOMIC_ADD\]](#), page 129, [Section 9.36 \[ATOMIC_OR\]](#), page 135, [Section 9.38 \[ATOMIC_XOR\]](#), page 136

9.39 BACKTRACE — Show a backtrace

Description:

BACKTRACE shows a backtrace at an arbitrary place in user code. Program execution continues normally afterwards. The backtrace information is printed to the unit corresponding to ERROR_UNIT in ISO_FORTRAN_ENV.

Standard: GNU Extension

Class: Subroutine

Syntax: CALL BACKTRACE

Arguments:

None

See also: [Section 9.2 \[ABORT\]](#), page 111

9.40 BESSEL_J0 — Bessel function of the first kind of order 0

Description:

BESSEL_J0(X) computes the Bessel function of the first kind of order 0 of X. This function is available under the name BESJ0 as a GNU extension.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: **RESULT = BESSEL_J0(X)**

Arguments:

X The type shall be **REAL**.

Return value:

The return value is of type **REAL** and lies in the range $-0.4027... \leq Bessel(0, x) \leq 1$. It has the same kind as **X**.

Example:

```
program test_besj0
  real(8) :: x = 0.0_8
  x = bessej0(x)
end program test_besj0
```

Specific names:

Name	Argument	Return type	Standard
DBESJ0(X)	REAL(8) X	REAL(8)	GNU extension

9.41 BESSEL_J1 — Bessel function of the first kind of order 1

Description:

BESSEL_J1(X) computes the Bessel function of the first kind of order 1 of **X**. This function is available under the name **BESJ1** as a GNU extension.

Standard: Fortran 2008

Class: Elemental function

Syntax: **RESULT = BESSEL_J1(X)**

Arguments:

X The type shall be **REAL**.

Return value:

The return value is of type **REAL** and lies in the range $-0.5818... \leq Bessel(0, x) \leq 0.5818$. It has the same kind as **X**.

Example:

```
program test_besj1
  real(8) :: x = 1.0_8
  x = bessej1(x)
end program test_besj1
```

Specific names:

Name	Argument	Return type	Standard
DBESJ1(X)	REAL(8) X	REAL(8)	GNU extension

9.42 BESSEL_JN — Bessel function of the first kind

Description:

BESSEL_JN(N, X) computes the Bessel function of the first kind of order **N** of **X**. This function is available under the name **BESJN** as a GNU extension. If **N** and **X** are arrays, their ranks and shapes shall conform.

BESSEL_JN(N1, N2, X) returns an array with the Bessel functions of the first kind of the orders **N1** to **N2**.

Standard: Fortran 2008 and later, negative N is allowed as GNU extension

Class: Elemental function, except for the transformational function `BESSEL_JN(N1, N2, X)`

Syntax:

```
RESULT = BESSEL_JN(N, X)
RESULT = BESSEL_JN(N1, N2, X)
```

Arguments:

N	Shall be a scalar or an array of type <code>INTEGER</code> .
$N1$	Shall be a non-negative scalar of type <code>INTEGER</code> .
$N2$	Shall be a non-negative scalar of type <code>INTEGER</code> .
X	Shall be a scalar or an array of type <code>REAL</code> ; for <code>BESSEL_JN(N1, N2, X)</code> it shall be scalar.

Return value:

The return value is a scalar of type `REAL`. It has the same kind as X .

Note: The transformational function uses a recurrence algorithm which might, for some values of X , lead to different results than calls to the elemental function.

Example:

```
program test_besjn
  real(8) :: x = 1.0_8
  x = bessej_jn(5,x)
end program test_besjn
```

Specific names:

Name	Argument	Return type	Standard
<code>DBESJN(N, X)</code>	<code>INTEGER N</code> <code>REAL(8) X</code>	<code>REAL(8)</code>	GNU extension

9.43 BESSEL_Y0 — Bessel function of the second kind of order 0

Description:

`BESSEL_Y0(X)` computes the Bessel function of the second kind of order 0 of X . This function is available under the name `BESY0` as a GNU extension.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = BESSEL_Y0(X)`

Arguments:

X	The type shall be <code>REAL</code> .
-----	---------------------------------------

Return value:

The return value is of type `REAL`. It has the same kind as X .

Example:

```

program test_besy0
  real(8) :: x = 0.0_8
  x = besseli_0(x)
end program test_besy0

```

Specific names:

Name	Argument	Return type	Standard
DBESY0(X)	REAL(8) X	REAL(8)	GNU extension

9.44 BESSEL_Y1 — Bessel function of the second kind of order 1

Description:

BESSEL_Y1(X) computes the Bessel function of the second kind of order 1 of X. This function is available under the name BESY1 as a GNU extension.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = BESSEL_Y1(X)

Arguments:

X The type shall be REAL.

Return value:

The return value is of type REAL. It has the same kind as X.

Example:

```

program test_besy1
  real(8) :: x = 1.0_8
  x = besseli_1(x)
end program test_besy1

```

Specific names:

Name	Argument	Return type	Standard
DBESY1(X)	REAL(8) X	REAL(8)	GNU extension

9.45 BESSEL_YN — Bessel function of the second kind

Description:

BESSEL_YN(N, X) computes the Bessel function of the second kind of order N of X. This function is available under the name BESYN as a GNU extension. If N and X are arrays, their ranks and shapes shall conform.

BESSEL_YN(N1, N2, X) returns an array with the Bessel functions of the first kind of the orders N1 to N2.

Standard: Fortran 2008 and later, negative N is allowed as GNU extension

Class: Elemental function, except for the transformational function BESSEL_YN(N1, N2, X)

Syntax:

```

RESULT = BESSEL_YN(N, X)
RESULT = BESSEL_YN(N1, N2, X)

```

Arguments:

<i>N</i>	Shall be a scalar or an array of type <code>INTEGER</code> .
<i>N1</i>	Shall be a non-negative scalar of type <code>INTEGER</code> .
<i>N2</i>	Shall be a non-negative scalar of type <code>INTEGER</code> .
<i>X</i>	Shall be a scalar or an array of type <code>REAL</code> ; for <code>BESSEL_YN(N1, N2, X)</code> it shall be scalar.

Return value:

The return value is a scalar of type `REAL`. It has the same kind as *X*.

Note: The transformational function uses a recurrence algorithm which might, for some values of *X*, lead to different results than calls to the elemental function.

Example:

```

program test_besyn
  real(8) :: x = 1.0_8
  x = besse1_yn(5,x)
end program test_besyn

```

Specific names:

Name	Argument	Return type	Standard
<code>DBESYN(N,X)</code>	<code>INTEGER N</code> <code>REAL(8) X</code>	<code>REAL(8)</code>	GNU extension

9.46 BGE — Bitwise greater than or equal to

Description:

Determines whether an integral is a bitwise greater than or equal to another.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = BGE(I, J)`

Arguments:

<i>I</i>	Shall be of <code>INTEGER</code> type.
<i>J</i>	Shall be of <code>INTEGER</code> type, and of the same kind as <i>I</i> .

Return value:

The return value is of type `LOGICAL` and of the default kind.

See also: [Section 9.47 \[BGT\]](#), page 141, [Section 9.49 \[BLE\]](#), page 142, [Section 9.50 \[BLT\]](#), page 143

9.47 BGT — Bitwise greater than

Description:

Determines whether an integral is a bitwise greater than another.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = BGT(I, J)`

Arguments:

I Shall be of INTEGER type.
J Shall be of INTEGER type, and of the same kind as *I*.

Return value:

The return value is of type LOGICAL and of the default kind.

See also: Section 9.46 [BGE], page 141, Section 9.49 [BLE], page 142, Section 9.50 [BLT], page 143

9.48 BIT_SIZE — Bit size inquiry function

Description:

BIT_SIZE(*I*) returns the number of bits (integer precision plus sign bit) represented by the type of *I*. The result of BIT_SIZE(*I*) is independent of the actual value of *I*.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: RESULT = BIT_SIZE(*I*)

Arguments:

I The type shall be INTEGER.

Return value:

The return value is of type INTEGER

Example:

```
program test_bit_size
  integer :: i = 123
  integer :: size
  size = bit_size(i)
  print *, size
end program test_bit_size
```

9.49 BLE — Bitwise less than or equal to

Description:

Determines whether an integral is a bitwise less than or equal to another.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = BLE(*I*, *J*)

Arguments:

I Shall be of INTEGER type.
J Shall be of INTEGER type, and of the same kind as *I*.

Return value:

The return value is of type LOGICAL and of the default kind.

See also: Section 9.47 [BGT], page 141, Section 9.46 [BGE], page 141, Section 9.50 [BLT], page 143

9.50 BLT — Bitwise less than

Description:

Determines whether an integral is a bitwise less than another.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = BLT(I, J)

Arguments:

I Shall be of INTEGER type.
J Shall be of INTEGER type, and of the same kind as *I*.

Return value:

The return value is of type LOGICAL and of the default kind.

See also: [Section 9.46 \[BGE\], page 141](#), [Section 9.47 \[BGT\], page 141](#), [Section 9.49 \[BLE\], page 142](#)

9.51 BTEST — Bit test function

Description:

BTEST(I,POS) returns logical .TRUE. if the bit at *POS* in *I* is set. The counting of the bits starts at 0.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = BTEST(I, POS)

Arguments:

I The type shall be INTEGER.
POS The type shall be INTEGER.

Return value:

The return value is of type LOGICAL

Example:

```
program test_btest
  integer :: i = 32768 + 1024 + 64
  integer :: pos
  logical :: bool
  do pos=0,16
    bool = btest(i, pos)
    print *, pos, bool
  end do
end program test_btest
```

Specific names:

Name	Argument	Return type	Standard
BTEST(I,POS)	INTEGER I,POS	LOGICAL	F95 and later
BBTEST(I,POS)	INTEGER(1) I,POS	LOGICAL(1)	GNU extension

BITEST(I,POS)	INTEGER(2) I,POS	LOGICAL(2)	GNU extension
BJTEST(I,POS)	INTEGER(4) I,POS	LOGICAL(4)	GNU extension
BKTEST(I,POS)	INTEGER(8) I,POS	LOGICAL(8)	GNU extension

9.52 C_ASSOCIATED — Status of a C pointer

Description:

`C_ASSOCIATED(c_ptr_1[, c_ptr_2])` determines the status of the C pointer `c_ptr_1` or if `c_ptr_1` is associated with the target `c_ptr_2`.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: `RESULT = C_ASSOCIATED(c_ptr_1[, c_ptr_2])`

Arguments:

`c_ptr_1` Scalar of the type `C_PTR` or `C_FUNPTR`.
`c_ptr_2` (Optional) Scalar of the same type as `c_ptr_1`.

Return value:

The return value is of type `LOGICAL`; it is `.false.` if either `c_ptr_1` is a C NULL pointer or if `c_ptr_1` and `c_ptr_2` point to different addresses.

Example:

```
subroutine association_test(a,b)
  use iso_c_binding, only: c_associated, c_loc, c_ptr
  implicit none
  real, pointer :: a
  type(c_ptr) :: b
  if(c_associated(b, c_loc(a))) &
    stop 'b and a do not point to same target'
end subroutine association_test
```

See also: [Section 9.56 \[C_LOC\], page 146](#), [Section 9.55 \[C_FUNLOC\], page 146](#)

9.53 C_F_POINTER — Convert C into Fortran pointer

Description:

`C_F_POINTER(CPTR, FPTR[, SHAPE])` assigns the target of the C pointer `CPTR` to the Fortran pointer `FPTR` and specifies its shape.

Standard: Fortran 2003 and later

Class: Subroutine

Syntax: `CALL C_F_POINTER(CPTR, FPTR[, SHAPE])`

Arguments:

`CPTR` scalar of the type `C_PTR`. It is `INTENT(IN)`.
`FPTR` pointer interoperable with `cptr`. It is `INTENT(OUT)`.

SHAPE (Optional) Rank-one array of type `INTEGER` with `INTENT(IN)`. It shall be present if and only if *fptr* is an array. The size must be equal to the rank of *fptr*.

Example:

```
program main
  use iso_c_binding
  implicit none
  interface
    subroutine my_routine(p) bind(c,name='myC_func')
      import :: c_ptr
      type(c_ptr), intent(out) :: p
    end subroutine
  end interface
  type(c_ptr) :: cptr
  real, pointer :: a(:)
  call my_routine(cptr)
  call c_f_pointer(cptr, a, [12])
end program main
```

See also: [Section 9.56 \[C-LOC\], page 146](#), [Section 9.54 \[C_F_PROCPOINTER\], page 145](#)

9.54 C_F_PROCPOINTER — Convert C into Fortran procedure pointer

Description:

`C_F_PROCPOINTER(CPTR, FPTR)` Assign the target of the C function pointer *CPTR* to the Fortran procedure pointer *FPTR*.

Standard: Fortran 2003 and later

Class: Subroutine

Syntax: `CALL C_F_PROCPOINTER(cptr, fptr)`

Arguments:

CPTR scalar of the type `C_FUNPTR`. It is `INTENT(IN)`.
FPTR procedure pointer interoperable with *cptr*. It is `INTENT(OUT)`.

Example:

```
program main
  use iso_c_binding
  implicit none
  abstract interface
    function func(a)
      import :: c_float
      real(c_float), intent(in) :: a
      real(c_float) :: func
    end function
  end interface
  interface
    function getIterFunc() bind(c,name="getIterFunc")
      import :: c_funptr
      type(c_funptr) :: getIterFunc
    end function
  end interface
  type(c_funptr) :: cfunptr
```

```

    procedure(func), pointer :: myFunc
    cfunptr = getIterFunc()
    call c_f_procpointer(cfunptr, myFunc)
end program main

```

See also: [Section 9.56 \[C_LOC\], page 146](#), [Section 9.53 \[C_F_POINTER\], page 144](#)

9.55 C_FUNLOC — Obtain the C address of a procedure

Description:

C_FUNLOC(x) determines the C address of the argument.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: RESULT = C_FUNLOC(x)

Arguments:

x Interoperable function or pointer to such function.

Return value:

The return value is of type C_FUNPTR and contains the C address of the argument.

Example:

```

module x
  use iso_c_binding
  implicit none
contains
  subroutine sub(a) bind(c)
    real(c_float) :: a
    a = sqrt(a)+5.0
  end subroutine sub
end module x
program main
  use iso_c_binding
  use x
  implicit none
  interface
    subroutine my_routine(p) bind(c,name='myC_func')
      import :: c_funptr
      type(c_funptr), intent(in) :: p
    end subroutine
  end interface
  call my_routine(c_funloc(sub))
end program main

```

See also: [Section 9.52 \[C_ASSOCIATED\], page 144](#), [Section 9.56 \[C_LOC\], page 146](#), [Section 9.53 \[C_F_POINTER\], page 144](#), [Section 9.54 \[C_F_PROCPOINTER\], page 145](#)

9.56 C_LOC — Obtain the C address of an object

Description:

C_LOC(X) determines the C address of the argument.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: `RESULT = C_LOC(X)`

Arguments:

`X` Shall have either the `POINTER` or `TARGET` attribute. It shall not be a coindexed object. It shall either be a variable with interoperable type and kind type parameters, or be a scalar, nonpolymorphic variable with no length type parameters.

Return value:

The return value is of type `C_PTR` and contains the C address of the argument.

Example:

```
subroutine association_test(a,b)
  use iso_c_binding, only: c_associated, c_loc, c_ptr
  implicit none
  real, pointer :: a
  type(c_ptr) :: b
  if(c_associated(b, c_loc(a))) &
    stop 'b and a do not point to same target'
end subroutine association_test
```

See also: [Section 9.52 \[C_ASSOCIATED\]](#), page 144, [Section 9.55 \[C_FUNLOC\]](#), page 146, [Section 9.53 \[C_F_POINTER\]](#), page 144, [Section 9.54 \[C_F_PROCPONTER\]](#), page 145

9.57 C_SIZEOF — Size in bytes of an expression

Description:

`C_SIZEOF(X)` calculates the number of bytes of storage the expression `X` occupies.

Standard: Fortran 2008

Class: Inquiry function of the module `ISO_C_BINDING`

Syntax: `N = C_SIZEOF(X)`

Arguments:

`X` The argument shall be an interoperable data entity.

Return value:

The return value is of type integer and of the system-dependent kind `C_SIZE_T` (from the `ISO_C_BINDING` module). Its value is the number of bytes occupied by the argument. If the argument has the `POINTER` attribute, the number of bytes of the storage area pointed to is returned. If the argument is of a derived type with `POINTER` or `ALLOCATABLE` components, the return value does not account for the sizes of the data pointed to by these components.

Example:

```

use iso_c_binding
integer(c_int) :: i
real(c_float) :: r, s(5)
print *, (c_sizeof(s)/c_sizeof(r) == 5)
end

```

The example will print `.TRUE.` unless you are using a platform where default REAL variables are unusually padded.

See also: [Section 9.254 \[SIZEOF\], page 267](#), [Section 9.261 \[STORAGE_SIZE\], page 271](#)

9.58 CEILING — Integer ceiling function

Description:

CEILING(A) returns the least integer greater than or equal to A.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = CEILING(A [, KIND])

Arguments:

A	The type shall be REAL.
KIND	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type INTEGER(KIND) if KIND is present and a default-kind INTEGER otherwise.

Example:

```

program test_ceiling
  real :: x = 63.29
  real :: y = -63.59
  print *, ceiling(x) ! returns 64
  print *, ceiling(y) ! returns -63
end program test_ceiling

```

See also: [Section 9.109 \[FLOOR\], page 182](#), [Section 9.207 \[NINT\], page 240](#)

9.59 CHAR — Character conversion function

Description:

CHAR(I [, KIND]) returns the character represented by the integer I.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: RESULT = CHAR(I [, KIND])

Arguments:

I	The type shall be INTEGER.
KIND	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type CHARACTER(1)

Example:

```

program test_char
  integer :: i = 74
  character(1) :: c
  c = char(i)
  print *, i, c ! returns 'J'
end program test_char

```

Specific names:

Name	Argument	Return type	Standard
CHAR(I)	INTEGER I	CHARACTER(LEN=1)	F77 and later

Note: See [Section 9.143 \[ICHAR\]](#), page 203 for a discussion of converting between numerical values and formatted string representations.

See also: [Section 9.5 \[ACHAR\]](#), page 113, [Section 9.135 \[IACHAR\]](#), page 198, [Section 9.143 \[ICHAR\]](#), page 203

9.60 CHDIR — Change working directory

Description:

Change current working directory to a specified path.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```

CALL CHDIR(NAME [, STATUS])
STATUS = CHDIR(NAME)

```

Arguments:

<i>NAME</i>	The type shall be CHARACTER of default kind and shall specify a valid path within the file system.
<i>STATUS</i>	(Optional) INTEGER status flag of the default kind. Returns 0 on success, and a system specific and nonzero error code otherwise.

Example:

```

PROGRAM test_chdir
  CHARACTER(len=255) :: path
  CALL getcwd(path)
  WRITE(*,*) TRIM(path)
  CALL chdir("/tmp")
  CALL getcwd(path)
  WRITE(*,*) TRIM(path)
END PROGRAM

```

See also: [Section 9.124 \[GETCWD\]](#), page 192

9.61 CHMOD — Change access permissions of files

Description:

CHMOD changes the permissions of a file.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL CHMOD(NAME, MODE[, STATUS])
STATUS = CHMOD(NAME, MODE)
```

Arguments:

<i>NAME</i>	Scalar CHARACTER of default kind with the file name. Trailing blanks are ignored unless the character <code>achar(0)</code> is present, then all characters up to and excluding <code>achar(0)</code> are used as the file name.
<i>MODE</i>	Scalar CHARACTER of default kind giving the file permission. <i>MODE</i> uses the same syntax as the <code>chmod</code> utility as defined by the POSIX standard. The argument shall either be a string of a nonnegative octal number or a symbolic mode.
<i>STATUS</i>	(optional) scalar INTEGER , which is 0 on success and nonzero otherwise.

Return value:

In either syntax, *STATUS* is set to 0 on success and nonzero otherwise.

Example: CHMOD as subroutine

```
program chmod_test
  implicit none
  integer :: status
  call chmod('test.dat','u+x',status)
  print *, 'Status: ', status
end program chmod_test
```

CHMOD as function:

```
program chmod_test
  implicit none
  integer :: status
  status = chmod('test.dat','u+x')
  print *, 'Status: ', status
end program chmod_test
```

9.62 CMLX — Complex conversion function

Description:

CMLX(*X* [, *Y* [, *KIND*]]) returns a complex number where *X* is converted to the real component. If *Y* is present it is converted to the imaginary component.

If *Y* is not present then the imaginary component is set to 0.0. If *X* is complex then *Y* must not be present.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = CMPLX(X [, Y [, KIND]])`

Arguments:

<i>X</i>	The type may be <code>INTEGER</code> , <code>REAL</code> , or <code>COMPLEX</code> .
<i>Y</i>	(Optional; only allowed if <i>X</i> is not <code>COMPLEX</code> .) May be <code>INTEGER</code> or <code>REAL</code> .
<i>KIND</i>	(Optional) An <code>INTEGER</code> initialization expression indicating the kind parameter of the result.

Return value:

The return value is of `COMPLEX` type, with a kind equal to *KIND* if it is specified. If *KIND* is not specified, the result is of the default `COMPLEX` kind, regardless of the kinds of *X* and *Y*.

Example:

```
program test_cmplx
  integer :: i = 42
  real :: x = 3.14
  complex :: z
  z = cmplx(i, x)
  print *, z, cmplx(x)
end program test_cmplx
```

See also: [Section 9.71 \[COMPLEX\]](#), page 157

9.63 CO_BROADCAST — Copy a value to all images the current set of images

Description:

`CO_BROADCAST` copies the value of argument *A* on the image with image index `SOURCE_IMAGE` to all images in the current team. *A* becomes defined as if by intrinsic assignment. If the execution was successful and *STAT* is present, it is assigned the value zero. If the execution failed, *STAT* gets assigned a nonzero value and, if present, *ERRMSG* gets assigned a value describing the occurred error.

Standard: Technical Specification (TS) 18508 or later

Class: Collective subroutine

Syntax: `CALL CO_BROADCAST(A, SOURCE_IMAGE [, STAT, ERRMSG])`

Arguments:

<i>A</i>	<code>INTENT(INOUT)</code> argument; shall have the same dynamic type and type parameters on all images of the current team. If it is an array, it shall have the same shape on all images.
<i>SOURCE_IMAGE</i>	Scalar integer expression. It shall have the same value on all images and refer to an image of the current team.

STAT (optional) a scalar integer variable
ERRMSG (optional) a scalar character variable

Example:

```
program test
  integer :: val(3)
  if (this_image() == 1) then
    val = [1, 5, 3]
  end if
  call co_broadcast (val, source_image=1)
  print *, this_image, ":", val
end program test
```

See also: [Section 9.64 \[CO_MAX\], page 152](#), [Section 9.65 \[CO_MIN\], page 153](#),
[Section 9.67 \[CO_SUM\], page 155](#), [Section 9.66 \[CO_REDUCE\], page 153](#)

9.64 CO_MAX — Maximal value on the current set of images

Description:

CO_MAX determines element-wise the maximal value of *A* on all images of the current team. If *RESULT_IMAGE* is present, the maximum values are returned in *A* on the specified image only and the value of *A* on the other images become undefined. If *RESULT_IMAGE* is not present, the value is returned on all images. If the execution was successful and *STAT* is present, it is assigned the value zero. If the execution failed, *STAT* gets assigned a nonzero value and, if present, *ERRMSG* gets assigned a value describing the occurred error.

Standard: Technical Specification (TS) 18508 or later

Class: Collective subroutine

Syntax: CALL CO_MAX(*A* [, *RESULT_IMAGE*, *STAT*, *ERRMSG*])

Arguments:

A shall be an integer, real or character variable, which has the same type and type parameters on all images of the team.
RESULT_IMAGE (optional) a scalar integer expression; if present, it shall have the same the same value on all images and refer to an image of the current team.
STAT (optional) a scalar integer variable
ERRMSG (optional) a scalar character variable

Example:

```
program test
  integer :: val
  val = this_image ()
  call co_max (val, result_image=1)
  if (this_image() == 1) then
    write(*,*) "Maximal value", val ! prints num_images()
  end if
end program test
```

See also: [Section 9.65 \[CO_MIN\], page 153](#), [Section 9.67 \[CO_SUM\], page 155](#),
[Section 9.66 \[CO_REDUCE\], page 153](#), [Section 9.63 \[CO_BROADCAST\], page 151](#)

9.65 CO_MIN — Minimal value on the current set of images

Description:

CO_MIN determines element-wise the minimal value of *A* on all images of the current team. If *RESULT_IMAGE* is present, the minimal values are returned in *A* on the specified image only and the value of *A* on the other images become undefined. If *RESULT_IMAGE* is not present, the value is returned on all images. If the execution was successful and *STAT* is present, it is assigned the value zero. If the execution failed, *STAT* gets assigned a nonzero value and, if present, *ERRMSG* gets assigned a value describing the occurred error.

Standard: Technical Specification (TS) 18508 or later

Class: Collective subroutine

Syntax: CALL CO_MIN(*A* [, *RESULT_IMAGE*, *STAT*, *ERRMSG*])

Arguments:

<i>A</i>	shall be an integer, real or character variable, which has the same type and type parameters on all images of the team.
<i>RESULT_IMAGE</i>	(optional) a scalar integer expression; if present, it shall have the same the same value on all images and refer to an image of the current team.
<i>STAT</i>	(optional) a scalar integer variable
<i>ERRMSG</i>	(optional) a scalar character variable

Example:

```

program test
  integer :: val
  val = this_image ()
  call co_min (val, result_image=1)
  if (this_image() == 1) then
    write(*,*) "Minimal value", val ! prints 1
  end if
end program test

```

See also: [Section 9.64 \[CO_MAX\], page 152](#), [Section 9.67 \[CO_SUM\], page 155](#), [Section 9.66 \[CO_REDUCE\], page 153](#), [Section 9.63 \[CO_BROADCAST\], page 151](#)

9.66 CO_REDUCE — Reduction of values on the current set of images

Description:

CO_REDUCE determines element-wise the reduction of the value of *A* on all images of the current team. The pure function passed as *OPERATOR* is used to pairwise reduce the values of *A* by passing either the value of *A* of different images or the result values of such a reduction as argument. If *A* is an array, the deduction is done element wise. If *RESULT_IMAGE* is present, the result values are returned in *A* on the specified image only and the value of *A* on the other images become undefined. If *RESULT_IMAGE* is not present, the value is returned on all images. If the execution was successful and *STAT* is present,

it is assigned the value zero. If the execution failed, *STAT* gets assigned a nonzero value and, if present, *ERRMSG* gets assigned a value describing the occurred error.

Standard: Technical Specification (TS) 18508 or later

Class: Collective subroutine

Syntax: `CALL CO_REDUCE(A, OPERATOR, [, RESULT_IMAGE, STAT, ERRMSG])`

Arguments:

<i>A</i>	is an <code>INTENT(INOUT)</code> argument and shall be nonpolymorphic. If it is allocatable, it shall be allocated; if it is a pointer, it shall be associated. <i>A</i> shall have the same type and type parameters on all images of the team; if it is an array, it shall have the same shape on all images.
<i>OPERATOR</i>	pure function with two scalar nonallocatable arguments, which shall be nonpolymorphic and have the same type and type parameters as <i>A</i> . The function shall return a nonallocatable scalar of the same type and type parameters as <i>A</i> . The function shall be the same on all images and with regards to the arguments mathematically commutative and associative. Note that <i>OPERATOR</i> may not be an elemental function, unless it is an intrinsic function.
<i>RESULT_IMAGE</i>	(optional) a scalar integer expression; if present, it shall have the same the same value on all images and refer to an image of the current team.
<i>STAT</i>	(optional) a scalar integer variable
<i>ERRMSG</i>	(optional) a scalar character variable

Example:

```

program test
  integer :: val
  val = this_image ()
  call co_reduce (val, result_image=1, operator=myprod)
  if (this_image() == 1) then
    write(*,*) "Product value", val ! prints num_images() factorial
  end if
contains
  pure function myprod(a, b)
    integer, value :: a, b
    integer :: myprod
    myprod = a * b
  end function myprod
end program test

```

Note: While the rules permit in principle an intrinsic function, none of the intrinsics in the standard fulfill the criteria of having a specific function, which takes two arguments of the same type and returning that type as result.

See also: [Section 9.65 \[CO_MIN\], page 153](#), [Section 9.64 \[CO_MAX\], page 152](#), [Section 9.67 \[CO-SUM\], page 155](#), [Section 9.63 \[CO-BROADCAST\], page 151](#)

9.67 CO_SUM — Sum of values on the current set of images

Description:

CO_SUM sums up the values of each element of *A* on all images of the current team. If *RESULT_IMAGE* is present, the summed-up values are returned in *A* on the specified image only and the value of *A* on the other images become undefined. If *RESULT_IMAGE* is not present, the value is returned on all images. If the execution was successful and *STAT* is present, it is assigned the value zero. If the execution failed, *STAT* gets assigned a nonzero value and, if present, *ERRMSG* gets assigned a value describing the occurred error.

Standard: Technical Specification (TS) 18508 or later

Class: Collective subroutine

Syntax: CALL CO_MIN(*A* [, *RESULT_IMAGE*, *STAT*, *ERRMSG*])

Arguments:

<i>A</i>	shall be an integer, real or complex variable, which has the same type and type parameters on all images of the team.
<i>RESULT_IMAGE</i>	(optional) a scalar integer expression; if present, it shall have the same the same value on all images and refer to an image of the current team.
<i>STAT</i>	(optional) a scalar integer variable
<i>ERRMSG</i>	(optional) a scalar character variable

Example:

```

program test
  integer :: val
  val = this_image ()
  call co_sum (val, result_image=1)
  if (this_image() == 1) then
    write(*,*) "The sum is ", val ! prints (n**2 + n)/2, with n = num_images()
  end if
end program test

```

See also: [Section 9.64 \[CO_MAX\]](#), page 152, [Section 9.65 \[CO_MIN\]](#), page 153, [Section 9.66 \[CO_REDUCE\]](#), page 153, [Section 9.63 \[CO_BROADCAST\]](#), page 151

9.68 COMMAND_ARGUMENT_COUNT — Get number of command line arguments

Description:

COMMAND_ARGUMENT_COUNT returns the number of arguments passed on the command line when the containing program was invoked.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: RESULT = COMMAND_ARGUMENT_COUNT()

Arguments:

None

Return value:

The return value is an `INTEGER` of default kind.

Example:

```

program test_command_argument_count
  integer :: count
  count = command_argument_count()
  print *, count
end program test_command_argument_count

```

See also: [Section 9.122 \[GET_COMMAND\]](#), [page 191](#), [Section 9.123 \[GET_COMMAND_ARGUMENT\]](#), [page 191](#)

9.69 COMPILER_OPTIONS — Options passed to the compiler

Description:

`COMPILER_OPTIONS` returns a string with the options used for compiling.

Standard: Fortran 2008

Class: Inquiry function of the module `ISO_FORTRAN_ENV`

Syntax: `STR = COMPILER_OPTIONS()`

Arguments:

None.

Return value:

The return value is a default-kind string with system-dependent length. It contains the compiler flags used to compile the file, which called the `COMPILER_OPTIONS` intrinsic.

Example:

```

use iso_fortran_env
print '(4a)', 'This file was compiled by ', &
           compiler_version(), ' using the options ', &
           compiler_options()
end

```

See also: [Section 9.70 \[COMPILER_VERSION\]](#), [page 156](#), [Section 10.1 \[ISO_FORTRAN_ENV\]](#), [page 287](#)

9.70 COMPILER_VERSION — Compiler version string

Description:

`COMPILER_VERSION` returns a string with the name and the version of the compiler.

Standard: Fortran 2008

Class: Inquiry function of the module `ISO_FORTRAN_ENV`

Syntax: `STR = COMPILER_VERSION()`

Arguments:

None.

Return value:

The return value is a default-kind string with system-dependent length. It contains the name of the compiler and its version number.

Example:

```

use iso_fortran_env
print '(4a)', 'This file was compiled by ', &
    compiler_version(), ' using the options ', &
    compiler_options()
end

```

See also: [Section 9.69 \[COMPILER_OPTIONS\]](#), page 156, [Section 10.1 \[ISO_FORTRAN_ENV\]](#), page 287

9.71 COMPLEX — Complex conversion function

Description:

COMPLEX(*X*, *Y*) returns a complex number where *X* is converted to the real component and *Y* is converted to the imaginary component.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = COMPLEX(*X*, *Y*)

Arguments:

<i>X</i>	The type may be INTEGER or REAL.
<i>Y</i>	The type may be INTEGER or REAL.

Return value:

If *X* and *Y* are both of INTEGER type, then the return value is of default COMPLEX type.

If *X* and *Y* are of REAL type, or one is of REAL type and one is of INTEGER type, then the return value is of COMPLEX type with a kind equal to that of the REAL argument with the highest precision.

Example:

```

program test_complex
  integer :: i = 42
  real :: x = 3.14
  print *, complex(i, x)
end program test_complex

```

See also: [Section 9.62 \[CMPLX\]](#), page 150

9.72 CONJG — Complex conjugate function

Description:

CONJG(*Z*) returns the conjugate of *Z*. If *Z* is (*x*, *y*) then the result is (*x*, *-y*)

Standard: Fortran 77 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: Z = CONJG(*Z*)

Arguments:

Z The type shall be `COMPLEX`.

Return value:

The return value is of type `COMPLEX`.

Example:

```
program test_conjg
  complex :: z = (2.0, 3.0)
  complex(8) :: dz = (2.71_8, -3.14_8)
  z = conjg(z)
  print *, z
  dz = dconjg(dz)
  print *, dz
end program test_conjg
```

Specific names:

Name	Argument	Return type	Standard
<code>CONJG(Z)</code>	<code>COMPLEX Z</code>	<code>COMPLEX</code>	GNU extension
<code>DCONJG(Z)</code>	<code>COMPLEX(8) Z</code>	<code>COMPLEX(8)</code>	GNU extension

9.73 COS — Cosine function

Description:

`COS(X)` computes the cosine of *X*.

Standard: Fortran 77 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = COS(X)`

Arguments:

X The type shall be `REAL` or `COMPLEX`.

Return value:

The return value is of the same type and kind as *X*. The real part of the result is in radians. If *X* is of the type `REAL`, the return value lies in the range $-1 \leq \cos(x) \leq 1$.

Example:

```
program test_cos
  real :: x = 0.0
  x = cos(x)
end program test_cos
```

Specific names:

Name	Argument	Return type	Standard
<code>COS(X)</code>	<code>REAL(4) X</code>	<code>REAL(4)</code>	Fortran 77 and later
<code>DCOS(X)</code>	<code>REAL(8) X</code>	<code>REAL(8)</code>	Fortran 77 and later
<code>CCOS(X)</code>	<code>COMPLEX(4) X</code>	<code>COMPLEX(4)</code>	Fortran 77 and later
<code>ZCOS(X)</code>	<code>COMPLEX(8) X</code>	<code>COMPLEX(8)</code>	GNU extension
<code>CDCOS(X)</code>	<code>COMPLEX(8) X</code>	<code>COMPLEX(8)</code>	GNU extension

See also: Inverse function: [Section 9.6 \[ACOS\]](#), page 114 Degrees function: [Section 9.74 \[COSD\]](#), page 159

9.74 COSD — Cosine function, degrees

Description:

`COSD(X)` computes the cosine of X in degrees.

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘`-fdec-math`’.

Class: Elemental function

Syntax: `RESULT = COSD(X)`

Arguments:

X The type shall be `REAL` or `COMPLEX`.

Return value:

The return value is of the same type and kind as X . The real part of the result is in degrees. If X is of the type `REAL`, the return value lies in the range $-1 \leq \text{cosd}(x) \leq 1$.

Example:

```
program test_cosd
  real :: x = 0.0
  x = cosd(x)
end program test_cosd
```

Specific names:

Name	Argument	Return type	Standard
<code>COSD(X)</code>	<code>REAL(4) X</code>	<code>REAL(4)</code>	GNU Extension
<code>DCOSD(X)</code>	<code>REAL(8) X</code>	<code>REAL(8)</code>	GNU Extension
<code>CCOSD(X)</code>	<code>COMPLEX(4) X</code>	<code>COMPLEX(4)</code>	GNU Extension
<code>ZCOSD(X)</code>	<code>COMPLEX(8) X</code>	<code>COMPLEX(8)</code>	GNU extension
<code>CDCOSD(X)</code>	<code>COMPLEX(8) X</code>	<code>COMPLEX(8)</code>	GNU extension

See also: Inverse function: [Section 9.7 \[ACOSD\]](#), page 115 Radians function: [Section 9.73 \[COS\]](#), page 158

9.75 COSH — Hyperbolic cosine function

Description:

`COSH(X)` computes the hyperbolic cosine of X .

Standard: Fortran 77 and later, for a complex argument Fortran 2008 or later

Class: Elemental function

Syntax: `X = COSH(X)`

Arguments:

X The type shall be `REAL` or `COMPLEX`.

Return value:

The return value has same type and kind as X . If X is complex, the imaginary part of the result is in radians. If X is `REAL`, the return value has a lower bound of one, $\cosh(x) \geq 1$.

Example:

```
program test_cosh
  real(8) :: x = 1.0_8
  x = cosh(x)
end program test_cosh
```

Specific names:

Name	Argument	Return type	Standard
COSH(X)	REAL(4) X	REAL(4)	Fortran 77 and later
DCOSH(X)	REAL(8) X	REAL(8)	Fortran 77 and later

See also: Inverse function: [Section 9.8 \[ACOSH\]](#), page 115

9.76 COTAN — Cotangent function

Description:

COTAN(X) computes the cotangent of X. Equivalent to COS(x) divided by SIN(x), or 1 / TAN(x).

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘-fdec-math’.

Class: Elemental function

Syntax: RESULT = COTAN(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has same type and kind as X, and its value is in radians.

Example:

```
program test_cotan
  real(8) :: x = 0.165_8
  x = cotan(x)
end program test_cotan
```

Specific names:

Name	Argument	Return type	Standard
COTAN(X)	REAL(4) X	REAL(4)	GNU Extension
DCOTAN(X)	REAL(8) X	REAL(8)	GNU Extension

See also: Converse function: [Section 9.266 \[TAN\]](#), page 274 Degrees function: [Section 9.77 \[COTAND\]](#), page 160

9.77 COTAND — Cotangent function, degrees

Description:

COTAND(X) computes the cotangent of X in degrees. Equivalent to COSD(x) divided by SIND(x), or 1 / TAND(x).

Standard: GNU Extension, enabled with ‘`-fdec-math`’.

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Class: Elemental function

Syntax: `RESULT = COTAND(X)`

Arguments:

`X` The type shall be `REAL` or `COMPLEX`.

Return value:

The return value has same type and kind as `X`, and its value is in degrees.

Example:

```
program test_cotand
  real(8) :: x = 0.165_8
  x = cotand(x)
end program test_cotand
```

Specific names:

Name	Argument	Return type	Standard
<code>COTAND(X)</code>	<code>REAL(4) X</code>	<code>REAL(4)</code>	GNU Extension
<code>DCOTAND(X)</code>	<code>REAL(8) X</code>	<code>REAL(8)</code>	GNU Extension

See also: Converse function: [Section 9.267 \[TAND\]](#), [page 275](#) Radians function: [Section 9.76 \[COTAN\]](#), [page 160](#)

9.78 COUNT — Count function

Description:

Counts the number of `.TRUE.` elements in a logical `MASK`, or, if the `DIM` argument is supplied, counts the number of elements along each row of the array in the `DIM` direction. If the array has zero size, or all of the elements of `MASK` are `.FALSE.`, then the result is 0.

Standard: Fortran 95 and later, with `KIND` argument Fortran 2003 and later

Class: Transformational function

Syntax: `RESULT = COUNT(MASK [, DIM, KIND])`

Arguments:

`MASK` The type shall be `LOGICAL`.
`DIM` (Optional) The type shall be `INTEGER`.
`KIND` (Optional) An `INTEGER` initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind `KIND`. If `KIND` is absent, the return value is of default integer kind. If `DIM` is present, the result is an array with a rank one less than the rank of `ARRAY`, and a size corresponding to the shape of `ARRAY` with the `DIM` dimension removed.

Example:

```

program test_count
  integer, dimension(2,3) :: a, b
  logical, dimension(2,3) :: mask
  a = reshape( (/ 1, 2, 3, 4, 5, 6 /), (/ 2, 3 /))
  b = reshape( (/ 0, 7, 3, 4, 5, 8 /), (/ 2, 3 /))
  print '(3i3)', a(1,:)
  print '(3i3)', a(2,:)
  print *
  print '(3i3)', b(1,:)
  print '(3i3)', b(2,:)
  print *
  mask = a.ne.b
  print '(3l3)', mask(1,:)
  print '(3l3)', mask(2,:)
  print *
  print '(3i3)', count(mask)
  print *
  print '(3i3)', count(mask, 1)
  print *
  print '(3i3)', count(mask, 2)
end program test_count

```

9.79 CPU_TIME — CPU elapsed time in seconds

Description:

Returns a **REAL** value representing the elapsed CPU time in seconds. This is useful for testing segments of code to determine execution time.

If a time source is available, time will be reported with microsecond resolution. If no time source is available, *TIME* is set to -1.0.

Note that *TIME* may contain a, system dependent, arbitrary offset and may not start with 0.0. For **CPU_TIME**, the absolute value is meaningless, only differences between subsequent calls to this subroutine, as shown in the example below, should be used.

Standard: Fortran 95 and later

Class: Subroutine

Syntax: **CALL CPU_TIME(TIME)**

Arguments:

TIME The type shall be **REAL** with **INTENT(OUT)**.

Return value:

None

Example:

```

program test_cpu_time
  real :: start, finish
  call cpu_time(start)
  ! put code to test here
  call cpu_time(finish)
  print '("Time = ",f6.3," seconds.")',finish-start
end program test_cpu_time

```

See also: [Section 9.265 \[SYSTEM_CLOCK\]](#), page 273, [Section 9.82 \[DATE_AND_TIME\]](#), page 164

9.80 CSHIFT — Circular shift elements of an array

Description:

CSHIFT(*ARRAY*, SHIFT [, *DIM*]) performs a circular shift on elements of *ARRAY* along the dimension of *DIM*. If *DIM* is omitted it is taken to be 1. *DIM* is a scalar of type INTEGER in the range of $1 \leq DIM \leq n$ where *n* is the rank of *ARRAY*. If the rank of *ARRAY* is one, then all elements of *ARRAY* are shifted by *SHIFT* places. If rank is greater than one, then all complete rank one sections of *ARRAY* along the given dimension are shifted. Elements shifted out one end of each rank one section are shifted back in the other end.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = CSHIFT(*ARRAY*, SHIFT [, *DIM*])

Arguments:

<i>ARRAY</i>	Shall be an array of any type.
<i>SHIFT</i>	The type shall be INTEGER.
<i>DIM</i>	The type shall be INTEGER.

Return value:

Returns an array of same type and rank as the *ARRAY* argument.

Example:

```

program test_cshift
  integer, dimension(3,3) :: a
  a = reshape( (/ 1, 2, 3, 4, 5, 6, 7, 8, 9 /), (/ 3, 3 /))
  print '(3i3)', a(1,:)
  print '(3i3)', a(2,:)
  print '(3i3)', a(3,:)
  a = cshift(a, SHIFT=(/1, 2, -1/), DIM=2)
  print *
  print '(3i3)', a(1,:)
  print '(3i3)', a(2,:)
  print '(3i3)', a(3,:)
end program test_cshift

```

9.81 CTIME — Convert a time into a string

Description:

CTIME converts a system time value, such as returned by [Section 9.271 \[TIME8\]](#), [page 278](#), to a string. The output will be of the form 'Sat Aug 19 18:13:14 1995'.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```

CALL CTIME(TIME, RESULT).
RESULT = CTIME(TIME).

```

Arguments:

<i>TIME</i>	The type shall be of type <code>INTEGER</code> .
<i>RESULT</i>	The type shall be of type <code>CHARACTER</code> and of default kind. It is an <code>INTENT(OUT)</code> argument. If the length of this variable is too short for the time and date string to fit completely, it will be blank on procedure return.

Return value:

The converted date and time as a string.

Example:

```

program test_ctime
  integer(8) :: i
  character(len=30) :: date
  i = time8()

  ! Do something, main part of the program

  call ctime(i,date)
  print *, 'Program was started on ', date
end program test_ctime

```

See Also: [Section 9.82 \[DATE_AND_TIME\]](#), page 164, [Section 9.131 \[GMTIME\]](#), page 196, [Section 9.184 \[LTIME\]](#), page 226, [Section 9.270 \[TIME\]](#), page 277, [Section 9.271 \[TIMES\]](#), page 278

9.82 DATE_AND_TIME — Date and time subroutine

Description:

`DATE_AND_TIME`(`DATE`, `TIME`, `ZONE`, `VALUES`) gets the corresponding date and time information from the real-time system clock. *DATE* is `INTENT(OUT)` and has form `ccyymmdd`. *TIME* is `INTENT(OUT)` and has form `hhmmss.sss`. *ZONE* is `INTENT(OUT)` and has form `(+)-hhmm`, representing the difference with respect to Coordinated Universal Time (UTC). Unavailable time and date parameters return blanks.

VALUES is `INTENT(OUT)` and provides the following:

<code>VALUE(1):</code>	The year
<code>VALUE(2):</code>	The month
<code>VALUE(3):</code>	The day of the month
<code>VALUE(4):</code>	Time difference with UTC in minutes
<code>VALUE(5):</code>	The hour of the day
<code>VALUE(6):</code>	The minutes of the hour
<code>VALUE(7):</code>	The seconds of the minute
<code>VALUE(8):</code>	The milliseconds of the second

Standard: Fortran 95 and later

Class: Subroutine

Syntax: `CALL DATE_AND_TIME([DATE, TIME, ZONE, VALUES])`

Arguments:

<i>DATE</i>	(Optional) The type shall be CHARACTER(LEN=8) or larger, and of default kind.
<i>TIME</i>	(Optional) The type shall be CHARACTER(LEN=10) or larger, and of default kind.
<i>ZONE</i>	(Optional) The type shall be CHARACTER(LEN=5) or larger, and of default kind.
<i>VALUES</i>	(Optional) The type shall be INTEGER(8).

Return value:

None

Example:

```

program test_time_and_date
  character(8)  :: date
  character(10) :: time
  character(5)  :: zone
  integer,dimension(8) :: values
  ! using keyword arguments
  call date_and_time(date,time,zone,values)
  call date_and_time(DATE=date,ZONE=zone)
  call date_and_time(TIME=time)
  call date_and_time(VALUES=values)
  print '(a,2x,a,2x,a)', date, time, zone
  print '(8i5)', values
end program test_time_and_date

```

See also: [Section 9.79 \[CPU_TIME\], page 162](#), [Section 9.265 \[SYSTEM_CLOCK\], page 273](#)

9.83 DBLE — Double conversion function

Description:

DBLE(A) Converts A to double precision real type.

Standard: Fortran 77 and later*Class:* Elemental function*Syntax:* RESULT = DBLE(A)*Arguments:*

A The type shall be INTEGER, REAL, or COMPLEX.

Return value:

The return value is of type double precision real.

Example:

```

program test_dble
  real      :: x = 2.18
  integer   :: i = 5
  complex   :: z = (2.3,1.14)
  print *, dble(x), dble(i), dble(z)
end program test_dble

```

See also: [Section 9.229 \[REAL\], page 252](#)

9.84 DCMPLX — Double complex conversion function

Description:

DCMPLX(X [,Y]) returns a double complex number where X is converted to the real component. If Y is present it is converted to the imaginary component. If Y is not present then the imaginary component is set to 0.0. If X is complex then Y must not be present.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = DCMPLX(X [, Y])

Arguments:

X	The type may be INTEGER, REAL, or COMPLEX.
Y	(Optional if X is not COMPLEX.) May be INTEGER or REAL.

Return value:

The return value is of type COMPLEX(8)

Example:

```

program test_dcplx
  integer :: i = 42
  real :: x = 3.14
  complex :: z
  z = cmplx(i, x)
  print *, dcplx(i)
  print *, dcplx(x)
  print *, dcplx(z)
  print *, dcplx(x,i)
end program test_dcplx

```

9.85 DIGITS — Significant binary digits function

Description:

DIGITS(X) returns the number of significant binary digits of the internal model representation of X. For example, on a system using a 32-bit floating point representation, a default real number would likely return 24.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: RESULT = DIGITS(X)

Arguments:

X	The type may be INTEGER or REAL.
---	----------------------------------

Return value:

The return value is of type INTEGER.

Example:

```

program test_digits
  integer :: i = 12345
  real :: x = 3.143

```

```

      real(8) :: y = 2.33
      print *, digits(i)
      print *, digits(x)
      print *, digits(y)
end program test_digits

```

9.86 DIM — Positive difference

Description:

DIM(X,Y) returns the difference X-Y if the result is positive; otherwise returns zero.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: RESULT = DIM(X, Y)

Arguments:

X	The type shall be INTEGER or REAL
Y	The type shall be the same type and kind as X.

Return value:

The return value is of type INTEGER or REAL.

Example:

```

program test_dim
  integer :: i
  real(8) :: x
  i = dim(4, 15)
  x = dim(4.345_8, 2.111_8)
  print *, i
  print *, x
end program test_dim

```

Specific names:

Name	Argument	Return type	Standard
DIM(X,Y)	REAL(4) X, Y	REAL(4)	Fortran 77 and later
IDIM(X,Y)	INTEGER(4) X, Y	INTEGER(4)	Fortran 77 and later
DDIM(X,Y)	REAL(8) X, Y	REAL(8)	Fortran 77 and later

9.87 DOT_PRODUCT — Dot product function

Description:

DOT_PRODUCT(VECTOR_A, VECTOR_B) computes the dot product multiplication of two vectors VECTOR_A and VECTOR_B. The two vectors may be either numeric or logical and must be arrays of rank one and of equal size. If the vectors are INTEGER or REAL, the result is SUM(VECTOR_A*VECTOR_B). If the vectors are COMPLEX, the result is SUM(CONJG(VECTOR_A)*VECTOR_B). If the vectors are LOGICAL, the result is ANY(VECTOR_A .AND. VECTOR_B).

Standard: Fortran 95 and later

Class: Transformational function

Syntax: **RESULT = DOT_PRODUCT(VECTOR_A, VECTOR_B)**

Arguments:

VECTOR_A The type shall be numeric or **LOGICAL**, rank 1.
VECTOR_B The type shall be numeric if **VECTOR_A** is of numeric type
 or **LOGICAL** if **VECTOR_A** is of type **LOGICAL**. **VECTOR_B**
 shall be a rank-one array.

Return value:

If the arguments are numeric, the return value is a scalar of numeric type, **INTEGER**, **REAL**, or **COMPLEX**. If the arguments are **LOGICAL**, the return value is **.TRUE.** or **.FALSE..**

Example:

```
program test_dot_prod
  integer, dimension(3) :: a, b
  a = (/ 1, 2, 3 /)
  b = (/ 4, 5, 6 /)
  print '(3i3)', a
  print *
  print '(3i3)', b
  print *
  print *, dot_product(a,b)
end program test_dot_prod
```

9.88 DPROD — Double product function

Description:

DPROD(X,Y) returns the product **X*Y**.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: **RESULT = DPROD(X, Y)**

Arguments:

X The type shall be **REAL**.
Y The type shall be **REAL**.

Return value:

The return value is of type **REAL(8)**.

Example:

```
program test_dprod
  real :: x = 5.2
  real :: y = 2.3
  real(8) :: d
  d = dprod(x,y)
  print *, d
end program test_dprod
```

Specific names:

Name	Argument	Return type	Standard
DPROD(X,Y)	REAL(4) X, Y	REAL(8)	Fortran 77 and later

9.89 DREAL — Double real part function

Description:

DREAL(*Z*) returns the real part of complex variable *Z*.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = DREAL(*A*)

Arguments:

A The type shall be COMPLEX(8).

Return value:

The return value is of type REAL(8).

Example:

```
program test_dreal
  complex(8) :: z = (1.3_8,7.2_8)
  print *, dreal(z)
end program test_dreal
```

See also: [Section 9.11 \[AIMAG\], page 117](#)

9.90 DSHIFTL — Combined left shift

Description:

DSHIFTL(*I*, *J*, SHIFT) combines bits of *I* and *J*. The rightmost *SHIFT* bits of the result are the leftmost *SHIFT* bits of *J*, and the remaining bits are the rightmost bits of *I*.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = DSHIFTL(*I*, *J*, SHIFT)

Arguments:

<i>I</i>	Shall be of type INTEGER or a BOZ constant.
<i>J</i>	Shall be of type INTEGER or a BOZ constant. If both <i>I</i> and <i>J</i> have integer type, then they shall have the same kind type parameter. <i>I</i> and <i>J</i> shall not both be BOZ constants.
<i>SHIFT</i>	Shall be of type INTEGER. It shall be nonnegative. If <i>I</i> is not a BOZ constant, then <i>SHIFT</i> shall be less than or equal to BIT_SIZE(<i>I</i>); otherwise, <i>SHIFT</i> shall be less than or equal to BIT_SIZE(<i>J</i>).

Return value:

If either *I* or *J* is a BOZ constant, it is first converted as if by the intrinsic function INT to an integer type with the kind type parameter of the other.

See also: [Section 9.91 \[DSHIFTR\], page 170](#)

9.91 DSHIFTR — Combined right shift

Description:

DSHIFTR(*I*, *J*, *SHIFT*) combines bits of *I* and *J*. The leftmost *SHIFT* bits of the result are the rightmost *SHIFT* bits of *I*, and the remaining bits are the leftmost bits of *J*.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = DSHIFTR(*I*, *J*, *SHIFT*)

Arguments:

<i>I</i>	Shall be of type <code>INTEGER</code> or a <code>BOZ</code> constant.
<i>J</i>	Shall be of type <code>INTEGER</code> or a <code>BOZ</code> constant. If both <i>I</i> and <i>J</i> have integer type, then they shall have the same kind type parameter. <i>I</i> and <i>J</i> shall not both be <code>BOZ</code> constants.
<i>SHIFT</i>	Shall be of type <code>INTEGER</code> . It shall be nonnegative. If <i>I</i> is not a <code>BOZ</code> constant, then <i>SHIFT</i> shall be less than or equal to <code>BIT_SIZE(I)</code> ; otherwise, <i>SHIFT</i> shall be less than or equal to <code>BIT_SIZE(J)</code> .

Return value:

If either *I* or *J* is a `BOZ` constant, it is first converted as if by the intrinsic function `INT` to an integer type with the kind type parameter of the other.

See also: [Section 9.90 \[DSHIFTL\], page 169](#)

9.92 DTIME — Execution time subroutine (or function)

Description:

DTIME(*VALUES*, *TIME*) initially returns the number of seconds of runtime since the start of the process's execution in *TIME*. *VALUES* returns the user and system components of this time in *VALUES*(1) and *VALUES*(2) respectively. *TIME* is equal to *VALUES*(1) + *VALUES*(2).

Subsequent invocations of DTIME return values accumulated since the previous invocation.

On some systems, the underlying timings are represented using types with sufficiently small limits that overflows (wrap around) are possible, such as 32-bit types. Therefore, the values returned by this intrinsic might be, or become, negative, or numerically less than previous values, during a single run of the compiled program.

Please note, that this implementation is thread safe if used within OpenMP directives, i.e., its state will be consistent while called from multiple threads. However, if DTIME is called from multiple threads, the result is still the time since the last invocation. This may not give the intended results. If possible, use `CPU_TIME` instead.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

VALUES and *TIME* are *INTENT(OUT)* and provide the following:

<i>VALUES</i> (1):	User time in seconds.
<i>VALUES</i> (2):	System time in seconds.
<i>TIME</i> :	Run time since start in seconds.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL DTIME(VALUES, TIME).
TIME = DTIME(VALUES), (not recommended).
```

Arguments:

<i>VALUES</i>	The type shall be <i>REAL(4)</i> , <i>DIMENSION(2)</i> .
<i>TIME</i>	The type shall be <i>REAL(4)</i> .

Return value:

Elapsed time in seconds since the last invocation or since the start of program execution if not called before.

Example:

```
program test_dtime
  integer(8) :: i, j
  real, dimension(2) :: tarray
  real :: result
  call dtime(tarray, result)
  print *, result
  print *, tarray(1)
  print *, tarray(2)
  do i=1,100000000    ! Just a delay
    j = i * i - i
  end do
  call dtime(tarray, result)
  print *, result
  print *, tarray(1)
  print *, tarray(2)
end program test_dtime
```

See also: [Section 9.79 \[CPU_TIME\]](#), page 162

9.93 EOSHIFT — End-off shift elements of an array

Description:

EOSHIFT(*ARRAY*, *SHIFT*[, *BOUNDARY*, *DIM*]) performs an end-off shift on elements of *ARRAY* along the dimension of *DIM*. If *DIM* is omitted it is taken to be 1. *DIM* is a scalar of type *INTEGER* in the range of $1 \leq DIM \leq n$ where *n* is the rank of *ARRAY*. If the rank of *ARRAY* is one, then all elements of *ARRAY* are shifted by *SHIFT* places. If rank is greater than one, then all complete rank one sections of *ARRAY* along the given dimension are shifted. Elements shifted out one end of each rank one section are dropped. If *BOUNDARY* is present then the corresponding value of from *BOUNDARY* is copied back in the other end. If *BOUNDARY* is not present then the following are copied in depending on the type of *ARRAY*.

<i>Array Type</i>	<i>Boundary Value</i>
Numeric	0 of the type and kind of <i>ARRAY</i> .
Logical	.FALSE..
Character(<i>len</i>)	<i>len</i> blanks.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: `RESULT = EOSHIFT(ARRAY, SHIFT [, BOUNDARY, DIM])`

Arguments:

<i>ARRAY</i>	May be any type, not scalar.
<i>SHIFT</i>	The type shall be <code>INTEGER</code> .
<i>BOUNDARY</i>	Same type as <i>ARRAY</i> .
<i>DIM</i>	The type shall be <code>INTEGER</code> .

Return value:

Returns an array of same type and rank as the *ARRAY* argument.

Example:

```

program test_eoshift
  integer, dimension(3,3) :: a
  a = reshape( (/ 1, 2, 3, 4, 5, 6, 7, 8, 9 /), (/ 3, 3 /))
  print '(3i3)', a(1,:)
  print '(3i3)', a(2,:)
  print '(3i3)', a(3,:)
  a = EOSHIFT(a, SHIFT=(/1, 2, 1/), BOUNDARY=-5, DIM=2)
  print *
  print '(3i3)', a(1,:)
  print '(3i3)', a(2,:)
  print '(3i3)', a(3,:)
end program test_eoshift

```

9.94 EPSILON — Epsilon function

Description:

`EPSILON(X)` returns the smallest number *E* of the same kind as *X* such that $1 + E > 1$.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: `RESULT = EPSILON(X)`

Arguments:

<i>X</i>	The type shall be <code>REAL</code> .
----------	---------------------------------------

Return value:

The return value is of same type as the argument.

Example:

```

program test_epsilon
  real :: x = 3.143
  real(8) :: y = 2.33
  print *, EPSILON(x)
  print *, EPSILON(y)
end program test_epsilon

```

9.95 ERF — Error function

Description:

ERF(X) computes the error function of X.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = ERF(X)

Arguments:

X The type shall be REAL.

Return value:

The return value is of type REAL, of the same kind as X and lies in the range $-1 \leq \text{erf}(x) \leq 1$.

Example:

```
program test_erf
  real(8) :: x = 0.17_8
  x = erf(x)
end program test_erf
```

Specific names:

Name	Argument	Return type	Standard
DERF(X)	REAL(8) X	REAL(8)	GNU extension

9.96 ERFC — Error function

Description:

ERFC(X) computes the complementary error function of X.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = ERFC(X)

Arguments:

X The type shall be REAL.

Return value:

The return value is of type REAL and of the same kind as X. It lies in the range $0 \leq \text{erfc}(x) \leq 2$.

Example:

```
program test_erfc
  real(8) :: x = 0.17_8
  x = erfc(x)
end program test_erfc
```

Specific names:

Name	Argument	Return type	Standard
DERFC(X)	REAL(8) X	REAL(8)	GNU extension

9.97 ERFC_SCALED — Error function

Description:

ERFC_SCALED(*X*) computes the exponentially-scaled complementary error function of *X*.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = ERFC_SCALED(*X*)

Arguments:

X The type shall be REAL.

Return value:

The return value is of type REAL and of the same kind as *X*.

Example:

```
program test_erfc_scaled
  real(8) :: x = 0.17_8
  x = erfc_scaled(x)
end program test_erfc_scaled
```

9.98 ETIME — Execution time subroutine (or function)

Description:

ETIME(*VALUES*, *TIME*) returns the number of seconds of runtime since the start of the process's execution in *TIME*. *VALUES* returns the user and system components of this time in *VALUES*(1) and *VALUES*(2) respectively. *TIME* is equal to *VALUES*(1) + *VALUES*(2).

On some systems, the underlying timings are represented using types with sufficiently small limits that overflows (wrap around) are possible, such as 32-bit types. Therefore, the values returned by this intrinsic might be, or become, negative, or numerically less than previous values, during a single run of the compiled program.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

VALUES and *TIME* are INTENT(OUT) and provide the following:

<i>VALUES</i> (1):	User time in seconds.
<i>VALUES</i> (2):	System time in seconds.
<i>TIME</i> :	Run time since start in seconds.

Standard: GNU extension

Class: Subroutine, function

Syntax:

CALL ETIME(*VALUES*, *TIME*).
TIME = ETIME(*VALUES*), (not recommended).

Arguments:

VALUES The type shall be REAL(4), DIMENSION(2).
TIME The type shall be REAL(4).

Return value:

Elapsed time in seconds since the start of program execution.

Example:

```

program test_etime
  integer(8) :: i, j
  real, dimension(2) :: tarray
  real :: result
  call ETIME(tarray, result)
  print *, result
  print *, tarray(1)
  print *, tarray(2)
  do i=1,100000000    ! Just a delay
    j = i * i - i
  end do
  call ETIME(tarray, result)
  print *, result
  print *, tarray(1)
  print *, tarray(2)
end program test_etime

```

See also: [Section 9.79 \[CPU_TIME\], page 162](#)

9.99 EVENT_QUERY — Query whether a coarray event has occurred

Description:

EVENT_QUERY assigns the number of events to *COUNT* which have been posted to the *EVENT* variable and not yet been removed by calling EVENT_WAIT. When *STAT* is present and the invocation was successful, it is assigned the value 0. If it is present and the invocation has failed, it is assigned a positive value and *COUNT* is assigned the value -1 .

Standard: TS 18508 or later

Class: subroutine

Syntax: CALL EVENT_QUERY (EVENT, COUNT [, STAT])

Arguments:

<i>EVENT</i>	(intent(IN)) Scalar of type EVENT_TYPE, defined in ISO_FORTRAN_ENV; shall not be coindexed.
<i>COUNT</i>	(intent(out)) Scalar integer with at least the precision of default integer.
<i>STAT</i>	(optional) Scalar default-kind integer variable.

Example:

```

program atomic
  use iso_fortran_env
  implicit none
  type(event_type) :: event_value_has_been_set[*]
  integer :: cnt
  if (this_image() == 1) then
    call event_query (event_value_has_been_set, cnt)
    if (cnt > 0) write(*,*) "Value has been set"
  end if
end program atomic

```

```

elseif (this_image() == 2) then
  event post (event_value_has_been_set[1])
end if
end program atomic

```

9.100 EXECUTE_COMMAND_LINE — Execute a shell command

Description:

EXECUTE_COMMAND_LINE runs a shell command, synchronously or asynchronously.

The **COMMAND** argument is passed to the shell and executed (The shell is **sh** on Unix systems, and **cmd.exe** on Windows.). If **WAIT** is present and has the value false, the execution of the command is asynchronous if the system supports it; otherwise, the command is executed synchronously using the C library's **system** call.

The three last arguments allow the user to get status information. After synchronous execution, **EXITSTAT** contains the integer exit code of the command, as returned by **system**. **CMDSTAT** is set to zero if the command line was executed (whatever its exit status was). **CMDMSG** is assigned an error message if an error has occurred.

Note that the **system** function need not be thread-safe. It is the responsibility of the user to ensure that **system** is not called concurrently.

For asynchronous execution on supported targets, the POSIX **posix_spawn** or **fork** functions are used. Also, a signal handler for the **SIGCHLD** signal is installed.

Standard: Fortran 2008 and later

Class: Subroutine

Syntax: CALL EXECUTE_COMMAND_LINE(COMMAND [, WAIT, EXITSTAT, CMDSTAT, CMDMSG])

Arguments:

COMMAND	Shall be a default CHARACTER scalar.
WAIT	(Optional) Shall be a default LOGICAL scalar.
EXITSTAT	(Optional) Shall be an INTEGER of the default kind.
CMDSTAT	(Optional) Shall be an INTEGER of the default kind.
CMDMSG	(Optional) Shall be an CHARACTER scalar of the default kind.

Example:

```

program test_exec
  integer :: i

  call execute_command_line ("external_prog.exe", exitstat=i)
  print *, "Exit status of external_prog.exe was ", i

  call execute_command_line ("reindex_files.exe", wait=.false.)
  print *, "Now reindexing files in the background"

end program test_exec

```


Note:

Because this intrinsic is implemented in terms of the `system` function call, its behavior with respect to signaling is processor dependent. In particular, on POSIX-compliant systems, the SIGINT and SIGQUIT signals will be ignored, and the SIGCHLD will be blocked. As such, if the parent process is terminated, the child process might not be terminated alongside.

See also: [Section 9.264 \[SYSTEM\], page 273](#)

9.101 EXIT — Exit the program with status.

Description:

EXIT causes immediate termination of the program with status. If status is omitted it returns the canonical *success* for the system. All Fortran I/O units are closed.

Standard: GNU extension

Class: Subroutine

Syntax: CALL EXIT([STATUS])

Arguments:

STATUS Shall be an INTEGER of the default kind.

Return value:

STATUS is passed to the parent process on exit.

Example:

```
program test_exit
  integer :: STATUS = 0
  print *, 'This program is going to exit.'
  call EXIT(STATUS)
end program test_exit
```

See also: [Section 9.2 \[ABORT\], page 111](#), [Section 9.163 \[KILL\], page 215](#)

9.102 EXP — Exponential function

Description:

EXP(X) computes the base *e* exponential of X.

Standard: Fortran 77 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = EXP(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has same type and kind as X.

Example:

```

program test_exp
  real :: x = 1.0
  x = exp(x)
end program test_exp

```

Specific names:

Name	Argument	Return type	Standard
EXP(X)	REAL(4) X	REAL(4)	Fortran 77 and later
DEXP(X)	REAL(8) X	REAL(8)	Fortran 77 and later
CEXP(X)	COMPLEX(4) X	COMPLEX(4)	Fortran 77 and later
ZEXP(X)	COMPLEX(8) X	COMPLEX(8)	GNU extension
CDEXP(X)	COMPLEX(8) X	COMPLEX(8)	GNU extension

9.103 EXPONENT — Exponent function

Description:

EXPONENT(X) returns the value of the exponent part of X. If X is zero the value returned is zero.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = EXPONENT(X)

Arguments:

X The type shall be REAL.

Return value:

The return value is of type default INTEGER.

Example:

```

program test_exponent
  real :: x = 1.0
  integer :: i
  i = exponent(x)
  print *, i
  print *, exponent(0.0)
end program test_exponent

```

9.104 EXTENDS_TYPE_OF — Query dynamic type for extension

Description:

Query dynamic type for extension.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: RESULT = EXTENDS_TYPE_OF(A, MOLD)

Arguments:

A Shall be an object of extensible declared type or unlimited polymorphic.

MOLD Shall be an object of extensible declared type or unlimited polymorphic.

Return value:

The return value is a scalar of type default logical. It is true if and only if the dynamic type of A is an extension type of the dynamic type of MOLD.

See also: [Section 9.235 \[SAME-TYPE-AS\]](#), page 256

9.105 FDATE — Get the current time as a string

Description:

FDATE(DATE) returns the current date (using the same format as [Section 9.81 \[CTIME\]](#), page 163) in DATE. It is equivalent to CALL CTIME(DATE, TIME()). This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL FDATE(DATE).
DATE = FDATE().
```

Arguments:

DATE	The type shall be of type CHARACTER of the default kind. It is an INTENT(OUT) argument. If the length of this variable is too short for the date and time string to fit completely, it will be blank on procedure return.
------	---

Return value:

The current date and time as a string.

Example:

```
program test_fdate
  integer(8) :: i, j
  character(len=30) :: date
  call fdate(date)
  print *, 'Program started on ', date
  do i = 1, 100000000 ! Just a delay
    j = i * i - i
  end do
  call fdate(date)
  print *, 'Program ended on ', date
end program test_fdate
```

See also: [Section 9.82 \[DATE-AND-TIME\]](#), page 164, [Section 9.81 \[CTIME\]](#), page 163

9.106 FGET — Read a single character in stream mode from stdin

Description:

Read a single character in stream mode from stdin by bypassing normal formatted output. Stream I/O should not be mixed with normal record-oriented (formatted or unformatted) I/O on the same unit; the results are unpredictable.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Note that the `FGET` intrinsic is provided for backwards compatibility with g77. GNU Fortran provides the Fortran 2003 Stream facility. Programmers should consider the use of new stream IO feature in new code for future portability. See also [Section 4.1 \[Fortran 2003 status\]](#), page 39.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL FGET(C [, STATUS])
STATUS = FGET(C)
```

Arguments:

<i>C</i>	The type shall be <code>CHARACTER</code> and of default kind.
<i>STATUS</i>	(Optional) status flag of type <code>INTEGER</code> . Returns 0 on success, -1 on end-of-file, and a system specific positive error code otherwise.

Example:

```
PROGRAM test_fget
  INTEGER, PARAMETER :: strlen = 100
  INTEGER :: status, i = 1
  CHARACTER(len=strlen) :: str = ""

  WRITE (*,*) 'Enter text:'
  DO
    CALL fget(str(i:i), status)
    if (status /= 0 .OR. i > strlen) exit
    i = i + 1
  END DO
  WRITE (*,*) TRIM(str)
END PROGRAM
```

See also: [Section 9.107 \[FGETC\]](#), page 180, [Section 9.112 \[FPUT\]](#), page 184, [Section 9.113 \[FPUTC\]](#), page 185

9.107 FGETC — Read a single character in stream mode

Description:

Read a single character in stream mode by bypassing normal formatted output. Stream I/O should not be mixed with normal record-oriented (formatted or unformatted) I/O on the same unit; the results are unpredictable.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Note that the `FGET` intrinsic is provided for backwards compatibility with g77. GNU Fortran provides the Fortran 2003 Stream facility. Programmers should consider the use of new stream IO feature in new code for future portability. See also [Section 4.1 \[Fortran 2003 status\]](#), page 39.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL FGETC(UNIT, C [, STATUS])
STATUS = FGETC(UNIT, C)
```

Arguments:

<i>UNIT</i>	The type shall be <code>INTEGER</code> .
<i>C</i>	The type shall be <code>CHARACTER</code> and of default kind.
<i>STATUS</i>	(Optional) status flag of type <code>INTEGER</code> . Returns 0 on success, -1 on end-of-file and a system specific positive error code otherwise.

Example:

```
PROGRAM test_fgetc
  INTEGER :: fd = 42, status
  CHARACTER :: c

  OPEN(UNIT=fd, FILE="/etc/passwd", ACTION="READ", STATUS = "OLD")
  DO
    CALL fgetc(fd, c, status)
    IF (status /= 0) EXIT
    call fput(c)
  END DO
  CLOSE(UNIT=fd)
END PROGRAM
```

See also: [Section 9.106 \[FGET\], page 179](#), [Section 9.112 \[FPUT\], page 184](#), [Section 9.113 \[FPUTC\], page 185](#)

9.108 FINDLOC — Search an array for a value

Description:

Determines the location of the element in the array with the value given in the *VALUE* argument, or, if the *DIM* argument is supplied, determines the locations of the maximum element along each row of the array in the *DIM* direction. If *MASK* is present, only the elements for which *MASK* is `.TRUE.` are considered. If more than one element in the array has the value *VALUE*, the location returned is that of the first such element in array element order if the *BACK* is not present or if it is `.FALSE.`. If *BACK* is true, the location returned is that of the last such element. If the array has zero size, or all of the elements of *MASK* are `.FALSE.`, then the result is an array of zeroes. Similarly, if *DIM* is supplied and all of the elements of *MASK* along a given row are zero, the result value for that row is zero.

Standard: Fortran 2008 and later.

Class: Transformational function

Syntax:

```
RESULT = FINDLOC(ARRAY, VALUE, DIM [, MASK] [,KIND] [,BACK])
RESULT = FINDLOC(ARRAY, VALUE, [, MASK] [,KIND] [,BACK])
```

Arguments:

<i>ARRAY</i>	Shall be an array of intrinsic type.
<i>VALUE</i>	A scalar of intrinsic type which is in type conformance with <i>ARRAY</i> .
<i>DIM</i>	(Optional) Shall be a scalar of type <i>INTEGER</i> , with a value between one and the rank of <i>ARRAY</i> , inclusive. It may not be an optional dummy argument.
<i>KIND</i>	(Optional) An <i>INTEGER</i> initialization expression indicating the kind parameter of the result.
<i>BACK</i>	(Optional) A scalar of type <i>LOGICAL</i> .

Return value:

If *DIM* is absent, the result is a rank-one array with a length equal to the rank of *ARRAY*. If *DIM* is present, the result is an array with a rank one less than the rank of *ARRAY*, and a size corresponding to the size of *ARRAY* with the *DIM* dimension removed. If *DIM* is present and *ARRAY* has a rank of one, the result is a scalar. If the optional argument *KIND* is present, the result is an integer of kind *KIND*, otherwise it is of default kind.

See also: [Section 9.191 \[MAXLOC\], page 230](#), [Section 9.199 \[MINLOC\], page 234](#)

9.109 FLOOR — Integer floor function

Description:

FLOOR(A) returns the greatest integer less than or equal to *X*.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: **RESULT = FLOOR(A [, KIND])**

Arguments:

<i>A</i>	The type shall be <i>REAL</i> .
<i>KIND</i>	(Optional) An <i>INTEGER</i> initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type *INTEGER(KIND)* if *KIND* is present and of default-kind *INTEGER* otherwise.

Example:

```

program test_floor
  real :: x = 63.29
  real :: y = -63.59
  print *, floor(x) ! returns 63
  print *, floor(y) ! returns -64
end program test_floor

```

See also: [Section 9.58 \[CEILING\], page 148](#), [Section 9.207 \[NINT\], page 240](#)

9.110 FLUSH — Flush I/O unit(s)

Description:

Flushes Fortran unit(s) currently open for output. Without the optional argument, all units are flushed, otherwise just the unit specified.

Standard: GNU extension

Class: Subroutine

Syntax: CALL FLUSH(UNIT)

Arguments:

UNIT (Optional) The type shall be INTEGER.

Note: Beginning with the Fortran 2003 standard, there is a FLUSH statement that should be preferred over the FLUSH intrinsic.

The FLUSH intrinsic and the Fortran 2003 FLUSH statement have identical effect: they flush the runtime library's I/O buffer so that the data becomes visible to other processes. This does not guarantee that the data is committed to disk.

On POSIX systems, you can request that all data is transferred to the storage device by calling the `fsync` function, with the POSIX file descriptor of the I/O unit as argument (retrieved with GNU intrinsic `FNUM`). The following example shows how:

```
! Declare the interface for POSIX fsync function
interface
  function fsync (fd) bind(c,name="fsync")
    use iso_c_binding, only: c_int
    integer(c_int), value :: fd
    integer(c_int) :: fsync
  end function fsync
end interface

! Variable declaration
integer :: ret

! Opening unit 10
open (10,file="foo")

! ...
! Perform I/O on unit 10
! ...

! Flush and sync
flush(10)
ret = fsync(fnum(10))

! Handle possible error
if (ret /= 0) stop "Error calling FSYNC"
```

9.111 FNUM — File number function

Description:

FNUM(UNIT) returns the POSIX file descriptor number corresponding to the open Fortran I/O unit UNIT.

Standard: GNU extension

Class: Function

Syntax: RESULT = FNUM(UNIT)

Arguments:
 UNIT The type shall be INTEGER.

Return value:
 The return value is of type INTEGER

Example:

```
program test_fnum
  integer :: i
  open (unit=10, status = "scratch")
  i = fnum(10)
  print *, i
  close (10)
end program test_fnum
```

9.112 FPUT — Write a single character in stream mode to stdout

Description:

Write a single character in stream mode to stdout by bypassing normal formatted output. Stream I/O should not be mixed with normal record-oriented (formatted or unformatted) I/O on the same unit; the results are unpredictable. This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Note that the FGET intrinsic is provided for backwards compatibility with g77. GNU Fortran provides the Fortran 2003 Stream facility. Programmers should consider the use of new stream IO feature in new code for future portability. See also [Section 4.1 \[Fortran 2003 status\]](#), page 39.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL FPUT(C [, STATUS])
STATUS = FPUT(C)
```

Arguments:
 C The type shall be CHARACTER and of default kind.
 STATUS (Optional) status flag of type INTEGER. Returns 0 on success, -1 on end-of-file and a system specific positive error code otherwise.

Example:

```
PROGRAM test_fput
  CHARACTER(len=10) :: str = "gfortran"
  INTEGER :: i
  DO i = 1, len_trim(str)
```



```

        CALL fput(str(i:i))
    END DO
END PROGRAM

```

See also: [Section 9.113 \[FPUTC\]](#), page 185, [Section 9.106 \[FGET\]](#), page 179, [Section 9.107 \[FGETC\]](#), page 180

9.113 FPUTC — Write a single character in stream mode

Description:

Write a single character in stream mode by bypassing normal formatted output. Stream I/O should not be mixed with normal record-oriented (formatted or unformatted) I/O on the same unit; the results are unpredictable.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Note that the FGET intrinsic is provided for backwards compatibility with g77. GNU Fortran provides the Fortran 2003 Stream facility. Programmers should consider the use of new stream IO feature in new code for future portability. See also [Section 4.1 \[Fortran 2003 status\]](#), page 39.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```

CALL FPUTC(UNIT, C [, STATUS])
STATUS = FPUTC(UNIT, C)

```

Arguments:

<i>UNIT</i>	The type shall be <code>INTEGER</code> .
<i>C</i>	The type shall be <code>CHARACTER</code> and of default kind.
<i>STATUS</i>	(Optional) status flag of type <code>INTEGER</code> . Returns 0 on success, -1 on end-of-file and a system specific positive error code otherwise.

Example:

```

PROGRAM test_fputc
  CHARACTER(len=10) :: str = "gfortran"
  INTEGER :: fd = 42, i

  OPEN(UNIT = fd, FILE = "out", ACTION = "WRITE", STATUS="NEW")
  DO i = 1, len_trim(str)
    CALL fputc(fd, str(i:i))
  END DO
  CLOSE(fd)
END PROGRAM

```

See also: [Section 9.112 \[FPUT\]](#), page 184, [Section 9.106 \[FGET\]](#), page 179, [Section 9.107 \[FGETC\]](#), page 180

9.114 FRACTION — Fractional part of the model representation

Description:

FRACTION(X) returns the fractional part of the model representation of X.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: Y = FRACTION(X)

Arguments:

X The type of the argument shall be a REAL.

Return value:

The return value is of the same type and kind as the argument. The fractional part of the model representation of X is returned; it is $X * \text{RADIX}(X)^{-(\text{EXPONENT}(X))}$.

Example:

```
program test_fraction
  real :: x
  x = 178.1387e-4
  print *, fraction(x), x * radix(x)**(-exponent(x))
end program test_fraction
```

9.115 FREE — Frees memory

Description:

Frees memory previously allocated by MALLOC. The FREE intrinsic is an extension intended to be used with Cray pointers, and is provided in GNU Fortran to allow user to compile legacy code. For new code using Fortran 95 pointers, the memory de-allocation intrinsic is DEALLOCATE.

Standard: GNU extension

Class: Subroutine

Syntax: CALL FREE(PTR)

Arguments:

PTR The type shall be INTEGER. It represents the location of the memory that should be de-allocated.

Return value:

None

Example: See MALLOC for an example.

See also: [Section 9.185 \[MALLOC\]](#), page 227

9.116 FSEEK — Low level file positioning subroutine

Description:

Moves UNIT to the specified OFFSET. If WHENCE is set to 0, the OFFSET is taken as an absolute value SEEK_SET, if set to 1, OFFSET is taken to be

relative to the current position `SEEK_CUR`, and if set to 2 relative to the end of the file `SEEK_END`. On error, `STATUS` is set to a nonzero value. If `STATUS` the seek fails silently.

This intrinsic routine is not fully backwards compatible with g77. In g77, the `FSEEK` takes a statement label instead of a `STATUS` variable. If `FSEEK` is used in old code, change

```
CALL FSEEK(UNIT, OFFSET, WHENCE, *label)
```

to

```
INTEGER :: status
CALL FSEEK(UNIT, OFFSET, WHENCE, status)
IF (status /= 0) GOTO label
```

Please note that GNU Fortran provides the Fortran 2003 Stream facility. Programmers should consider the use of new stream IO feature in new code for future portability. See also [Section 4.1 \[Fortran 2003 status\]](#), page 39.

Standard: GNU extension

Class: Subroutine

Syntax: `CALL FSEEK(UNIT, OFFSET, WHENCE[, STATUS])`

Arguments:

<i>UNIT</i>	Shall be a scalar of type <code>INTEGER</code> .
<i>OFFSET</i>	Shall be a scalar of type <code>INTEGER</code> .
<i>WHENCE</i>	Shall be a scalar of type <code>INTEGER</code> . Its value shall be either 0, 1 or 2.
<i>STATUS</i>	(Optional) shall be a scalar of type <code>INTEGER(4)</code> .

Example:

```
PROGRAM test_fseek
  INTEGER, PARAMETER :: SEEK_SET = 0, SEEK_CUR = 1, SEEK_END = 2
  INTEGER :: fd, offset, ierr

  ierr = 0
  offset = 5
  fd = 10

  OPEN(UNIT=fd, FILE="fseek.test")
  CALL FSEEK(fd, offset, SEEK_SET, ierr) ! move to OFFSET
  print *, FTELL(fd), ierr

  CALL FSEEK(fd, 0, SEEK_END, ierr)      ! move to end
  print *, FTELL(fd), ierr

  CALL FSEEK(fd, 0, SEEK_SET, ierr)      ! move to beginning
  print *, FTELL(fd), ierr

  CLOSE(UNIT=fd)
END PROGRAM
```

See also: [Section 9.118 \[FTELL\]](#), page 188

9.117 FSTAT — Get file status

Description:

FSTAT is identical to [Section 9.260 \[STAT\], page 270](#), except that information about an already opened file is obtained.

The elements in `VALUES` are the same as described by [Section 9.260 \[STAT\], page 270](#).

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL FSTAT(UNIT, VALUES [, STATUS])
STATUS = FSTAT(UNIT, VALUES)
```

Arguments:

<code>UNIT</code>	An open I/O unit number of type <code>INTEGER</code> .
<code>VALUES</code>	The type shall be <code>INTEGER(4)</code> , <code>DIMENSION(13)</code> .
<code>STATUS</code>	(Optional) status flag of type <code>INTEGER(4)</code> . Returns 0 on success and a system specific error code otherwise.

Example: See [Section 9.260 \[STAT\], page 270](#) for an example.

See also: To stat a link: [Section 9.183 \[LSTAT\], page 225](#), to stat a file: [Section 9.260 \[STAT\], page 270](#)

9.118 FTELL — Current stream position

Description:

Retrieves the current position within an open file.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL FTELL(UNIT, OFFSET)
OFFSET = FTELL(UNIT)
```

Arguments:

<code>OFFSET</code>	Shall of type <code>INTEGER</code> .
<code>UNIT</code>	Shall of type <code>INTEGER</code> .

Return value:

In either syntax, `OFFSET` is set to the current offset of unit number `UNIT`, or to `-1` if the unit is not currently open.

Example:

```

PROGRAM test_ftell
  INTEGER :: i
  OPEN(10, FILE="temp.dat")
  CALL ftell(10,i)
  WRITE(*,*) i
END PROGRAM

```

See also: [Section 9.116 \[FSEEK\]](#), page 186

9.119 GAMMA — Gamma function

Description:

GAMMA(X) computes Gamma (Γ) of X. For positive, integer values of X the Gamma function simplifies to the factorial function $\Gamma(x) = (x - 1)!$.

$$\Gamma(x) = \int_0^{\infty} t^{x-1} e^{-t} dt$$

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: X = GAMMA(X)

Arguments:

X Shall be of type REAL and neither zero nor a negative integer.

Return value:

The return value is of type REAL of the same kind as X.

Example:

```

program test_gamma
  real :: x = 1.0
  x = gamma(x) ! returns 1.0
end program test_gamma

```

Specific names:

Name	Argument	Return type	Standard
GAMMA(X)	REAL(4) X	REAL(4)	GNU Extension
DGAMMA(X)	REAL(8) X	REAL(8)	GNU Extension

See also: Logarithm of the Gamma function: [Section 9.179 \[LOG_GAMMA\]](#), page 223

9.120 GERROR — Get last system error message

Description:

Returns the system error message corresponding to the last system error. This resembles the functionality of `strerror(3)` in C.

Standard: GNU extension

Class: Subroutine

Syntax: CALL GERROR(RESULT)

Arguments:

RESULT Shall of type CHARACTER and of default

Example:

```
PROGRAM test_gerror
  CHARACTER(len=100) :: msg
  CALL gerror(msg)
  WRITE(*,*) msg
END PROGRAM
```

See also: [Section 9.146 \[IERRNO\]](#), page 205, [Section 9.215 \[PERROR\]](#), page 245

9.121 GETARG — Get command line arguments

Description:

Retrieve the *POS*-th argument that was passed on the command line when the containing program was invoked.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.123 \[GET_COMMAND_ARGUMENT\]](#), page 191 intrinsic defined by the Fortran 2003 standard.

Standard: GNU extension

Class: Subroutine

Syntax: CALL GETARG(*POS*, *VALUE*)

Arguments:

<i>POS</i>	Shall be of type <code>INTEGER</code> and not wider than the default integer kind; $POS \geq 0$
<i>VALUE</i>	Shall be of type <code>CHARACTER</code> and of default kind.
<i>VALUE</i>	Shall be of type <code>CHARACTER</code> .

Return value:

After GETARG returns, the *VALUE* argument holds the *POS*th command line argument. If *VALUE* cannot hold the argument, it is truncated to fit the length of *VALUE*. If there are less than *POS* arguments specified at the command line, *VALUE* will be filled with blanks. If $POS = 0$, *VALUE* is set to the name of the program (on systems that support this feature).

Example:

```
PROGRAM test_getarg
  INTEGER :: i
  CHARACTER(len=32) :: arg

  DO i = 1, iargc()
    CALL getarg(i, arg)
    WRITE (*,*) arg
  END DO
END PROGRAM
```

See also: GNU Fortran 77 compatibility function: [Section 9.139 \[IARGC\]](#), page 201

Fortran 2003 functions and subroutines: [Section 9.122 \[GET_COMMAND\]](#), page 191, [Section 9.123 \[GET_COMMAND_ARGUMENT\]](#), page 191, [Section 9.68 \[COMMAND_ARGUMENT_COUNT\]](#), page 155

9.122 GET_COMMAND — Get the entire command line

Description:

Retrieve the entire command line that was used to invoke the program.

Standard: Fortran 2003 and later

Class: Subroutine

Syntax: CALL GET_COMMAND([COMMAND, LENGTH, STATUS])

Arguments:

COMMAND (Optional) shall be of type CHARACTER and of default kind.
LENGTH (Optional) Shall be of type INTEGER and of default kind.
STATUS (Optional) Shall be of type INTEGER and of default kind.

Return value:

If **COMMAND** is present, stores the entire command line that was used to invoke the program in **COMMAND**. If **LENGTH** is present, it is assigned the length of the command line. If **STATUS** is present, it is assigned 0 upon success of the command, -1 if **COMMAND** is too short to store the command line, or a positive value in case of an error.

Example:

```
PROGRAM test_get_command
  CHARACTER(len=255) :: cmd
  CALL get_command(cmd)
  WRITE (*,*) TRIM(cmd)
END PROGRAM
```

See also: [Section 9.123 \[GET_COMMAND_ARGUMENT\]](#), page 191, [Section 9.68 \[COMMAND_ARGUMENT_COUNT\]](#), page 155

9.123 GET_COMMAND_ARGUMENT — Get command line arguments

Description:

Retrieve the **NUMBER**-th argument that was passed on the command line when the containing program was invoked.

Standard: Fortran 2003 and later

Class: Subroutine

Syntax: CALL GET_COMMAND_ARGUMENT(NUMBER [, VALUE, LENGTH, STATUS])

Arguments:

NUMBER Shall be a scalar of type INTEGER and of default kind, $NUMBER \geq 0$
VALUE (Optional) Shall be a scalar of type CHARACTER and of default kind.
LENGTH (Optional) Shall be a scalar of type INTEGER and of default kind.
STATUS (Optional) Shall be a scalar of type INTEGER and of default kind.

Return value:

After `GET_COMMAND_ARGUMENT` returns, the *VALUE* argument holds the *NUMBER*-th command line argument. If *VALUE* cannot hold the argument, it is truncated to fit the length of *VALUE*. If there are less than *NUMBER* arguments specified at the command line, *VALUE* will be filled with blanks. If *NUMBER* = 0, *VALUE* is set to the name of the program (on systems that support this feature). The *LENGTH* argument contains the length of the *NUMBER*-th command line argument. If the argument retrieval fails, *STATUS* is a positive number; if *VALUE* contains a truncated command line argument, *STATUS* is -1; and otherwise the *STATUS* is zero.

Example:

```
PROGRAM test_get_command_argument
  INTEGER :: i
  CHARACTER(len=32) :: arg

  i = 0
  DO
    CALL get_command_argument(i, arg)
    IF (LEN_TRIM(arg) == 0) EXIT

    WRITE (*,*) TRIM(arg)
    i = i+1
  END DO
END PROGRAM
```

See also: [Section 9.122 \[GET_COMMAND\]](#), page 191, [Section 9.68 \[COMMAND_ARGUMENT_COUNT\]](#), page 155

9.124 GETCWD — Get current working directory

Description:

Get current working directory.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL GETCWD(C [, STATUS])
STATUS = GETCWD(C)
```

Arguments:

<i>C</i>	The type shall be <code>CHARACTER</code> and of default kind.
<i>STATUS</i>	(Optional) status flag. Returns 0 on success, a system specific and nonzero error code otherwise.

Example:

```
PROGRAM test_getcwd
  CHARACTER(len=255) :: cwd
  CALL getcwd(cwd)
  WRITE(*,*) TRIM(cwd)
END PROGRAM
```


See also: [Section 9.60 \[CHDIR\]](#), page 149

9.125 GETENV — Get an environmental variable

Description:

Get the *VALUE* of the environmental variable *NAME*.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.126 \[GET_ENVIRONMENT_VARIABLE\]](#), page 193 intrinsic defined by the Fortran 2003 standard.

Note that **GETENV** need not be thread-safe. It is the responsibility of the user to ensure that the environment is not being updated concurrently with a call to the **GETENV** intrinsic.

Standard: GNU extension

Class: Subroutine

Syntax: `CALL GETENV(NAME, VALUE)`

Arguments:

<i>NAME</i>	Shall be of type CHARACTER and of default kind.
<i>VALUE</i>	Shall be of type CHARACTER and of default kind.

Return value:

Stores the value of *NAME* in *VALUE*. If *VALUE* is not large enough to hold the data, it is truncated. If *NAME* is not set, *VALUE* will be filled with blanks.

Example:

```
PROGRAM test_getenv
  CHARACTER(len=255) :: homedir
  CALL getenv("HOME", homedir)
  WRITE (*,*) TRIM(homedir)
END PROGRAM
```

See also: [Section 9.126 \[GET_ENVIRONMENT_VARIABLE\]](#), page 193

9.126 GET_ENVIRONMENT_VARIABLE — Get an environmental variable

Description:

Get the *VALUE* of the environmental variable *NAME*.

Note that **GET_ENVIRONMENT_VARIABLE** need not be thread-safe. It is the responsibility of the user to ensure that the environment is not being updated concurrently with a call to the **GET_ENVIRONMENT_VARIABLE** intrinsic.

Standard: Fortran 2003 and later

Class: Subroutine

Syntax: `CALL GET_ENVIRONMENT_VARIABLE(NAME[, VALUE, LENGTH, STATUS, TRIM_NAME])`

Arguments:

<i>NAME</i>	Shall be a scalar of type <code>CHARACTER</code> and of default kind.
<i>VALUE</i>	(Optional) Shall be a scalar of type <code>CHARACTER</code> and of default kind.
<i>LENGTH</i>	(Optional) Shall be a scalar of type <code>INTEGER</code> and of default kind.
<i>STATUS</i>	(Optional) Shall be a scalar of type <code>INTEGER</code> and of default kind.
<i>TRIM_NAME</i>	(Optional) Shall be a scalar of type <code>LOGICAL</code> and of default kind.

Return value:

Stores the value of *NAME* in *VALUE*. If *VALUE* is not large enough to hold the data, it is truncated. If *NAME* is not set, *VALUE* will be filled with blanks. Argument *LENGTH* contains the length needed for storing the environment variable *NAME* or zero if it is not present. *STATUS* is -1 if *VALUE* is present but too short for the environment variable; it is 1 if the environment variable does not exist and 2 if the processor does not support environment variables; in all other cases *STATUS* is zero. If *TRIM_NAME* is present with the value `.FALSE.`, the trailing blanks in *NAME* are significant; otherwise they are not part of the environment variable name.

Example:

```
PROGRAM test_getenv
  CHARACTER(len=255) :: homedir
  CALL get_environment_variable("HOME", homedir)
  WRITE (*,*) TRIM(homedir)
END PROGRAM
```

9.127 GETGID — Group ID function

Description:

Returns the numerical group ID of the current process.

Standard: GNU extension

Class: Function

Syntax: `RESULT = GETGID()`

Return value:

The return value of `GETGID` is an `INTEGER` of the default kind.

Example: See `GETPID` for an example.

See also: [Section 9.129 \[GETPID\]](#), page 195, [Section 9.130 \[GETUID\]](#), page 195

9.128 GETLOG — Get login name

Description:

Gets the username under which the program is running.

Standard: GNU extension

Class: Subroutine

Syntax: CALL GETLOG(C)

Arguments:

C Shall be of type CHARACTER and of default kind.

Return value:

Stores the current user name in *LOGIN*. (On systems where POSIX functions `geteuid` and `getpwuid` are not available, and the `getlogin` function is not implemented either, this will return a blank string.)

Example:

```
PROGRAM TEST_GETLOG
  CHARACTER(32) :: login
  CALL GETLOG(login)
  WRITE(*,*) login
END PROGRAM
```

See also: [Section 9.130 \[GETUID\]](#), page 195

9.129 GETPID — Process ID function

Description:

Returns the numerical process identifier of the current process.

Standard: GNU extension

Class: Function

Syntax: RESULT = GETPID()

Return value:

The return value of GETPID is an INTEGER of the default kind.

Example:

```
program info
  print *, "The current process ID is ", getpid()
  print *, "Your numerical user ID is ", getuid()
  print *, "Your numerical group ID is ", getgid()
end program info
```

See also: [Section 9.127 \[GETGID\]](#), page 194, [Section 9.130 \[GETUID\]](#), page 195

9.130 GETUID — User ID function

Description:

Returns the numerical user ID of the current process.

Standard: GNU extension

Class: Function

Syntax: RESULT = GETUID()

Return value:

The return value of GETUID is an INTEGER of the default kind.

Example: See GETPID for an example.

See also: [Section 9.129 \[GETPID\]](#), page 195, [Section 9.128 \[GETLOG\]](#), page 194

9.131 GMTIME — Convert time to GMT info

Description:

Given a system time value *TIME* (as provided by the [Section 9.270 \[TIME\]](#), [page 277](#) intrinsic), fills *VALUES* with values extracted from it appropriate to the UTC time zone (Universal Coordinated Time, also known in some countries as GMT, Greenwich Mean Time), using `gmtime(3)`.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.82 \[DATE_AND_TIME\]](#), [page 164](#) intrinsic defined by the Fortran 95 standard.

Standard: GNU extension

Class: Subroutine

Syntax: `CALL GMTIME(TIME, VALUES)`

Arguments:

<i>TIME</i>	An <code>INTEGER</code> scalar expression corresponding to a system time, with <code>INTENT(IN)</code> .
<i>VALUES</i>	A default <code>INTEGER</code> array with 9 elements, with <code>INTENT(OUT)</code> .

Return value:

The elements of *VALUES* are assigned as follows:

1. Seconds after the minute, range 0–59 or 0–61 to allow for leap seconds
2. Minutes after the hour, range 0–59
3. Hours past midnight, range 0–23
4. Day of month, range 1–31
5. Number of months since January, range 0–11
6. Years since 1900
7. Number of days since Sunday, range 0–6
8. Days since January 1, range 0–365
9. Daylight savings indicator: positive if daylight savings is in effect, zero if not, and negative if the information is not available.

See also: [Section 9.82 \[DATE_AND_TIME\]](#), [page 164](#), [Section 9.81 \[CTIME\]](#), [page 163](#), [Section 9.184 \[LTIME\]](#), [page 226](#), [Section 9.270 \[TIME\]](#), [page 277](#), [Section 9.271 \[TIME8\]](#), [page 278](#)

9.132 HOSTNM — Get system host name

Description:

Retrieves the host name of the system on which the program is running.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL HOSTNM(C [, STATUS])
STATUS = HOSTNM(NAME)
```

Arguments:

C Shall be of type **CHARACTER** and of default kind.
STATUS (Optional) status flag of type **INTEGER**. Returns 0 on success,
or a system specific error code otherwise.

Return value:

In either syntax, *NAME* is set to the current hostname if it can be obtained, or to a blank string otherwise.

9.133 HUGE — Largest number of a kind

Description:

HUGE(*X*) returns the largest number that is not an infinity in the model of the type of *X*.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: RESULT = HUGE(*X*)

Arguments:

X Shall be of type **REAL** or **INTEGER**.

Return value:

The return value is of the same type and kind as *X*

Example:

```
program test_huge_tiny
  print *, huge(0), huge(0.0), huge(0.0d0)
  print *, tiny(0.0), tiny(0.0d0)
end program test_huge_tiny
```

9.134 HYPOT — Euclidean distance function

Description:

HYPOT(*X*,*Y*) is the Euclidean distance function. It is equal to $\sqrt{X^2 + Y^2}$, without undue underflow or overflow.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = HYPOT(*X*, *Y*)

Arguments:

X The type shall be **REAL**.
Y The type and kind type parameter shall be the same as *X*.

Return value:

The return value has the same type and kind type parameter as *X*.

Example:

```
program test_hypot
  real(4) :: x = 1.e0_4, y = 0.5e0_4
  x = hypot(x,y)
end program test_hypot
```

9.135 IACHAR — Code in ASCII collating sequence

Description:

IACHAR(C) returns the code for the ASCII character in the first character position of C.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: RESULT = IACHAR(C [, KIND])

Arguments:

<i>C</i>	Shall be a scalar CHARACTER, with INTENT(IN)
<i>KIND</i>	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type INTEGER and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Example:

```
program test_iachar
  integer i
  i = iachar(' ')
end program test_iachar
```

Note: See [Section 9.143 \[ICHAR\]](#), page 203 for a discussion of converting between numerical values and formatted string representations.

See also: [Section 9.5 \[ACHAR\]](#), page 113, [Section 9.59 \[CHAR\]](#), page 148, [Section 9.143 \[ICHAR\]](#), page 203

9.136 IALL — Bitwise AND of array elements

Description:

Reduces with bitwise AND the elements of *ARRAY* along dimension *DIM* if the corresponding element in *MASK* is TRUE.

Standard: Fortran 2008 and later

Class: Transformational function

Syntax:

```
RESULT = IALL(ARRAY[, MASK])
RESULT = IALL(ARRAY, DIM[, MASK])
```

Arguments:

<i>ARRAY</i>	Shall be an array of type INTEGER
--------------	-----------------------------------

DIM (Optional) shall be a scalar of type **INTEGER** with a value in the range from 1 to *n*, where *n* equals the rank of *ARRAY*.

MASK (Optional) shall be of type **LOGICAL** and either be a scalar or an array of the same shape as *ARRAY*.

Return value:

The result is of the same type as *ARRAY*.

If *DIM* is absent, a scalar with the bitwise **ALL** of all elements in *ARRAY* is returned. Otherwise, an array of rank *n*-1, where *n* equals the rank of *ARRAY*, and a shape similar to that of *ARRAY* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_iall
  INTEGER(1) :: a(2)

  a(1) = b'00100100'
  a(2) = b'01101010'

  ! prints 00100000
  PRINT '(b8.8)', IALL(a)
END PROGRAM
```

See also: [Section 9.138 \[IANY\], page 200](#), [Section 9.153 \[IPARITY\], page 209](#), [Section 9.137 \[IAND\], page 199](#)

9.137 IAND — Bitwise logical and

Description:

Bitwise logical AND.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: **RESULT = IAND(I, J)**

Arguments:

I The type shall be **INTEGER** or a boz-literal-constant.

J The type shall be **INTEGER** with the same kind type parameter as *I* or a boz-literal-constant. *I* and *J* shall not both be boz-literal-constants.

Return value:

The return type is **INTEGER** with the kind type parameter of the arguments. A boz-literal-constant is converted to an **INTEGER** with the kind type parameter of the other argument as-if a call to [Section 9.149 \[INT\], page 207](#) occurred.

Example:

```
PROGRAM test_iand
  INTEGER :: a, b
  DATA a / Z'F' /, b / Z'3' /
  WRITE (*,*) IAND(a, b)
END PROGRAM
```

Specific names:

Name	Argument	Return type	Standard
IAND(A)	INTEGER A	INTEGER	Fortran 95 and later
BIAND(A)	INTEGER(1) A	INTEGER(1)	GNU extension
IIAND(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JIAND(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KIAND(A)	INTEGER(8) A	INTEGER(8)	GNU extension

See also: [Section 9.152 \[IOR\], page 208](#), [Section 9.145 \[IEOR\], page 204](#), [Section 9.141 \[IBITS\], page 202](#), [Section 9.142 \[IBSET\], page 202](#), [Section 9.140 \[IBCLR\], page 201](#), [Section 9.209 \[NOT\], page 241](#)

9.138 IANY — Bitwise OR of array elements

Description:

Reduces with bitwise OR (inclusive or) the elements of *ARRAY* along dimension *DIM* if the corresponding element in *MASK* is *TRUE*.

Standard: Fortran 2008 and later

Class: Transformational function

Syntax:

```
RESULT = IANY(ARRAY[, MASK])
RESULT = IANY(ARRAY, DIM[, MASK])
```

Arguments:

ARRAY Shall be an array of type *INTEGER*

DIM (Optional) shall be a scalar of type *INTEGER* with a value in the range from 1 to n, where n equals the rank of *ARRAY*.

MASK (Optional) shall be of type *LOGICAL* and either be a scalar or an array of the same shape as *ARRAY*.

Return value:

The result is of the same type as *ARRAY*.

If *DIM* is absent, a scalar with the bitwise OR of all elements in *ARRAY* is returned. Otherwise, an array of rank n-1, where n equals the rank of *ARRAY*, and a shape similar to that of *ARRAY* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_iany
  INTEGER(1) :: a(2)

  a(1) = b'00100100'
  a(2) = b'01101010'

  ! prints 01101110
  PRINT '(b8.8)', IANY(a)
END PROGRAM
```

See also: [Section 9.153 \[IPARITY\], page 209](#), [Section 9.136 \[IALL\], page 198](#), [Section 9.152 \[IOR\], page 208](#)

9.139 IARGC — Get the number of command line arguments

Description:

IARGC returns the number of arguments passed on the command line when the containing program was invoked.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.68 \[COMMAND_ARGUMENT_COUNT\]](#), [page 155](#) intrinsic defined by the Fortran 2003 standard.

Standard: GNU extension

Class: Function

Syntax: RESULT = IARGC()

Arguments:

None.

Return value:

The number of command line arguments, type INTEGER(4).

Example: See [Section 9.121 \[GETARG\]](#), [page 190](#)

See also: GNU Fortran 77 compatibility subroutine: [Section 9.121 \[GETARG\]](#), [page 190](#)
Fortran 2003 functions and subroutines: [Section 9.122 \[GET_COMMAND\]](#), [page 191](#), [Section 9.123 \[GET_COMMAND_ARGUMENT\]](#), [page 191](#), [Section 9.68 \[COMMAND_ARGUMENT_COUNT\]](#), [page 155](#)

9.140 IBCLR — Clear bit

Description:

IBCLR returns the value of *I* with the bit at position *POS* set to zero.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = IBCLR(*I*, *POS*)

Arguments:

I The type shall be INTEGER.
POS The type shall be INTEGER.

Return value:

The return value is of type INTEGER and of the same kind as *I*.

Specific names:

Name	Argument	Return type	Standard
IBCLR(A)	INTEGER A	INTEGER	Fortran 95 and later
BBCLR(A)	INTEGER(1) A	INTEGER(1)	GNU extension
IIBCLR(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JIBCLR(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KIBCLR(A)	INTEGER(8) A	INTEGER(8)	GNU extension

See also: [Section 9.141 \[IBITS\]](#), page 202, [Section 9.142 \[IBSET\]](#), page 202, [Section 9.137 \[IAND\]](#), page 199, [Section 9.152 \[IOR\]](#), page 208, [Section 9.145 \[IEOR\]](#), page 204, [Section 9.204 \[MVBITS\]](#), page 238

9.141 IBITS — Bit extraction

Description:

IBITS extracts a field of length *LEN* from *I*, starting from bit position *POS* and extending left for *LEN* bits. The result is right-justified and the remaining bits are zeroed. The value of *POS*+*LEN* must be less than or equal to the value *BIT_SIZE(I)*.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = IBITS(I, POS, LEN)`

Arguments:

<i>I</i>	The type shall be <code>INTEGER</code> .
<i>POS</i>	The type shall be <code>INTEGER</code> .
<i>LEN</i>	The type shall be <code>INTEGER</code> .

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

Specific names:

Name	Argument	Return type	Standard
<code>IBITS(A)</code>	<code>INTEGER A</code>	<code>INTEGER</code>	Fortran 95 and later
<code>BITS(A)</code>	<code>INTEGER(1) A</code>	<code>INTEGER(1)</code>	GNU extension
<code>IIBITS(A)</code>	<code>INTEGER(2) A</code>	<code>INTEGER(2)</code>	GNU extension
<code>JIBITS(A)</code>	<code>INTEGER(4) A</code>	<code>INTEGER(4)</code>	GNU extension
<code>KIBITS(A)</code>	<code>INTEGER(8) A</code>	<code>INTEGER(8)</code>	GNU extension

See also: [Section 9.48 \[BIT_SIZE\]](#), page 142, [Section 9.140 \[IBCLR\]](#), page 201, [Section 9.142 \[IBSET\]](#), page 202, [Section 9.137 \[IAND\]](#), page 199, [Section 9.152 \[IOR\]](#), page 208, [Section 9.145 \[IEOR\]](#), page 204

9.142 IBSET — Set bit

Description:

IBSET returns the value of *I* with the bit at position *POS* set to one.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = IBSET(I, POS)`

Arguments:

<i>I</i>	The type shall be <code>INTEGER</code> .
<i>POS</i>	The type shall be <code>INTEGER</code> .

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

Specific names:

Name	Argument	Return type	Standard
IBSET(A)	INTEGER A	INTEGER	Fortran 95 and later
BBSET(A)	INTEGER(1) A	INTEGER(1)	GNU extension
IIBSET(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JIBSET(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KIBSET(A)	INTEGER(8) A	INTEGER(8)	GNU extension

See also: Section 9.140 [IBCLR], page 201, Section 9.141 [IBITS], page 202, Section 9.137 [IAND], page 199, Section 9.152 [IOR], page 208, Section 9.145 [IEOR], page 204, Section 9.204 [MVBITS], page 238

9.143 ICHAR — Character-to-integer conversion function

Description:

ICHAR(C) returns the code for the character in the first character position of C in the system's native character set. The correspondence between characters and their codes is not necessarily the same across different GNU Fortran implementations.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: RESULT = ICHAR(C [, KIND])

Arguments:

<i>C</i>	Shall be a scalar CHARACTER, with INTENT(IN)
<i>KIND</i>	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type INTEGER and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Example:

```
program test_ichar
  integer i
  i = ichar(' ')
end program test_ichar
```

Specific names:

Name	Argument	Return type	Standard
ICHAR(C)	CHARACTER C	INTEGER(4)	Fortran 77 and later

Note: No intrinsic exists to convert between a numeric value and a formatted character string representation – for instance, given the CHARACTER value '154', obtaining an INTEGER or REAL value with the value 154, or vice versa. Instead, this functionality is provided by internal-file I/O, as in the following example:

```
program read_val
  integer value
  character(len=10) string, string2
  string = '154'
```

```

      ! Convert a string to a numeric value
      read (string,'(I10)') value
      print *, value

      ! Convert a value to a formatted string
      write (string2,'(I10)') value
      print *, string2
    end program read_val

```

See also: [Section 9.5 \[ACHAR\]](#), page 113, [Section 9.59 \[CHAR\]](#), page 148, [Section 9.135 \[IACHAR\]](#), page 198

9.144 IDATE — Get current local time subroutine (day/month/year)

Description:

IDATE(VALUE) Fills *VALUES* with the numerical values at the current local time. The day (in the range 1-31), month (in the range 1-12), and year appear in elements 1, 2, and 3 of *VALUES*, respectively. The year has four significant digits.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.82 \[DATE_AND_TIME\]](#), page 164 intrinsic defined by the Fortran 95 standard.

Standard: GNU extension

Class: Subroutine

Syntax: CALL IDATE(VALUE)

Arguments:

VALUES The type shall be INTEGER, DIMENSION(3) and the kind shall be the default integer kind.

Return value:

Does not return anything.

Example:

```

    program test_idate
      integer, dimension(3) :: tarray
      call idate(tarray)
      print *, tarray(1)
      print *, tarray(2)
      print *, tarray(3)
    end program test_idate

```

See also: [Section 9.82 \[DATE_AND_TIME\]](#), page 164

9.145 IEOR — Bitwise logical exclusive or

Description:

IEOR returns the bitwise Boolean exclusive-OR of *I* and *J*.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = IEOR(I, J)`

Arguments:

I The type shall be `INTEGER` or a boz-literal-constant.
J The type shall be `INTEGER` with the same kind type parameter as *I* or a boz-literal-constant. *I* and *J* shall not both be boz-literal-constants.

Return value:

The return type is `INTEGER` with the kind type parameter of the arguments. A boz-literal-constant is converted to an `INTEGER` with the kind type parameter of the other argument as-if a call to [Section 9.149 \[INT\]](#), page 207 occurred.

Specific names:

Name	Argument	Return type	Standard
<code>IEOR(A)</code>	<code>INTEGER A</code>	<code>INTEGER</code>	Fortran 95 and later
<code>BIEOR(A)</code>	<code>INTEGER(1) A</code>	<code>INTEGER(1)</code>	GNU extension
<code>IIIEOR(A)</code>	<code>INTEGER(2) A</code>	<code>INTEGER(2)</code>	GNU extension
<code>JIEOR(A)</code>	<code>INTEGER(4) A</code>	<code>INTEGER(4)</code>	GNU extension
<code>KIEOR(A)</code>	<code>INTEGER(8) A</code>	<code>INTEGER(8)</code>	GNU extension

See also: [Section 9.152 \[IOR\]](#), page 208, [Section 9.137 \[IAND\]](#), page 199, [Section 9.141 \[IBITS\]](#), page 202, [Section 9.142 \[IBSET\]](#), page 202, [Section 9.140 \[IBCLR\]](#), page 201, [Section 9.209 \[NOT\]](#), page 241

9.146 IERRNO — Get the last system error number

Description:

Returns the last system error number, as given by the C `errno` variable.

Standard: GNU extension

Class: Function

Syntax: `RESULT = IERRNO()`

Arguments:

None.

Return value:

The return value is of type `INTEGER` and of the default integer kind.

See also: [Section 9.215 \[PERROR\]](#), page 245

9.147 IMAGE_INDEX — Function that converts a cosubscript to an image index

Description:

Returns the image index belonging to a cosubscript.

Standard: Fortran 2008 and later

Class: Inquiry function.

Syntax: **RESULT = IMAGE_INDEX(COARRAY, SUB)**

Arguments: None.

COARRAY Coarray of any type.
SUB default integer rank-1 array of a size equal to the corank of **COARRAY**.

Return value:

Scalar default integer with the value of the image index which corresponds to the cosubscripts. For invalid cosubscripts the result is zero.

Example:

```
INTEGER :: array[2,-1:4,8,*]
! Writes 28 (or 0 if there are fewer than 28 images)
WRITE (*,*) IMAGE_INDEX (array, [2,0,3,1])
```

See also: Section 9.269 [THIS_IMAGE], page 276, Section 9.211 [NUM_IMAGES], page 242

9.148 INDEX — Position of a substring within a string

Description:

Returns the position of the start of the first occurrence of string *SUBSTRING* as a substring in *STRING*, counting from one. If *SUBSTRING* is not present in *STRING*, zero is returned. If the *BACK* argument is present and true, the return value is the start of the last occurrence rather than the first.

Standard: Fortran 77 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: **RESULT = INDEX(STRING, SUBSTRING [, BACK [, KIND]])**

Arguments:

STRING Shall be a scalar **CHARACTER**, with **INTENT(IN)**
SUBSTRING Shall be a scalar **CHARACTER**, with **INTENT(IN)**
BACK (Optional) Shall be a scalar **LOGICAL**, with **INTENT(IN)**
KIND (Optional) An **INTEGER** initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type **INTEGER** and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Specific names:

Name	Argument	Return type	Standard
INDEX(STRING, SUBSTRING)	CHARACTER	INTEGER(4)	Fortran 77 and later

See also: Section 9.237 [SCAN], page 256, Section 9.283 [VERIFY], page 284

9.149 INT — Convert to integer type

Description:

Convert to integer type

Standard: Fortran 77 and later

Class: Elemental function

Syntax: RESULT = INT(A [, KIND])

Arguments:

A Shall be of type INTEGER, REAL, or COMPLEX.
 KIND (Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

These functions return a INTEGER variable or array under the following rules:

- (A) If A is of type INTEGER, INT(A) = A
- (B) If A is of type REAL and $|A| < 1$, INT(A) equals 0. If $|A| \geq 1$, then INT(A) is the integer whose magnitude is the largest integer that does not exceed the magnitude of A and whose sign is the same as the sign of A.
- (C) If A is of type COMPLEX, rule B is applied to the real part of A.

Example:

```
program test_int
  integer :: i = 42
  complex :: z = (-3.7, 1.0)
  print *, int(i)
  print *, int(z), int(z,8)
end program
```

Specific names:

Name	Argument	Return type	Standard
INT(A)	REAL(4) A	INTEGER	Fortran 77 and later
IFIX(A)	REAL(4) A	INTEGER	Fortran 77 and later
IDINT(A)	REAL(8) A	INTEGER	Fortran 77 and later

9.150 INT2 — Convert to 16-bit integer type

Description:

Convert to a KIND=2 integer type. This is equivalent to the standard INT intrinsic with an optional argument of KIND=2, and is only included for backwards compatibility.

The SHORT intrinsic is equivalent to INT2.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = INT2(A)

Arguments:

A Shall be of type `INTEGER`, `REAL`, or `COMPLEX`.

Return value:

The return value is a `INTEGER(2)` variable.

See also: [Section 9.149 \[INT\]](#), page 207, [Section 9.151 \[INT8\]](#), page 208, [Section 9.181 \[LONG\]](#), page 224

9.151 INT8 — Convert to 64-bit integer type

Description:

Convert to a `KIND=8` integer type. This is equivalent to the standard `INT` intrinsic with an optional argument of `KIND=8`, and is only included for backwards compatibility.

Standard: GNU extension

Class: Elemental function

Syntax: `RESULT = INT8(A)`

Arguments:

A Shall be of type `INTEGER`, `REAL`, or `COMPLEX`.

Return value:

The return value is a `INTEGER(8)` variable.

See also: [Section 9.149 \[INT\]](#), page 207, [Section 9.150 \[INT2\]](#), page 207, [Section 9.181 \[LONG\]](#), page 224

9.152 IOR — Bitwise logical or

Description:

`IOR` returns the bitwise Boolean inclusive-OR of *I* and *J*.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = IOR(I, J)`

Arguments:

I The type shall be `INTEGER` or a boz-literal-constant.
J The type shall be `INTEGER` with the same kind type parameter as *I* or a boz-literal-constant. *I* and *J* shall not both be boz-literal-constants.

Return value:

The return type is `INTEGER` with the kind type parameter of the arguments. A boz-literal-constant is converted to an `INTEGER` with the kind type parameter of the other argument as-if a call to [Section 9.149 \[INT\]](#), page 207 occurred.

Specific names:

Name	Argument	Return type	Standard
------	----------	-------------	----------

IOR(A)	INTEGER A	INTEGER	Fortran 95 and later
BIOR(A)	INTEGER(1) A	INTEGER(1)	GNU extension
IIOR(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JIOR(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KIOR(A)	INTEGER(8) A	INTEGER(8)	GNU extension

See also: [Section 9.145 \[IEOR\]](#), page 204, [Section 9.137 \[IAND\]](#), page 199, [Section 9.141 \[IBITS\]](#), page 202, [Section 9.142 \[IBSET\]](#), page 202, [Section 9.140 \[IBCLR\]](#), page 201, [Section 9.209 \[NOT\]](#), page 241

9.153 IPARITY — Bitwise XOR of array elements

Description:

Reduces with bitwise XOR (exclusive or) the elements of *ARRAY* along dimension *DIM* if the corresponding element in *MASK* is *TRUE*.

Standard: Fortran 2008 and later

Class: Transformational function

Syntax:

```
RESULT = IPARITY(ARRAY[, MASK])
RESULT = IPARITY(ARRAY, DIM[, MASK])
```

Arguments:

<i>ARRAY</i>	Shall be an array of type <i>INTEGER</i>
<i>DIM</i>	(Optional) shall be a scalar of type <i>INTEGER</i> with a value in the range from 1 to n, where n equals the rank of <i>ARRAY</i> .
<i>MASK</i>	(Optional) shall be of type <i>LOGICAL</i> and either be a scalar or an array of the same shape as <i>ARRAY</i> .

Return value:

The result is of the same type as *ARRAY*.

If *DIM* is absent, a scalar with the bitwise XOR of all elements in *ARRAY* is returned. Otherwise, an array of rank n-1, where n equals the rank of *ARRAY*, and a shape similar to that of *ARRAY* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_iparity
  INTEGER(1) :: a(2)

  a(1) = b'00100100'
  a(2) = b'01101010'

  ! prints 01001110
  PRINT '(b8.8)', IPARITY(a)
END PROGRAM
```

See also: [Section 9.138 \[IANY\]](#), page 200, [Section 9.136 \[IALL\]](#), page 198, [Section 9.145 \[IEOR\]](#), page 204, [Section 9.214 \[PARITY\]](#), page 244

9.154 IRAND — Integer pseudo-random number

Description:

IRAND(FLAG) returns a pseudo-random number from a uniform distribution between 0 and a system-dependent limit (which is in most cases 2147483647). If *FLAG* is 0, the next number in the current sequence is returned; if *FLAG* is 1, the generator is restarted by CALL SRAND(0); if *FLAG* has any other value, it is used as a new seed with SRAND.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. It implements a simple modulo generator as provided by g77. For new code, one should consider the use of [Section 9.225 \[RANDOM_NUMBER\]](#), page 250 as it implements a superior algorithm.

Standard: GNU extension

Class: Function

Syntax: RESULT = IRAND(I)

Arguments:

I Shall be a scalar INTEGER of kind 4.

Return value:

The return value is of INTEGER(kind=4) type.

Example:

```
program test_irand
  integer,parameter :: seed = 86456

  call srand(seed)
  print *, irand(), irand(), irand(), irand()
  print *, irand(seed), irand(), irand(), irand()
end program test_irand
```

9.155 IS_CONTIGUOUS — Test whether an array is contiguous

Description:

IS_CONTIGUOUS tests whether an array is contiguous.

Standard: Fortran 2008 and later

Class: Inquiry function

Syntax: RESULT = IS_CONTIGUOUS(ARRAY)

Arguments:

ARRAY Shall be an array of any type.

Return value:

Returns a LOGICAL of the default kind, which .TRUE. if *ARRAY* is contiguous and false otherwise.

Example:

```
program test
  integer :: a(10)
```

```

a = [1,2,3,4,5,6,7,8,9,10]
call sub (a)      ! every element, is contiguous
call sub (a(::2)) ! every other element, is noncontiguous
contains
subroutine sub (x)
  integer :: x(:)
  if (is_contiguous (x)) then
    write (*,*) 'X is contiguous'
  else
    write (*,*) 'X is not contiguous'
  end if
end subroutine sub
end program test

```

9.156 IS_IOSTAT_END — Test for end-of-file value

Description:

IS_IOSTAT_END tests whether an variable has the value of the I/O status “end of file”. The function is equivalent to comparing the variable with the IOSTAT_END parameter of the intrinsic module ISO_FORTRAN_ENV.

Standard: Fortran 2003 and later

Class: Elemental function

Syntax: RESULT = IS_IOSTAT_END(I)

Arguments:

I Shall be of the type INTEGER.

Return value:

Returns a LOGICAL of the default kind, which .TRUE. if *I* has the value which indicates an end of file condition for IOSTAT= specifiers, and is .FALSE. otherwise.

Example:

```

PROGRAM iostat
  IMPLICIT NONE
  INTEGER :: stat, i
  OPEN(88, FILE='test.dat')
  READ(88, *, IOSTAT=stat) i
  IF(IS_IOSTAT_END(stat)) STOP 'END OF FILE'
END PROGRAM

```

9.157 IS_IOSTAT_EOR — Test for end-of-record value

Description:

IS_IOSTAT_EOR tests whether an variable has the value of the I/O status “end of record”. The function is equivalent to comparing the variable with the IOSTAT_EOR parameter of the intrinsic module ISO_FORTRAN_ENV.

Standard: Fortran 2003 and later

Class: Elemental function

Syntax: RESULT = IS_IOSTAT_EOR(I)

Arguments:

I Shall be of the type `INTEGER`.

Return value:

Returns a `LOGICAL` of the default kind, which `.TRUE.` if *I* has the value which indicates an end of file condition for `IOSTAT=` specifiers, and is `.FALSE.` otherwise.

Example:

```
PROGRAM iostat
  IMPLICIT NONE
  INTEGER :: stat, i(50)
  OPEN(88, FILE='test.dat', FORM='UNFORMATTED')
  READ(88, IOSTAT=stat) i
  IF(IS_IOSTAT_EOR(stat)) STOP 'END OF RECORD'
END PROGRAM
```

9.158 ISATTY — Whether a unit is a terminal device.

Description:

Determine whether a unit is connected to a terminal device.

Standard: GNU extension

Class: Function

Syntax: `RESULT = ISATTY(UNIT)`

Arguments:

UNIT Shall be a scalar `INTEGER`.

Return value:

Returns `.TRUE.` if the *UNIT* is connected to a terminal device, `.FALSE.` otherwise.

Example:

```
PROGRAM test_isatty
  INTEGER(kind=1) :: unit
  DO unit = 1, 10
    write(*,*) isatty(unit=unit)
  END DO
END PROGRAM
```

See also: [Section 9.277 \[TTYNAM\]](#), page 281

9.159 ISHFT — Shift bits

Description:

`ISHFT` returns a value corresponding to *I* with all of the bits shifted *SHIFT* places. A value of *SHIFT* greater than zero corresponds to a left shift, a value of zero corresponds to no shift, and a value less than zero corresponds to a right shift. If the absolute value of *SHIFT* is greater than `BIT_SIZE(I)`, the value is undefined. Bits shifted out from the left end or right end are lost; zeros are shifted in from the opposite end.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = ISHFT(I, SHIFT)`

Arguments:

I The type shall be `INTEGER`.
SHIFT The type shall be `INTEGER`.

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

Specific names:

Name	Argument	Return type	Standard
<code>ISHFT(A)</code>	<code>INTEGER A</code>	<code>INTEGER</code>	Fortran 95 and later
<code>BSHFT(A)</code>	<code>INTEGER(1) A</code>	<code>INTEGER(1)</code>	GNU extension
<code>IISHFT(A)</code>	<code>INTEGER(2) A</code>	<code>INTEGER(2)</code>	GNU extension
<code>JISHFT(A)</code>	<code>INTEGER(4) A</code>	<code>INTEGER(4)</code>	GNU extension
<code>KISHFT(A)</code>	<code>INTEGER(8) A</code>	<code>INTEGER(8)</code>	GNU extension

See also: [Section 9.160 \[ISHFTC\], page 213](#)

9.160 ISHFTC — Shift bits circularly

Description:

`ISHFTC` returns a value corresponding to *I* with the rightmost *SIZE* bits shifted circularly *SHIFT* places; that is, bits shifted out one end are shifted into the opposite end. A value of *SHIFT* greater than zero corresponds to a left shift, a value of zero corresponds to no shift, and a value less than zero corresponds to a right shift. The absolute value of *SHIFT* must be less than *SIZE*. If the *SIZE* argument is omitted, it is taken to be equivalent to `BIT_SIZE(I)`.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: `RESULT = ISHFTC(I, SHIFT [, SIZE])`

Arguments:

I The type shall be `INTEGER`.
SHIFT The type shall be `INTEGER`.
SIZE (Optional) The type shall be `INTEGER`; the value must be greater than zero and less than or equal to `BIT_SIZE(I)`.

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

Specific names:

Name	Argument	Return type	Standard
<code>ISHFTC(A)</code>	<code>INTEGER A</code>	<code>INTEGER</code>	Fortran 95 and later
<code>BSHFTC(A)</code>	<code>INTEGER(1) A</code>	<code>INTEGER(1)</code>	GNU extension
<code>IISHFTC(A)</code>	<code>INTEGER(2) A</code>	<code>INTEGER(2)</code>	GNU extension
<code>JISHFTC(A)</code>	<code>INTEGER(4) A</code>	<code>INTEGER(4)</code>	GNU extension
<code>KISHFTC(A)</code>	<code>INTEGER(8) A</code>	<code>INTEGER(8)</code>	GNU extension

See also: [Section 9.159 \[ISHFT\]](#), page 212

9.161 ISNAN — Test for a NaN

Description:

ISNAN tests whether a floating-point value is an IEEE Not-a-Number (NaN).

Standard: GNU extension

Class: Elemental function

Syntax: ISNAN(X)

Arguments:

X Variable of the type REAL.

Return value:

Returns a default-kind LOGICAL. The returned value is TRUE if X is a NaN and FALSE otherwise.

Example:

```
program test_nan
  implicit none
  real :: x
  x = -1.0
  x = sqrt(x)
  if (isnan(x)) stop '"x" is a NaN'
end program test_nan
```

9.162 ITIME — Get current local time subroutine (hour/minutes/seconds)

Description:

ITIME(VALUE) Fills *VALUES* with the numerical values at the current local time. The hour (in the range 1-24), minute (in the range 1-60), and seconds (in the range 1-60) appear in elements 1, 2, and 3 of *VALUES*, respectively.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.82 \[DATE_AND_TIME\]](#), page 164 intrinsic defined by the Fortran 95 standard.

Standard: GNU extension

Class: Subroutine

Syntax: CALL ITIME(VALUE)

Arguments:

VALUES The type shall be INTEGER, DIMENSION(3) and the kind shall be the default integer kind.

Return value:

Does not return anything.

Example:

```

program test_itime
  integer, dimension(3) :: tarray
  call itime(tarray)
  print *, tarray(1)
  print *, tarray(2)
  print *, tarray(3)
end program test_itime

```

See also: [Section 9.82 \[DATE_AND_TIME\]](#), page 164

9.163 KILL — Send a signal to a process

Description:

Standard: Sends the signal specified by *SIG* to the process *PID*. See `kill(2)`.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Class: Subroutine, function

Syntax:

```

CALL KILL(PID, SIG [, STATUS])
STATUS = KILL(PID, SIG)

```

Arguments:

<i>PID</i>	Shall be a scalar <code>INTEGER</code> with <code>INTENT(IN)</code> .
<i>SIG</i>	Shall be a scalar <code>INTEGER</code> with <code>INTENT(IN)</code> .
<i>STATUS</i>	[Subroutine](Optional) Shall be a scalar <code>INTEGER</code> . Returns 0 on success; otherwise a system-specific error code is returned.
<i>STATUS</i>	[Function] The kind type parameter is that of <code>pid</code> . Returns 0 on success; otherwise a system-specific error code is returned.

See also: [Section 9.2 \[ABORT\]](#), page 111, [Section 9.101 \[EXIT\]](#), page 177

9.164 KIND — Kind of an entity

Description:

`KIND(X)` returns the kind value of the entity *X*.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: `K = KIND(X)`

Arguments:

<i>X</i>	Shall be of type <code>LOGICAL</code> , <code>INTEGER</code> , <code>REAL</code> , <code>COMPLEX</code> or <code>CHARACTER</code> .
----------	---

Return value:

The return value is a scalar of type `INTEGER` and of the default integer kind.

Example:

```

program test_kind
  integer,parameter :: kc = kind(' ')

```

```
integer,parameter :: kl = kind(.true.)

print *, "The default character kind is ", kc
print *, "The default logical kind is ", kl
end program test_kind
```

9.165 LBOUND — Lower dimension bounds of an array

Description:

Returns the lower bounds of an array, or a single lower bound along the *DIM* dimension.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Inquiry function

Syntax: `RESULT = LBOUND(ARRAY [, DIM [, KIND]])`

Arguments:

<i>ARRAY</i>	Shall be an array, of any type.
<i>DIM</i>	(Optional) Shall be a scalar <code>INTEGER</code> .
<i>KIND</i>	(Optional) An <code>INTEGER</code> initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind. If *DIM* is absent, the result is an array of the lower bounds of *ARRAY*. If *DIM* is present, the result is a scalar corresponding to the lower bound of the array along that dimension. If *ARRAY* is an expression rather than a whole array or array structure component, or if it has a zero extent along the relevant dimension, the lower bound is taken to be 1.

See also: [Section 9.278 \[UBOUND\], page 281](#), [Section 9.166 \[LCOBUND\], page 216](#)

9.166 LCOBOUND — Lower codimension bounds of an array

Description:

Returns the lower bounds of a coarray, or a single lower cobound along the *DIM* codimension.

Standard: Fortran 2008 and later

Class: Inquiry function

Syntax: `RESULT = LCOBOUND(COARRAY [, DIM [, KIND]])`

Arguments:

<i>ARRAY</i>	Shall be an coarray, of any type.
<i>DIM</i>	(Optional) Shall be a scalar <code>INTEGER</code> .
<i>KIND</i>	(Optional) An <code>INTEGER</code> initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind `KIND`. If `KIND` is absent, the return value is of default integer kind. If `DIM` is absent, the result is an array of the lower cobounds of `COARRAY`. If `DIM` is present, the result is a scalar corresponding to the lower cobound of the array along that codimension.

See also: [Section 9.279 \[UCOBOUND\]](#), page 282, [Section 9.165 \[LBOUND\]](#), page 216

9.167 LEADZ — Number of leading zero bits of an integer

Description:

LEADZ returns the number of leading zero bits of an integer.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = LEADZ(I)`

Arguments:

`I` Shall be of type `INTEGER`.

Return value:

The type of the return value is the default `INTEGER`. If all the bits of `I` are zero, the result value is `BIT_SIZE(I)`.

Example:

```
PROGRAM test_leadz
  WRITE (*,*) BIT_SIZE(1) ! prints 32
  WRITE (*,*) LEADZ(1)    ! prints 31
END PROGRAM
```

See also: [Section 9.48 \[BIT_SIZE\]](#), page 142, [Section 9.273 \[TRAILZ\]](#), page 279, [Section 9.216 \[POPCNT\]](#), page 245, [Section 9.217 \[POPPAR\]](#), page 246

9.168 LEN — Length of a character entity

Description:

Returns the length of a character string. If `STRING` is an array, the length of an element of `STRING` is returned. Note that `STRING` need not be defined when this intrinsic is invoked, since only the length, not the content, of `STRING` is needed.

Standard: Fortran 77 and later, with `KIND` argument Fortran 2003 and later

Class: Inquiry function

Syntax: `L = LEN(STRING [, KIND])`

Arguments:

`STRING` Shall be a scalar or array of type `CHARACTER`, with `INTENT(IN)`
`KIND` (Optional) An `INTEGER` initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Specific names:

Name	Argument	Return type	Standard
<code>LEN(String)</code>	<code>CHARACTER</code>	<code>INTEGER</code>	Fortran 77 and later

See also: [Section 9.169 \[LEN_TRIM\], page 218](#), [Section 9.9 \[ADJUSTL\], page 116](#), [Section 9.10 \[ADJUSTR\], page 116](#)

9.169 `LEN_TRIM` — Length of a character entity without trailing blank characters

Description:

Returns the length of a character string, ignoring any trailing blanks.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: `RESULT = LEN_TRIM(String [, KIND])`

Arguments:

String Shall be a scalar of type `CHARACTER`, with `INTENT(IN)`
KIND (Optional) An `INTEGER` initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

See also: [Section 9.168 \[LEN\], page 217](#), [Section 9.9 \[ADJUSTL\], page 116](#), [Section 9.10 \[ADJUSTR\], page 116](#)

9.170 `LGE` — Lexical greater than or equal

Description:

Determines whether one string is lexically greater than or equal to another string, where the two strings are interpreted as containing ASCII character codes. If the String A and String B are not the same length, the shorter is compared as if spaces were appended to it to form a value that has the same length as the longer.

In general, the lexical comparison intrinsics `LGE`, `LGT`, `LLE`, and `LLT` differ from the corresponding intrinsic operators `.GE.`, `.GT.`, `.LE.`, and `.LT.`, in that the latter use the processor's character ordering (which is not ASCII on some targets), whereas the former always use the ASCII ordering.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = LGE(String_A, String_B)`

Arguments:

STRING_A Shall be of default **CHARACTER** type.
STRING_B Shall be of default **CHARACTER** type.

Return value:

Returns **.TRUE.** if *STRING_A* **>=** *STRING_B*, and **.FALSE.** otherwise, based on the ASCII ordering.

Specific names:

Name	Argument	Return type	Standard
LGE (<i>STRING_A</i> , <i>STRING_B</i>)	CHARACTER	LOGICAL	Fortran 77 and later

See also: [Section 9.171 \[LGT\], page 219](#), [Section 9.173 \[LLE\], page 220](#), [Section 9.174 \[LLT\], page 221](#)

9.171 LGT — Lexical greater than

Description:

Determines whether one string is lexically greater than another string, where the two strings are interpreted as containing ASCII character codes. If the String A and String B are not the same length, the shorter is compared as if spaces were appended to it to form a value that has the same length as the longer.

In general, the lexical comparison intrinsics **LGE**, **LGT**, **LLE**, and **LLT** differ from the corresponding intrinsic operators **.GE.**, **.GT.**, **.LE.**, and **.LT.**, in that the latter use the processor's character ordering (which is not ASCII on some targets), whereas the former always use the ASCII ordering.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: **RESULT = LGT**(*STRING_A*, *STRING_B*)

Arguments:

STRING_A Shall be of default **CHARACTER** type.
STRING_B Shall be of default **CHARACTER** type.

Return value:

Returns **.TRUE.** if *STRING_A* **>** *STRING_B*, and **.FALSE.** otherwise, based on the ASCII ordering.

Specific names:

Name	Argument	Return type	Standard
LGT (<i>STRING_A</i> , <i>STRING_B</i>)	CHARACTER	LOGICAL	Fortran 77 and later

See also: [Section 9.170 \[LGE\], page 218](#), [Section 9.173 \[LLE\], page 220](#), [Section 9.174 \[LLT\], page 221](#)

9.172 LINK — Create a hard link

Description:

Makes a (hard) link from file *PATH1* to *PATH2*. A null character (`CHAR(0)`) can be used to mark the end of the names in *PATH1* and *PATH2*; otherwise, trailing blanks in the file names are ignored. If the *STATUS* argument is supplied, it contains 0 on success or a nonzero error code upon return; see `link(2)`.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL LINK(PATH1, PATH2 [, STATUS])
STATUS = LINK(PATH1, PATH2)
```

Arguments:

<i>PATH1</i>	Shall be of default <code>CHARACTER</code> type.
<i>PATH2</i>	Shall be of default <code>CHARACTER</code> type.
<i>STATUS</i>	(Optional) Shall be of default <code>INTEGER</code> type.

See also: [Section 9.263 \[SYMLNK\], page 272](#), [Section 9.281 \[UNLINK\], page 282](#)

9.173 LLE — Lexical less than or equal

Description:

Determines whether one string is lexically less than or equal to another string, where the two strings are interpreted as containing ASCII character codes. If the String A and String B are not the same length, the shorter is compared as if spaces were appended to it to form a value that has the same length as the longer.

In general, the lexical comparison intrinsics `LGE`, `LGT`, `LLE`, and `LLT` differ from the corresponding intrinsic operators `.GE.`, `.GT.`, `.LE.`, and `.LT.`, in that the latter use the processor's character ordering (which is not ASCII on some targets), whereas the former always use the ASCII ordering.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = LLE(STRING_A, STRING_B)`

Arguments:

<i>STRING_A</i>	Shall be of default <code>CHARACTER</code> type.
<i>STRING_B</i>	Shall be of default <code>CHARACTER</code> type.

Return value:

Returns `.TRUE.` if `STRING_A <= STRING_B`, and `.FALSE.` otherwise, based on the ASCII ordering.

Specific names:

Name	Argument	Return type	Standard
LLE(<i>STRING_A</i> , <i>STRING_B</i>)	CHARACTER	LOGICAL	Fortran 77 and later

See also: [Section 9.170 \[LGE\]](#), page 218, [Section 9.171 \[LGT\]](#), page 219, [Section 9.174 \[LLT\]](#), page 221

9.174 LLT — Lexical less than

Description:

Determines whether one string is lexically less than another string, where the two strings are interpreted as containing ASCII character codes. If the String A and String B are not the same length, the shorter is compared as if spaces were appended to it to form a value that has the same length as the longer.

In general, the lexical comparison intrinsics LGE, LGT, LLE, and LLT differ from the corresponding intrinsic operators `.GE.`, `.GT.`, `.LE.`, and `.LT.`, in that the latter use the processor's character ordering (which is not ASCII on some targets), whereas the former always use the ASCII ordering.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = LLT(STRING_A, STRING_B)`

Arguments:

STRING_A Shall be of default CHARACTER type.
STRING_B Shall be of default CHARACTER type.

Return value:

Returns `.TRUE.` if `STRING_A < STRING_B`, and `.FALSE.` otherwise, based on the ASCII ordering.

Specific names:

Name	Argument	Return type	Standard
LLT(<i>STRING_A</i> , <i>STRING_B</i>)	CHARACTER	LOGICAL	Fortran 77 and later

See also: [Section 9.170 \[LGE\]](#), page 218, [Section 9.171 \[LGT\]](#), page 219, [Section 9.173 \[LLE\]](#), page 220

9.175 LNBLNK — Index of the last non-blank character in a string

Description:

Returns the length of a character string, ignoring any trailing blanks. This is identical to the standard `LEN_TRIM` intrinsic, and is only included for backwards compatibility.

Standard: GNU extension

Class: Elemental function

Syntax: `RESULT = LNBLNK (STRING)`

Arguments:
 `STRING` Shall be a scalar of type `CHARACTER`, with `INTENT(IN)`

Return value:
 The return value is of `INTEGER(kind=4)` type.

See also: Section 9.148 [`INDEX` intrinsic], page 206, Section 9.169 [`LEN_TRIM`], page 218

9.176 LOC — Returns the address of a variable

Description:
 `LOC(X)` returns the address of `X` as an integer.

Standard: GNU extension

Class: Inquiry function

Syntax: `RESULT = LOC(X)`

Arguments:
 `X` Variable of any type.

Return value:
 The return value is of type `INTEGER`, with a `KIND` corresponding to the size (in bytes) of a memory address on the target machine.

Example:

```
program test_loc
  integer :: i
  real :: r
  i = loc(r)
  print *, i
end program test_loc
```

9.177 LOG — Natural logarithm function

Description:
 `LOG(X)` computes the natural logarithm of `X`, i.e. the logarithm to the base e .

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = LOG(X)`

Arguments:
 `X` The type shall be `REAL` or `COMPLEX`.

Return value:
 The return value is of type `REAL` or `COMPLEX`. The kind type parameter is the same as `X`. If `X` is `COMPLEX`, the imaginary part ω is in the range $-\pi < \omega \leq \pi$.

Example:

```

program test_log
  real(8) :: x = 2.7182818284590451_8
  complex :: z = (1.0, 2.0)
  x = log(x)      ! will yield (approximately) 1
  z = log(z)
end program test_log

```

Specific names:

Name	Argument	Return type	Standard
ALOG(X)	REAL(4) X	REAL(4)	f95, gnu
DLOG(X)	REAL(8) X	REAL(8)	f95, gnu
CLOG(X)	COMPLEX(4) X	COMPLEX(4)	f95, gnu
ZLOG(X)	COMPLEX(8) X	COMPLEX(8)	f95, gnu
CDLOG(X)	COMPLEX(8) X	COMPLEX(8)	f95, gnu

9.178 LOG10 — Base 10 logarithm function

Description:

LOG10(X) computes the base 10 logarithm of X.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: RESULT = LOG10(X)

Arguments:

X The type shall be REAL.

Return value:

The return value is of type REAL or COMPLEX. The kind type parameter is the same as X.

Example:

```

program test_log10
  real(8) :: x = 10.0_8
  x = log10(x)
end program test_log10

```

Specific names:

Name	Argument	Return type	Standard
ALOG10(X)	REAL(4) X	REAL(4)	Fortran 95 and later
DLOG10(X)	REAL(8) X	REAL(8)	Fortran 95 and later

9.179 LOG_GAMMA — Logarithm of the Gamma function

Description:

LOG_GAMMA(X) computes the natural logarithm of the absolute value of the Gamma (Γ) function.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: X = LOG_GAMMA(X)

Arguments:

X Shall be of type **REAL** and neither zero nor a negative integer.

Return value:

The return value is of type **REAL** of the same kind as *X*.

Example:

```
program test_log_gamma
  real :: x = 1.0
  x = lgamma(x) ! returns 0.0
end program test_log_gamma
```

Specific names:

Name	Argument	Return type	Standard
LGAMMA(<i>X</i>)	REAL(4) <i>X</i>	REAL(4)	GNU Extension
ALGAMA(<i>X</i>)	REAL(4) <i>X</i>	REAL(4)	GNU Extension
DLGAMA(<i>X</i>)	REAL(8) <i>X</i>	REAL(8)	GNU Extension

See also: Gamma function: [Section 9.119 \[GAMMA\]](#), page 189

9.180 LOGICAL — Convert to logical type

Description:

Converts one kind of **LOGICAL** variable to another.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = LOGICAL(*L* [, *KIND*])

Arguments:

L The type shall be **LOGICAL**.
KIND (Optional) An **INTEGER** initialization expression indicating the kind parameter of the result.

Return value:

The return value is a **LOGICAL** value equal to *L*, with a kind corresponding to *KIND*, or of the default logical kind if *KIND* is not given.

See also: [Section 9.149 \[INT\]](#), page 207, [Section 9.229 \[REAL\]](#), page 252, [Section 9.62 \[CMPLX\]](#), page 150

9.181 LONG — Convert to integer type

Description:

Convert to a *KIND*=4 integer type, which is the same size as a C long integer. This is equivalent to the standard **INT** intrinsic with an optional argument of *KIND*=4, and is only included for backwards compatibility.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = LONG(*A*)

Arguments:

A Shall be of type INTEGER, REAL, or COMPLEX.

Return value:

The return value is a INTEGER(4) variable.

See also: [Section 9.149 \[INT\]](#), page 207, [Section 9.150 \[INT2\]](#), page 207, [Section 9.151 \[INT8\]](#), page 208

9.182 LSHIFT — Left shift bits*Description:*

LSHIFT returns a value corresponding to *I* with all of the bits shifted left by *SHIFT* places. *SHIFT* shall be nonnegative and less than or equal to BIT_SIZE(*I*), otherwise the result value is undefined. Bits shifted out from the left end are lost; zeros are shifted in from the opposite end.

This function has been superseded by the ISHFT intrinsic, which is standard in Fortran 95 and later, and the SHIFTL intrinsic, which is standard in Fortran 2008 and later.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = LSHIFT(*I*, *SHIFT*)

Arguments:

I The type shall be INTEGER.

SHIFT The type shall be INTEGER.

Return value:

The return value is of type INTEGER and of the same kind as *I*.

See also: [Section 9.159 \[ISHFT\]](#), page 212, [Section 9.160 \[ISHFTC\]](#), page 213, [Section 9.234 \[RSHIFT\]](#), page 255, [Section 9.245 \[SHIFTA\]](#), page 262, [Section 9.246 \[SHIFTL\]](#), page 262, [Section 9.247 \[SHIFTR\]](#), page 262

9.183 LSTAT — Get file status*Description:*

LSTAT is identical to [Section 9.260 \[STAT\]](#), page 270, except that if path is a symbolic link, then the link itself is stat'd, not the file that it refers to.

The elements in VALUES are the same as described by [Section 9.260 \[STAT\]](#), page 270.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL LSTAT(NAME, VALUES [, STATUS])
STATUS = LSTAT(NAME, VALUES)
```

Arguments:

<i>NAME</i>	The type shall be CHARACTER of the default kind, a valid path within the file system.
<i>VALUES</i>	The type shall be INTEGER(4) , DIMENSION(13) .
<i>STATUS</i>	(Optional) status flag of type INTEGER(4) . Returns 0 on success and a system specific error code otherwise.

Example: See [Section 9.260 \[STAT\]](#), page 270 for an example.

See also: To stat an open file: [Section 9.117 \[FSTAT\]](#), page 188, to stat a file: [Section 9.260 \[STAT\]](#), page 270

9.184 LTIME — Convert time to local time info

Description:

Given a system time value *TIME* (as provided by the [Section 9.270 \[TIME\]](#), page 277 intrinsic), fills *VALUES* with values extracted from it appropriate to the local time zone using `localtime(3)`.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. In new code, programmers should consider the use of the [Section 9.82 \[DATE_AND_TIME\]](#), page 164 intrinsic defined by the Fortran 95 standard.

Standard: GNU extension

Class: Subroutine

Syntax: `CALL LTIME(TIME, VALUES)`

Arguments:

<i>TIME</i>	An INTEGER scalar expression corresponding to a system time, with INTENT(IN) .
<i>VALUES</i>	A default INTEGER array with 9 elements, with INTENT(OUT) .

Return value:

The elements of *VALUES* are assigned as follows:

1. Seconds after the minute, range 0–59 or 0–61 to allow for leap seconds
2. Minutes after the hour, range 0–59
3. Hours past midnight, range 0–23
4. Day of month, range 1–31
5. Number of months since January, range 0–11
6. Years since 1900
7. Number of days since Sunday, range 0–6
8. Days since January 1, range 0–365
9. Daylight savings indicator: positive if daylight savings is in effect, zero if not, and negative if the information is not available.

See also: [Section 9.82 \[DATE_AND_TIME\]](#), page 164, [Section 9.81 \[CTIME\]](#), page 163, [Section 9.131 \[GMTIME\]](#), page 196, [Section 9.270 \[TIME\]](#), page 277, [Section 9.271 \[TIME8\]](#), page 278

9.185 MALLOC — Allocate dynamic memory

Description:

MALLOC(SIZE) allocates *SIZE* bytes of dynamic memory and returns the address of the allocated memory. The MALLOC intrinsic is an extension intended to be used with Cray pointers, and is provided in GNU Fortran to allow the user to compile legacy code. For new code using Fortran 95 pointers, the memory allocation intrinsic is ALLOCATE.

Standard: GNU extension

Class: Function

Syntax: PTR = MALLOC(SIZE)

Arguments:

SIZE The type shall be INTEGER.

Return value:

The return value is of type INTEGER(K), with *K* such that variables of type INTEGER(K) have the same size as C pointers (`sizeof(void *)`).

Example: The following example demonstrates the use of MALLOC and FREE with Cray pointers.

```

program test_malloc
  implicit none
  integer i
  real*8 x(*), z
  pointer(ptr_x,x)

  ptr_x = malloc(20*8)
  do i = 1, 20
    x(i) = sqrt(1.0d0 / i)
  end do
  z = 0
  do i = 1, 20
    z = z + x(i)
    print *, z
  end do
  call free(ptr_x)
end program test_malloc

```

See also: [Section 9.115 \[FREE\]](#), page 186

9.186 MASKL — Left justified mask

Description:

MASKL(I[, KIND]) has its leftmost *I* bits set to 1, and the remaining bits set to 0.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = MASKL(I[, KIND])

Arguments:

I Shall be of type INTEGER.
KIND Shall be a scalar constant expression of type INTEGER.

Return value:

The return value is of type INTEGER. If *KIND* is present, it specifies the kind value of the return type; otherwise, it is of the default integer kind.

See also: [Section 9.187 \[MASKR\], page 228](#)

9.187 MASKR — Right justified mask

Description:

MASKR(*I* [, *KIND*]) has its rightmost *I* bits set to 1, and the remaining bits set to 0.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = MASKR(*I* [, *KIND*])

Arguments:

I Shall be of type INTEGER.
KIND Shall be a scalar constant expression of type INTEGER.

Return value:

The return value is of type INTEGER. If *KIND* is present, it specifies the kind value of the return type; otherwise, it is of the default integer kind.

See also: [Section 9.186 \[MASKL\], page 227](#)

9.188 MATMUL — matrix multiplication

Description:

Performs a matrix multiplication on numeric or logical arguments.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = MATMUL(MATRIX_A, MATRIX_B)

Arguments:

MATRIX_A An array of INTEGER, REAL, COMPLEX, or LOGICAL type, with a rank of one or two.
MATRIX_B An array of INTEGER, REAL, or COMPLEX type if *MATRIX_A* is of a numeric type; otherwise, an array of LOGICAL type. The rank shall be one or two, and the first (or only) dimension of *MATRIX_B* shall be equal to the last (or only) dimension of *MATRIX_A*. *MATRIX_A* and *MATRIX_B* shall not both be rank one arrays.

Return value:

The matrix product of *MATRIX_A* and *MATRIX_B*. The type and kind of the result follow the usual type and kind promotion rules, as for the `*` or `.AND.` operators.

See also:

9.189 MAX — Maximum value of an argument list

Description:

Returns the argument with the largest (most positive) value.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = MAX(A1, A2 [, A3 [, ...]])`

Arguments:

A1 The type shall be `INTEGER` or `REAL`.
A2, A3, ... An expression of the same type and kind as *A1*. (As a GNU extension, arguments of different kinds are permitted.)

Return value:

The return value corresponds to the maximum value among the arguments, and has the same type and kind as the first argument.

Specific names:

Name	Argument	Return type	Standard
<code>MAX0(A1)</code>	<code>INTEGER(4) A1</code>	<code>INTEGER(4)</code>	Fortran 77 and later
<code>AMAX0(A1)</code>	<code>INTEGER(4) A1</code>	<code>REAL(MAX(X))</code>	Fortran 77 and later
<code>MAX1(A1)</code>	<code>REAL A1</code>	<code>INT(MAX(X))</code>	Fortran 77 and later
<code>AMAX1(A1)</code>	<code>REAL(4) A1</code>	<code>REAL(4)</code>	Fortran 77 and later
<code>DMAX1(A1)</code>	<code>REAL(8) A1</code>	<code>REAL(8)</code>	Fortran 77 and later

See also: [Section 9.191 \[MAXLOC\], page 230](#) [Section 9.192 \[MAXVAL\], page 231](#),
[Section 9.197 \[MIN\], page 233](#)

9.190 MAXEXPONENT — Maximum exponent of a real kind

Description:

`MAXEXPONENT(X)` returns the maximum exponent in the model of the type of *X*.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: `RESULT = MAXEXPONENT(X)`

Arguments:

X Shall be of type `REAL`.

Return value:

The return value is of type `INTEGER` and of the default integer kind.

Example:

```

program exponents
  real(kind=4) :: x
  real(kind=8) :: y

  print *, minexponent(x), maxexponent(x)
  print *, minexponent(y), maxexponent(y)
end program exponents

```

9.191 MAXLOC — Location of the maximum value within an array

Description:

Determines the location of the element in the array with the maximum value, or, if the *DIM* argument is supplied, determines the locations of the maximum element along each row of the array in the *DIM* direction. If *MASK* is present, only the elements for which *MASK* is *.TRUE.* are considered. If more than one element in the array has the maximum value, the location returned is that of the first such element in array element order if the *BACK* is not present, or is false; if *BACK* is true, the location returned is that of the last such element. If the array has zero size, or all of the elements of *MASK* are *.FALSE.*, then the result is an array of zeroes. Similarly, if *DIM* is supplied and all of the elements of *MASK* along a given row are zero, the result value for that row is zero.

Standard: Fortran 95 and later; *ARRAY* of *CHARACTER* and the *KIND* argument are available in Fortran 2003 and later. The *BACK* argument is available in Fortran 2008 and later.

Class: Transformational function

Syntax:

```

RESULT = MAXLOC(ARRAY, DIM [, MASK] [,KIND] [,BACK])
RESULT = MAXLOC(ARRAY [, MASK] [,KIND] [,BACK])

```

Arguments:

<i>ARRAY</i>	Shall be an array of type <i>INTEGER</i> or <i>REAL</i> .
<i>DIM</i>	(Optional) Shall be a scalar of type <i>INTEGER</i> , with a value between one and the rank of <i>ARRAY</i> , inclusive. It may not be an optional dummy argument.
<i>MASK</i>	Shall be an array of type <i>LOGICAL</i> , and conformable with <i>ARRAY</i> .
<i>KIND</i>	(Optional) An <i>INTEGER</i> initialization expression indicating the kind parameter of the result.
<i>BACK</i>	(Optional) A scalar of type <i>LOGICAL</i> .

Return value:

If *DIM* is absent, the result is a rank-one array with a length equal to the rank of *ARRAY*. If *DIM* is present, the result is an array with a rank one less than the rank of *ARRAY*, and a size corresponding to the size of *ARRAY* with the *DIM* dimension removed. If *DIM* is present and *ARRAY* has a rank of one,

the result is a scalar. If the optional argument *KIND* is present, the result is an integer of kind *KIND*, otherwise it is of default kind.

See also: [Section 9.108 \[FINDLOC\]](#), page 181, [Section 9.189 \[MAX\]](#), page 229, [Section 9.192 \[MAXVAL\]](#), page 231

9.192 MAXVAL — Maximum value of an array

Description:

Determines the maximum value of the elements in an array value, or, if the *DIM* argument is supplied, determines the maximum value along each row of the array in the *DIM* direction. If *MASK* is present, only the elements for which *MASK* is *.TRUE.* are considered. If the array has zero size, or all of the elements of *MASK* are *.FALSE.*, then the result is *-HUGE(ARRAY)* if *ARRAY* is numeric, or a string of nulls if *ARRAY* is of character type.

Standard: Fortran 95 and later

Class: Transformational function

Syntax:

```
RESULT = MAXVAL(ARRAY, DIM [, MASK])
RESULT = MAXVAL(ARRAY [, MASK])
```

Arguments:

<i>ARRAY</i>	Shall be an array of type <i>INTEGER</i> or <i>REAL</i> .
<i>DIM</i>	(Optional) Shall be a scalar of type <i>INTEGER</i> , with a value between one and the rank of <i>ARRAY</i> , inclusive. It may not be an optional dummy argument.
<i>MASK</i>	Shall be an array of type <i>LOGICAL</i> , and conformable with <i>ARRAY</i> .

Return value:

If *DIM* is absent, or if *ARRAY* has a rank of one, the result is a scalar. If *DIM* is present, the result is an array with a rank one less than the rank of *ARRAY*, and a size corresponding to the size of *ARRAY* with the *DIM* dimension removed. In all cases, the result is of the same type and kind as *ARRAY*.

See also: [Section 9.189 \[MAX\]](#), page 229, [Section 9.191 \[MAXLOC\]](#), page 230

9.193 MCLOCK — Time function

Description:

Returns the number of clock ticks since the start of the process, based on the function `clock(3)` in the C standard library.

This intrinsic is not fully portable, such as to systems with 32-bit *INTEGER* types but supporting times wider than 32 bits. Therefore, the values returned by this intrinsic might be, or become, negative, or numerically less than previous values, during a single run of the compiled program.

Standard: GNU extension

Class: Function

Syntax: RESULT = MCLOCK()

Return value:

The return value is a scalar of type INTEGER(4), equal to the number of clock ticks since the start of the process, or -1 if the system does not support clock(3).

See also: Section 9.81 [CTIME], page 163, Section 9.131 [GMTIME], page 196, Section 9.184 [LTIME], page 226, Section 9.193 [MCLOCK], page 231, Section 9.270 [TIME], page 277

9.194 MCLOCK8 — Time function (64-bit)

Description:

Returns the number of clock ticks since the start of the process, based on the function clock(3) in the C standard library.

Warning: this intrinsic does not increase the range of the timing values over that returned by clock(3). On a system with a 32-bit clock(3), MCLOCK8 will return a 32-bit value, even though it is converted to a 64-bit INTEGER(8) value. That means overflows of the 32-bit value can still occur. Therefore, the values returned by this intrinsic might be or become negative or numerically less than previous values during a single run of the compiled program.

Standard: GNU extension

Class: Function

Syntax: RESULT = MCLOCK8()

Return value:

The return value is a scalar of type INTEGER(8), equal to the number of clock ticks since the start of the process, or -1 if the system does not support clock(3).

See also: Section 9.81 [CTIME], page 163, Section 9.131 [GMTIME], page 196, Section 9.184 [LTIME], page 226, Section 9.193 [MCLOCK], page 231, Section 9.271 [TIME8], page 278

9.195 MERGE — Merge variables

Description:

Select values from two arrays according to a logical mask. The result is equal to *TSOURCE* if *MASK* is .TRUE., or equal to *FSOURCE* if it is .FALSE..

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = MERGE(TSOURCE, FSOURCE, MASK)

Arguments:

TSOURCE May be of any type.

FSOURCE Shall be of the same type and type parameters as *TSOURCE*.
MASK Shall be of type LOGICAL.

Return value:

The result is of the same type and type parameters as *TSOURCE*.

9.196 MERGE_BITS — Merge of bits under mask

Description:

MERGE_BITS(*I*, *J*, *MASK*) merges the bits of *I* and *J* as determined by the mask. The *i*-th bit of the result is equal to the *i*-th bit of *I* if the *i*-th bit of *MASK* is 1; it is equal to the *i*-th bit of *J* otherwise.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = MERGE_BITS(*I*, *J*, *MASK*)

Arguments:

I Shall be of type INTEGER or a boz-literal-constant.
J Shall be of type INTEGER with the same kind type parameter as *I* or a boz-literal-constant. *I* and *J* shall not both be boz-literal-constants.
MASK Shall be of type INTEGER or a boz-literal-constant and of the same kind as *I*.

Return value:

The result is of the same type and kind as *I*.

9.197 MIN — Minimum value of an argument list

Description:

Returns the argument with the smallest (most negative) value.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: RESULT = MIN(*A1*, *A2* [, *A3*, ...])

Arguments:

A1 The type shall be INTEGER or REAL.
A2, *A3*, ... An expression of the same type and kind as *A1*. (As a GNU extension, arguments of different kinds are permitted.)

Return value:

The return value corresponds to the maximum value among the arguments, and has the same type and kind as the first argument.

Specific names:

Name	Argument	Return type	Standard
MINO(<i>A1</i>)	INTEGER(4) <i>A1</i>	INTEGER(4)	Fortran 77 and later
AMINO(<i>A1</i>)	INTEGER(4) <i>A1</i>	REAL(4)	Fortran 77 and later

MIN1(A1)	REAL A1	INTEGER(4)	Fortran 77 and later
AMIN1(A1)	REAL(4) A1	REAL(4)	Fortran 77 and later
DMIN1(A1)	REAL(8) A1	REAL(8)	Fortran 77 and later

See also: [Section 9.189 \[MAX\]](#), page 229, [Section 9.199 \[MINLOC\]](#), page 234, [Section 9.200 \[MINVAL\]](#), page 235

9.198 MINEXPONENT — Minimum exponent of a real kind

Description:

MINEXPONENT(X) returns the minimum exponent in the model of the type of X.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: RESULT = MINEXPONENT(X)

Arguments:

X Shall be of type REAL.

Return value:

The return value is of type INTEGER and of the default integer kind.

Example: See MAXEXPONENT for an example.

9.199 MINLOC — Location of the minimum value within an array

Description:

Determines the location of the element in the array with the minimum value, or, if the *DIM* argument is supplied, determines the locations of the minimum element along each row of the array in the *DIM* direction. If *MASK* is present, only the elements for which *MASK* is *.TRUE.* are considered. If more than one element in the array has the minimum value, the location returned is that of the first such element in array element order if the *BACK* is not present, or is false; if *BACK* is true, the location returned is that of the last such element. If the array has zero size, or all of the elements of *MASK* are *.FALSE.*, then the result is an array of zeroes. Similarly, if *DIM* is supplied and all of the elements of *MASK* along a given row are zero, the result value for that row is zero.

Standard: Fortran 95 and later; *ARRAY* of *CHARACTER* and the *KIND* argument are available in Fortran 2003 and later. The *BACK* argument is available in Fortran 2008 and later.

Class: Transformational function

Syntax:

```
RESULT = MINLOC(ARRAY, DIM [, MASK] [,KIND] [,BACK])
RESULT = MINLOC(ARRAY [, MASK] [,KIND] [,BACK])
```

Arguments:

ARRAY Shall be an array of type INTEGER, REAL or CHARACTER.

<i>DIM</i>	(Optional) Shall be a scalar of type INTEGER , with a value between one and the rank of <i>ARRAY</i> , inclusive. It may not be an optional dummy argument.
<i>MASK</i>	Shall be an array of type LOGICAL , and conformable with <i>ARRAY</i> .
<i>KIND</i>	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.
<i>BACK</i>	(Optional) A scalar of type LOGICAL .

Return value:

If *DIM* is absent, the result is a rank-one array with a length equal to the rank of *ARRAY*. If *DIM* is present, the result is an array with a rank one less than the rank of *ARRAY*, and a size corresponding to the size of *ARRAY* with the *DIM* dimension removed. If *DIM* is present and *ARRAY* has a rank of one, the result is a scalar. If the optional argument *KIND* is present, the result is an integer of kind *KIND*, otherwise it is of default kind.

See also: [Section 9.108 \[FINDLOC\], page 181](#), [Section 9.197 \[MIN\], page 233](#), [Section 9.200 \[MINVAL\], page 235](#)

9.200 MINVAL — Minimum value of an array

Description:

Determines the minimum value of the elements in an array value, or, if the *DIM* argument is supplied, determines the minimum value along each row of the array in the *DIM* direction. If *MASK* is present, only the elements for which *MASK* is **.TRUE.** are considered. If the array has zero size, or all of the elements of *MASK* are **.FALSE.**, then the result is **HUGE(ARRAY)** if *ARRAY* is numeric, or a string of **CHAR(255)** characters if *ARRAY* is of character type.

Standard: Fortran 95 and later

Class: Transformational function

Syntax:

```
RESULT = MINVAL(ARRAY, DIM [, MASK])
RESULT = MINVAL(ARRAY [, MASK])
```

Arguments:

<i>ARRAY</i>	Shall be an array of type INTEGER or REAL .
<i>DIM</i>	(Optional) Shall be a scalar of type INTEGER , with a value between one and the rank of <i>ARRAY</i> , inclusive. It may not be an optional dummy argument.
<i>MASK</i>	Shall be an array of type LOGICAL , and conformable with <i>ARRAY</i> .

Return value:

If *DIM* is absent, or if *ARRAY* has a rank of one, the result is a scalar. If *DIM* is present, the result is an array with a rank one less than the rank of *ARRAY*, and a size corresponding to the size of *ARRAY* with the *DIM* dimension removed. In all cases, the result is of the same type and kind as *ARRAY*.

See also: [Section 9.197 \[MIN\]](#), page 233, [Section 9.199 \[MINLOC\]](#), page 234

9.201 MOD — Remainder function

Description:

MOD(A,P) computes the remainder of the division of A by P.

Standard: Fortran 77 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = MOD(A, P)

Arguments:

A Shall be a scalar of type **INTEGER** or **REAL**.
P Shall be a scalar of the same type and kind as A and not equal to zero.

Return value:

The return value is the result of $A - (\text{INT}(A/P) * P)$. The type and kind of the return value is the same as that of the arguments. The returned value has the same sign as A and a magnitude less than the magnitude of P.

Example:

```
program test_mod
  print *, mod(17,3)
  print *, mod(17.5,5.5)
  print *, mod(17.5d0,5.5)
  print *, mod(17.5,5.5d0)

  print *, mod(-17,3)
  print *, mod(-17.5,5.5)
  print *, mod(-17.5d0,5.5)
  print *, mod(-17.5,5.5d0)

  print *, mod(17,-3)
  print *, mod(17.5,-5.5)
  print *, mod(17.5d0,-5.5)
  print *, mod(17.5,-5.5d0)
end program test_mod
```

Specific names:

Name	Arguments	Return type	Standard
MOD(A,P)	INTEGER A,P	INTEGER	Fortran 95 and later
AMOD(A,P)	REAL(4) A,P	REAL(4)	Fortran 95 and later
DMOD(A,P)	REAL(8) A,P	REAL(8)	Fortran 95 and later
BMOD(A,P)	INTEGER(1) A,P	INTEGER(1)	GNU extension
IMOD(A,P)	INTEGER(2) A,P	INTEGER(2)	GNU extension
JMOD(A,P)	INTEGER(4) A,P	INTEGER(4)	GNU extension
KMOD(A,P)	INTEGER(8) A,P	INTEGER(8)	GNU extension

See also: [Section 9.202 \[MODULO\], page 237](#)

9.202 MODULO — Modulo function

Description:

MODULO(A,P) computes the *A* modulo *P*.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = MODULO(A, P)

Arguments:

<i>A</i>	Shall be a scalar of type INTEGER or REAL.
<i>P</i>	Shall be a scalar of the same type and kind as <i>A</i> . It shall not be zero.

Return value:

The type and kind of the result are those of the arguments.

If *A* and *P* are of type INTEGER:

MODULO(A,P) has the value *R* such that $A=Q*P+R$, where *Q* is an integer and *R* is between 0 (inclusive) and *P* (exclusive).

If *A* and *P* are of type REAL:

MODULO(A,P) has the value of $A - \text{FLOOR}(A / P) * P$.

The returned value has the same sign as *P* and a magnitude less than the magnitude of *P*.

Example:

```
program test_modulo
  print *, modulo(17,3)
  print *, modulo(17.5,5.5)

  print *, modulo(-17,3)
  print *, modulo(-17.5,5.5)

  print *, modulo(17,-3)
  print *, modulo(17.5,-5.5)
end program
```

See also: [Section 9.201 \[MOD\], page 236](#)

9.203 MOVE_ALLOC — Move allocation from one object to another

Description:

MOVE_ALLOC(FROM, TO) moves the allocation from *FROM* to *TO*. *FROM* will become deallocated in the process.

Standard: Fortran 2003 and later

Class: Pure subroutine

Syntax: CALL MOVE_ALLOC(FROM, TO)

Arguments:

FROM ALLOCATABLE, INTENT(INOUT), may be of any type and kind.
TO ALLOCATABLE, INTENT(OUT), shall be of the same type, kind
 and rank as *FROM*.

Return value:

None

Example:

```

program test_move_alloc
  integer, allocatable :: a(:), b(:)

  allocate(a(3))
  a = [ 1, 2, 3 ]
  call move_alloc(a, b)
  print *, allocated(a), allocated(b)
  print *, b
end program test_move_alloc

```

9.204 MVBITS — Move bits from one integer to another

Description:

Moves *LEN* bits from positions *FROMPOS* through *FROMPOS+LEN-1* of *FROM* to positions *TOPOS* through *TOPOS+LEN-1* of *TO*. The portion of argument *TO* not affected by the movement of bits is unchanged. The values of *FROMPOS+LEN-1* and *TOPOS+LEN-1* must be less than *BIT_SIZE(FROM)*.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental subroutine

Syntax: CALL MVBITS(*FROM*, *FROMPOS*, *LEN*, *TO*, *TOPOS*)

Arguments:

FROM The type shall be INTEGER.
FROMPOS The type shall be INTEGER.
LEN The type shall be INTEGER.
TO The type shall be INTEGER, of the same kind as *FROM*.
TOPOS The type shall be INTEGER.

Specific names:

Name	Argument	Return type	Standard
MVBITS(A)	INTEGER A	INTEGER	Fortran 95 and later
BMVBITS(A)	INTEGER(1) A	INTEGER(1)	GNU extension
IMVBITS(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JMVBITS(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KMVBITS(A)	INTEGER(8) A	INTEGER(8)	GNU extension

See also: [Section 9.140 \[IBCLR\], page 201](#), [Section 9.142 \[IBSET\], page 202](#), [Section 9.141 \[IBITS\], page 202](#), [Section 9.137 \[IAND\], page 199](#), [Section 9.152 \[IOR\], page 208](#), [Section 9.145 \[IEOR\], page 204](#)

9.205 NEAREST — Nearest representable number

Description:

NEAREST(X, S) returns the processor-representable number nearest to X in the direction indicated by the sign of S.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = NEAREST(X, S)

Arguments:

X	Shall be of type REAL.
S	Shall be of type REAL and not equal to zero.

Return value:

The return value is of the same type as X. If S is positive, NEAREST returns the processor-representable number greater than X and nearest to it. If S is negative, NEAREST returns the processor-representable number smaller than X and nearest to it.

Example:

```
program test_nearest
  real :: x, y
  x = nearest(42.0, 1.0)
  y = nearest(42.0, -1.0)
  write (*,"(3(G20.15))") x, y, x - y
end program test_nearest
```

9.206 NEW_LINE — New line character

Description:

NEW_LINE(C) returns the new-line character.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: RESULT = NEW_LINE(C)

Arguments:

C	The argument shall be a scalar or array of the type CHARACTER.
---	--

Return value:

Returns a CHARACTER scalar of length one with the new-line character of the same kind as parameter C.

Example:

```
program newline
  implicit none
  write(*,'(A)') 'This is record 1.'//NEW_LINE('A')// 'This is record 2.'
end program newline
```

9.207 NINT — Nearest whole number

Description:

NINT(A) rounds its argument to the nearest whole number.

Standard: Fortran 77 and later, with *KIND* argument Fortran 90 and later

Class: Elemental function

Syntax: RESULT = NINT(A [, KIND])

Arguments:

A The type of the argument shall be **REAL**.
KIND (Optional) An **INTEGER** initialization expression indicating the kind parameter of the result.

Return value:

Returns *A* with the fractional portion of its magnitude eliminated by rounding to the nearest whole number and with its sign preserved, converted to an **INTEGER** of the default kind.

Example:

```
program test_nint
  real(4) x4
  real(8) x8
  x4 = 1.234E0_4
  x8 = 4.321_8
  print *, nint(x4), idnint(x8)
end program test_nint
```

Specific names:

Name	Argument	Return Type	Standard
NINT(A)	REAL(4) A	INTEGER	Fortran 95 and later
IDNINT(A)	REAL(8) A	INTEGER	Fortran 95 and later

See also: [Section 9.58 \[CEILING\], page 148](#), [Section 9.109 \[FLOOR\], page 182](#)

9.208 NORM2 — Euclidean vector norms

Description:

Calculates the Euclidean vector norm (L_2 norm) of of *ARRAY* along dimension *DIM*.

Standard: Fortran 2008 and later

Class: Transformational function

Syntax:

RESULT = NORM2(ARRAY[, DIM])

Arguments:

ARRAY Shall be an array of type **REAL**
DIM (Optional) shall be a scalar of type **INTEGER** with a value in the range from 1 to n, where n equals the rank of *ARRAY*.

Return value:

The result is of the same type as *ARRAY*.

If *DIM* is absent, a scalar with the square root of the sum of all elements in *ARRAY* squared is returned. Otherwise, an array of rank $n-1$, where n equals the rank of *ARRAY*, and a shape similar to that of *ARRAY* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_sum
  REAL :: x(5) = [ real :: 1, 2, 3, 4, 5 ]
  print *, NORM2(x)  ! = sqrt(55.) ~ 7.416
END PROGRAM
```

9.209 NOT — Logical negation

Description:

NOT returns the bitwise Boolean inverse of *I*.

Standard: Fortran 95 and later, has overloads that are GNU extensions

Class: Elemental function

Syntax: RESULT = NOT(I)

Arguments:

I The type shall be INTEGER.

Return value:

The return type is INTEGER, of the same kind as the argument.

Specific names:

Name	Argument	Return type	Standard
NOT(A)	INTEGER A	INTEGER	Fortran 95 and later
BNOT(A)	INTEGER(1) A	INTEGER(1)	GNU extension
INOT(A)	INTEGER(2) A	INTEGER(2)	GNU extension
JNOT(A)	INTEGER(4) A	INTEGER(4)	GNU extension
KNOT(A)	INTEGER(8) A	INTEGER(8)	GNU extension

See also: [Section 9.137 \[IAND\], page 199](#), [Section 9.145 \[IEOR\], page 204](#), [Section 9.152 \[IOR\], page 208](#), [Section 9.141 \[IBITS\], page 202](#), [Section 9.142 \[IBSET\], page 202](#), [Section 9.140 \[IBCLR\], page 201](#)

9.210 NULL — Function that returns an disassociated pointer

Description:

Returns a disassociated pointer.

If *MOLD* is present, a disassociated pointer of the same type is returned, otherwise the type is determined by context.

In Fortran 95, *MOLD* is optional. Please note that Fortran 2003 includes cases where it is required.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: PTR => NULL([MOLD])

Arguments:
 MOLD (Optional) shall be a pointer of any association status and of any type.

Return value:
 A disassociated pointer.

Example:

```
REAL, POINTER, DIMENSION(:) :: VEC => NULL ()
```

See also: [Section 9.22 \[ASSOCIATED\]](#), page 124

9.211 NUM_IMAGES — Function that returns the number of images

Description:
 Returns the number of images.

Standard: Fortran 2008 and later. With *DISTANCE* or *FAILED* argument, Technical Specification (TS) 18508 or later

Class: Transformational function

Syntax: RESULT = NUM_IMAGES(DISTANCE, FAILED)

Arguments:
 DISTANCE (optional, intent(in)) Nonnegative scalar integer
 FAILED (optional, intent(in)) Scalar logical expression

Return value:
 Scalar default-kind integer. If *DISTANCE* is not present or has value 0, the number of images in the current team is returned. For values smaller or equal distance to the initial team, it returns the number of images index on the ancestor team which has a distance of *DISTANCE* from the invoking team. If *DISTANCE* is larger than the distance to the initial team, the number of images of the initial team is returned. If *FAILED* is not present the total number of images is returned; if it has the value *.TRUE.*, the number of failed images is returned, otherwise, the number of images which do have not the failed status.

Example:

```
INTEGER :: value[*]
INTEGER :: i
value = THIS_IMAGE()
SYNC ALL
IF (THIS_IMAGE() == 1) THEN
  DO i = 1, NUM_IMAGES()
    WRITE(*,'(2(a,i0))') 'value[' , i, ']' is ', value[i]
  END DO
END IF
```

See also: [Section 9.269 \[THIS_IMAGE\]](#), page 276, [Section 9.147 \[IMAGE_INDEX\]](#), page 205

9.212 OR — Bitwise logical OR

Description:

Bitwise logical OR.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. For integer arguments, programmers should consider the use of the [Section 9.152 \[IOR\], page 208](#) intrinsic defined by the Fortran standard.

Standard: GNU extension

Class: Function

Syntax: `RESULT = OR(I, J)`

Arguments:

<i>I</i>	The type shall be either a scalar <code>INTEGER</code> type or a scalar <code>LOGICAL</code> type or a boz-literal-constant.
<i>J</i>	The type shall be the same as the type of <i>I</i> or a boz-literal-constant. <i>I</i> and <i>J</i> shall not both be boz-literal-constants. If either <i>I</i> and <i>J</i> is a boz-literal-constant, then the other argument must be a scalar <code>INTEGER</code> .

Return value:

The return type is either a scalar `INTEGER` or a scalar `LOGICAL`. If the kind type parameters differ, then the smaller kind type is implicitly converted to larger kind, and the return has the larger kind. A boz-literal-constant is converted to an `INTEGER` with the kind type parameter of the other argument as-if a call to [Section 9.149 \[INT\], page 207](#) occurred.

Example:

```
PROGRAM test_or
  LOGICAL :: T = .TRUE., F = .FALSE.
  INTEGER :: a, b
  DATA a / Z'F' /, b / Z'3' /

  WRITE (*,*) OR(T, T), OR(T, F), OR(F, T), OR(F, F)
  WRITE (*,*) OR(a, b)
END PROGRAM
```

See also: Fortran 95 elemental function: [Section 9.152 \[IOR\], page 208](#)

9.213 PACK — Pack an array into an array of rank one

Description:

Stores the elements of *ARRAY* in an array of rank one.

The beginning of the resulting array is made up of elements whose *MASK* equals `TRUE`. Afterwards, positions are filled with elements taken from *VECTOR*.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: `RESULT = PACK(ARRAY, MASK[, VECTOR])`

Arguments:

ARRAY Shall be an array of any type.

MASK Shall be an array of type **LOGICAL** and of the same size as **ARRAY**. Alternatively, it may be a **LOGICAL** scalar.

VECTOR (Optional) shall be an array of the same type as **ARRAY** and of rank one. If present, the number of elements in **VECTOR** shall be equal to or greater than the number of true elements in **MASK**. If **MASK** is scalar, the number of elements in **VECTOR** shall be equal to or greater than the number of elements in **ARRAY**.

Return value:

The result is an array of rank one and the same type as that of **ARRAY**. If **VECTOR** is present, the result size is that of **VECTOR**, the number of **TRUE** values in **MASK** otherwise.

Example: Gathering nonzero elements from an array:

```
PROGRAM test_pack_1
  INTEGER :: m(6)
  m = (/ 1, 0, 0, 0, 5, 0 /)
  WRITE(*, FMT="(6(I0, ' '))") pack(m, m /= 0) ! "1 5"
END PROGRAM
```

Gathering nonzero elements from an array and appending elements from **VECTOR**:

```
PROGRAM test_pack_2
  INTEGER :: m(4)
  m = (/ 1, 0, 0, 2 /)
  WRITE(*, FMT="(4(I0, ' '))") pack(m, m /= 0, (/ 0, 0, 3, 4 /)) ! "1 2 3 4"■
END PROGRAM
```

See also: [Section 9.282 \[UNPACK\], page 283](#)

9.214 PARITY — Reduction with exclusive OR

Description:

Calculates the parity, i.e. the reduction using **.XOR.**, of **MASK** along dimension **DIM**.

Standard: Fortran 2008 and later

Class: Transformational function

Syntax:

```
RESULT = PARITY(MASK[, DIM])
```

Arguments:

LOGICAL Shall be an array of type **LOGICAL**

DIM (Optional) shall be a scalar of type **INTEGER** with a value in the range from 1 to n, where n equals the rank of **MASK**.

Return value:

The result is of the same type as **MASK**.

If *DIM* is absent, a scalar with the parity of all elements in *MASK* is returned, i.e. true if an odd number of elements is `.true.` and false otherwise. If *DIM* is present, an array of rank $n - 1$, where n equals the rank of *ARRAY*, and a shape similar to that of *MASK* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_sum
  LOGICAL :: x(2) = [ .true., .false. ]
  print *, PARITY(x) ! prints "T" (true).
END PROGRAM
```

9.215 PERROR — Print system error message

Description:

Prints (on the C `stderr` stream) a newline-terminated error message corresponding to the last system error. This is prefixed by *STRING*, a colon and a space. See `perror(3)`.

Standard: GNU extension

Class: Subroutine

Syntax: CALL PERROR(STRING)

Arguments:

STRING A scalar of type CHARACTER and of the default kind.

See also: [Section 9.146 \[IERRNO\]](#), page 205

9.216 POPCNT — Number of bits set

Description:

POPCNT(I) returns the number of bits set ('1' bits) in the binary representation of I.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = POPCNT(I)

Arguments:

I Shall be of type INTEGER.

Return value:

The return value is of type INTEGER and of the default integer kind.

See also: [Section 9.217 \[POPPAR\]](#), page 246, [Section 9.167 \[LEADZ\]](#), page 217, [Section 9.273 \[TRAILZ\]](#), page 279

Example:

```
program test_population
  print *, popcnt(127),        poppar(127)
  print *, popcnt(huge(0_4)), poppar(huge(0_4))
  print *, popcnt(huge(0_8)), poppar(huge(0_8))
end program test_population
```

9.217 POPPAR — Parity of the number of bits set

Description:

POPPAR(I) returns parity of the integer I, i.e. the parity of the number of bits set ('1' bits) in the binary representation of I. It is equal to 0 if I has an even number of bits set, and 1 for an odd number of '1' bits.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = POPPAR(I)

Arguments:

I Shall be of type INTEGER.

Return value:

The return value is of type INTEGER and of the default integer kind.

See also: Section 9.216 [POPCNT], page 245, Section 9.167 [LEADZ], page 217,
Section 9.273 [TRAILZ], page 279

Example:

```
program test_population
  print *, popcnt(127),      poppar(127)
  print *, popcnt(huge(0_4)), poppar(huge(0_4))
  print *, popcnt(huge(0_8)), poppar(huge(0_8))
end program test_population
```

9.218 PRECISION — Decimal precision of a real kind

Description:

PRECISION(X) returns the decimal precision in the model of the type of X.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: RESULT = PRECISION(X)

Arguments:

X Shall be of type REAL or COMPLEX.

Return value:

The return value is of type INTEGER and of the default integer kind.

See also: Section 9.242 [SELECTED_REAL_KIND], page 259, Section 9.227 [RANGE],
page 251

Example:

```
program prec_and_range
  real(kind=4) :: x(2)
  complex(kind=8) :: y

  print *, precision(x), range(x)
  print *, precision(y), range(y)
end program prec_and_range
```

9.219 PRESENT — Determine whether an optional dummy argument is specified

Description:

Determines whether an optional dummy argument is present.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: `RESULT = PRESENT(A)`

Arguments:

A May be of any type and may be a pointer, scalar or array value, or a dummy procedure. It shall be the name of an optional dummy argument accessible within the current sub-routine or function.

Return value:

Returns either `TRUE` if the optional argument *A* is present, or `FALSE` otherwise.

Example:

```
PROGRAM test_present
  WRITE(*,*) f(), f(42)      ! "F T"
CONTAINS
  LOGICAL FUNCTION f(x)
    INTEGER, INTENT(IN), OPTIONAL :: x
    f = PRESENT(x)
  END FUNCTION
END PROGRAM
```

9.220 PRODUCT — Product of array elements

Description:

Multiplies the elements of *ARRAY* along dimension *DIM* if the corresponding element in *MASK* is `TRUE`.

Standard: Fortran 95 and later

Class: Transformational function

Syntax:

```
RESULT = PRODUCT(ARRAY[, MASK])
RESULT = PRODUCT(ARRAY, DIM[, MASK])
```

Arguments:

ARRAY Shall be an array of type `INTEGER`, `REAL` or `COMPLEX`.
DIM (Optional) shall be a scalar of type `INTEGER` with a value in the range from 1 to *n*, where *n* equals the rank of *ARRAY*.
MASK (Optional) shall be of type `LOGICAL` and either be a scalar or an array of the same shape as *ARRAY*.

Return value:

The result is of the same type as *ARRAY*.

If *DIM* is absent, a scalar with the product of all elements in *ARRAY* is returned. Otherwise, an array of rank *n*-1, where *n* equals the rank of *ARRAY*, and a shape similar to that of *ARRAY* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_product
  INTEGER :: x(5) = (/ 1, 2, 3, 4 ,5 /)
  print *, PRODUCT(x)           ! all elements, product = 120
  print *, PRODUCT(x, MASK=MOD(x, 2)==1) ! odd elements, product = 15
END PROGRAM
```

See also: [Section 9.262 \[SUM\]](#), page 272

9.221 RADIX — Base of a model number

Description:

RADIX(X) returns the base of the model representing the entity *X*.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: **RESULT = RADIX(X)**

Arguments:

X Shall be of type **INTEGER** or **REAL**

Return value:

The return value is a scalar of type **INTEGER** and of the default integer kind.

See also: [Section 9.242 \[SELECTED_REAL_KIND\]](#), page 259

Example:

```
program test_radix
  print *, "The radix for the default integer kind is", radix(0)
  print *, "The radix for the default real kind is", radix(0.0)
end program test_radix
```

9.222 RAN — Real pseudo-random number

Description:

For compatibility with HP FORTRAN 77/iX, the **RAN** intrinsic is provided as an alias for **RAND**. See [Section 9.223 \[RAND\]](#), page 248 for complete documentation.

Standard: GNU extension

Class: Function

See also: [Section 9.223 \[RAND\]](#), page 248, [Section 9.225 \[RANDOM_NUMBER\]](#), page 250

9.223 RAND — Real pseudo-random number

Description:

RAND(FLAG) returns a pseudo-random number from a uniform distribution between 0 and 1. If *FLAG* is 0, the next number in the current sequence is returned; if *FLAG* is 1, the generator is restarted by **CALL SRAND(0)**; if *FLAG* has any other value, it is used as a new seed with **SRAND**.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. It implements a simple modulo generator as provided by g77. For new code, one should consider the use of [Section 9.225 \[RANDOM_NUMBER\]](#), [page 250](#) as it implements a superior algorithm.

Standard: GNU extension

Class: Function

Syntax: RESULT = RAND(I)

Arguments:

I Shall be a scalar INTEGER of kind 4.

Return value:

The return value is of REAL type and the default kind.

Example:

```
program test_rand
  integer,parameter :: seed = 86456

  call srand(seed)
  print *, rand(), rand(), rand(), rand()
  print *, rand(seed), rand(), rand(), rand()
end program test_rand
```

See also: [Section 9.259 \[SRAND\]](#), [page 269](#), [Section 9.225 \[RANDOM_NUMBER\]](#), [page 250](#)

9.224 RANDOM_INIT — Initialize a pseudo-random number generator

Description:

Initializes the state of the pseudorandom number generator used by RANDOM_NUMBER.

Standard: Fortran 2018

Class: Subroutine

Syntax: CALL RANDOM_INIT(REPEATABLE, IMAGE_DISTINCT)

Arguments:

REPEATABLE Shall be a scalar with a LOGICAL type, and it is INTENT(IN). If it is `.true.`, the seed is set to a processor-dependent value that is the same each time RANDOM_INIT is called from the same image. The term “same image” means a single instance of program execution. The sequence of random numbers is different for repeated execution of the program. If it is `.false.`, the seed is set to a processor-dependent value.

IMAGE_DISTINCT Shall be a scalar with a LOGICAL type, and it is INTENT(IN). If it is `.true.`, the seed is set to a processor-dependent value that is distinct from the seed set by a call to RANDOM_INIT in another image. If it is `.false.`, the seed is set to a value that does not depend on which image called RANDOM_INIT.

Example:

```

program test_random_seed
  implicit none
  real x(3), y(3)
  call random_init(.true., .true.)
  call random_number(x)
  call random_init(.true., .true.)
  call random_number(y)
  ! x and y are the same sequence
  if (any(x /= y)) call abort
end program test_random_seed

```

See also: [Section 9.225 \[RANDOM_NUMBER\]](#), page 250, [Section 9.226 \[RANDOM_SEED\]](#), page 250

9.225 RANDOM_NUMBER — Pseudo-random number

Description:

Returns a single pseudorandom number or an array of pseudorandom numbers from the uniform distribution over the range $0 \leq x < 1$.

The runtime-library implements the xorshift1024* random number generator (RNG). This generator has a period of $2^{1024} - 1$, and when using multiple threads up to 2^{512} threads can each generate 2^{512} random numbers before any aliasing occurs.

Note that in a multi-threaded program (e.g. using OpenMP directives), each thread will have its own random number state. For details of the seeding procedure, see the documentation for the `RANDOM_SEED` intrinsic.

Standard: Fortran 95 and later

Class: Subroutine

Syntax: `RANDOM_NUMBER(HARVEST)`

Arguments:

`HARVEST` Shall be a scalar or an array of type `REAL`.

Example:

```

program test_random_number
  REAL :: r(5,5)
  CALL RANDOM_NUMBER(r)
end program

```

See also: [Section 9.226 \[RANDOM_SEED\]](#), page 250, [Section 9.224 \[RANDOM_INIT\]](#), page 249

9.226 RANDOM_SEED — Initialize a pseudo-random number sequence

Description:

Restarts or queries the state of the pseudorandom number generator used by `RANDOM_NUMBER`.

If `RANDOM_SEED` is called without arguments, it is seeded with random data retrieved from the operating system.

As an extension to the Fortran standard, the GFortran `RANDOM_NUMBER` supports multiple threads. Each thread in a multi-threaded program has its own seed. When `RANDOM_SEED` is called either without arguments or with the *PUT* argument, the given seed is copied into a master seed as well as the seed of the current thread. When a new thread uses `RANDOM_NUMBER` for the first time, the seed is copied from the master seed, and forwarded $N * 2^{512}$ steps to guarantee that the random stream does not alias any other stream in the system, where N is the number of threads that have used `RANDOM_NUMBER` so far during the program execution.

Standard: Fortran 95 and later

Class: Subroutine

Syntax: `CALL RANDOM_SEED([SIZE, PUT, GET])`

Arguments:

<i>SIZE</i>	(Optional) Shall be a scalar and of type default <code>INTEGER</code> , with <code>INTENT(OUT)</code> . It specifies the minimum size of the arrays used with the <i>PUT</i> and <i>GET</i> arguments.
<i>PUT</i>	(Optional) Shall be an array of type default <code>INTEGER</code> and rank one. It is <code>INTENT(IN)</code> and the size of the array must be larger than or equal to the number returned by the <i>SIZE</i> argument.
<i>GET</i>	(Optional) Shall be an array of type default <code>INTEGER</code> and rank one. It is <code>INTENT(OUT)</code> and the size of the array must be larger than or equal to the number returned by the <i>SIZE</i> argument.

Example:

```

program test_random_seed
  implicit none
  integer, allocatable :: seed(:)
  integer :: n

  call random_seed(size = n)
  allocate(seed(n))
  call random_seed(get=seed)
  write (*, *) seed
end program test_random_seed

```

See also: [Section 9.225 \[RANDOM_NUMBER\]](#), page 250, [Section 9.224 \[RANDOM_INIT\]](#), page 249

9.227 RANGE — Decimal exponent range

Description:

`RANGE(X)` returns the decimal exponent range in the model of the type of `X`.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: `RESULT = RANGE(X)`

Arguments:

X Shall be of type `INTEGER`, `REAL` or `COMPLEX`.

Return value:

The return value is of type `INTEGER` and of the default integer kind.

See also: [Section 9.242 \[SELECTED-REAL-KIND\]](#), page 259, [Section 9.218 \[PRECISION\]](#), page 246

Example: See `PRECISION` for an example.

9.228 RANK — Rank of a data object

Description:

`RANK(A)` returns the rank of a scalar or array data object.

Standard: Technical Specification (TS) 29113

Class: Inquiry function

Syntax: `RESULT = RANK(A)`

Arguments:

A can be of any type

Return value:

The return value is of type `INTEGER` and of the default integer kind. For arrays, their rank is returned; for scalars zero is returned.

Example:

```

program test_rank
  integer :: a
  real, allocatable :: b(:, :)

  print *, rank(a), rank(b) ! Prints: 0 2
end program test_rank

```

9.229 REAL — Convert to real type

Description:

`REAL(A [, KIND])` converts its argument *A* to a real type. The `REALPART` function is provided for compatibility with `g77`, and its use is strongly discouraged.

Standard: Fortran 77 and later

Class: Elemental function

Syntax:

```

RESULT = REAL(A [, KIND])
RESULT = REALPART(Z)

```

Arguments:

A Shall be `INTEGER`, `REAL`, or `COMPLEX`.
KIND (Optional) An `INTEGER` initialization expression indicating the kind parameter of the result.

Return value:

These functions return a **REAL** variable or array under the following rules:

- (A) **REAL(A)** is converted to a default real type if *A* is an integer or real variable.
- (B) **REAL(A)** is converted to a real type with the kind type parameter of *A* if *A* is a complex variable.
- (C) **REAL(A, KIND)** is converted to a real type with kind type parameter *KIND* if *A* is a complex, integer, or real variable.

Example:

```
program test_real
  complex :: x = (1.0, 2.0)
  print *, real(x), real(x,8), realpart(x)
end program test_real
```

Specific names:

Name	Argument	Return type	Standard
FLOAT(A)	INTEGER(4)	REAL(4)	Fortran 77 and later
DFLOAT(A)	INTEGER(4)	REAL(8)	GNU extension
FLOATI(A)	INTEGER(2)	REAL(4)	GNU extension
FLOATJ(A)	INTEGER(4)	REAL(4)	GNU extension
FLOATK(A)	INTEGER(8)	REAL(4)	GNU extension
SNGL(A)	INTEGER(8)	REAL(4)	Fortran 77 and later

See also: [Section 9.83 \[DBLE\]](#), page 165

9.230 RENAME — Rename a file

Description:

Renames a file from file *PATH1* to *PATH2*. A null character (**CHAR(0)**) can be used to mark the end of the names in *PATH1* and *PATH2*; otherwise, trailing blanks in the file names are ignored. If the *STATUS* argument is supplied, it contains 0 on success or a nonzero error code upon return; see **rename(2)**.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL RENAME(PATH1, PATH2 [, STATUS])
STATUS = RENAME(PATH1, PATH2)
```

Arguments:

PATH1 Shall be of default **CHARACTER** type.
PATH2 Shall be of default **CHARACTER** type.
STATUS (Optional) Shall be of default **INTEGER** type.

See also: [Section 9.172 \[LINK\]](#), page 220

9.231 REPEAT — Repeated string concatenation

Description:

Concatenates *NCOPIES* copies of a string.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: `RESULT = REPEAT(STRING, NCOPIES)`

Arguments:

STRING Shall be scalar and of type CHARACTER.
NCOPIES Shall be scalar and of type INTEGER.

Return value:

A new scalar of type CHARACTER built up from *NCOPIES* copies of *STRING*.

Example:

```
program test_repeat
  write(*,*) repeat("x", 5)   ! "xxxxx"
end program
```

9.232 RESHAPE — Function to reshape an array

Description:

Reshapes *SOURCE* to correspond to *SHAPE*. If necessary, the new array may be padded with elements from *PAD* or permuted as defined by *ORDER*.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: `RESULT = RESHAPE(SOURCE, SHAPE[, PAD, ORDER])`

Arguments:

SOURCE Shall be an array of any type.
SHAPE Shall be of type INTEGER and an array of rank one. Its values must be positive or zero.
PAD (Optional) shall be an array of the same type as *SOURCE*.
ORDER (Optional) shall be of type INTEGER and an array of the same shape as *SHAPE*. Its values shall be a permutation of the numbers from 1 to n, where n is the size of *SHAPE*. If *ORDER* is absent, the natural ordering shall be assumed.

Return value:

The result is an array of shape *SHAPE* with the same type as *SOURCE*.

Example:

```
PROGRAM test_reshape
  INTEGER, DIMENSION(4) :: x
  WRITE(*,*) SHAPE(x)                ! prints "4"
  WRITE(*,*) SHAPE(RESHAPE(x, (/2, 2/))) ! prints "2 2"
END PROGRAM
```

See also: [Section 9.244 \[SHAPE\]](#), page 261

9.233 RRSPACING — Reciprocal of the relative spacing

Description:

RRSPACING(*X*) returns the reciprocal of the relative spacing of model numbers near *X*.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = RRSPACING(*X*)

Arguments:

X Shall be of type REAL.

Return value:

The return value is of the same type and kind as *X*. The value returned is equal to ABS(FRACTION(*X*)) * FLOAT(RADIX(*X*))**DIGITS(*X*).

See also: [Section 9.256 \[SPACING\]](#), page 268

9.234 RSHIFT — Right shift bits

Description:

RSHIFT returns a value corresponding to *I* with all of the bits shifted right by *SHIFT* places. *SHIFT* shall be nonnegative and less than or equal to BIT_SIZE(*I*), otherwise the result value is undefined. Bits shifted out from the right end are lost. The fill is arithmetic: the bits shifted in from the left end are equal to the leftmost bit, which in two's complement representation is the sign bit.

This function has been superseded by the SHIFTA intrinsic, which is standard in Fortran 2008 and later.

Standard: GNU extension

Class: Elemental function

Syntax: RESULT = RSHIFT(*I*, *SHIFT*)

Arguments:

I The type shall be INTEGER.
SHIFT The type shall be INTEGER.

Return value:

The return value is of type INTEGER and of the same kind as *I*.

See also: [Section 9.159 \[ISHFT\]](#), page 212, [Section 9.160 \[ISHFTC\]](#), page 213, [Section 9.182 \[LSHIFT\]](#), page 225, [Section 9.245 \[SHIFTA\]](#), page 262, [Section 9.247 \[SHIFTR\]](#), page 262, [Section 9.246 \[SHIFTL\]](#), page 262

9.235 SAME_TYPE_AS — Query dynamic types for equality

Description:

Query dynamic types for equality.

Standard: Fortran 2003 and later

Class: Inquiry function

Syntax: RESULT = SAME_TYPE_AS(A, B)

Arguments:

<i>A</i>	Shall be an object of extensible declared type or unlimited polymorphic.
<i>B</i>	Shall be an object of extensible declared type or unlimited polymorphic.

Return value:

The return value is a scalar of type default logical. It is true if and only if the dynamic type of A is the same as the dynamic type of B.

See also: [Section 9.104 \[EXTENDS_TYPE_OF\]](#), page 178

9.236 SCALE — Scale a real value

Description:

SCALE(X,I) returns $X * RADIX(X)^{**I}$.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = SCALE(X, I)

Arguments:

<i>X</i>	The type of the argument shall be a REAL .
<i>I</i>	The type of the argument shall be a INTEGER .

Return value:

The return value is of the same type and kind as X. Its value is $X * RADIX(X)^{**I}$.

Example:

```

program test_scale
  real :: x = 178.1387e-4
  integer :: i = 5
  print *, scale(x,i), x*radix(x)**i
end program test_scale

```

9.237 SCAN — Scan a string for the presence of a set of characters

Description:

Scans a *STRING* for any of the characters in a *SET* of characters.

If *BACK* is either absent or equals **FALSE**, this function returns the position of the leftmost character of *STRING* that is in *SET*. If *BACK* equals **TRUE**, the rightmost position is returned. If no character of *SET* is found in *STRING*, the result is zero.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: `RESULT = SCAN(STRING, SET[, BACK [, KIND]])`

Arguments:

<i>STRING</i>	Shall be of type CHARACTER .
<i>SET</i>	Shall be of type CHARACTER .
<i>BACK</i>	(Optional) shall be of type LOGICAL .
<i>KIND</i>	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type **INTEGER** and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Example:

```
PROGRAM test_scan
  WRITE(*,*) SCAN("FORTRAN", "AO")           ! 2, found 'O'
  WRITE(*,*) SCAN("FORTRAN", "AO", .TRUE.)    ! 6, found 'A'
  WRITE(*,*) SCAN("FORTRAN", "C++")          ! 0, found none
END PROGRAM
```

See also: [Section 9.148 \[INDEX intrinsic\]](#), page 206, [Section 9.283 \[VERIFY\]](#), page 284

9.238 SECNDS — Time function

Description:

SECNDS(X) gets the time in seconds from the real-time system clock. *X* is a reference time, also in seconds. If this is zero, the time in seconds from midnight is returned. This function is non-standard and its use is discouraged.

Standard: GNU extension

Class: Function

Syntax: `RESULT = SECNDS (X)`

Arguments:

<i>T</i>	Shall be of type REAL(4) .
<i>X</i>	Shall be of type REAL(4) .

Return value:

None

Example:

```
program test_secnds
  integer :: i
  real(4) :: t1, t2
  print *, secnds (0.0)    ! seconds since midnight
```

```

      t1 = secnds (0.0)      ! reference time
      do i = 1, 10000000    ! do something
      end do
      t2 = secnds (t1)      ! elapsed time
      print *, "Something took ", t2, " seconds."
    end program test_secnds

```

9.239 SECOND — CPU time function

Description:

Returns a `REAL(4)` value representing the elapsed CPU time in seconds. This provides the same functionality as the standard `CPU_TIME` intrinsic, and is only included for backwards compatibility.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```

CALL SECOND(TIME)
TIME = SECOND()

```

Arguments:

TIME Shall be of type `REAL(4)`.

Return value:

In either syntax, *TIME* is set to the process's current runtime in seconds.

See also: [Section 9.79 \[CPU_TIME\]](#), page 162

9.240 SELECTED_CHAR_KIND — Choose character kind

Description:

`SELECTED_CHAR_KIND(NAME)` returns the kind value for the character set named *NAME*, if a character set with such a name is supported, or `-1` otherwise. Currently, supported character sets include “ASCII” and “DEFAULT”, which are equivalent, and “ISO_10646” (Universal Character Set, UCS-4) which is commonly known as Unicode.

Standard: Fortran 2003 and later

Class: Transformational function

Syntax: `RESULT = SELECTED_CHAR_KIND(NAME)`

Arguments:

NAME Shall be a scalar and of the default character type.

Example:

```

program character_kind
  use iso_fortran_env
  implicit none

```

```

integer, parameter :: ascii = selected_char_kind ("ascii")
integer, parameter :: ucs4  = selected_char_kind ('ISO_10646')

character(kind=ascii, len=26) :: alphabet
character(kind=ucs4,  len=30) :: hello_world

alphabet = ascii_"abcdefghijklmnopqrstuvwxyz"
hello_world = ucs4_'Hello World and Ni Hao -- ' &
              // char (int (z'4F60'), ucs4)      &
              // char (int (z'597D'), ucs4)

write (*,*) alphabet

open (output_unit, encoding='UTF-8')
write (*,*) trim (hello_world)
end program character_kind

```

9.241 SELECTED_INT_KIND — Choose integer kind

Description:

SELECTED_INT_KIND(R) return the kind value of the smallest integer type that can represent all values ranging from -10^R (exclusive) to 10^R (exclusive). If there is no integer kind that accommodates this range, SELECTED_INT_KIND returns -1 .

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = SELECTED_INT_KIND(R)

Arguments:

R Shall be a scalar and of type INTEGER.

Example:

```

program large_integers
  integer,parameter :: k5 = selected_int_kind(5)
  integer,parameter :: k15 = selected_int_kind(15)
  integer(kind=k5) :: i5
  integer(kind=k15) :: i15

  print *, huge(i5), huge(i15)

  ! The following inequalities are always true
  print *, huge(i5) >= 10_k5**5-1
  print *, huge(i15) >= 10_k15**15-1
end program large_integers

```

9.242 SELECTED_REAL_KIND — Choose real kind

Description:

SELECTED_REAL_KIND(P,R) returns the kind value of a real data type with decimal precision of at least P digits, exponent range of at least R, and with a radix of RADIX.

Standard: Fortran 95 and later, with RADIX Fortran 2008 or later

Class: Transformational function

Syntax: `RESULT = SELECTED_REAL_KIND([P, R, RADIX])`

Arguments:

P (Optional) shall be a scalar and of type `INTEGER`.

R (Optional) shall be a scalar and of type `INTEGER`.

RADIX (Optional) shall be a scalar and of type `INTEGER`.

Before Fortran 2008, at least one of the arguments *R* or *P* shall be present; since Fortran 2008, they are assumed to be zero if absent.

Return value:

`SELECTED_REAL_KIND` returns the value of the kind type parameter of a real data type with decimal precision of at least *P* digits, a decimal exponent range of at least *R*, and with the requested *RADIX*. If the *RADIX* parameter is absent, real kinds with any radix can be returned. If more than one real data type meet the criteria, the kind of the data type with the smallest decimal precision is returned. If no real data type matches the criteria, the result is

- 1 if the processor does not support a real data type with a precision greater than or equal to *P*, but the *R* and *RADIX* requirements can be fulfilled
- 2 if the processor does not support a real type with an exponent range greater than or equal to *R*, but *P* and *RADIX* are fulfillable
- 3 if *RADIX* but not *P* and *R* requirements are fulfillable
- 4 if *RADIX* and either *P* or *R* requirements are fulfillable
- 5 if there is no real type with the given *RADIX*

See also: [Section 9.218 \[PRECISION\], page 246](#), [Section 9.227 \[RANGE\], page 251](#), [Section 9.221 \[RADIX\], page 248](#)

Example:

```
program real_kinds
  integer,parameter :: p6 = selected_real_kind(6)
  integer,parameter :: p10r100 = selected_real_kind(10,100)
  integer,parameter :: r400 = selected_real_kind(r=400)
  real(kind=p6) :: x
  real(kind=p10r100) :: y
  real(kind=r400) :: z

  print *, precision(x), range(x)
  print *, precision(y), range(y)
  print *, precision(z), range(z)
end program real_kinds
```

9.243 SET_EXPONENT — Set the exponent of the model

Description:

`SET_EXPONENT(X, I)` returns the real number whose fractional part is that of *X* and whose exponent part is *I*.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: `RESULT = SET_EXPONENT(X, I)`

Arguments:

<i>X</i>	Shall be of type <code>REAL</code> .
<i>I</i>	Shall be of type <code>INTEGER</code> .

Return value:

The return value is of the same type and kind as *X*. The real number whose fractional part is that of *X* and whose exponent part if *I* is returned; it is `FRACTION(X) * RADIX(X)**I`.

Example:

```
PROGRAM test_setexp
  REAL :: x = 178.1387e-4
  INTEGER :: i = 17
  PRINT *, SET_EXPONENT(x, i), FRACTION(x) * RADIX(x)**i
END PROGRAM
```

9.244 SHAPE — Determine the shape of an array

Description:

Determines the shape of an array.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Inquiry function

Syntax: `RESULT = SHAPE(SOURCE [, KIND])`

Arguments:

<i>SOURCE</i>	Shall be an array or scalar of any type. If <i>SOURCE</i> is a pointer it must be associated and allocatable arrays must be allocated.
<i>KIND</i>	(Optional) An <code>INTEGER</code> initialization expression indicating the kind parameter of the result.

Return value:

An `INTEGER` array of rank one with as many elements as *SOURCE* has dimensions. The elements of the resulting array correspond to the extent of *SOURCE* along the respective dimensions. If *SOURCE* is a scalar, the result is the rank one array of size zero. If *KIND* is absent, the return value has the default integer kind otherwise the specified kind.

Example:

```
PROGRAM test_shape
  INTEGER, DIMENSION(-1:1, -1:2) :: A
  WRITE(*,*) SHAPE(A)           ! (/ 3, 4 /)
  WRITE(*,*) SIZE(SHAPE(42))    ! (/ /)
END PROGRAM
```

See also: [Section 9.232 \[RESHAPE\], page 254](#), [Section 9.253 \[SIZE\], page 266](#)

9.245 SHIFTA — Right shift with fill

Description:

SHIFTA returns a value corresponding to *I* with all of the bits shifted right by *SHIFT* places. *SHIFT* shall be nonnegative and less than or equal to `BIT_SIZE(I)`, otherwise the result value is undefined. Bits shifted out from the right end are lost. The fill is arithmetic: the bits shifted in from the left end are equal to the leftmost bit, which in two's complement representation is the sign bit.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = SHIFTA(I, SHIFT)`

Arguments:

<i>I</i>	The type shall be <code>INTEGER</code> .
<i>SHIFT</i>	The type shall be <code>INTEGER</code> .

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

See also: [Section 9.246 \[SHIFTL\], page 262](#), [Section 9.247 \[SHIFTR\], page 262](#)

9.246 SHIFTL — Left shift

Description:

SHIFTL returns a value corresponding to *I* with all of the bits shifted left by *SHIFT* places. *SHIFT* shall be nonnegative and less than or equal to `BIT_SIZE(I)`, otherwise the result value is undefined. Bits shifted out from the left end are lost, and bits shifted in from the right end are set to 0.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = SHIFTL(I, SHIFT)`

Arguments:

<i>I</i>	The type shall be <code>INTEGER</code> .
<i>SHIFT</i>	The type shall be <code>INTEGER</code> .

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

See also: [Section 9.245 \[SHIFTA\], page 262](#), [Section 9.247 \[SHIFTR\], page 262](#)

9.247 SHIFTR — Right shift

Description:

SHIFTR returns a value corresponding to *I* with all of the bits shifted right by *SHIFT* places. *SHIFT* shall be nonnegative and less than or equal to `BIT_SIZE(I)`, otherwise the result value is undefined. Bits shifted out from the right end are lost, and bits shifted in from the left end are set to 0.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: `RESULT = SHIFTR(I, SHIFT)`

Arguments:

I The type shall be `INTEGER`.

SHIFT The type shall be `INTEGER`.

Return value:

The return value is of type `INTEGER` and of the same kind as *I*.

See also: [Section 9.245 \[SHIFTA\]](#), page 262, [Section 9.246 \[SHIFTL\]](#), page 262

9.248 SIGN — Sign copying function

Description:

`SIGN(A,B)` returns the value of *A* with the sign of *B*.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = SIGN(A, B)`

Arguments:

A Shall be of type `INTEGER` or `REAL`

B Shall be of the same type and kind as *A*

Return value:

The kind of the return value is that of *A* and *B*. If $B \geq 0$ then the result is `ABS(A)`, else it is `-ABS(A)`.

Example:

```
program test_sign
  print *, sign(-12,1)
  print *, sign(-12,0)
  print *, sign(-12,-1)

  print *, sign(-12.,1.)
  print *, sign(-12.,0.)
  print *, sign(-12.,-1.)
end program test_sign
```

Specific names:

Name	Arguments	Return type	Standard
<code>SIGN(A,B)</code>	<code>REAL(4) A, B</code>	<code>REAL(4)</code>	f77, gnu
<code>ISIGN(A,B)</code>	<code>INTEGER(4) A,</code> <code>B</code>	<code>INTEGER(4)</code>	f77, gnu
<code>DSIGN(A,B)</code>	<code>REAL(8) A, B</code>	<code>REAL(8)</code>	f77, gnu

9.249 SIGNAL — Signal handling subroutine (or function)

Description:

`SIGNAL(NUMBER, HANDLER [, STATUS])` causes external subroutine *HANDLER* to be executed with a single integer argument when signal *NUMBER* occurs. If *HANDLER* is an integer, it can be used to turn off handling of signal *NUMBER* or revert to its default action. See `signal(2)`.

If `SIGNAL` is called as a subroutine and the *STATUS* argument is supplied, it is set to the value returned by `signal(2)`.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL SIGNAL(NUMBER, HANDLER [, STATUS])
STATUS = SIGNAL(NUMBER, HANDLER)
```

Arguments:

<i>NUMBER</i>	Shall be a scalar integer, with <code>INTENT(IN)</code>
<i>HANDLER</i>	Signal handler (<code>INTEGER FUNCTION</code> or <code>SUBROUTINE</code>) or dummy/global <code>INTEGER</code> scalar. <code>INTEGER</code> . It is <code>INTENT(IN)</code> .
<i>STATUS</i>	(Optional) <i>STATUS</i> shall be a scalar integer. It has <code>INTENT(OUT)</code> .

Return value:

The `SIGNAL` function returns the value returned by `signal(2)`.

Example:

```
program test_signal
  intrinsic signal
  external handler_print

  call signal (12, handler_print)
  call signal (10, 1)

  call sleep (30)
end program test_signal
```

9.250 SIN — Sine function

Description:

`SIN(X)` computes the sine of *X*.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: `RESULT = SIN(X)`

Arguments:

<i>X</i>	The type shall be <code>REAL</code> or <code>COMPLEX</code> .
----------	---

Return value:

The return value has same type and kind as *X*.

Example:

```
program test_sin
  real :: x = 0.0
  x = sin(x)
end program test_sin
```

Specific names:

Name	Argument	Return type	Standard
SIN(X)	REAL(4) X	REAL(4)	f77, gnu
DSIN(X)	REAL(8) X	REAL(8)	f95, gnu
CSIN(X)	COMPLEX(4) X	COMPLEX(4)	f95, gnu
ZSIN(X)	COMPLEX(8) X	COMPLEX(8)	f95, gnu
CDSIN(X)	COMPLEX(8) X	COMPLEX(8)	f95, gnu

See also: Inverse function: [Section 9.19 \[ASIN\], page 123](#) Degrees function: [Section 9.251 \[SIND\], page 265](#)

9.251 SIND — Sine function, degrees

Description:

SIND(X) computes the sine of X in degrees.

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘-fdec-math’.

Class: Elemental function

Syntax: RESULT = SIND(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has same type and kind as X, and its value is in degrees.

Example:

```
program test_sind
  real :: x = 0.0
  x = sind(x)
end program test_sind
```

Specific names:

Name	Argument	Return type	Standard
SIND(X)	REAL(4) X	REAL(4)	GNU Extension
DSIND(X)	REAL(8) X	REAL(8)	GNU Extension
CSIND(X)	COMPLEX(4) X	COMPLEX(4)	GNU Extension
ZSIND(X)	COMPLEX(8) X	COMPLEX(8)	GNU Extension
CDSIND(X)	COMPLEX(8) X	COMPLEX(8)	GNU Extension

See also: Inverse function: [Section 9.20 \[ASIND\], page 123](#) Radians function: [Section 9.250 \[SIN\], page 264](#)

9.252 SINH — Hyperbolic sine function

Description:

SINH(X) computes the hyperbolic sine of X.

Standard: Fortran 95 and later, for a complex argument Fortran 2008 or later

Class: Elemental function

Syntax: RESULT = SINH(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has same type and kind as X.

Example:

```
program test_sinh
  real(8) :: x = - 1.0_8
  x = sinh(x)
end program test_sinh
```

Specific names:

Name	Argument	Return type	Standard
SINH(X)	REAL(4) X	REAL(4)	Fortran 95 and later
DSINH(X)	REAL(8) X	REAL(8)	Fortran 95 and later

See also: [Section 9.21 \[ASINH\]](#), [page 124](#)

9.253 SIZE — Determine the size of an array

Description:

Determine the extent of *ARRAY* along a specified dimension *DIM*, or the total number of elements in *ARRAY* if *DIM* is absent.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Inquiry function

Syntax: RESULT = SIZE(ARRAY[, DIM [, KIND]])

Arguments:

ARRAY Shall be an array of any type. If *ARRAY* is a pointer it must be associated and allocatable arrays must be allocated.

DIM (Optional) shall be a scalar of type *INTEGER* and its value shall be in the range from 1 to n, where n equals the rank of *ARRAY*.

KIND (Optional) An *INTEGER* initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type *INTEGER* and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Example:

```

PROGRAM test_size
  WRITE(*,*) SIZE((/ 1, 2 /))    ! 2
END PROGRAM

```

See also: [Section 9.244 \[SHAPE\]](#), page 261, [Section 9.232 \[RESHAPE\]](#), page 254

9.254 SIZEOF — Size in bytes of an expression

Description:

SIZEOF(X) calculates the number of bytes of storage the expression X occupies.

Standard: GNU extension

Class: Inquiry function

Syntax: N = SIZEOF(X)

Arguments:

X The argument shall be of any type, rank or shape.

Return value:

The return value is of type integer and of the system-dependent kind *C_SIZE_T* (from the *ISO_C_BINDING* module). Its value is the number of bytes occupied by the argument. If the argument has the *POINTER* attribute, the number of bytes of the storage area pointed to is returned. If the argument is of a derived type with *POINTER* or *ALLOCATABLE* components, the return value does not account for the sizes of the data pointed to by these components. If the argument is polymorphic, the size according to the dynamic type is returned. The argument may not be a procedure or procedure pointer. Note that the code assumes for arrays that those are contiguous; for contiguous arrays, it returns the storage or an array element multiplied by the size of the array.

Example:

```

integer :: i
real :: r, s(5)
print *, (sizeof(s)/sizeof(r) == 5)
end

```

The example will print *.TRUE.* unless you are using a platform where default *REAL* variables are unusually padded.

See also: [Section 9.57 \[C_SIZEOF\]](#), page 147, [Section 9.261 \[STORAGE_SIZE\]](#), page 271

9.255 SLEEP — Sleep for the specified number of seconds

Description:

Calling this subroutine causes the process to pause for *SECONDS* seconds.

Standard: GNU extension

Class: Subroutine

Syntax: CALL SLEEP(SECONDS)

Arguments:

SECONDS The type shall be of default *INTEGER*.

Example:

```
program test_sleep
  call sleep(5)
end
```

9.256 SPACING — Smallest distance between two numbers of a given type

Description:

Determines the distance between the argument *X* and the nearest adjacent number of the same type.

Standard: Fortran 95 and later

Class: Elemental function

Syntax: RESULT = SPACING(*X*)

Arguments:

X Shall be of type REAL.

Return value:

The result is of the same type as the input argument *X*.

Example:

```
PROGRAM test_spacing
  INTEGER, PARAMETER :: SGL = SELECTED_REAL_KIND(p=6, r=37)
  INTEGER, PARAMETER :: DBL = SELECTED_REAL_KIND(p=13, r=200)

  WRITE(*,*) spacing(1.0_SGL)      ! "1.1920929E-07"          on i686
  WRITE(*,*) spacing(1.0_DBL)      ! "2.220446049250313E-016" on i686
END PROGRAM
```

See also: [Section 9.233 \[RRSPACING\]](#), page 255

9.257 SPREAD — Add a dimension to an array

Description:

Replicates a *SOURCE* array *NCOPIES* times along a specified dimension *DIM*.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = SPREAD(*SOURCE*, *DIM*, *NCOPIES*)

Arguments:

SOURCE Shall be a scalar or an array of any type and a rank less than seven.

DIM Shall be a scalar of type INTEGER with a value in the range from 1 to n+1, where n equals the rank of *SOURCE*.

NCOPIES Shall be a scalar of type INTEGER.

Return value:

The result is an array of the same type as *SOURCE* and has rank n+1 where n equals the rank of *SOURCE*.

Example:

```
PROGRAM test_spread
  INTEGER :: a = 1, b(2) = (/ 1, 2 /)
  WRITE(*,*) SPREAD(A, 1, 2)      ! "1 1"
  WRITE(*,*) SPREAD(B, 1, 2)      ! "1 1 2 2"
END PROGRAM
```

See also: [Section 9.282 \[UNPACK\]](#), page 283

9.258 SQRT — Square-root function

Description:

SQRT(X) computes the square root of X.

Standard: Fortran 77 and later

Class: Elemental function

Syntax: RESULT = SQRT(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value is of type REAL or COMPLEX. The kind type parameter is the same as X.

Example:

```
program test_sqrt
  real(8) :: x = 2.0_8
  complex :: z = (1.0, 2.0)
  x = sqrt(x)
  z = sqrt(z)
end program test_sqrt
```

Specific names:

Name	Argument	Return type	Standard
SQRT(X)	REAL(4) X	REAL(4)	Fortran 95 and later
DSQRT(X)	REAL(8) X	REAL(8)	Fortran 95 and later
CSQRT(X)	COMPLEX(4) X	COMPLEX(4)	Fortran 95 and later
ZSQRT(X)	COMPLEX(8) X	COMPLEX(8)	GNU extension
CDSQRT(X)	COMPLEX(8) X	COMPLEX(8)	GNU extension

9.259 SRAND — Reinitialize the random number generator

Description:

SRAND reinitializes the pseudo-random number generator called by RAND and IRAND. The new seed used by the generator is specified by the required argument SEED.

Standard: GNU extension

Class: Subroutine

Syntax: CALL SRAND(SEED)

Arguments:

SEED Shall be a scalar `INTEGER(kind=4)`.

Return value:

Does not return anything.

Example: See `RAND` and `IRAND` for examples.

Notes: The Fortran standard specifies the intrinsic subroutines `RANDOM_SEED` to initialize the pseudo-random number generator and `RANDOM_NUMBER` to generate pseudo-random numbers. These subroutines should be used in new codes.

Please note that in GNU Fortran, these two sets of intrinsics (`RAND`, `IRAND` and `SRAND` on the one hand, `RANDOM_NUMBER` and `RANDOM_SEED` on the other hand) access two independent pseudo-random number generators.

See also: [Section 9.223 \[RAND\], page 248](#), [Section 9.226 \[RANDOM_SEED\], page 250](#),
[Section 9.225 \[RANDOM_NUMBER\], page 250](#)

9.260 STAT — Get file status

Description:

This function returns information about a file. No permissions are required on the file itself, but execute (search) permission is required on all of the directories in path that lead to the file.

The elements that are obtained and stored in the array `VALUES`:

<code>VALUES(1)</code>	Device ID
<code>VALUES(2)</code>	Inode number
<code>VALUES(3)</code>	File mode
<code>VALUES(4)</code>	Number of links
<code>VALUES(5)</code>	Owner's uid
<code>VALUES(6)</code>	Owner's gid
<code>VALUES(7)</code>	ID of device containing directory entry for file (0 if not available)
<code>VALUES(8)</code>	File size (bytes)
<code>VALUES(9)</code>	Last access time
<code>VALUES(10)</code>	Last modification time
<code>VALUES(11)</code>	Last file status change time
<code>VALUES(12)</code>	Preferred I/O block size (-1 if not available)
<code>VALUES(13)</code>	Number of blocks allocated (-1 if not available)

Not all these elements are relevant on all systems. If an element is not relevant, it is returned as 0.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL STAT(NAME, VALUES [, STATUS])
```

```
STATUS = STAT(NAME, VALUES)
```

Arguments:

<i>NAME</i>	The type shall be CHARACTER , of the default kind and a valid path within the file system.
<i>VALUES</i>	The type shall be INTEGER(4) , DIMENSION(13) .
<i>STATUS</i>	(Optional) status flag of type INTEGER(4) . Returns 0 on success and a system specific error code otherwise.

Example:

```
PROGRAM test_stat
  INTEGER, DIMENSION(13) :: buff
  INTEGER :: status

  CALL STAT("/etc/passwd", buff, status)

  IF (status == 0) THEN
    WRITE (*, FMT="(Device ID:', T30, I19)") buff(1)
    WRITE (*, FMT="(Inode number:', T30, I19)") buff(2)
    WRITE (*, FMT="(File mode (octal):', T30, I19)") buff(3)
    WRITE (*, FMT="(Number of links:', T30, I19)") buff(4)
    WRITE (*, FMT="(Owner's uid:', T30, I19)") buff(5)
    WRITE (*, FMT="(Owner's gid:', T30, I19)") buff(6)
    WRITE (*, FMT="(Device where located:', T30, I19)") buff(7)
    WRITE (*, FMT="(File size:', T30, I19)") buff(8)
    WRITE (*, FMT="(Last access time:', T30, A19)") CTIME(buff(9))
    WRITE (*, FMT="(Last modification time', T30, A19)") CTIME(buff(10))
    WRITE (*, FMT="(Last status change time:', T30, A19)") CTIME(buff(11))
    WRITE (*, FMT="(Preferred block size:', T30, I19)") buff(12)
    WRITE (*, FMT="(No. of blocks allocated:', T30, I19)") buff(13)
  END IF
END PROGRAM
```

See also: To stat an open file: [Section 9.117 \[FSTAT\]](#), page 188, to stat a link: [Section 9.183 \[LSTAT\]](#), page 225

9.261 STORAGE_SIZE — Storage size in bits

Description:

Returns the storage size of argument *A* in bits.

Standard: Fortran 2008 and later

Class: Inquiry function

Syntax: **RESULT = STORAGE_SIZE(A [, KIND])**

Arguments:

<i>A</i>	Shall be a scalar or array of any type.
<i>KIND</i>	(Optional) shall be a scalar integer constant expression.

Return Value:

The result is a scalar integer with the kind type parameter specified by *KIND* (or default integer type if *KIND* is missing). The result value is the size expressed in bits for an element of an array that has the dynamic type and type parameters of *A*.

See also: [Section 9.57 \[C_SIZEOF\]](#), page 147, [Section 9.254 \[SIZEOF\]](#), page 267

9.262 SUM — Sum of array elements

Description:

Adds the elements of *ARRAY* along dimension *DIM* if the corresponding element in *MASK* is `TRUE`.

Standard: Fortran 95 and later

Class: Transformational function

Syntax:

```
RESULT = SUM(ARRAY[, MASK])
RESULT = SUM(ARRAY, DIM[, MASK])
```

Arguments:

<i>ARRAY</i>	Shall be an array of type <code>INTEGER</code> , <code>REAL</code> or <code>COMPLEX</code> .
<i>DIM</i>	(Optional) shall be a scalar of type <code>INTEGER</code> with a value in the range from 1 to <i>n</i> , where <i>n</i> equals the rank of <i>ARRAY</i> .
<i>MASK</i>	(Optional) shall be of type <code>LOGICAL</code> and either be a scalar or an array of the same shape as <i>ARRAY</i> .

Return value:

The result is of the same type as *ARRAY*.

If *DIM* is absent, a scalar with the sum of all elements in *ARRAY* is returned. Otherwise, an array of rank *n*-1, where *n* equals the rank of *ARRAY*, and a shape similar to that of *ARRAY* with dimension *DIM* dropped is returned.

Example:

```
PROGRAM test_sum
  INTEGER :: x(5) = (/ 1, 2, 3, 4, 5 /)
  print *, SUM(x)                ! all elements, sum = 15
  print *, SUM(x, MASK=MOD(x, 2)==1) ! odd elements, sum = 9
END PROGRAM
```

See also: [Section 9.220 \[PRODUCT\]](#), page 247

9.263 SYMLNK — Create a symbolic link

Description:

Makes a symbolic link from file *PATH1* to *PATH2*. A null character (`CHAR(0)`) can be used to mark the end of the names in *PATH1* and *PATH2*; otherwise, trailing blanks in the file names are ignored. If the *STATUS* argument is supplied, it contains 0 on success or a nonzero error code upon return; see `symlink(2)`. If the system does not supply `symlink(2)`, `ENOSYS` is returned.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL SYMLNK(PATH1, PATH2 [, STATUS])
STATUS = SYMLNK(PATH1, PATH2)
```

Arguments:

PATH1 Shall be of default CHARACTER type.
PATH2 Shall be of default CHARACTER type.
STATUS (Optional) Shall be of default INTEGER type.

See also: [Section 9.172 \[LINK\], page 220](#), [Section 9.281 \[UNLINK\], page 282](#)

9.264 SYSTEM — Execute a shell command*Description:*

Passes the command *COMMAND* to a shell (see `system(3)`). If argument *STATUS* is present, it contains the value returned by `system(3)`, which is presumably 0 if the shell command succeeded. Note that which shell is used to invoke the command is system-dependent and environment-dependent.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Note that the `system` function need not be thread-safe. It is the responsibility of the user to ensure that `system` is not called concurrently.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL SYSTEM(COMMAND [, STATUS])
STATUS = SYSTEM(COMMAND)
```

Arguments:

COMMAND Shall be of default CHARACTER type.
STATUS (Optional) Shall be of default INTEGER type.

See also: [Section 9.100 \[EXECUTE_COMMAND_LINE\], page 176](#), which is part of the Fortran 2008 standard and should be considered in new code for future portability.

9.265 SYSTEM_CLOCK — Time function*Description:*

Determines the *COUNT* of a processor clock since an unspecified time in the past modulo *COUNT_MAX*, *COUNT_RATE* determines the number of clock ticks per second. If the platform supports a monotonic clock, that clock is used and can, depending on the platform clock implementation, provide up to nanosecond resolution. If a monotonic clock is not available, the implementation falls back to a realtime clock.

COUNT_RATE is system dependent and can vary depending on the kind of the arguments. For *kind=4* arguments (and smaller integer kinds), *COUNT* represents milliseconds, while for *kind=8* arguments (and larger integer kinds),

COUNT typically represents micro- or nanoseconds depending on resolution of the underlying platform clock. *COUNT_MAX* usually equals *HUGE(COUNT_MAX)*. Note that the millisecond resolution of the *kind=4* version implies that the *COUNT* will wrap around in roughly 25 days. In order to avoid issues with the wrap around and for more precise timing, please use the *kind=8* version.

If there is no clock, or querying the clock fails, *COUNT* is set to *-HUGE(COUNT)*, and *COUNT_RATE* and *COUNT_MAX* are set to zero.

When running on a platform using the GNU C library (glibc) version 2.16 or older, or a derivative thereof, the high resolution monotonic clock is available only when linking with the *rt* library. This can be done explicitly by adding the *-lrt* flag when linking the application, but is also done implicitly when using OpenMP.

On the Windows platform, the version with *kind=4* arguments uses the *GetTickCount* function, whereas the *kind=8* version uses *QueryPerformanceCounter* and *QueryPerformanceCounterFrequency*. For more information, and potential caveats, please see the platform documentation.

Standard: Fortran 95 and later

Class: Subroutine

Syntax: `CALL SYSTEM_CLOCK([COUNT, COUNT_RATE, COUNT_MAX])`

Arguments:

COUNT (Optional) shall be a scalar of type *INTEGER* with *INTENT(OUT)*.

COUNT_RATE (Optional) shall be a scalar of type *INTEGER* or *REAL*, with *INTENT(OUT)*.

COUNT_MAX (Optional) shall be a scalar of type *INTEGER* with *INTENT(OUT)*.

Example:

```
PROGRAM test_system_clock
  INTEGER :: count, count_rate, count_max
  CALL SYSTEM_CLOCK(count, count_rate, count_max)
  WRITE(*,*) count, count_rate, count_max
END PROGRAM
```

See also: [Section 9.82 \[DATE_AND_TIME\]](#), page 164, [Section 9.79 \[CPU_TIME\]](#), page 162

9.266 TAN — Tangent function

Description:

TAN(X) computes the tangent of *X*.

Standard: Fortran 77 and later, for a complex argument Fortran 2008 or later

Class: Elemental function

Syntax: `RESULT = TAN(X)`

Arguments:

X The type shall be `REAL` or `COMPLEX`.

Return value:

The return value has same type and kind as X , and its value is in radians.

Example:

```
program test_tan
  real(8) :: x = 0.165_8
  x = tan(x)
end program test_tan
```

Specific names:

Name	Argument	Return type	Standard
<code>TAN(X)</code>	<code>REAL(4) X</code>	<code>REAL(4)</code>	Fortran 95 and later
<code>DTAN(X)</code>	<code>REAL(8) X</code>	<code>REAL(8)</code>	Fortran 95 and later

See also: Inverse function: [Section 9.23 \[ATAN\], page 126](#) Degrees function: [Section 9.267 \[TAND\], page 275](#)

9.267 TAND — Tangent function, degrees

Description:

`TAND(X)` computes the tangent of X in degrees.

This function is for compatibility only and should be avoided in favor of standard constructs wherever possible.

Standard: GNU Extension, enabled with ‘`-fdec-math`’.

Class: Elemental function

Syntax: `RESULT = TAND(X)`

Arguments:

X The type shall be `REAL` or `COMPLEX`.

Return value:

The return value has same type and kind as X , and its value is in degrees.

Example:

```
program test_tand
  real(8) :: x = 0.165_8
  x = tand(x)
end program test_tand
```

Specific names:

Name	Argument	Return type	Standard
<code>TAND(X)</code>	<code>REAL(4) X</code>	<code>REAL(4)</code>	GNU Extension
<code>DTAND(X)</code>	<code>REAL(8) X</code>	<code>REAL(8)</code>	GNU Extension

See also: Inverse function: [Section 9.24 \[ATAND\], page 126](#) Radians function: [Section 9.266 \[TAN\], page 274](#)

9.268 TANH — Hyperbolic tangent function

Description:

TANH(X) computes the hyperbolic tangent of X.

Standard: Fortran 77 and later, for a complex argument Fortran 2008 or later

Class: Elemental function

Syntax: X = TANH(X)

Arguments:

X The type shall be REAL or COMPLEX.

Return value:

The return value has same type and kind as X. If X is complex, the imaginary part of the result is in radians. If X is REAL, the return value lies in the range $-1 \leq \tanh(x) \leq 1$.

Example:

```
program test_tanh
  real(8) :: x = 2.1_8
  x = tanh(x)
end program test_tanh
```

Specific names:

Name	Argument	Return type	Standard
TANH(X)	REAL(4) X	REAL(4)	Fortran 95 and later
DTANH(X)	REAL(8) X	REAL(8)	Fortran 95 and later

See also: [Section 9.27 \[ATANH\], page 129](#)

9.269 THIS_IMAGE — Function that returns the cosubscript index of this image

Description:

Returns the cosubscript for this image.

Standard: Fortran 2008 and later. With *DISTANCE* argument, Technical Specification (TS) 18508 or later

Class: Transformational function

Syntax:

```
RESULT = THIS_IMAGE()
RESULT = THIS_IMAGE(DISTANCE)
RESULT = THIS_IMAGE(COARRAY [, DIM])
```

Arguments:

DISTANCE (optional, intent(in)) Nonnegative scalar integer (not permitted together with *COARRAY*).

COARRAY Coarray of any type (optional; if *DIM* present, required).

DIM default integer scalar (optional). If present, *DIM* shall be between one and the corank of *COARRAY*.

Return value:

Default integer. If *COARRAY* is not present, it is scalar; if *DISTANCE* is not present or has value 0, its value is the image index on the invoking image for the current team, for values smaller or equal distance to the initial team, it returns the image index on the ancestor team which has a distance of *DISTANCE* from the invoking team. If *DISTANCE* is larger than the distance to the initial team, the image index of the initial team is returned. Otherwise when the *COARRAY* is present, if *DIM* is not present, a rank-1 array with corank elements is returned, containing the cosubscripts for *COARRAY* specifying the invoking image. If *DIM* is present, a scalar is returned, with the value of the *DIM* element of *THIS_IMAGE(COARRAY)*.

Example:

```

INTEGER :: value[*]
INTEGER :: i
value = THIS_IMAGE()
SYNC ALL
IF (THIS_IMAGE() == 1) THEN
  DO i = 1, NUM_IMAGES()
    WRITE(*,'(2(a,i0))') 'value[' , i, ']' is ', value[i]
  END DO
END IF

! Check whether the current image is the initial image
IF (THIS_IMAGE(HUGE(1)) /= THIS_IMAGE())
  error stop "something is rotten here"
```

See also: [Section 9.211 \[NUM_IMAGES\]](#), page 242, [Section 9.147 \[IMAGE_INDEX\]](#), page 205

9.270 TIME — Time function

Description:

Returns the current time encoded as an integer (in the manner of the function `time(3)` in the C standard library). This value is suitable for passing to [Section 9.81 \[CTIME\]](#), page 163, [Section 9.131 \[GMTIME\]](#), page 196, and [Section 9.184 \[LTIME\]](#), page 226.

This intrinsic is not fully portable, such as to systems with 32-bit `INTEGER` types but supporting times wider than 32 bits. Therefore, the values returned by this intrinsic might be, or become, negative, or numerically less than previous values, during a single run of the compiled program.

See [Section 9.271 \[TIME8\]](#), page 278, for information on a similar intrinsic that might be portable to more GNU Fortran implementations, though to fewer Fortran compilers.

Standard: GNU extension

Class: Function

Syntax: `RESULT = TIME()`

Return value:

The return value is a scalar of type `INTEGER(4)`.

See also: Section 9.82 [DATE_AND_TIME], page 164, Section 9.81 [CTIME], page 163,
 Section 9.131 [GMTIME], page 196, Section 9.184 [LTIME], page 226,
 Section 9.193 [MCLOCK], page 231, Section 9.271 [TIME8], page 278

9.271 TIME8 — Time function (64-bit)

Description:

Returns the current time encoded as an integer (in the manner of the function `time(3)` in the C standard library). This value is suitable for passing to Section 9.81 [CTIME], page 163, Section 9.131 [GMTIME], page 196, and Section 9.184 [LTIME], page 226.

Warning: this intrinsic does not increase the range of the timing values over that returned by `time(3)`. On a system with a 32-bit `time(3)`, `TIME8` will return a 32-bit value, even though it is converted to a 64-bit `INTEGER(8)` value. That means overflows of the 32-bit value can still occur. Therefore, the values returned by this intrinsic might be or become negative or numerically less than previous values during a single run of the compiled program.

Standard: GNU extension

Class: Function

Syntax: `RESULT = TIME8()`

Return value:

The return value is a scalar of type `INTEGER(8)`.

See also: Section 9.82 [DATE_AND_TIME], page 164, Section 9.81 [CTIME], page 163,
 Section 9.131 [GMTIME], page 196, Section 9.184 [LTIME], page 226,
 Section 9.194 [MCLOCK8], page 232, Section 9.270 [TIME], page 277

9.272 TINY — Smallest positive number of a real kind

Description:

`TINY(X)` returns the smallest positive (non zero) number in the model of the type of `X`.

Standard: Fortran 95 and later

Class: Inquiry function

Syntax: `RESULT = TINY(X)`

Arguments:

`X` Shall be of type `REAL`.

Return value:

The return value is of the same type and kind as `X`

Example: See `HUGE` for an example.

9.273 TRAILZ — Number of trailing zero bits of an integer

Description:

TRAILZ returns the number of trailing zero bits of an integer.

Standard: Fortran 2008 and later

Class: Elemental function

Syntax: RESULT = TRAILZ(I)

Arguments:

I Shall be of type INTEGER.

Return value:

The type of the return value is the default INTEGER. If all the bits of I are zero, the result value is BIT_SIZE(I).

Example:

```
PROGRAM test_trailz
  WRITE (*,*) TRAILZ(8)  ! prints 3
END PROGRAM
```

See also: [Section 9.48 \[BIT_SIZE\], page 142](#), [Section 9.167 \[LEADZ\], page 217](#), [Section 9.217 \[POPPAR\], page 246](#), [Section 9.216 \[POPCNT\], page 245](#)

9.274 TRANSFER — Transfer bit patterns

Description:

Interprets the bitwise representation of *SOURCE* in memory as if it is the representation of a variable or array of the same type and type parameters as *MOLD*.

This is approximately equivalent to the C concept of *casting* one type to another.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: RESULT = TRANSFER(SOURCE, MOLD[, SIZE])

Arguments:

SOURCE Shall be a scalar or an array of any type.
MOLD Shall be a scalar or an array of any type.
SIZE (Optional) shall be a scalar of type INTEGER.

Return value:

The result has the same type as *MOLD*, with the bit level representation of *SOURCE*. If *SIZE* is present, the result is a one-dimensional array of length *SIZE*. If *SIZE* is absent but *MOLD* is an array (of any size or shape), the result is a one-dimensional array of the minimum length needed to contain the entirety of the bitwise representation of *SOURCE*. If *SIZE* is absent and *MOLD* is a scalar, the result is a scalar.

If the bitwise representation of the result is longer than that of *SOURCE*, then the leading bits of the result correspond to those of *SOURCE* and any trailing bits are filled arbitrarily.

When the resulting bit representation does not correspond to a valid representation of a variable of the same type as *MOLD*, the results are undefined, and subsequent operations on the result cannot be guaranteed to produce sensible behavior. For example, it is possible to create LOGICAL variables for which *VAR* and *.NOT.VAR* both appear to be true.

Example:

```
PROGRAM test_transfer
  integer :: x = 2143289344
  print *, transfer(x, 1.0)    ! prints "NaN" on i686
END PROGRAM
```

9.275 TRANSPOSE — Transpose an array of rank two

Description:

Transpose an array of rank two. Element (i, j) of the result has the value *MATRIX*(j, i), for all i, j.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: *RESULT* = TRANSPOSE(*MATRIX*)

Arguments:

MATRIX Shall be an array of any type and have a rank of two.

Return value:

The result has the same type as *MATRIX*, and has shape (/ *m*, *n* /) if *MATRIX* has shape (/ *n*, *m* /).

9.276 TRIM — Remove trailing blank characters of a string

Description:

Removes trailing blank characters of a string.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: *RESULT* = TRIM(*STRING*)

Arguments:

STRING Shall be a scalar of type CHARACTER.

Return value:

A scalar of type CHARACTER which length is that of *STRING* less the number of trailing blanks.

Example:

```
PROGRAM test_trim
  CHARACTER(len=10), PARAMETER :: s = "GFORTRAN  "
  WRITE(*,*) LEN(s), LEN(TRIM(s)) ! "10 8", with/without trailing blanks
END PROGRAM
```

See also: [Section 9.9 \[ADJUSTL\], page 116](#), [Section 9.10 \[ADJUSTR\], page 116](#)

9.277 TTYNAM — Get the name of a terminal device.

Description:

Get the name of a terminal device. For more information, see `ttynam(3)`.
This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL TTYNAM(UNIT, NAME)
NAME = TTYNAM(UNIT)
```

Arguments:

<i>UNIT</i>	Shall be a scalar <code>INTEGER</code> .
<i>NAME</i>	Shall be of type <code>CHARACTER</code> .

Example:

```
PROGRAM test_ttynam
  INTEGER :: unit
  DO unit = 1, 10
    IF (isatty(unit=unit)) write(*,*) ttynam(unit)
  END DO
END PROGRAM
```

See also: [Section 9.158 \[ISATTY\], page 212](#)

9.278 UBOUND — Upper dimension bounds of an array

Description:

Returns the upper bounds of an array, or a single upper bound along the *DIM* dimension.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Inquiry function

Syntax: `RESULT = UBOUND(ARRAY [, DIM [, KIND]])`

Arguments:

<i>ARRAY</i>	Shall be an array, of any type.
<i>DIM</i>	(Optional) Shall be a scalar <code>INTEGER</code> .
<i>KIND</i>	(Optional) An <code>INTEGER</code> initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind. If *DIM* is absent, the result is an array of the upper bounds of *ARRAY*. If *DIM* is present, the result is a scalar corresponding to the upper bound of the array along that dimension. If *ARRAY* is an expression rather than a whole array or array structure component, or if it has a zero extent along the relevant dimension, the upper bound is taken to be the number of elements along the relevant dimension.

See also: [Section 9.165 \[LBOUND\]](#), page 216, [Section 9.166 \[LCOBUND\]](#), page 216

9.279 UCBOUND — Upper codimension bounds of an array

Description:

Returns the upper cobounds of a coarray, or a single upper cobound along the *DIM* codimension.

Standard: Fortran 2008 and later

Class: Inquiry function

Syntax: `RESULT = UCBOUND(COARRAY [, DIM [, KIND]])`

Arguments:

<i>ARRAY</i>	Shall be an coarray, of any type.
<i>DIM</i>	(Optional) Shall be a scalar <code>INTEGER</code> .
<i>KIND</i>	(Optional) An <code>INTEGER</code> initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type `INTEGER` and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind. If *DIM* is absent, the result is an array of the lower cobounds of *COARRAY*. If *DIM* is present, the result is a scalar corresponding to the lower cobound of the array along that codimension.

See also: [Section 9.166 \[LCOBUND\]](#), page 216, [Section 9.165 \[LBOUND\]](#), page 216

9.280 UMASK — Set the file creation mask

Description:

Sets the file creation mask to *MASK*. If called as a function, it returns the old value. If called as a subroutine and argument *OLD* if it is supplied, it is set to the old value. See `umask(2)`.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL UMASK(MASK [, OLD])
OLD = UMASK(MASK)
```

Arguments:

<i>MASK</i>	Shall be a scalar of type <code>INTEGER</code> .
<i>OLD</i>	(Optional) Shall be a scalar of type <code>INTEGER</code> .

9.281 UNLINK — Remove a file from the file system

Description:

Unlinks the file *PATH*. A null character (`CHAR(0)`) can be used to mark the end of the name in *PATH*; otherwise, trailing blanks in the file name are ignored. If

the *STATUS* argument is supplied, it contains 0 on success or a nonzero error code upon return; see `unlink(2)`.

This intrinsic is provided in both subroutine and function forms; however, only one form can be used in any given program unit.

Standard: GNU extension

Class: Subroutine, function

Syntax:

```
CALL UNLINK(PATH [, STATUS])
STATUS = UNLINK(PATH)
```

Arguments:

PATH Shall be of default CHARACTER type.
STATUS (Optional) Shall be of default INTEGER type.

See also: [Section 9.172 \[LINK\], page 220](#), [Section 9.263 \[SYMLNK\], page 272](#)

9.282 UNPACK — Unpack an array of rank one into an array

Description:

Store the elements of *VECTOR* in an array of higher rank.

Standard: Fortran 95 and later

Class: Transformational function

Syntax: `RESULT = UNPACK(VECTOR, MASK, FIELD)`

Arguments:

VECTOR Shall be an array of any type and rank one. It shall have at least as many elements as *MASK* has TRUE values.
MASK Shall be an array of type LOGICAL.
FIELD Shall be of the same type as *VECTOR* and have the same shape as *MASK*.

Return value:

The resulting array corresponds to *FIELD* with TRUE elements of *MASK* replaced by values from *VECTOR* in array element order.

Example:

```
PROGRAM test_unpack
  integer :: vector(2) = (/1,1/)
  logical :: mask(4) = (/ .TRUE., .FALSE., .FALSE., .TRUE. /)
  integer :: field(2,2) = 0, unity(2,2)

  ! result: unity matrix
  unity = unpack(vector, reshape(mask, (/2,2/)), field)
END PROGRAM
```

See also: [Section 9.213 \[PACK\], page 243](#), [Section 9.257 \[SPREAD\], page 268](#)

9.283 VERIFY — Scan a string for characters not a given set

Description:

Verifies that all the characters in *STRING* belong to the set of characters in *SET*.

If *BACK* is either absent or equals **FALSE**, this function returns the position of the leftmost character of *STRING* that is not in *SET*. If *BACK* equals **TRUE**, the rightmost position is returned. If all characters of *STRING* are found in *SET*, the result is zero.

Standard: Fortran 95 and later, with *KIND* argument Fortran 2003 and later

Class: Elemental function

Syntax: `RESULT = VERIFY(STRING, SET[, BACK [, KIND]])`

Arguments:

<i>STRING</i>	Shall be of type CHARACTER .
<i>SET</i>	Shall be of type CHARACTER .
<i>BACK</i>	(Optional) shall be of type LOGICAL .
<i>KIND</i>	(Optional) An INTEGER initialization expression indicating the kind parameter of the result.

Return value:

The return value is of type **INTEGER** and of kind *KIND*. If *KIND* is absent, the return value is of default integer kind.

Example:

```

PROGRAM test_verify
  WRITE(*,*) VERIFY("FORTRAN", "AO")           ! 1, found 'F'
  WRITE(*,*) VERIFY("FORTRAN", "FOO")           ! 3, found 'R'
  WRITE(*,*) VERIFY("FORTRAN", "C++")           ! 1, found 'F'
  WRITE(*,*) VERIFY("FORTRAN", "C++", .TRUE.)   ! 7, found 'N'
  WRITE(*,*) VERIFY("FORTRAN", "FORTRAN")       ! 0' found none
END PROGRAM

```

See also: [Section 9.237 \[SCAN\], page 256](#), [Section 9.148 \[INDEX intrinsic\], page 206](#)

9.284 XOR — Bitwise logical exclusive OR

Description:

Bitwise logical exclusive or.

This intrinsic routine is provided for backwards compatibility with GNU Fortran 77. For integer arguments, programmers should consider the use of the [Section 9.145 \[IEOR\], page 204](#) intrinsic and for logical arguments the **.NEQV.** operator, which are both defined by the Fortran standard.

Standard: GNU extension

Class: Function

Syntax: `RESULT = XOR(I, J)`

Arguments:

- I* The type shall be either a scalar `INTEGER` type or a scalar `LOGICAL` type or a boz-literal-constant.
- J* The type shall be the same as the type of *I* or a boz-literal-constant. *I* and *J* shall not both be boz-literal-constants. If either *I* and *J* is a boz-literal-constant, then the other argument must be a scalar `INTEGER`.

Return value:

The return type is either a scalar `INTEGER` or a scalar `LOGICAL`. If the kind type parameters differ, then the smaller kind type is implicitly converted to larger kind, and the return has the larger kind. A boz-literal-constant is converted to an `INTEGER` with the kind type parameter of the other argument as-if a call to [Section 9.149 \[INT\]](#), page 207 occurred.

Example:

```
PROGRAM test_xor
  LOGICAL :: T = .TRUE., F = .FALSE.
  INTEGER :: a, b
  DATA a / Z'F' /, b / Z'3' /

  WRITE (*,*) XOR(T, T), XOR(T, F), XOR(F, T), XOR(F, F)
  WRITE (*,*) XOR(a, b)
END PROGRAM
```

See also: Fortran 95 elemental function: [Section 9.145 \[IEOR\]](#), page 204

10 Intrinsic Modules

10.1 ISO_FORTRAN_ENV

Standard: Fortran 2003 and later, except when otherwise noted

The `ISO_FORTRAN_ENV` module provides the following scalar default-integer named constants:

`ATOMIC_INT_KIND:`

Default-kind integer constant to be used as kind parameter when defining integer variables used in atomic operations. (Fortran 2008 or later.)

`ATOMIC_LOGICAL_KIND:`

Default-kind integer constant to be used as kind parameter when defining logical variables used in atomic operations. (Fortran 2008 or later.)

`CHARACTER_KINDS:`

Default-kind integer constant array of rank one containing the supported kind parameters of the `CHARACTER` type. (Fortran 2008 or later.)

`CHARACTER_STORAGE_SIZE:`

Size in bits of the character storage unit.

`ERROR_UNIT:`

Identifies the preconnected unit used for error reporting.

`FILE_STORAGE_SIZE:`

Size in bits of the file-storage unit.

`INPUT_UNIT:`

Identifies the preconnected unit identified by the asterisk (*) in `READ` statement.

`INT8, INT16, INT32, INT64:`

Kind type parameters to specify an `INTEGER` type with a storage size of 16, 32, and 64 bits. It is negative if a target platform does not support the particular kind. (Fortran 2008 or later.)

`INTEGER_KINDS:`

Default-kind integer constant array of rank one containing the supported kind parameters of the `INTEGER` type. (Fortran 2008 or later.)

`IOSTAT_END:`

The value assigned to the variable passed to the `IOSTAT=` specifier of an input/output statement if an end-of-file condition occurred.

`IOSTAT_EOR:`

The value assigned to the variable passed to the `IOSTAT=` specifier of an input/output statement if an end-of-record condition occurred.

`IOSTAT_INQUIRE_INTERNAL_UNIT:`

Scalar default-integer constant, used by `INQUIRE` for the `IOSTAT=` specifier to denote an that a unit number identifies an internal unit. (Fortran 2008 or later.)

NUMERIC_STORAGE_SIZE:

The size in bits of the numeric storage unit.

LOGICAL_KINDS:

Default-kind integer constant array of rank one containing the supported kind parameters of the **LOGICAL** type. (Fortran 2008 or later.)

OUTPUT_UNIT:

Identifies the preconnected unit identified by the asterisk (*) in **WRITE** statement.

REAL32, REAL64, REAL128:

Kind type parameters to specify a **REAL** type with a storage size of 32, 64, and 128 bits. It is negative if a target platform does not support the particular kind. (Fortran 2008 or later.)

REAL_KINDS:

Default-kind integer constant array of rank one containing the supported kind parameters of the **REAL** type. (Fortran 2008 or later.)

STAT_LOCKED:

Scalar default-integer constant used as **STAT=** return value by **LOCK** to denote that the lock variable is locked by the executing image. (Fortran 2008 or later.)

STAT_LOCKED_OTHER_IMAGE:

Scalar default-integer constant used as **STAT=** return value by **UNLOCK** to denote that the lock variable is locked by another image. (Fortran 2008 or later.)

STAT_STOPPED_IMAGE:

Positive, scalar default-integer constant used as **STAT=** return value if the argument in the statement requires synchronisation with an image, which has initiated the termination of the execution. (Fortran 2008 or later.)

STAT_FAILED_IMAGE:

Positive, scalar default-integer constant used as **STAT=** return value if the argument in the statement requires communication with an image, which has is in the failed state. (TS 18508 or later.)

STAT_UNLOCKED:

Scalar default-integer constant used as **STAT=** return value by **UNLOCK** to denote that the lock variable is unlocked. (Fortran 2008 or later.)

The module provides the following derived type:

LOCK_TYPE:

Derived type with private components to be use with the **LOCK** and **UNLOCK** statement. A variable of its type has to be always declared as coarray and may not appear in a variable-definition context. (Fortran 2008 or later.)

The module also provides the following intrinsic procedures: [Section 9.69 \[COMPILER_OPTIONS\]](#), page 156 and [Section 9.70 \[COMPILER_VERSION\]](#), page 156.

10.2 ISO_C_BINDING

Standard: Fortran 2003 and later, GNU extensions

The following intrinsic procedures are provided by the module; their definition can be found in the section Intrinsic Procedures of this manual.

C_ASSOCIATED
C_F_POINTER
C_F_PROCPTR
C_FUNLOC
C_LOC
C_SIZEOF

The ISO_C_BINDING module provides the following named constants of type default integer, which can be used as KIND type parameters.

In addition to the integer named constants required by the Fortran 2003 standard and C_PTRDIFF_T of TS 29113, GNU Fortran provides as an extension named constants for the 128-bit integer types supported by the C compiler: C_INT128_T, C_INT_LEAST128_T, C_INT_FAST128_T. Furthermore, if __float128 is supported in C, the named constants C_FLOAT128, C_FLOAT128_COMPLEX are defined.

Fortran Type	Named constant	C type	Extension
INTEGER	C_INT	int	
INTEGER	C_SHORT	short int	
INTEGER	C_LONG	long int	
INTEGER	C_LONG_LONG	long long int	
INTEGER	C_SIGNED_CHAR	signed char/unsigned char	
INTEGER	C_SIZE_T	size_t	
INTEGER	C_INT8_T	int8_t	
INTEGER	C_INT16_T	int16_t	
INTEGER	C_INT32_T	int32_t	
INTEGER	C_INT64_T	int64_t	
INTEGER	C_INT128_T	int128_t	Ext.
INTEGER	C_INT_LEAST8_T	int_least8_t	
INTEGER	C_INT_LEAST16_T	int_least16_t	
INTEGER	C_INT_LEAST32_T	int_least32_t	
INTEGER	C_INT_LEAST64_T	int_least64_t	
INTEGER	C_INT_LEAST128_T	int_least128_t	Ext.
INTEGER	C_INT_FAST8_T	int_fast8_t	
INTEGER	C_INT_FAST16_T	int_fast16_t	
INTEGER	C_INT_FAST32_T	int_fast32_t	
INTEGER	C_INT_FAST64_T	int_fast64_t	
INTEGER	C_INT_FAST128_T	int_fast128_t	Ext.
INTEGER	C_INTMAX_T	intmax_t	
INTEGER	C_INTPTR_T	intptr_t	
INTEGER	C_PTRDIFF_T	ptrdiff_t	TS 29113
REAL	C_FLOAT	float	
REAL	C_DOUBLE	double	

REAL	C_LONG_DOUBLE	long double	
REAL	C_FLOAT128	__float128	Ext.
COMPLEX	C_FLOAT_COMPLEX	float _Complex	
COMPLEX	C_DOUBLE_COMPLEX	double _Complex	
COMPLEX	C_LONG_DOUBLE_COMPLEX	long double _Complex	
REAL	C_FLOAT128_COMPLEX	__float128 _Complex	Ext.
LOGICAL	C_BOOL	_Bool	
CHARACTER	C_CHAR	char	

Additionally, the following parameters of type `CHARACTER(KIND=C_CHAR)` are defined.

Name	C definition	Value
C_NULL_CHAR	null character	'\0'
C_ALERT	alert	'\a'
C_BACKSPACE	backspace	'\b'
C_FORM_FEED	form feed	'\f'
C_NEW_LINE	new line	'\n'
C_CARRIAGE_	carriage return	'\r'
RETURN		
C_HORIZONTAL_	horizontal tab	'\t'
TAB		
C_VERTICAL_TAB	vertical tab	'\v'

Moreover, the following two named constants are defined:

Name	Type
C_NULL_PTR	C_PTR
C_NULL_FUNPTR	C_FUNPTR

Both are equivalent to the value `NULL` in C.

10.3 IEEE modules: IEEE_EXCEPTIONS, IEEE_ARITHMETIC, and IEEE_FEATURES

Standard: Fortran 2003 and later

The `IEEE_EXCEPTIONS`, `IEEE_ARITHMETIC`, and `IEEE_FEATURES` intrinsic modules provide support for exceptions and IEEE arithmetic, as defined in Fortran 2003 and later standards, and the IEC 60559:1989 standard (*Binary floating-point arithmetic for microprocessor systems*). These modules are only provided on the following supported platforms:

- i386 and x86_64 processors
- platforms which use the GNU C Library (glibc)
- platforms with support for SysV/386 routines for floating point interface (including Solaris and BSDs)
- platforms with the AIX OS

For full compliance with the Fortran standards, code using the `IEEE_EXCEPTIONS` or `IEEE_ARITHMETIC` modules should be compiled with the following options: `-fno-unsafe-math-optimizations -frounding-math -fsignaling-nans`.

10.4 OpenMP Modules OMP_LIB and OMP_LIB_KINDS

Standard: OpenMP Application Program Interface v4.5

The OpenMP Fortran runtime library routines are provided both in a form of two Fortran 90 modules, named `OMP_LIB` and `OMP_LIB_KINDS`, and in a form of a Fortran `include` file named `'omp_lib.h'`. The procedures provided by `OMP_LIB` can be found in the [Section “Introduction”](#) in *GNU Offloading and Multi Processing Runtime Library* manual, the named constants defined in the modules are listed below.

For details refer to the actual [OpenMP Application Program Interface v4.5](#).

`OMP_LIB_KINDS` provides the following scalar default-integer named constants:

```
omp_lock_kind
omp_nest_lock_kind
omp_proc_bind_kind
omp_sched_kind
```

`OMP_LIB` provides the scalar default-integer named constant `openmp_version` with a value of the form `yyyymm`, where `yyyy` is the year and `mm` the month of the OpenMP version; for OpenMP v4.5 the value is 201511.

The following scalar integer named constants of the kind `omp_sched_kind`:

```
omp_sched_static
omp_sched_dynamic
omp_sched_guided
omp_sched_auto
```

And the following scalar integer named constants of the kind `omp_proc_bind_kind`:

```
omp_proc_bind_false
omp_proc_bind_true
omp_proc_bind_master
omp_proc_bind_close
omp_proc_bind_spread
```

10.5 OpenACC Module OPENACC

Standard: OpenACC Application Programming Interface v2.0

The OpenACC Fortran runtime library routines are provided both in a form of a Fortran 90 module, named `OPENACC`, and in form of a Fortran `include` file named `'openacc_lib.h'`. The procedures provided by `OPENACC` can be found in the [Section “Introduction”](#) in *GNU Offloading and Multi Processing Runtime Library* manual, the named constants defined in the modules are listed below.

For details refer to the actual [OpenACC Application Programming Interface v2.0](#).

`OPENACC` provides the scalar default-integer named constant `openacc_version` with a value of the form `yyyymm`, where `yyyy` is the year and `mm` the month of the OpenACC version; for OpenACC v2.0 the value is 201306.

Contributing

Free software is only possible if people contribute to efforts to create it. We're always in need of more people helping out with ideas and comments, writing documentation and contributing code.

If you want to contribute to GNU Fortran, have a look at the long lists of projects you can take on. Some of these projects are small, some of them are large; some are completely orthogonal to the rest of what is happening on GNU Fortran, but others are “mainstream” projects in need of enthusiastic hackers. All of these projects are important! We will eventually get around to the things here, but they are also things doable by someone who is willing and able.

Contributors to GNU Fortran

Most of the parser was hand-crafted by *Andy Vaught*, who is also the initiator of the whole project. Thanks Andy! Most of the interface with GCC was written by *Paul Brook*.

The following individuals have contributed code and/or ideas and significant help to the GNU Fortran project (in alphabetical order):

- Janne Blomqvist
- Steven Bosscher
- Paul Brook
- Tobias Burnus
- François-Xavier Coudert
- Bud Davis
- Jerry DeLisle
- Erik Edelman
- Bernhard Fischer
- Daniel Franke
- Richard Guenther
- Richard Henderson
- Katherine Holcomb
- Jakub Jelinek
- Niels Kristian Bech Jensen
- Steven Johnson
- Steven G. Kargl
- Thomas Koenig
- Asher Langton
- H. J. Lu
- Toon Moene
- Brooks Moses
- Andrew Pinski
- Tim Prince

- Christopher D. Rickett
- Richard Sandiford
- Tobias Schlüter
- Roger Sayle
- Paul Thomas
- Andy Vaught
- Feng Wang
- Janus Weil
- Daniel Kraft

The following people have contributed bug reports, smaller or larger patches, and much needed feedback and encouragement for the GNU Fortran project:

- Bill Clodius
- Dominique d’Humières
- Kate Hedstrom
- Erik Schnetter
- Joost VandeVondele

Many other individuals have helped debug, test and improve the GNU Fortran compiler over the past few years, and we welcome you to do the same! If you already have done so, and you would like to see your name listed in the list above, please contact us.

Projects

Help build the test suite

Solicit more code for donation to the test suite: the more extensive the testsuite, the smaller the risk of breaking things in the future! We can keep code private on request.

Bug hunting/squishing

Find bugs and write more test cases! Test cases are especially very welcome, because it allows us to concentrate on fixing bugs instead of isolating them. Going through the bugzilla database at <https://gcc.gnu.org/bugzilla/> to reduce testcases posted there and add more information (for example, for which version does the testcase work, for which versions does it fail?) is also very helpful.

Proposed Extensions

Here’s a list of proposed extensions for the GNU Fortran compiler, in no particular order. Most of these are necessary to be fully compatible with existing Fortran compilers, but they are not part of the official J3 Fortran 95 standard.

Compiler extensions:

- User-specified alignment rules for structures.
- Automatically extend single precision constants to double.

- Compile code that conserves memory by dynamically allocating common and module storage either on stack or heap.
- Compile flag to generate code for array conformance checking (suggest -CC).
- User control of symbol names (underscores, etc).
- Compile setting for maximum size of stack frame size before spilling parts to static or heap.
- Flag to force local variables into static space.
- Flag to force local variables onto stack.

Environment Options

- Pluggable library modules for random numbers, linear algebra. LA should use BLAS calling conventions.
- Environment variables controlling actions on arithmetic exceptions like overflow, underflow, precision loss—Generate NaN, abort, default. action.
- Set precision for fp units that support it (i387).
- Variable for setting fp rounding mode.
- Variable to fill uninitialized variables with a user-defined bit pattern.
- Environment variable controlling filename that is opened for that unit number.
- Environment variable to clear/trash memory being freed.
- Environment variable to control tracing of allocations and frees.
- Environment variable to display allocated memory at normal program end.
- Environment variable for filename for * IO-unit.
- Environment variable for temporary file directory.
- Environment variable forcing standard output to be line buffered (Unix).

GNU General Public License

Version 3, 29 June 2007

Copyright © 2007 Free Software Foundation, Inc. <http://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software; it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, although the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS

0. Definitions.

“This License” refers to version 3 of the GNU General Public License.

“Copyright” also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

“The Program” refers to any copyrightable work licensed under this License. Each licensee is addressed as “you”. “Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users’ Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- a. The work must carry prominent notices stating that you modified it, and giving a relevant date.
- b. The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to "keep intact all notices".
- c. You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- d. If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an "aggregate" if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation's users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- a. Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.
- b. Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.
- c. Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- d. Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.
- e. Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source.

The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- a. Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- b. Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- c. Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or

- d. Limiting the use for publicity purposes of names of licensors or authors of the material; or
- e. Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or
- f. Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance.

However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party’s predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A “contributor” is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor’s “contributor version”.

A contributor’s “essential patent claims” are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, “control” includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor’s essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a “patent license” is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To “grant” such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so

available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. “Knowingly relying” means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient’s use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is “discriminatory” if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others’ Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively state the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

one line to give the program's name and a brief idea of what it does.
Copyright (C) year name of author

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <http://www.gnu.org/licenses/>.

Also add information on how to contact you by electronic and paper mail.

If the program does terminal interaction, make it output a short notice like this when it starts in an interactive mode:

program Copyright (C) year name of author
This program comes with ABSOLUTELY NO WARRANTY; for details type 'show w'.
This is free software, and you are welcome to redistribute it
under certain conditions; type 'show c' for details.

The hypothetical commands ‘show w’ and ‘show c’ should show the appropriate parts of the General Public License. Of course, your program’s commands might be different; for a GUI interface, you would use an “about box”.

You should also get your employer (if you work as a programmer) or school, if any, to sign a “copyright disclaimer” for the program, if necessary. For more information on this, and how to apply and follow the GNU GPL, see <http://www.gnu.org/licenses/>.

The GNU General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Lesser General Public License instead of this License. But first, please read <http://www.gnu.org/philosophy/why-not-lgpl.html>.

GNU Free Documentation License

Version 1.3, 3 November 2008

Copyright © 2000, 2001, 2002, 2007, 2008 Free Software Foundation, Inc.

<http://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document *free* in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or non-commercially. Secondly, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of “copyleft”, which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The “Document”, below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as “you”. You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A “Modified Version” of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A “Secondary Section” is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document’s overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The “Invariant Sections” are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released

under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The “Cover Texts” are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A “Transparent” copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not “Transparent” is called “Opaque”.

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The “Title Page” means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, “Title Page” means the text near the most prominent appearance of the work’s title, preceding the beginning of the body of the text.

The “publisher” means any person or entity that distributes copies of the Document to the public.

A section “Entitled XYZ” means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as “Acknowledgements”, “Dedications”, “Endorsements”, or “History”.) To “Preserve the Title” of such a section when you modify the Document means that it remains a section “Entitled XYZ” according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any,

be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.

- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.
- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
- L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their

titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements."

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an “aggregate” if the copyright resulting from the compilation is not used to limit the legal rights of the compilation’s users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document’s Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled “Acknowledgements”, “Dedications”, or “History”, the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, or distribute it is void, and will automatically terminate your rights under this License.

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, receipt of a copy of some or all of the same material does not give you any rights to use it.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License “or any later version” applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation. If the Document specifies that a proxy can decide which future versions of this License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Document.

11. RELICENSING

“Massive Multiauthor Collaboration Site” (or “MMC Site”) means any World Wide Web server that publishes copyrightable works and also provides prominent facilities for anybody to edit those works. A public wiki that anybody can edit is an example of such a server. A “Massive Multiauthor Collaboration” (or “MMC”) contained in the site means any set of copyrightable works thus published on the MMC site.

“CC-BY-SA” means the Creative Commons Attribution-Share Alike 3.0 license published by Creative Commons Corporation, a not-for-profit corporation with a principal place of business in San Francisco, California, as well as future copyleft versions of that license published by that same organization.

“Incorporate” means to publish or republish a Document, in whole or in part, as part of another Document.

An MMC is “eligible for relicensing” if it is licensed under this License, and if all works that were first published under this License somewhere other than this MMC, and subsequently incorporated in whole or in part into the MMC, (1) had no cover texts or invariant sections, and (2) were thus incorporated prior to November 1, 2008.

The operator of an MMC Site may republish an MMC contained in the site under CC-BY-SA on the same site at any time before August 1, 2009, provided the MMC is eligible for relicensing.

ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

```
Copyright (C)  year  your name.
Permission is granted to copy, distribute and/or modify this document
under the terms of the GNU Free Documentation License, Version 1.3
or any later version published by the Free Software Foundation;
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover
Texts. A copy of the license is included in the section entitled ‘‘GNU
Free Documentation License’’.
```

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the “with...Texts.” line with this:

```
with the Invariant Sections being list their titles, with
the Front-Cover Texts being list, and with the Back-Cover Texts
being list.
```

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Funding Free Software

If you want to have more free software a few years from now, it makes sense for you to help encourage people to contribute funds for its development. The most effective approach known is to encourage commercial redistributors to donate.

Users of free software systems can boost the pace of development by encouraging for-a-fee distributors to donate part of their selling price to free software developers—the Free Software Foundation, and others.

The way to convince distributors to do this is to demand it and expect it from them. So when you compare distributors, judge them partly by how much they give to free software development. Show distributors they must compete to be the one who gives the most.

To make this approach work, you must insist on numbers that you can compare, such as, “We will donate ten dollars to the Frobnitz project for each disk sold.” Don’t be satisfied with a vague promise, such as “A portion of the profits are donated,” since it doesn’t give a basis for comparison.

Even a precise fraction “of the profits from this disk” is not very meaningful, since creative accounting and unrelated business decisions can greatly alter what fraction of the sales price counts as profit. If the price you pay is \$50, ten percent of the profit is probably less than a dollar; it might be a few cents, or nothing at all.

Some redistributors do development work themselves. This is useful too; but to keep everyone honest, you need to inquire how much they do, and what kind. Some kinds of development make much more long-term difference than others. For example, maintaining a separate version of a program contributes very little; maintaining the standard version of a program for the whole community contributes much. Easy new ports contribute little, since someone else would surely do them; difficult ports such as adding a new CPU to the GNU Compiler Collection contribute more; major new features or packages contribute the most.

By establishing the idea that supporting further development is “the proper thing to do” when distributing free software for a fee, we can assure a steady flow of resources into making more free software.

Copyright © 1994 Free Software Foundation, Inc.

Verbatim copying and redistribution of this section is permitted without royalty; alteration is not permitted.

Option Index

gfortran's command line options are indexed here without any initial '-' or '--'. Where an option has both positive and negative forms (such as -foption and -fno-option), relevant entries in the manual are indexed under the most appropriate form; it may sometimes be useful to look up both forms.

A

A-predicate=answer 14
 Apredicate=answer..... 14

B

backslash..... 9

C

c-prototypes 30
 c-prototypes-external..... 30
 CC..... 15
 cpp..... 13
 C..... 14

D

dD..... 13
 dI..... 13
 dM..... 13
 Dname..... 15
 Dname=definition..... 15
 dN..... 13
 dU..... 13

F

faggressive-function-elimination 29
 falign-commons..... 29
 fall-intrinsics 8
 fblas-matmul-limit 28
 fbounds-check 26
 fcheck 25
 fcheck-array-temporaries 26
 fcoarray..... 25
 fconvert=conversion 22
 fcray-pointer 10
 fd-lines-as-code..... 9
 fd-lines-as-comments 9
 fdec..... 9
 fdec-include 9
 fdec-intrinsic-ints 9
 fdec-math..... 9
 fdec-static..... 9
 fdec-structure 9
 fdefault-double-8..... 12
 fdefault-integer-8 11
 fdefault-real-10..... 11

fdefault-real-16..... 11
 fdefault-real-8..... 11
 fdollar-ok..... 9
 fdump-fortran-global..... 20
 fdump-fortran-optimized..... 20
 fdump-fortran-original..... 20
 fdump-parse-tree..... 20
 fexternal-blas..... 28
 ff2c 23
 ffixed-form..... 8
 ffixed-line-length-n..... 10
 ffpe-summary=list..... 21
 ffpe-trap=list..... 21
 ffree-form..... 8
 ffree-line-length-n..... 10
 fimPLICIT-none..... 10
 finit-character..... 28
 finit-derived..... 28
 finit-integer 28
 finit-local-zero..... 28
 finit-logical 28
 finit-real..... 28
 finline-matmul-limit..... 28
 finteger-4-integer-8..... 12
 fintrinsic-modules-path dir 22
 fmax-array-constructor..... 27
 fmax-errors=n 16
 fmax-identifier-length=n 10
 fmax-stack-var-size 27
 fmax-subrecord-length=length..... 23
 fmodule-private..... 10
 fno-automatic 23
 fno-backtrace 21
 fno-protect-parens 29
 fno-underscoring..... 24
 fopenacc..... 10
 fopenmp..... 11
 fpack-derived 27
 fpad-source 10
 fpp..... 13
 frange-check..... 11
 freal-4-real-10..... 12
 freal-4-real-16..... 12
 freal-4-real-8..... 12
 freal-8-real-10..... 12
 freal-8-real-16..... 12
 freal-8-real-4..... 12
 frealloc-lhs 29
 frecord-marker=length..... 22

frecursive.....	28
frepack-arrays.....	27
frontend-loop-interchange.....	30
frontend-optimize.....	30
fsecond-underscore.....	25
fshort-enums.....	27, 40
fsign-zero.....	23
fstack-arrays.....	27
fsyntax-only.....	16
ftest-forall-temp.....	12
fworking-directory.....	13

H

H.....	15
--------	----

I

Idir.....	22
idirafter <i>dir</i>	13
imultilib <i>dir</i>	14
iprefix <i>prefix</i>	14
iquote <i>dir</i>	14
isysroot <i>dir</i>	14
isystem <i>dir</i>	14

J

Jdir.....	22
-----------	----

M

Mdir.....	22
-----------	----

N

nostdinc.....	14
---------------	----

P

pedantic.....	16
pedantic-errors.....	16
P.....	15

S

static-libgfortran.....	22
-------------------------	----

std=std option.....	12
---------------------	----

T

tail-call-workaround.....	26
---------------------------	----

U

Uname.....	15
undef.....	14

W

Waliasing.....	16
Walign-commons.....	19
Wall.....	16
Wampersand.....	17
Wargument-mismatch.....	17
Warray-temporaries.....	17
Wc-binding-type.....	17
Wcharacter-truncation.....	17
Wcompare-reals.....	19
Wconversion.....	17
Wconversion-extra.....	17
Wdo-subscript.....	20
Werror.....	20
Wextra.....	17
Wfrontend-loop-interchange.....	17
Wfunction-elimination.....	19
Wimplicit-interface.....	18
Wimplicit-procedure.....	18
Winteger-division.....	18
Wintrinsic-shadow.....	19
Wintrinsics-std.....	18
Wline-truncation.....	17
Wpedantic.....	16
Wreal-q-constant.....	18
Wrealloc-lhs.....	19
Wrealloc-lhs-all.....	19
Wsurprising.....	18
Wtabs.....	18
Wtgt-lifetime.....	19
Wundefined-do-loop.....	18
Wunderflow.....	18
Wunused-dummy-argument.....	19
Wunused-parameter.....	19
Wuse-without-only.....	19
Wzerotrip.....	20

Keyword Index

\$

\$ 9

%

%LOC 59

%REF 59

%VAL 59

&

& 17

[

[...] 40

—

_gfortran_set_args 81

_gfortran_set_convert 82

_gfortran_set_fpe 83

_gfortran_set_max_subrecord_length 84

_gfortran_set_options 81

_gfortran_set_record_marker 83

A

ABORT 111

absolute value 112

ABS 112

ACCESS 113

ACCESS='STREAM' I/O 40

ACHAR 113

ACOS 114

ACOSH 115

ACOSH 115

adjust string 116

ADJUSTL 116

ADJUSTR 116

AIMAG 117

AINT 117

ALARM 118

ALGAMA 223

aliasing 16

alignment of COMMON blocks 19, 29

all warnings 16

ALLOCATABLE components of derived types 40

ALLOCATABLE dummy arguments 40

ALLOCATABLE function results 40

ALLOCATED 120

allocation, moving 237

allocation, status 120

ALL 119

ALOG 222

ALOG10 223

AMAXO 229

AMAX1 229

AMINO 233

AMIN1 233

AMOD 236

AND 120

ANINT 121

ANY 122

area hyperbolic cosine 115

area hyperbolic sine 124

area hyperbolic tangent 129

argument list functions 59

arguments, to program 155, 190, 191, 201

array, add elements 272

array, AND 198

array, apply condition 119, 122

array, bounds checking 25

array, change dimensions 254

array, combine arrays 232

array, condition testing 119, 122

array, conditionally add elements 272

array, conditionally count elements 161

array, conditionally multiply elements 247

array, constructors 40

array, contiguity 210

array, count elements 266

array, duplicate dimensions 268

array, duplicate elements 268

array, element counting 161

array, gather elements 243

array, increase dimension 268, 283

array, indices of type real 54

array, location of maximum element 230

array, location of minimum element 234

array, lower bound 216

array, maximum value 231

array, merge arrays 232

array, minimum value 235

array, multiply elements 247

array, number of elements 161, 266

array, OR 200

array, packing 243

array, parity 209

array, permutation 163

array, product 247

array, reduce dimension 243

array, rotate 163

array, scatter elements 283

array, shape 261

array, shift 171

array, shift circularly 163

array, size	266	BESJN	138
array, sum	272	Bessel function, first kind	137, 138
array, transmogrify	254	Bessel function, second kind	139, 140
array, transpose	280	BESSEL_J0	137
array, unpacking	283	BESSEL_J1	138
array, upper bound	281	BESSEL_JN	138
array, XOR	209	BESSEL_Y0	139
ASCII collating sequence	113, 198	BESSEL_Y1	140
ASIN	123	BESSEL_YN	140
ASIND	123	BESY0	139
ASINH	124	BESY1	140
ASSOCIATED	124	BESYN	140
association status	124	BGE	141
association status, C pointer	144	BGT	141
asynchronous I/O	49	BIAND	199
ATAN	126	BIEOR	204
ATAN2	127	binary representation	245, 246
ATAN2D	128	BIOR	208
ATAND	126	BIT_SIZE	142
ATANH	129	BITEST	143
Atomic subroutine, add	129	bits set	245
Atomic subroutine, ADD with fetch	132	bits, AND of array elements	198
Atomic subroutine, AND	130	bits, clear	201
Atomic subroutine, AND with fetch	133	bits, extract	202
Atomic subroutine, compare and swap	131	bits, get	202
Atomic subroutine, define	131	bits, merge	233
Atomic subroutine, OR	135	bits, move	238, 279
Atomic subroutine, OR with fetch	133	bits, negate	241
Atomic subroutine, reference	136	bits, number of	142
Atomic subroutine, XOR	136	bits, OR of array elements	200
Atomic subroutine, XOR with fetch	134	bits, set	202
ATOMIC_ADD	129	bits, shift	212
ATOMIC_AND	130	bits, shift circular	213
ATOMIC_DEFINE	131	bits, shift left	225, 262
ATOMIC_FETCH_ADD	132	bits, shift right	255, 262
ATOMIC_FETCH_AND	133	bits, testing	143
ATOMIC_FETCH_OR	133	bits, unset	201
ATOMIC_FETCH_XOR	134	bits, XOR of array elements	209
ATOMIC_OR	135	bitwise comparison	141, 142, 143
ATOMIC_REF	136	bitwise logical and	120, 199
ATOMIC_XOR	136	bitwise logical exclusive or	204, 284
Authors	293	bitwise logical not	241
AUTOMATIC	65	bitwise logical or	208, 243
B			
BABS	112	BJTEST	143
backslash	9	BKTEST	143
BACKSPACE	59	BLE	142
backtrace	21, 137	BLT	143
BACKTRACE	137	BMOD	236
base 10 logarithm function	223	BMVBITS	238
BBCLR	201	BNOT	241
BBITS	202	bounds checking	25
BBSET	202	BOZ literal constants	53
BBTEST	143	BSHFT	212
BESJ0	137	BSHFTC	213
BESJ1	138	BTEST	143

C

- C_ASSOCIATED 144
- C_F_POINTER 144
- C_F_PROCPINTER 145
- C_FUNLOC 146
- C_LOC 146
- C_SIZEOF 147
- CABS 112
- calling convention 23
- CARRIAGECONTROL 67
- CCOS 158
- CCOSD 159
- CDABS 112
- CDCOS 158
- CDCOSD 159
- CDEXP 177
- CDLOG 222
- CDSIN 264
- CDSIND 265
- CDSQRT 269
- ceiling 121, 148
- CEILING 148
- CEXP 177
- character kind 258
- character set 9
- CHAR 148
- CHDIR 149
- checking array temporaries 25
- checking subscripts 25
- CHMOD 150
- clock ticks 231, 232, 273
- CLOG 222
- CMPLX 150
- CO_BROADCAST 151
- CO_MAX 152
- CO_MIN 153
- CO_REDUCE 153
- CO_SUM 155
- Coarray, _gfortran_caf_atomic_cas 106
- Coarray, _gfortran_caf_atomic_define 105
- Coarray, _gfortran_caf_atomic_op 107
- Coarray, _gfortran_caf_atomic_ref 106
- Coarray, _gfortran_caf_co_broadcast 107
- Coarray, _gfortran_caf_co_max 108
- Coarray, _gfortran_caf_co_min 108
- Coarray, _gfortran_caf_co_reduce 109
- Coarray, _gfortran_caf_co_sum 109
- Coarray, _gfortran_caf_deregister 93
- Coarray, _gfortran_caf_error_stop 105
- Coarray, _gfortran_caf_error_stop_str 105
- Coarray, _gfortran_caf_event_post 102
- Coarray, _gfortran_caf_event_query 103
- Coarray, _gfortran_caf_event_wait 102
- Coarray, _gfortran_caf_fail_image 105
- Coarray, _gfortran_caf_failed_images 91
- Coarray, _gfortran_caf_finish 90
- Coarray, _gfortran_caf_get 95
- Coarray, _gfortran_caf_get_by_ref 98
- Coarray, _gfortran_caf_image_status 91
- Coarray, _gfortran_caf_init 90
- Coarray, _gfortran_caf_is_present 94
- Coarray, _gfortran_caf_lock 101
- Coarray, _gfortran_caf_num_images 91
- Coarray, _gfortran_caf_register 92
- Coarray, _gfortran_caf_send 94
- Coarray, _gfortran_caf_send_by_ref 97
- Coarray, _gfortran_caf_sendget 96
- Coarray, _gfortran_caf_sendget_by_ref 99
- Coarray, _gfortran_caf_stopped_images 92
- Coarray, _gfortran_caf_sync_all 103
- Coarray, _gfortran_caf_sync_images 104
- Coarray, _gfortran_caf_sync_memory 104
- Coarray, _gfortran_caf_this_image 90
- Coarray, _gfortran_caf_unlock 101
- coarray, IMAGE_INDEX 205
- coarray, lower bound 216
- coarray, NUM_IMAGES 242
- coarray, THIS_IMAGE 276
- coarray, upper bound 282
- Coarrays 87
- coarrays 25
- code generation, conventions 23
- collating sequence, ASCII 113, 198
- Collectives, generic reduction 153
- Collectives, maximal value 152
- Collectives, minimal value 153
- Collectives, sum of values 155
- Collectives, value broadcasting 151
- command line 176
- command options 7
- command-line arguments 155, 190, 191, 201
- command-line arguments, number of 155, 201
- COMMAND_ARGUMENT_COUNT 155
- COMMON 70
- compiler flags inquiry function 156
- compiler, name and version 156
- COMPILER_OPTIONS 156
- COMPILER_VERSION 156
- complex conjugate 157
- Complex function 70
- complex numbers, conversion to 150, 157, 166
- complex numbers, imaginary part 117
- complex numbers, real part 169, 252
- COMPLEX 157
- Conditional compilation 2
- CONJG 157
- consistency, durability 47
- Contributing 293
- Contributors 293
- conversion 17
- conversion, to character 148
- conversion, to complex 150, 157, 166
- conversion, to integer .. 54, 198, 203, 207, 208, 224
- conversion, to logical 54, 224
- conversion, to real 165, 252
- conversion, to string 163

CONVERT specifier	57
core, dump	111
COS	158
COSD	159
COSH	159
cosine	158
cosine, degrees	159
cosine, hyperbolic	159
cosine, hyperbolic, inverse	115
cosine, inverse	114
cosine, inverse, degrees	115
COTAND	160
cotangent	160
cotangent, degrees	160
COTAN	160
COUNT	161
CPP	2, 13
CPU_TIME	162
Credits	293
CSHIFT	163
CSIN	264
CSIND	265
CSQRT	269
CTIME	163
current date	164, 179, 204
current time	164, 179, 214, 277, 278

D

DABS	112
DACOS	114
DACOSD	115
DACOSH	115
DASIN	123
DASIND	123
DASINH	124
DATAN	126
DATAN2	127
DATAN2D	128
DATAND	126
DATANH	129
date, current	164, 179, 204
DATE_AND_TIME	164
DBESJ0	137
DBESJ1	138
DBESJN	138
DBESY0	139
DBESY1	140
DBESYN	140
DBLE	165
DCMLPX	166
DCONJG	157
DCOS	158
DCOSD	159
DCOSH	159
DCOTAN	160
DCOTAND	160
DDIM	167

debugging information options	20
debugging, preprocessor	13
DECODE	69
delayed execution	118, 267
DEXP	177
DFLOAT	252
DGAMMA	189
dialect options	8
DIGITS	166
DIM	167
DIMAG	117
DINT	117
directive, INCLUDE	22
directory, options	22
directory, search paths for inclusion	22
division, modulo	237
division, remainder	236
DLGAMA	223
DLOG	222
DLOG10	223
DMAX1	229
DMIN1	233
DMOD	236
DNINT	121
dot product	167
DOT_PRODUCT	167
DPROD	168
DREAL	169
DSHIFTL	169
DSHIFTR	170
DSIGN	263
DSIN	264
DSIND	265
DSINH	266
DSQRT	269
DTAN	274
DTAND	275
DTANH	276
DTIME	170
dummy argument, unused	19

E

elapsed time	170, 257, 258
Elimination of functions with identical argument lists	29
ENCODE	69
ENUM statement	40
ENUMERATOR statement	40
environment variable	31, 33, 193
EOF	59
EOSHIFT	171
EPSILON	172
ERF	173
ERFC	173
ERFC_SCALED	174
error function	173
error function, complementary	173

error function, complementary,
 exponentially-scaled 174
 errors, limiting 16
 escape characters 9
 ETIME 174
 Euclidean distance 197
 Euclidean vector norm 240
 EVENT_QUERY 175
 Events, EVENT_QUERY 175
 EXECUTE_COMMAND_LINE 176
 EXIT 177
 exponent 69
 exponential function 177
 exponential function, inverse 222, 223
 expression size 147, 267
 EXP 177
 EXPONENT 178
 EXTENDS_TYPE_OF 178
 extensions 51
 extensions, implemented 51
 extensions, not implemented 69
 extra warnings 17

F

f2c calling convention 23, 25
 Factorial function 189
 FDATE 179
 FDL, GNU Free Documentation License 309
 FGET 179
 FGETC 180
 file format, fixed 8, 10
 file format, free 8, 10
 file operation, file number 183
 file operation, flush 183
 file operation, position 186, 188
 file operation, read character 179, 180
 file operation, seek 186
 file operation, write character 184, 185
 file system, access mode 113
 file system, change access mode 150
 file system, create link 220, 272
 file system, file creation mask 282
 file system, file status 188, 225, 270
 file system, hard link 220
 file system, remove file 282
 file system, rename file 253
 file system, soft link 272
 file, symbolic link 48
 file, unformatted sequential 48
 findloc 181
 FINDLOC 181
 flags inquiry function 156
 floating point, exponent 178
 floating point, fraction 186
 floating point, nearest different 239
 floating point, relative spacing 255, 268
 floating point, scale 256

floating point, set exponent 260
 FLOAT 252
 FLOATI 252
 FLOATJ 252
 FLOATK 252
 floor 117, 182
 FLOOR 182
 FLUSH 183
 FLUSH statement 40
 FNUM 183
 form feed whitespace 66
 FORMAT 70
 Fortran 77 3
 FPP 2
 FPUT 184
 FPUTC 185
 FRACTION 186
 FREE 186
 Front-end optimization 30
 FSEEK 186
 FSTAT 188
 FTELL 188
 function elimination 19

G

g77 3
 g77 calling convention 23, 25
 Gamma function 189
 Gamma function, logarithm of 223
 GAMMA 189
 GCC 2
 Generating C prototypes from external procedures
 30
 Generating C prototypes from Fortran BIND(C)
 entities 30
 GERROR 189
 GET_COMMAND 191
 GET_COMMAND_ARGUMENT 191
 GET_ENVIRONMENT_VARIABLE 193
 GETARG 190
 GETCWD 192
 GETENV 193
 GETGID 194
 GETLOG 194
 GETPID 195
 GETUID 195
 GMTIME 196
 GNU Compiler Collection 2
 GNU Fortran command options 7

H

Hollerith constants 55
 HOSTNM 196
 HUGE 197
 hyperbolic cosine 159
 hyperbolic function, cosine 159

hyperbolic function, cosine, inverse.....	115
hyperbolic function, sine.....	266
hyperbolic function, sine, inverse.....	124
hyperbolic function, tangent.....	276
hyperbolic function, tangent, inverse.....	129
hyperbolic sine.....	266
hyperbolic tangent.....	276
HYPOT.....	197

I

I/O item lists.....	53
I/O specifiers.....	67
IABS.....	112
IACHAR.....	198
IALL.....	198
IAND.....	199
IANY.....	200
IARGC.....	201
IBCLR.....	201
IBITS.....	202
IBSET.....	202
ICHAR.....	203
IDATE.....	204
IDIM.....	167
IDINT.....	207
IDNINT.....	240
IEEE, ISNAN.....	214
IEOR.....	204
IERRNO.....	205
IFIX.....	207
IIABS.....	112
IIAND.....	199
IIBCLR.....	201
IIBITS.....	202
IIBSET.....	202
IIEOR.....	204
IIOR.....	208
IISHFT.....	212
IISHFTC.....	213
images, cosubscript to image index conversion.....	205
images, index of this image.....	276
images, number of.....	242
IMAG.....	117
IMAGE_INDEX.....	205
IMAGPART.....	117
IMOD.....	236
IMPORT statement.....	40
IMVBITS.....	238
INCLUDE directive.....	22
inclusion, directory search paths for.....	22
INDEX.....	206
INOT.....	241
input/output, asynchronous.....	49
INT2.....	207
INT8.....	208
integer kind.....	259

Interoperability.....	73
intrinsic.....	19
intrinsic Modules.....	287
intrinsic procedures.....	111
intrinsic, integer.....	64
intrinsic, math.....	66
intrinsic, trigonometric functions.....	66
INT.....	207
inverse hyperbolic cosine.....	115
inverse hyperbolic sine.....	124
inverse hyperbolic tangent.....	129
IOMSG= specifier.....	40
IOR.....	208
IOSTAT, end of file.....	211
IOSTAT, end of record.....	211
IPARITY.....	209
IRAND.....	210
IS_IOSTAT_END.....	211
IS_IOSTAT_EOR.....	210, 211
ISATTY.....	212
ISHFT.....	212
ISHFTC.....	213
ISIGN.....	263
ISNAN.....	214
ISO_FORTRAN_ENV statement.....	40
ITIME.....	214

J

JIABS.....	112
JIAND.....	199
JIBCLR.....	201
JIBITS.....	202
JIBSET.....	202
JIEOR.....	204
JIOR.....	208
JISHFT.....	212
JISHFTC.....	213
JMOD.....	236
JMVBITS.....	238
JNOT.....	241

K

KIABS.....	112
KIAND.....	199
KIBCLR.....	201
KIBITS.....	202
KIBSET.....	202
KIEOR.....	204
KILL.....	215
kind.....	45, 215
kind, character.....	258
kind, integer.....	259
kind, old-style.....	51
kind, real.....	259
KIND.....	215
KIOR.....	208

KISHFT 212
 KISHFTC 213
 KMOD 236
 KMBITS 238
 KNOT 241

L

L2 vector norm 240
 language, dialect options 8
 LBOUND 216
 LCOBOUND 216
 LEADZ 217
 left shift, combined 169
 LEN 217
 LEN_TRIM 218
 lexical comparison of strings 218, 219, 220, 221
 LGAMMA 223
 LGE 218
 LGT 219
 libf2c calling convention 23, 25
 libgfortran initialization, set_args 81
 libgfortran initialization, set_convert 82
 libgfortran initialization, set_fpe 83
 libgfortran initialization, set_max_subrecord_length 84
 libgfortran initialization, set_options 81
 libgfortran initialization, set_record_marker 83
 limits, largest number 197
 limits, smallest number 278
 linking, static 22
 LINK 220
 LLE 220
 LLT 221
 LNBLNK 221
 location of a variable in memory 222
 LOC 66, 222
 LOG_GAMMA 223
 LOG10 223
 logarithm function 222
 logarithm function with base 10 223
 logarithm function, inverse 177
 logical and, bitwise 120, 199
 logical exclusive or, bitwise 204, 284
 logical not, bitwise 241
 logical or, bitwise 208, 243
 logical, bitwise 67
 logical, variable representation 45
 login name 194
 LOG 222
 LOGICAL 224
 LONG 224
 loop interchange, Fortran 30
 loop interchange, warning 17
 LSHIFT 225
 LSTAT 225
 LTIME 226

M

MALLOC 227
 MAP 62
 mask, left justified 227
 mask, right justified 228
 MASKL 227
 MASKR 228
 MATMUL 228
 matrix multiplication 228
 matrix, transpose 280
 MAX 229
 MAX, MIN, NaN 46
 MAX0 229
 MAX1 229
 MAXEXPONENT 229
 maximum value 229, 231
 MAXLOC 230
 MAXVAL 231
 MCLOCK 231
 MCLOCK8 232
 memory checking 25
 MERGE 232
 MERGE_BITS 233
 messages, error 15
 messages, warning 15
 MIN 233
 MINO 233
 MIN1 233
 MINEXPONENT 234
 minimum value 233, 235
 MINLOC 234
 MINVAL 235
 Mixed-language programming 73
 model representation, base 248
 model representation, epsilon 172
 model representation, largest number 197
 model representation, maximum exponent 229
 model representation, minimum exponent 234
 model representation, precision 246
 model representation, radix 248
 model representation, range 251
 model representation, significant digits 166
 model representation, smallest number 278
 module entities 10
 module search path 22
 modulo 237
 MOD 236
 MODULO 237
 MOVE_ALLOC 237
 moving allocation 237
 multiply array elements 247
 MVBITS 238

N

Namelist 52
 NAME 70
 natural logarithm function 222

NEAREST	239
NEW_LINE	239
newline	239
NINT	240
norm, Euclidean	240
NORM2	240
NOSHARED	67
NOT	241
NULL	241
NUM_IMAGES	242

O

open, action	48
OpenACC	10, 58
OpenMP	11, 58
operators, unary	54
operators, xor	67
options inquiry function	156
options, code generation	23
options, debugging	20
options, dialect	8
options, directory search	22
options, errors	15
options, Fortran dialect	8
options, gfortran command	7
options, linking	22
options, negative forms	7
options, preprocessor	13
options, real kind type promotion	12
options, run-time	23
options, runtime	22
options, warnings	15
OR	243
output, newline	239

P

PACK	243
PARAMETER	69
Parity	244
parity	246
PARITY	244
paths, search	22
PERROR	245
pointer checking	25
pointer, C address of pointers	145
pointer, C address of procedures	146
pointer, C association status	144
pointer, convert C to Fortran	144
pointer, Cray	55
pointer, cray	186, 227
pointer, disassociated	241
pointer, status	124, 241
POPCNT	245
POPPAR	246
positive difference	167
PRECISION	246

Preprocessing	2
preprocessing, assertion	14
preprocessing, define macros	15
preprocessing, include path	13, 14
preprocessing, keep comments	14, 15
preprocessing, no linemarkers	15
preprocessing, undefine macros	15
preprocessor	13
preprocessor, debugging	13
preprocessor, disable	13
preprocessor, enable	13
preprocessor, include file handling	2
preprocessor, working directory	13
PRESENT	247
private	10
procedure pointer, convert C to Fortran	146
process ID	195
product, double-precision	168
product, matrix	228
product, vector	167
PRODUCT	247
program termination	177
program termination, with core dump	111
PROTECTED statement	40

Q

Q edit descriptor	70
Q exponent-letter	53

R

radix, real	259
RADIX	248
random number generation	210, 248, 250
random number generation, initialization	249
random number generation, seeding	250, 269
RAND	248
RANDOM_INIT	249
RANDOM_NUMBER	250
RANDOM_SEED	250
range checking	25
RANGE	251
rank	252
RAN	248
RANK	252
re-association of parenthesized expressions	29
read character, stream mode	179, 180
READONLY	67
real kind	259
real number, exponent	178
real number, fraction	186
real number, nearest different	239
real number, relative spacing	255, 268
real number, scale	256
real number, set exponent	260
Reallocate the LHS in assignments	29

Reallocate the LHS in assignments, notification 19

REAL 252

REALPART 252

record marker 48

RECORD 59

Reduction, XOR 244

remainder 236

RENAME 253

repacking arrays 27

REPEAT 254

RESHAPE 254

REWIND 59

right shift, combined 170

root 269

rounding, ceiling 121, 148

rounding, floor 117, 182

rounding, nearest whole number 240

RRSPACING 255

RSHIFT 255

run-time checking 25

S

SAME_TYPE_AS 256

SAVE statement 23

SCALE 256

SCAN 256

search path 22

search paths, for included files 22

SECNDS 257

SECOND 258

seeding a random number generator 250, 269

SELECTED_CHAR_KIND 258

SELECTED_INT_KIND 259

SELECTED_REAL_KIND 259

sequential, unformatted 48

SET_EXPONENT 260

SHAPE 261

SHARE 67

SHARED 67

shift, left 169, 262

shift, right 170, 262

shift, right with fill 262

SHIFTA 262

SHIFTL 262

SHIFTR 262

SHORT 207

sign copying 263

SIGN 263

SIGNAL 264

SIND 265

sine 264

sine, degrees 265

sine, hyperbolic 266

sine, hyperbolic, inverse 124

sine, inverse 123

sine, inverse, degrees 123

SIN 264

SINH 266

size of a variable, in bits 142

size of an expression 147, 267

SIZE 266

SIZEOF 267

SLEEP 267

SNGL 252

SPACING 268

SPREAD 268

SQRT 269

square-root 269

SRAND 269

Standards 4

statement, ENUM 40

statement, ENUMERATOR 40

statement, FLUSH 40

statement, IMPORT 40

statement, ISO_FORTRAN_ENV 40

statement, PROTECTED 40

statement, SAVE 23

statement, USE, INTRINSIC 40

statement, VALUE 40

statement, VOLATILE 40

STAT 270

STATIC 65

storage size 271

STORAGE_SIZE 271

stream mode, read character 179, 180

stream mode, write character 184, 185

STREAM I/O 40

string, adjust left 116

string, adjust right 116

string, comparison 218, 219, 220, 221

string, concatenate 254

string, find missing set 284

string, find non-blank character 221

string, find subset 256

string, find substring 206

string, length 217

string, length, without trailing whitespace 218

string, remove trailing whitespace 280

string, repeat 254

strings, varying length 4

structure packing 27

STRUCTURE 59

subrecord 48

subscript checking 25

substring position 206

sum array elements 272

SUM 272

suppressing warnings 15

symbol names 9

symbol names, transforming 24, 25

symbol names, underscores 24, 25

SYMLNK 272

syntax checking 16

system, error handling 189, 205, 245

system, group ID	194
system, host name	196
system, login name	194
system, process ID	195
system, signal handling	264
system, system call	176, 273
system, terminal	212, 281
system, user ID	195
system, working directory	149, 192
SYSTEM	273
SYSTEM_CLOCK	273

T

tabulators	18
TAND	275
tangent	274
tangent, degrees	275
tangent, hyperbolic	276
tangent, hyperbolic, inverse	129
tangent, inverse	126, 127
tangent, inverse, degrees	126, 128
TAN	274
TANH	276
terminate program	177
terminate program, with core dump	111
THIS_IMAGE	276
thread-safety, threads	46
time, clock ticks	231, 232, 273
time, conversion to GMT info	196
time, conversion to local time info	226
time, conversion to string	163
time, current	164, 179, 214, 277, 278
time, elapsed	162, 170, 174, 257, 258
TIME	277
TIME8	278
TINY	278
TR 15581	40
trace	21
TRAILZ	279
TRANSFER	279
transforming symbol names	24, 25
transpose	280
TRANPOSE	280
trigonometric function, cosine	158
trigonometric function, cosine, degrees	159
trigonometric function, cosine, inverse	114
trigonometric function, cosine, inverse, degrees	115
trigonometric function, cotangent	160
trigonometric function, cotangent, degrees	160
trigonometric function, sine	264
trigonometric function, sine, degrees	265
trigonometric function, sine, inverse	123
trigonometric function, sine, inverse, degrees	123
trigonometric function, tangent	274
trigonometric function, tangent, degrees	275
trigonometric function, tangent, inverse	126, 127

trigonometric function, tangent, inverse, degrees	126, 128
TRIM	280
TTYNAM	281
type alias print	66
type cast	279

U

UBOUND	281
UCOBOUND	282
UMASK	282
underflow	18
underscore	24, 25
unformatted sequential	48
UNION	62
UNLINK	282
UNPACK	283
unused dummy argument	19
unused parameter	19
USE, INTRINSIC statement	40
user id	195

V

VALUE statement	40
variable attributes	65
Varying length character strings	4
Varying length strings	4
vector product	167
VERIFY	284
version of the compiler	156
VOLATILE	70
VOLATILE statement	40

W

warning, C binding type	17
warnings, aliasing	16
warnings, alignment of COMMON blocks	19
warnings, all	16
warnings, ampersand	17
warnings, argument mismatch	17
warnings, array temporaries	17
warnings, character truncation	17
warnings, conversion	17
warnings, division of integers	18
warnings, extra	17
warnings, function elimination	19
warnings, implicit interface	18
warnings, implicit procedure	18
warnings, integer division	18
warnings, interface mismatch	17
warnings, intrinsic	19
warnings, intrinsics of other standards	18
warnings, line truncation	17
warnings, loop interchange	17
warnings, non-standard intrinsics	18

warnings, parameter mismatch 17
warnings, **q** exponent-letter 18
warnings, suppressing 15
warnings, suspicious code 18
warnings, tabs 18
warnings, to errors 20
warnings, undefined do loop 18
warnings, underflow 18
warnings, unused dummy argument 19
warnings, unused parameter 19
warnings, use statements 19
write character, stream mode 184, 185

X

XOR 284
XOR reduction 244

Z

ZABS 112
ZCOS 158
ZCOSD 159
zero bits 217, 279
ZEXP 177
ZLOG 222
ZSIN 264
ZSIND 265
ZSQRT 269

